



Edinburgh 2014/15
 IFUM-1034-FT
 CERN-PH-TH/2013-253
 OUTP-14-11p
 CAVENDISH-HEP-14-11

Parton distributions for the LHC Run II

The NNPDF Collaboration:

Richard D. Ball^{1,2}, Valerio Bertone², Stefano Carrazza^{4,2},
 Christopher S. Deans¹, Luigi Del Debbio¹, Stefano Forte⁴, Alberto Guffanti⁵,
 Nathan P. Hartland¹, José I. Latorre³, Juan Rojo⁶ and Maria Ubiali⁷

¹ *The Higgs Centre for Theoretical Physics, University of Edinburgh,
 JCMB, KB, Mayfield Rd, Edinburgh EH9 3JZ, Scotland*

² *PH Department, TH Unit, CERN,
 CH-1211 Geneva 23, Switzerland*

³ *Departament d'Estructura i Constituents de la Matèria, Universitat de Barcelona,
 Diagonal 647, E-08028 Barcelona, Spain*

⁴ *Dipartimento di Fisica, Università di Milano and INFN, Sezione di Milano,
 Via Celoria 16, I-20133 Milano, Italy*

⁵ *Niels Bohr International Academy and Discovery Center,
 Niels Bohr Institute, University of Copenhagen,
 Blegdamsvej 17, DK-2100 Copenhagen, Denmark*

⁶ *Rudolf Peierls Centre for Theoretical Physics, 1 Keble Road,
 University of Oxford, OX1 3NP Oxford, UK*

⁷ *The Cavendish Laboratory, University of Cambridge,
 J.J. Thomson Avenue, CB3 0HE, UK*

Abstract

We present NNPDF3.0, the first set of parton distribution functions (PDFs) determined with a methodology validated by a closure test. NNPDF3.0 uses a global dataset including HERA-II deep-inelastic inclusive cross-sections, the combined HERA charm data, jet production from ATLAS and CMS, vector boson rapidity and transverse momentum distributions from ATLAS, CMS and LHCb, $W+c$ data from CMS and top quark pair production total cross sections from ATLAS and CMS. Results are based on LO, NLO and NNLO QCD theory and also include electroweak corrections. To validate our methodology, we show that PDFs determined from pseudo-data generated from a known underlying law correctly reproduce the statistical distributions expected on the basis of the assumed experimental uncertainties. This closure test ensures that our methodological uncertainties are negligible in comparison to the generic theoretical and experimental uncertainties of PDF determination. This enables us to determine with confidence PDFs at different perturbative orders and using a variety of experimental datasets ranging from HERA-only up to a global set including the latest LHC results, all using precisely the same validated methodology. We explore some of the phenomenological implications of our results for the upcoming 13 TeV Run of the LHC, in particular for Higgs production cross-sections.

Contents

1	Introduction	4
2	Experimental data	7
2.1	Overview	7
2.2	New experimental data in NNPDF3.0	8
2.3	Theoretical treatment	13
2.3.1	Computational tools	14
2.3.2	Approximate NNLO treatment of jets	15
2.3.3	Electroweak corrections	16
2.3.4	Treatment of heavy quarks	18
2.4	Construction of the dataset	19
2.4.1	Kinematic cuts	19
2.4.2	Treatment of correlated systematic uncertainties	21
3	Improved methodology	23
3.1	Code	23
3.1.1	FastKernel methodology	23
3.1.2	NNPDF++	24
3.2	PDF parametrization	26
3.2.1	Parametrization basis	26
3.2.2	Effective preprocessing exponents	28
3.2.3	Positivity constraints	29
3.3	Minimization algorithm	31
3.3.1	Genetic Algorithms	31
3.3.2	Stopping criterion	33
4	Closure testing	35
4.1	NNPDF closure testing	35
4.2	Validation of the training efficiency: Level 0 closure tests	39
4.2.1	Training methodology efficiency	39
4.2.2	Effective preprocessing exponents	44
4.3	PDF uncertainties in closure tests	45
4.3.1	PDF uncertainties: data, functional and extrapolation components	45
4.3.2	Quantitative validation of PDF uncertainties in closure tests	49
4.4	Validation of the closure test fits	51
4.4.1	Central values	51
4.4.2	PDF uncertainties: qualitative validation	51
4.4.3	PDF uncertainties: quantitative validation	56
4.4.4	Closure test validation using Bayesian reweighting	57
4.5	Robustness of the fitting methodology	60
4.5.1	Independence of the maximum number of GA generations	60
4.5.2	Dependence on the training fraction	60
4.5.3	Parametrization basis independence	62
4.5.4	Redundancy of the neural network architecture	63
4.5.5	Robustness with respect to the choice of input PDF set	64

5	The NNPDF3.0 PDF set	68
5.1	The NNPDF3.0 set of parton distributions	68
5.1.1	Fit quality	68
5.1.2	Parton distributions	70
5.1.3	Perturbative stability and theoretical uncertainties	73
5.2	Dependence on the dataset	78
5.2.1	Conservative PDFs from a consistent dataset	78
5.2.2	Impact of the new HERA and LHC data	82
5.2.3	Impact of jet data on the global fit	90
5.2.4	Nucleon strangeness	91
5.3	Stability	94
5.3.1	Impact of the NNPDF3.0 methodology	94
5.3.2	Constraints from positivity	96
5.3.3	Additive versus multiplicative systematics	97
5.3.4	Independence of preprocessing	99
5.3.5	Independence of the PDF fitting basis	99
5.4	Implications for LHC phenomenology	100
5.4.1	PDF luminosities	101
5.4.2	Implications for $\sqrt{s}=13$ TeV LHC processes	102
5.4.3	Higgs production in gluon fusion	106
5.4.4	New Physics particle production at high masses	107
6	Summary and outlook	111
A	QCD and weak corrections to vector boson production data	115
A.1	CMS double differential distributions	115
A.2	ATLAS high-mass Drell-Yan differential cross section	122
A.3	ATLAS Z rapidity distribution	122
A.4	LHCb Z forward rapidity distribution	123
B	Distance estimators	125

1 Introduction

Parton distribution functions (PDFs) are currently one of the major sources of uncertainty in processes at hadron colliders, specifically at the LHC. Ever since they have been quantified for the first time [1], it has been recognized that PDF uncertainties stem from three sources: the underlying data (which are affected by statistical and systematic errors), the theory which is used to describe them (which is typically based on the truncation of a perturbative expansion) and the procedure which is used to extract the PDFs from the data. This last source of uncertainty is the most elusive, and because of it, it was thought until very recently [2] that PDF uncertainties are not yet understood from a statistical point of view.

In a series of papers [3–13] we have introduced a methodology aimed at reducing as much as possible this procedural uncertainty, up to the point that it can be neglected in comparison to the remaining data and theory uncertainty. However, as data become more abundant and precise, and theoretical calculations more accurate, a full characterization of the procedural uncertainties becomes necessary.

In the present paper we construct for the first time a set of PDFs with an explicit characterization of procedural uncertainties. We do so by using a methodology which has been repeatedly suggested [2], and recently used to study specific sources of PDF uncertainty [14] though never (to the best of our knowledge) for full characterization: the so called “closure test”. The idea is to assume that the underlying PDFs are known, and use this assumed set of PDFs to generate artificial data. One may then set theoretical uncertainties to zero (by using the same theory to generate data and to analyze them), and fix the data uncertainty to any desired value. By determining PDFs from these artificial data, one can then tune the methodology until procedural uncertainties have been removed (or at least made small enough to be undetectable), and the ensuing PDFs faithfully reproduce the data uncertainty. Thus one may explicitly check that the output PDFs provide consistent and unbiased [15] estimators of the underlying truth, that confidence levels reproduce the correct coverage (e.g., that the true value is indeed within the 68% interval in 68% of cases), and so on. Of course, a full closure test requires verifying that the procedure is robust: so, for example, that the conclusion that there are no procedural uncertainties is independent of the assumed underlying set of PDFs.

This program will be fully realized in this paper, whose core is a discussion of the closure test itself. The main motivation for doing so at this time is, as mentioned, that data are more and more abundant and accurate and that more and more higher-order calculations have become available. In this work we will use a very wide dataset, which in particular includes, besides the data used in previous global PDF fits, all the most recent HERA deep-inelastic scattering data for inclusive and charm production cross-sections, essentially all the most recent single-inclusive jet data and W and Z rapidity distributions from the LHC, and also, for the first time in a global fit, $W+c$ production data, double-differential and high-mass Drell-Yan distributions, W transverse momentum distributions, and top total cross section data from ATLAS and CMS. For many of these data, higher-order calculations and/or fast computational interfaces have become available only recently. The dataset on which our PDF determination is based is summarized in Sect. 2, where in particular all new data included on top of those used for the previous NNPDF2.3 set [13] are discussed. The theoretical calculations and tools used for their description and inclusion in the PDF fit are summarized in Sect. 2.3.

The inclusion of these new data and processes in our PDF fitting code has required very substantial computational upgrades, including a migration of the code to C++. Also, whereas the basic principles of our methodology are the same as in our previous PDF determinations, namely, a Monte Carlo approach based on using neural networks as underlying interpolating

functions, several improvements have become necessary in order to ensure the full independence of procedural uncertainties which is required to pass the closure test. This includes a more general way to define the basis PDFs (needed in order to test basis-independence in the closure test); an improved way of making sure that the so-called preprocessing exponents which are introduced as part of the PDF parametrization do not bias the results of the fit; an improved version of the genetic algorithm used for PDF minimization; and an improved form of the cross-validation procedure which is used in order to determine the optimal fit (which in the presence of a redundant parametrization is not the absolute minimum of the figure of merit, since this would require fitting noise). All these methodological aspects are summarized in Sect. 3, where we also summarize the general structure and features of the new NNPDF C++ code.

The closure test procedure is presented in Sect. 4. We actually perform three different kinds of closure test, which we call respectively Level-0, Level-1 and Level-2. In the Level-0 closure test, perfect data are generated from the assumed underlying law, with no uncertainty. The test is successful if a perfect fit to the data can be produced, i.e., a fit with vanishing χ^2 . By showing that our PDFs pass this test we prove that our methodology is not limited by choice of basis functions, functional form, number of parameter, or minimization algorithm. Also, by looking at the remaining uncertainty on the final PDF we can determine what we call the “extrapolation uncertainty”, i.e., the uncertainty due to the fact that datapoints, even when infinitely precise, are not covering the whole of phase space. In a Level-1 closure test, data are generated with an uncertainty: this gives a full pseudodata set, which thus simulates a realistic situation (but with knowledge of the underlying physical law). However, instead of using the NNPDF Monte Carlo methodology, whereby pseudodata replicas are generated about the experimental data, we simply fit PDF replicas all to the same pseudodata. Despite the fact that the data are not fluctuating (i.e. the replicas are all fitted to the same pseudodata set), we find that results are affected by an uncertainty which we call “functional uncertainty”, and which is due to the fact that we are reconstructing a set of functions from a finite amount of discrete information. In a Level-2 closure test, once pseudodata are constructed, they are treated as real data according to the NNPDF methodology. We can then check that the closure test is successful with a wide variety of indicators, which include testing that the PDF replicas behave according to Bayes’ theorem upon the inclusion of new (pseudo)data. We finally check that our results are stable upon a wide variety of changes, such as different assumed underlying PDFs, different choices of basis, different size of the neural network used for parametrization, and so on. We also show how “self-closure” of the NNPDF3.0 set is succesful, by using it as the input of a Level-2 closure test.

Finally, we present the NNPDF3.0 set, which, as usual, we deliver at LO, NLO and NNLO in QCD theory, for a variety of values of α_s , and for a variety of values of the maximum value of active flavor numbers. Also, we provide PDFs determined using various subsets of our full dataset. This includes sets based on HERA data only (which may thus be compared to HERAPDF PDFs [16, 17], and may be used for applications where one wishes to use a PDF set based on a restrictive but consistent dataset); sets based on a “conservative” data subset, which is found by studying a specific statistical indicator introduced previously by us [18] which measures the consistency of each dataset with the remaining data, thereby allowing the removal data which are less compatible with the bulk of the dataset; sets based on the same data as our previous NNPDF2.3 dataset, thereby enabling us to separate the impact of new data from that of improved methodology when comparing NNPDF2.3 to NNPDF3.0; a set without LHC data, which is useful in order to gauge the impact of the latter; and sets where the HERA data are supplemented by all available data from either ATLAS or CMS, useful to compare with related studies presented by the LHC collaborations. All these datasets and their features are presented

and compared in Sect. 5, where we also discuss in light of our results specific issues of particular recent relevance in PDF determination: jet data and their impact on the gluon distribution, and the strange PDF which follows from both deep-inelastic and hadronic scattering data.

Finally, in Sect. 5.4 we briefly discuss the implications of the NNPDF3.0 set for LHC phenomenology, with special regard to Higgs production in gluon fusion, and a number of other representative processes, for a center-of-mass energy of 13 TeV. Information on delivery, usage, and future developments is collected in Sect. 6, while some technical results are relegated to appendices: in Appendix A we provide more details for the computation of QCD and electroweak corrections to LHC Drell-Yan data and in Appendix B we present the definitions of the distance estimators we use to compare between different sets of PDFs.

2 Experimental data

Here we present the new experimental data that have been added in the NNPDF3.0 analysis, and their treatment. After reviewing the NNPDF3.0 dataset, which includes most of the data on which the previous NNPDF2.3 set was based, we discuss each of the new datasets in turn. We then review our theoretical treatment, and specifically discuss the computational tools used to implement perturbative corrections, and particular issues related to NNLO QCD corrections to jet production, electroweak corrections, and the treatment of heavy quark mass effects. We finally review the way the total dataset is constructed, and in particular which data and which cuts are included at each perturbative order, and the treatment of experimental uncertainties, specifically the systematics.

2.1 Overview

The NNPDF2.3 fit [13] included fixed-target deep inelastic scattering (DIS) data from NMC [19, 20], BCDMS [21, 22] and SLAC [23]; the combined HERA-I DIS data set [24], HERA F_L [25] and the separated ZEUS and H1 F_2^c structure function data [26–32], some ZEUS HERA-II DIS cross-sections [33, 34]; CHORUS inclusive neutrino DIS [35], and NuTeV dimuon production data [36, 37]; fixed-target E605 [38] and E866 [39–41] Drell-Yan production data; CDF W asymmetry [42] and CDF [43] and D0 [44] Z rapidity distributions; CDF [45] and D0 [46] Run-II one-jet inclusive cross-sections; ATLAS [47], CMS [48] and LHCb [49] data on vector boson production; and ATLAS one-jet inclusive cross-sections [50] from the 2010 run.

In NNPDF3.0, several new datasets have been included. First, we have included all relevant inclusive cross-section measurements from HERA Run II [51–54]. This includes the complete set of inclusive measurements from H1 and ZEUS, as well as new H1 data at low Q^2 and high- y . In addition, we have replaced the separate H1 and ZEUS F_2^c data with the combined measurements of the charm production cross section σ_{cc}^{red} [55]. The separate HERA-II H1 and ZEUS data, as well as the inclusive HERA-I data, will be replaced with the combined HERA dataset as soon as it becomes available. However, we expect this replacement to have a small impact on the PDFs, since the neural network fit effectively performs a dataset combination. This was explicitly demonstrated for the combined HERA-I dataset in the NNPDF2.0 analysis [10].

We have included new vector boson production data from the LHC at a center-of-mass energy of 7 TeV from the 2010 and 2011 runs: the 5 fb $^{-1}$ high-mass Drell-Yan production [56] and the 31 pb $^{-1}$ transverse momentum distribution of W bosons data [57] from ATLAS; the 5 fb $^{-1}$ W muon asymmetry [58], the double differential Drell-Yan cross sections [59], and the differential distributions for W production in association with charm quarks from CMS [60]; and finally the 940 pb $^{-1}$ forward $Z \rightarrow ee$ production data from LHCb from the 2011 run [61]. In comparison to the NNPDF2.3 dataset, we have removed the CDF W asymmetry data [42], since more precise and cleaner data from the LHC (based on leptons rather than on the reconstructed W) covers the same region in x and extends their kinematical coverage.

Concerning LHC jet measurements, we have complemented the previously available ATLAS 7 TeV inclusive jet production data from the 2010 run [50] with the 7 TeV CMS 5 fb $^{-1}$ inclusive jet cross-sections from the 2011 run [62], and the ATLAS 2.76 TeV jet cross section data [63], supplemented by their correlations with the 7 TeV cross sections. In comparison to NNPDF2.3, we have removed the Tevatron D0 Run II inclusive jet cross-section measurements [46], which were obtained with the midpoint algorithm, which is infrared unsafe [64] and therefore cannot be used in conjunction with NNLO calculations. On the other hand, the CDF Run II inclusive jet measurements, based on the k_t algorithm, are retained in NNPDF3.0.

Finally, we have included total cross-section measurements for top quark pair production data from ATLAS and CMS at 7 and 8 TeV [65–70], for a total of six data points.

Besides extending the traditional dataset used in our previous global fit, NNPDF2.3, based on deep-inelastic scattering, Drell-Yan, and jet data, to include new information from HERA and the LHC (whose kinematic coverage and resolution is greatly expanded, specifically through the inclusive jet production and the high-mass and double-differential gauge boson production), the NNPDF3.0 dataset uses for the first time in a global fit many datasets whose importance has been repeatedly emphasized for constraining particular PDF combinations [71, 72]: correlated pairs of measurements at different center-of-mass energies (the ATLAS 2.76 and 7 TeV jet measurements) which provide a handle on PDFs due to the reduced experimental and theoretical systematics of cross-section ratios [73]; $W+c$ production which provides direct information on strangeness [74, 75]; top quark pair production which provides a handle on the gluon at large- x [76–78]; and the W p_T distribution which also helps to constrain the gluon and sea quarks in a wide range of x [79]. The inclusion of these data is especially powerful thanks to recent theoretical progress, specifically the completion of the computation of the NNLO corrections to the total cross-section for top quark pair production [80, 81] and the computation of NNLO corrections to jet production in the gluon channel [82].

For ease of reference, in Tables 1 (for DIS data) and 2 (for hadronic data) we summarize the complete list of data included in the NNPDF3.0 analysis. For each dataset we provide the corresponding published reference, the availability and treatment of systematics (to be discussed in Sect. 2.4.2 below), the number of data points before and after cuts at NLO and NNLO (to be discussed in Sect. 2.4.1 below). For the new data which were not in NNPDF2.3 we also give the kinematic coverage (which also flags the new datasets included for the first time in NNPDF3.0). Finally, in Tables 3 we also summarize the data which were in NNPDF2.3, but are no longer used in NNPDF3.0, for the reasons described above.

The kinematical coverage in the (x, Q^2) plane of the NNPDF3.0 dataset is shown in the scatter plot Fig. 1 (note that for hadronic data leading-order kinematics has been assumed for illustrative purposes, as discussed in [10]).

2.2 New experimental data in NNPDF3.0

We now examine in more detail the features of the new experimental data that has been included in NNPDF3.0. We begin with the new deep-inelastic scattering data. In this work we have supplemented the combined HERA-I dataset with the inclusion of all the relevant HERA-II inclusive cross-section measurements from H1 and ZEUS [51–54]. These data improve the statistical and systematic precision of medium- and high- Q^2 data from the HERA-I run, and thus provide valuable information for quarks at medium and large x . In addition, we have included low- Q^2 data from H1 that provides additional information on the small- x gluon and that could also be sensitive to potential deviations from fixed-order DGLAP at small- x .

From the H1 experiment, we have included the new high- Q^2 data from the HERA-II run [51], which covers the large Q^2 region, $60 \leq Q^2 \leq 5 \cdot 10^4 \text{ GeV}^2$, with improved statistical and systematic precision in comparison to Run-I. These data, taken at the default proton beam energy of $E_p = 920 \text{ GeV}$ used in most of the HERA-II run, have been supplemented with inclusive cross-section measurements performed at lower center-of-mass energies [52], obtained with proton beam energies of $E_p = 575 \text{ GeV}$ and $E_p = 460 \text{ GeV}$. These lower-energy measurements are the same ones used to determine the longitudinal structure function F_L . Therefore, we exclude the previously used F_L [25] data from the present fit to avoid any double counting, and for the same reason we do not include any other of the updated F_L extractions at HERA [84, 85]. For

Experiment	Dataset	Ref.	Sys. Unc.			N_{dat} no cuts (NLO/NNLO cuts)	Kinematics
NMC	NMC d/p	[19]	A	full		289 (132/132)	
	NMC $\sigma^{\text{NC},p}$	[20]	A	full		211 (224/224)	
SLAC	SLAC p	[23]	A	none	a	191 (37/37)	
	SLAC d	[23]	A	none	a	191 (37/37)	
BCDMS	BCDMS p	[21]	A	full	b	351 (333/333)	
	BCDMS d	[22]	A	full	b	254 (248/248)	
CHORUS	CHORUS ν	[35]	A	full	c	572 (431/431)	
	CHORUS $\bar{\nu}$	[35]	A	full	c	572 (431/431)	
NuTeV	NuTeV ν	[36, 37]	A	none		45 (41/41)	
	NuTeV $\bar{\nu}$	[36, 37]	A	none		44 (38/38)	
HERA-I	HERA-I NC e^+	[24]	M	full	d	434 (379/379)	
	HERA-I NC e^-	[24]	M	full	d	145 (145/145)	
	HERA-I CC e^+	[24]	M	full	d	34 (34/34)	
	HERA-I CC e^-	[24]	M	full	d	34 (34/34)	
ZEUS HERA-II	ZEUS-II NC e^-	[33]	M	full	e	90 (90/90)	
	ZEUS-II CC e^-	[34]	M	full	e	37 (37/37)	
	ZEUS-II NC e^+	[53]	M	full	f	90 (90/90)	$5 \cdot 10^{-3} \leq x \leq 0.40$
	ZEUS-II CC e^+	[54]	M	full	f	35 (35/35)	$200 \leq Q^2 \leq 3 \cdot 10^4 \text{ GeV}^2$ $7.8 \cdot 10^{-3} \leq x \leq 0.42$ $280 \leq Q^2 \leq 3 \cdot 10^4 \text{ GeV}^2$
H1 HERA-II	H1-II NC e^-	[51]	M	full	g	139 (139/139)	$2 \cdot 10^{-3} \leq x \leq 0.65$
	H1-II NC e^+	[51]	M	full	g	138 (138/138)	$120 \leq Q^2 \leq 4 \cdot 10^4 \text{ GeV}^2$ $2 \cdot 10^{-3} \leq x \leq 0.65$
	H1-II CC e^-	[51]	M	full	g	29 (29/29)	$120 \leq Q^2 \leq 4 \cdot 10^4 \text{ GeV}^2$ $8 \cdot 10^{-3} \leq x \leq 0.40$
	H1-II CC e^+	[51]	M	full	g	29 (29/29)	$300 \leq Q^2 \leq 3 \cdot 10^4 \text{ GeV}^2$ $8 \cdot 10^{-3} \leq x \leq 0.40$
	H1-II low Q^2	[52]	M	full		136 (124/124)	$300 \leq Q^2 \leq 3 \cdot 10^4 \text{ GeV}^2$ $2.8 \cdot 10^{-5} \leq x \leq 0.015$
	H1-II high y	[52]	M	full		55 (52/52)	$1.5 \leq Q^2 \leq 90 \text{ GeV}^2$ $2.9 \cdot 10^{-5} \leq x \leq 5 \cdot 10^{-3}$ $2.5 \leq Q^2 \leq 90 \text{ GeV}^2$
HERA σ_{NC}^e		[55]	M	full		52 (47/47)	$3 \cdot 10^{-5} \leq x \leq 0.05$ $2.5 \leq Q^2 \leq 2 \cdot 10^3 \text{ GeV}^2$

Table 1: List of all the deep-inelastic scattering data included in the NNPDF3.0 analysis. For each dataset we provide the corresponding reference in the third column. The next three columns provide information on the treatment of systematic uncertainties for each dataset. The fourth column specifies the treatment of systematic uncertainties, where M stands for multiplicative and A for additive. The fifth column states how the available experimental information on correlated uncertainties other than normalization uncertainties is provided by the experiment: “none” corresponds to datasets for which only the sum in quadrature of systematic uncertainties is provided, “cov” in case the covariance matrix is provided and “full” is the full breakup of systematics is provided. The sixth column gives information on whether datasets share systematics uncertainties among them: datasets that are marked with the same letter have common correlated systematics; see text for more details. Then, for each dataset, we provide the total number of data points N_{dat} available, as well as the number of data points left after kinematical cuts both in the NLO and in the NNLO fits (in parenthesis). For the experiments which are new in NNPDF3.0, we also include the information on their kinematical coverage in the last column.

Experiment	Dataset	Ref.	Sys. Unc.			N_{dat} no cuts (NLO/NNLO cuts)	Kinematics
DY E866	DY E866 d/p	[41]	M	none		15 (15/15)	
	DY E866 p	[39, 40]	M	none		184 (184/184)	
DY E605		[38]	M	none		119 (119/119)	
CDF	CDF Z rap	[43]	M	full	h	29 (29/29)	
	CDF Run-II k_t jets	[83]	M	full	h	76 (76/52)	
D0	D0 Z rap	[44]	M	full		28 (28/28)	
ATLAS	ATLAS W, Z 2010	[47]	M	full	i	30 (30/30)	
	ATLAS 7 TeV jets 2010	[50]	M	full	i,j	90 (90/9)	
	ATLAS 2.76 TeV jets	[63]	M	full	j	59 (59/3)	$20 \leq p_T^{\text{jet}} \leq 200$ GeV $0 \leq \eta^{\text{jet}} \leq 4.4$
	ATLAS high-mass DY	[56]	M	full		11 (5/5)	$116 \leq M_{ll} \leq 1500$ GeV
	ATLAS W p_T	[57]	M	full		11 (9/-)	$0 \leq p_T^W \leq 300$ GeV
CMS	CMS W electron asy	[48]	M	cov		11 (11/11)	
	CMS W muon asy	[58]	M	cov		11 (11/11)	$0 \leq \eta_l \leq 2.4$
	CMS jets 2011	[62]	M	full		133 (133/83)	$114 \leq p_T^{\text{jet}} \leq 2116$ GeV $0 \leq \eta^{\text{jet}} \leq 2.5$
	CMS $W + c$ total	[60]	M	cov		5 (5/5)	$0 \leq \eta_l \leq 2.1$
	CMS $W + c$ ratio	[60]	M	cov		5 (5/5)	$0 \leq \eta_l \leq 2.1$
	CMS 2D DY 2011	[59]	M	cov		124 (88/110)	$20 \leq M_{ll} \leq 1200$ GeV $0 \leq \eta_{ll} \leq 2.4$
LHCb	LHCb W rapidity	[49]	M	cov		10 (10/10)	
	LHCb Z rapidity	[61]	M	cov		9 (9/9)	$2.0 \leq \eta \leq 4.5$
$\sigma(t\bar{t})$	ATLAS $\sigma(t\bar{t})$	[65–67]	M	none		3 (3/3)	-
	CMS $\sigma(t\bar{t})$	[68–70]	M	none		3 (3/3)	-
Total						5179 (4276/4078)	

Table 2: Same as Table 1 for fixed-target Drell-Yan production, electroweak vector boson production data from the Tevatron and LHC data. Again we explicitly provide the kinematics only for the new experiments. The $t\bar{t}$ cross-sections are new in NNPDF3.0, but being total cross-sections no information on kinematics needs to be provided. In the bottom row of the table we give the number of data points in the global fit dataset, including also the DIS numbers from Table 1.

completeness, in NNPDF3.0 we also include the the high-inelasticity data that H1 extracted from their Run II measurements [52].

From the ZEUS experiment, in NNPDF2.3 we already included some HERA-II data [33, 34], for neutral- and charged-current DIS with an electron beam. We now also include neutral- and charged-current cross sections with a positron beam [54, 86], which have been published since. As in the case of H1, ZEUS Run II inclusive cross-sections exhibit reduced statistical and systematic uncertainties when compared to Run I data in the medium and large Q^2 region. For both H1 and ZEUS, we use the data averaged over lepton beam polarizations.

Turning to semi-inclusive measurements, in NNPDF3.0 we have replaced the separate charm structure function data F_2^c from the H1 and ZEUS experiments [26–32] with the combined HERA charm production cross-section dataset [55], as the cross-calibration between common systematics implies that the combined data is more accurate than the separate inclusion of the individual measurements. The combined HERA charm production cross-sections offer a handle

Experiment	Dataset	Ref.	Sys. Unc.	N_{dat} no cuts	Details
H1 F_2^c	H1 F_2^c 01	[30]	full	12	Superseded by combination
	H1 F_2^c 09	[31]	full	6	"
	H1 F_2^c 10	[32]	full	26	"
ZEUS F_2^c	ZEUS F_2^c 99	[26]	full	21	Superseded by combination
	ZEUS F_2^c 03	[27]	full	31	"
	ZEUS F_2^c 08	[28]	full	9	"
	ZEUS F_2^c 09	[29]	full	8	"
CDF	CDF W asymmetry	[42]	full	13	Lepton-level data from LHC
D0	D0 Run II cone jets	[46]	full	110	Infrared unsafe at NNLO

Table 3: Same as Table 1, this time for those experiments that were present in NNPDF2.3 but that have been excluded from NNPDF3.0. In the last column we provide information on why these experiments are not included anymore in NNPDF3.0.

on the small- x gluon [87], provide a unique testing ground for different treatments of heavy quark mass effects, and allow one to extract the running charm quark mass $m_c(m_c)$ with competitive uncertainties [88, 89].

Concerning vector boson production at the LHC, NNPDF2.3 already included some data taken at 7 TeV, in particular the CMS W electron asymmetry data from an integrated luminosity of 880 pb^{-1} [48], the ATLAS W^\pm and Z rapidity distributions from the 2010 run [47] and the LHCb forward W production data also from 2010 [49]. These datasets have now been supplemented with essentially all vector boson production data which have become available since.

From the ATLAS experiment, we include the high-mass Drell-Yan production data from the 2011 run [56], based on an integrated luminosity of 4.9 fb^{-1} . These data, presented in terms of the invariant mass of the electron pairs produced at an invariant mass larger than the Z peak, which extends to $M_{ll} = 1.5 \text{ TeV}$, can be used to constrain the large- x antiquarks. In addition, it was shown in Ref. [90] that high-mass Drell-Yan at the LHC can be used to constrain the photon PDFs of the proton, and it was indeed used there to construct the NNPDF2.3QED PDF set. We also include now the ATLAS measurement of the W boson transverse momentum distribution from the 2010 run of the LHC at $\sqrt{s} = 7 \text{ TeV}$ [57], corresponding to an integrated luminosity of 31 pb^{-1} . This data, unlike the 7 TeV ATLAS measurement of the Z boson transverse momentum distribution [91], is provided with all information on correlated uncertainties and has the potential to constrain the gluon and the light quark distributions in the medium- x region [79].

From the CMS experiment, we include the W muon asymmetry data based on the full statistics (5 fb^{-1}) of the 7 TeV run [58], which have substantially reduced statistical and systematic uncertainties. We also include the double-differential distributions for Drell-Yan production from low dilepton masses (from $M_{ll} = 20 \text{ GeV}$) to high dilepton masses (up to $M_{ll} = 1.5 \text{ TeV}$) in bins of the dilepton invariant mass and rapidity, from the full 2011 dataset [59]. The CMS data for the production of charm quarks associated to W bosons [60], which we also include, provide the absolute cross sections, differential on the lepton rapidity from the W decay η_l , as well as cross-section ratios W^+/W^- also binned in η_l . They are both included, as the former constrain the shape and overall normalization of the total strangeness $s + \bar{s}$ at $Q \sim M_W$ and the latter offer some handle on the strangeness asymmetry in the proton, $s - \bar{s}$. Data from this

NNPDF3.0 NLO dataset




Figure 1: The kinematical coverage in the (x, Q^2) plane of the NNPDF3.0 dataset. For hadronic data, leading-order kinematics have been assumed for illustrative purposes. The green stars mark the data already included in NNPDF2.3, while the circles correspond to experiments that are novel in NNPDF3.0.

same process are available from the ATLAS Collaboration [92], but are given at the hadron level and thus cannot be directly included in our fit (though they could be included by for example estimating a hadron-to-parton correction factor using MADGRAPH5_AMC@NLO).

Finally, we include the LHCb $Z \rightarrow ee$ rapidity distributions from the 2011 dataset [61], which are more precise than the previous data from the 2010 run. The forward kinematics of this data provide constraints on PDFs at smaller and large values of x than the vector boson production data from ATLAS and CMS. Further LHCb data from the 2011 run for Z boson rapidity distributions in the $\mu\mu$ channel [93] and for low mass Drell-Yan production [94] are still preliminary.

Concerning inclusive jet production from ATLAS and CMS, we include the CMS inclusive jet production measurement at 7 TeV from the full 5 fb^{-1} dataset [62], which has been provided with the full experimental covariance matrix, and which supersedes previous inclusive jet measurements from CMS [95]. This data has a large kinematical coverage: for example, in the central rapidity region, the CMS data reaches up to jet transverse momenta of more than 2 TeV, thus constraining the large- x quark and gluon PDFs [96,97]. From ATLAS, we include the new inclusive cross section measurement at $\sqrt{s} = 2.76 \text{ TeV}$ [63], which is provided with the full correlation matrix with the corresponding $\sqrt{s} = 7 \text{ TeV}$ measurement. Measuring the ratio of jet cross-sections at two different center of mass energies enhances the PDF sensitivity thanks to the

partial cancellation of theoretical (missing higher order corrections) and experimental (jet energy scale) systematic uncertainties [73]. On the other hand no LHC dijet data are included [98], since it is notoriously difficult to achieve a good description of these measurements [96].

Finally, we include six independent measurements of the total top quark pair production cross section from ATLAS and CMS, both at 7 TeV and at 8 TeV. These data provides information on the large- x gluon PDF, complementary to that provided by inclusive jet production. At 7 TeV we include the measurements in the dilepton channel, based on 0.70 fb^{-1} integrated luminosity from ATLAS [65] and on 2.3 fb^{-1} from CMS [68] and the measurements performed using lepton+jets events from ATLAS [66] and CMS [69]. At 8 TeV we have included the dilepton channel measurement corresponding to an integrated luminosity of 2.4 fb^{-1} by CMS [99] and the ATLAS analysis of the lepton+jets final state in a dataset corresponding to an integrated luminosity of 5.8 fb^{-1} [67].

2.3 Theoretical treatment

NNPDF3.0 PDFs are provided at LO, NLO and NNLO in perturbative QCD. While for most of the observables included in the fit NNLO QCD corrections are known, some observables are known only up to NLO, while for others only partial contributions to the full NNLO corrections have been calculated. Specifically, NNLO corrections are not available for two datasets included in our fit: the vector boson transverse momentum distribution and the $W+c$ rapidity distribution (since there are no NNLO calculations for V +jets and V +heavy quarks). For the jet inclusive cross section, only the gg -channel has been recently computed at NNLO [82, 100], while for the full cross section only an approximate NNLO prediction based on threshold resummation is available [101]. For all other observables included in the fit the cross sections are known up to NNLO.

The theoretical predictions for DIS observables have been implemented in the FASTKERNEL framework and thoroughly benchmarked [12, 102]. Drell-Yan cross sections, both for fixed target and for collider experiments, are computed at NNLO by using the local C -factors computed according to the procedure described in Ref. [12]: we define C -factors as the ratio of NNLO to NLO calculations for fixed NNLO PDFs, that is

$$C^{\text{nnlo}} \equiv \frac{\hat{\sigma}^{\text{nnlo}} \otimes \mathcal{L}^{\text{nlo}}}{\hat{\sigma}^{\text{nlo}} \otimes \mathcal{L}^{\text{nnlo}}} , \quad (1)$$

where $\hat{\sigma}$ is the partonic cross-section computed at either NNLO or NLO accuracy, and $\mathcal{L}^{\text{nnlo}}$ is the corresponding parton luminosity computed with a reference set of NNLO parton distributions. The tools used to compute the C -factors are described in Sect. 2.3.1. For a detailed description of the procedure that we adopt to include a subset of the inclusive jet data in our analysis, see Sect. 2.3.2.

Given that electroweak corrections can be relevant in the large invariant mass region covered by some of the experimental data included in our fit (see Ref. [103] for an overview), we provide EW corrections for all LHC vector boson production data. To include these corrections in our NLO and NNLO calculation, we compute C -factors, C^{ew} , defined analogously to Eq. (1), with the NNLO computation substituted by the NLO+EW one, and using NLO parton luminosities on both numerator and denominator. All details on their computation and implementation are provided in Sect. 2.3.3. The QCD NNLO and EW C -factors for LHC gauge boson production are listed in Appendix A.

In this work, we do not include all-order perturbative resummation of QCD corrections: these will be the object of a future separate study, possibly leading to the construction of dedicated

resummed sets. Also, we do not include nuclear corrections (which would be relevant for fixed-target deuterium DIS data and neutrino DIS data), given the substantial uncertainties in their modeling and the fact that dedicated studies show that their effect is small on the scale of the data uncertainties [104].

The way our dataset is constructed, and in particular which data are included at various perturbative orders, is discussed in Sect. 2.4.1 below.

2.3.1 Computational tools

The inclusion of perturbative corrections to hadronic processes in PDF fits requires the fast computation of the relevant cross-sections. Several fast interfaces have been developed to this purpose, including APPLGRID [105], which provides an interface to MCFM [106,107] and NLOJET++ [108], and FASTNLO [109,110], also interfaced to NLOJET++. Recently, a new fast interface has become available, namely AMCFast [111], a code which provides the complete automation of fast NLO QCD calculations for PDF fits, interfaced to MADGRAPH5_AMC@NLO [112], which in turn achieves the complete automation of the computations of tree-level and next-to-leading order cross sections and of their matching to parton shower simulations.

These tools have been used extensively in the present analysis. For the 7 TeV CMS jet data, we have used the FASTNLO calculation with central scales $\mu_F = \mu_R = p_T^{\text{jet}}$. For the 2.76 ATLAS jet data, we have used instead the APPLGRID calculation, which uses the same scale choices. For consistency, we use exactly the same settings of the calculation, including the central scales, as were used in the corresponding ATLAS 7 TeV inclusive jet analysis. The CDF Run II k_t jets have also been computed using the FASTNLO calculation with the same central scales as in the ATLAS and CMS case.

For electroweak vector boson production data, we have used the APPLGRID code interfaced to MCFM6.6 in all cases, with a consistent choice of electroweak parameters. We use the G_μ scheme, with $M_Z = 91.1876$ GeV, $M_W = 80.398$ GeV and $G_F = 1.16637 \cdot 10^{-5}$ GeV⁻² set as input parameters and α_e , $\sin\theta_W$ derived from those. We turn off the Narrow-Width approximation. For all rapidity distributions, we set $\mu_F = \mu_R = M_V$, with $V = W, Z$. For the W p_T distribution we set $\mu_F = \mu_R = M_W$, while in the case of CMS double-differential distribution we set the scales to the central value of the invariant mass bin. The MCFM6.6 calculations have been cross-checked with independent calculations of DYNLO [113–116] and FEWZ3.1 [117,118] at NLO, with the same settings, finding perfect agreement in all cases.

In the NNLO fits, the NNLO C -factors defined in Eq. (1) have been computed with FEWZ3.1 and cross-checked against DYNLO1.3. Very high statistics runs of all these codes have been necessary to achieve negligible integration error in all the data bins. In order to obtain smooth C -factors, the NNLO curves are smoothed with a high-degree polynomial interpolation. Notice that the difference between smoothed and original NNLO predictions is always within the Monte Carlo uncertainty of the code used to compute it. It turns out that these NNLO QCD corrections are sizable, especially for small invariant masses of the produced lepton pairs as it is shown on the right plot in Fig. 3: the NNLO C -factor for the CMS double differential Drell-Yan data at low M_{ll} is around 10%, independent of the dilepton rapidity. NNLO corrections are also important for the ATLAS high-mass Drell-Yan data, reaching almost 10% around $M_{ll} \sim 1$ TeV, see left plot of Fig. 3. Details of the computation of the EW corrections and their size are given below in Sect. 2.3.3.

For the computation of top quark pair production data at NLO, we again used APPLGRID interfaced to MCFM6.6. The NNLO C -factors have been computed by using the NNLO calculation of Ref. [81], as implemented in the TOP++ code [119]. Finally, we have used AMCFast in-

terfaced to MADGRAPH5_AMC@NLO to compute the Higgs rapidity distributions in gluon fusion at NLO with an unphysical boson of mass $m_h = \sqrt{5}$ GeV. As explained below in Sect. 3.2.1, this unphysical calculation has been used to enforce the positivity of cross-sections that depend on the small- x gluon. More details about Higgs production at NLO in the MADGRAPH5_AMC@NLO framework can be found in in Refs. [120, 121].

2.3.2 Approximate NNLO treatment of jets

NNLO corrections to jet production in the the gg -channel have become available recently [82, 100]. An approximate NNLO prediction based on threshold resummation (but including all partonic subchannels) was presented in [101]. Its accuracy can now be gauged against the existing part of the full calculation.

This was done recently in a systematic study [122], to which we refer for a more detailed treatment. An important consequence of the analysis of Ref. [122] is that some of the comparisons between the exact computations and the threshold approximation performed hitherto [101] were marred by the fact that the renormalization scale was set equal to the jet transverse momentum p_T in the threshold calculation, while the transverse momentum of the hardest jet p_T^{lead} was used instead for the exact calculation. It turns out that the perturbative expansion is rather better behaved, with smaller perturbative corrections, when the scale p_T is used, and also, that the regions of agreement are rather different if a consistent scale choice is made: specifically, when using p_T as a scale in the large rapidity region the accuracy of the threshold approximation appears to be very poor (while there is a better, accidental agreement when an inconsistent scale choice is made).

The percentage deviation between the exact and approximate calculation (with the renormalization scales set equal to p_T in the two cases) is shown in Fig. 2 in the (p_T, y) plane for the CMS 2011 [62], ATLAS 7 and 2.76 TeV [50, 63] and CDF Run II k_t [83] inclusive jet data included in our fit. A common NNLO reference PDF set, NNPDF2.3, is used in the two calculations. In Fig. 2, each of the entries of the contour plot correspond to one of the experimental data bins. It is apparent that the threshold approximation and the exact results are in reasonable agreement in the high- p_T and central- y regions. Note also that the disagreement between the two calculations can become very large for small transverse momenta and/or large rapidities.

In the kinematical region covered by the Tevatron and LHC jet cross-sections, see Fig. 2, the gluon channel is not necessarily dominant, however both theoretical arguments and the known NLO behavior suggest that the pattern of agreement or disagreement of the threshold and exact computations is similar in all partonic channels [122]. This result then suggests that the threshold approximation can be used provided all channels are included, and a sufficiently conservative cut is applied in order to keep only the regions in which the size of the deviation between the exact and threshold computations in the gg channel does not exceed a suitable low threshold. Reassuringly, these are also the regions in which NNLO corrections are not too large. For example, in the region in which the exact and threshold NNLO calculations in the gluon channel differ by less than 10%, the C -factor Eq. (1) (summing over all parton channels) is typically of order 15%, so the expected accuracy of the threshold approximation at NNLO is at the level of a few percent, smaller than the experimental uncertainties.

In NNPDF3.0, we thus follow the strategy of Ref. [122] and compute approximate NNLO C -factors, Eq. (1), using the threshold calculation, while restricting the fitted dataset to the region where, thanks to the comparison with the exact gg calculation, we know the former to be reliable. This leads to the set of cuts outlined below in Sect. 2.4.




Figure 2: Percentage difference $|\delta|$ between the exact and approximate gg -channel NNLO C -factors as a function of p_T and $|y|$ for the CMS, ATLAS 7 TeV and 2.76 TeV and CDF data that is included in NNPDF3.0. Each entry in the contour plot correspond to one of the experimental bins of the corresponding measurement. Differences larger than $|\delta| = 100\%$ are shown in red.

2.3.3 Electroweak corrections

Electroweak corrections, though generally small, may become large at high scales $Q^2 \gg M_V^2$. While this will certainly be an issue for future LHC data at higher center of mass energy, already for some high-mass data included in our analysis the high accuracy of the experimental measurements may require theoretical predictions at the percent level of precision, and the size of the EW corrections needs to be carefully assessed.

The NLO EW one-loop corrections are known [123–128] and have been implemented in several public codes such as HORACE [123] and ZGRAD2 [127, 128]. In FEWZ3.1 [117, 118] the NLO EW corrections are combined to the NNLO QCD corrections, using the complex mass scheme. This code allows the user to separate a gauge-invariant QED subset of the corrections from the full EW result. The QED subset includes initial-state QED radiation, final-state QED radiation (FSR) and the initial-final interference terms. The former cannot be consistently included without taking into account the photon PDF $\gamma(x, Q^2)$. Within the current uncertainties that affect the photon PDF, as determined in Ref. [90], initial-state QED corrections are compatible with zero for most of the data included in the analysis. In case they may be larger, we exclude the data from the fit, as explained in Sect. 2.4.1. As far as final state radiation (FSR) is concerned, when available, we use data from which it has already subtracted. These include all ATLAS and CMS electro-weak vector boson production data, which have been corrected for FSR (though the correction has not been applied to LHCb data).

We may thus consistently isolate and compute the weak component of the EW corrections with FEWZ3.1 and include it in our calculation via the computation of C -factors for all the

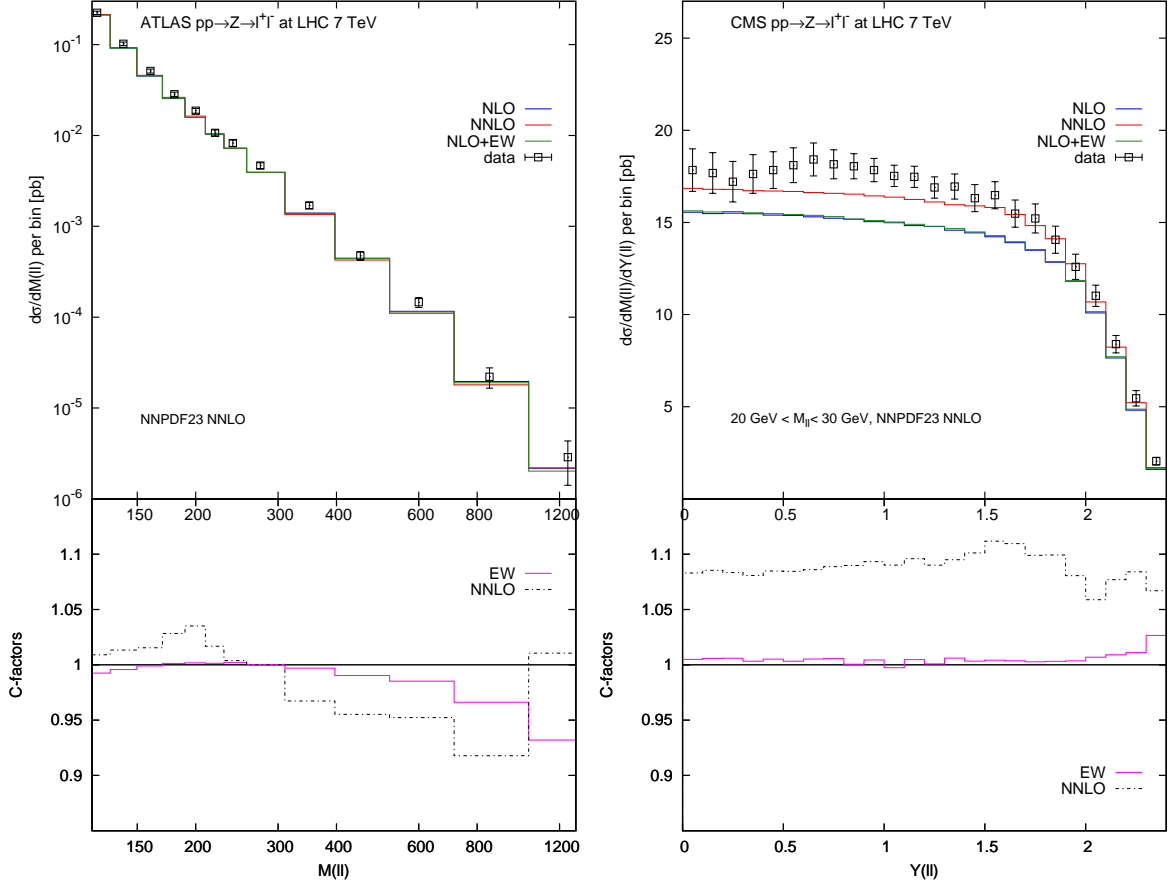


Figure 3: Left plot: the NLO, NNLO and NLO+EW predictions compared to the ATLAS high mass Drell-Yan distribution data as a function of the invariant mass of the dilepton system M_{ll} . As explained in the text, only virtual pure weak corrections have been included in the calculation (in particular no photon-initiated corrections are included). The three curves displayed have been computed with FEWZ3.1 with the same input PDF set, namely NNPDF2.3 with $n_f = 5$ and $\alpha_s(M_Z) = 0.118$. Right plot: same comparison for the CMS double differential Drell-Yan distribution as a function of the rapidity of the lepton pair in the lowest invariant mass bin, with $20 \leq M_{ll} \leq 30$ GeV. The bottom panels show the corresponding NNLO and EW C -factors that are applied to the NLO calculations.

electroweak gauge boson production. The size of the corrections is displayed in Fig. 3. We find that the effect of the pure EW corrections is negligible for most of the data in the Z peak region, of the order 1% or below. For the CMS double differential distributions, in the smallest invariant mass bin, we find that the pure EW corrections tend to increase the NLO and NNLO predictions by about 2%. At large invariant masses the pure EW corrections are rather large and negative, as expected from the results of Ref. [129]. This can be clearly seen in Fig. 3 for the ATLAS high-mass Drell-Yan data, where EW corrections reach $\mathcal{O}(7\%)$ in the last bin of the distribution corresponding to a dilepton invariant mass of $M_{ll} \in [1000, 1500]$. Note that for the last bin of invariant mass measured by CMS, $M_{ll} \in [200, 1500]$ GeV, EW effects are still moderate, since due to the steep fall-off of the cross-section, most of the events are from the region around $M_{ll} \sim 200$ GeV.

In summary, although the ATLAS high mass distribution is the only measurement for which EW corrections are required, for consistency we include them for all the Z/γ^* production data.

2.3.4 Treatment of heavy quarks

As in previous NNPDF analysis, in NNPDF3.0, heavy quark structure functions have been computed using the FONLL general-mass variable-flavor-number (GM-VFN) scheme [130]. In this scheme, the massive fixed-order calculation (in which the heavy quark is only counted in the final state) and resummed calculation (in which the heavy quark is treated as a massless parton) are consistently matched order by order. There is some latitude in deciding at which order the fixed-order massive result ought to be included when parton evolution is treated at a given perturbative order, and the FONLL method allows consistent matching in any case. Specifically, in a NLO computation one may decide to include massive contributions to structure functions up to $\mathcal{O}(\alpha_s)$, (on the grounds that this is the order at which the massless structure functions are also computed) or up to $\mathcal{O}(\alpha_s^2)$, (on the grounds that the massive structure function starts at $\mathcal{O}(\alpha_s)$). These are called respectively FONLL-A and FONLL-B schemes [130]. At NNLO, while a similar ambiguity would exist in principle, in practice the massive coefficient function can only be included up to $\mathcal{O}(\alpha_s^2)$ (FONLL-C) because the $\mathcal{O}(\alpha_s^3)$ massive result is now known (progress in this direction is reported e.g. in [131])

While the FONLL-A scheme was used for the NNPDF2.1 and NNPDF2.3 NLO PDF sets, we now adopt the FONLL-B scheme for NNPDF3.0 NLO, with FONLL-C used at NNLO. While this scheme is less systematic, in that when going to NLO to NNLO the massless computation goes up one order but the massive one does not, it has the advantage that massive terms at NLO are more accurate, thereby allowing for the inclusion of a somewhat wider set of data already at NLO, as small- x and Q^2 charm production data are affected by large $\mathcal{O}(\alpha_s^2)$ corrections, and cannot be described in the FONLL-A scheme. For the same reason, using FONLL-B allows for a more accurate description of the HERA inclusive cross-section data at small- x .

Heavy quark structure functions are computed using the expression which corresponds to the pole mass definition. Note that the distinction between pole and $\overline{\text{MS}}$ mass is only relevant at NNLO and beyond [132]. A full discussion of heavy quark mass dependence will be the subject of future work; for the time being we have checked that the dependence of our results on the value of the heavy quark masses is moderate, in agreement with previous findings in the framework of NNPDF2.1 [11]. In this paper, we adopt the following values for the heavy quark pole masses,

$$m_c = 1.275 \text{ GeV}, \quad m_b = 4.18 \text{ GeV}, \quad m_t = 173.07 \text{ GeV}, \quad (2)$$

which correspond to the current PDG values for the $\overline{\text{MS}}$ masses. Note that these values are

different from the ones used in NNPDF2.3, namely $m_c = \sqrt{2}$ GeV, $m_b = 4.75$ GeV and $m_t = 175$ GeV.

In NNPDF3.0 we take as a default the $n_f = 5$ scheme, in which the number of active flavors never exceeds $n_f = 5$ (i.e. in the fit the top quark is always treated as massive, never as a parton), though results will also be provided in the $n_f = 3$, $n_f = 4$ and $n_f = 6$ schemes. Note that in NNPDF2.1 and NNPDF2.3 the default was $n_f = 6$ (with the other cases also provided). While in previous PDF determinations the distinction between $n_f = 5$ and $n_f = 6$ was relevant only for delivery, as no data above threshold were available, now several jet data (especially for 2011 CMS inclusive jets, which has the highest reach in p_T) are above top threshold, as well as some CMS and ATLAS high-mass Drell-Yan data. Close to top threshold use of an $n_f = 5$ scheme might be advantageous because the top mass is treated exactly, while the loss of accuracy due to the fact that the $n_f = 5$ running of α_s differs from the exact $n_f = 6$ running [133] is a comparatively smaller effect (only visible for processes which start at high order in α_s , see e.g. [134]). Furthermore, most of the codes which we used for NNLO computations (specifically NLOJET++ and FEWZ) use an $n_f = 5$ scheme, and the same is true for many of the codes and interfaces used in the computation of LHC processes. With an ever increasing set of LHC data, the use of an $n_f = 5$ both in fitting, and as a default for PDF delivery appears thus to be more advantageous.

2.4 Construction of the dataset

We strive to include as much of the available experimental data as possible in our dataset. However, cuts have to be applied in order to ensure that only data for which the available perturbative treatment is adequate are included in the fit. These cuts are discussed in Sect. 2.4.1. Also, we include all available information on correlated uncertainties; however, there is some latitude in the treatment of correlated systematics, which we discuss in Sect. 2.4.2.

2.4.1 Kinematic cuts

As in previous NNPDF fits, we apply a cut in Q^2 and W^2 to fixed-target DIS data, in order to make sure that higher-twist corrections are not needed:

$$Q^2 \geq Q_{\min}^2 = 3.5 \text{ GeV}^2, \quad W^2 \geq W_{\min}^2 = 12.5 \text{ GeV}^2. \quad (3)$$

The stability of the fit with respect to these choices (and in particular the explicit check that they eliminate the need for higher twists) has been studied in detail in Ref. [104]. Low-scale DIS data carry less weight in our current fit than they did previously, because of the inclusion of a large amount of new HERA and LHC data. Therefore, the impact of these choices is even less important than it was previously. Note that all NNPDF fits include target-mass corrections, following the method of Ref. [6].

As discussed in Sect. 2.3 above, NNLO corrections are not available for the W p_T distribution and for $W + c$ production. As a consequence, ATLAS W p_T distribution data are included only in the NLO fit. However, we include the CMS $W + c$ distributions also in the NNLO fit, but with matrix elements computed up to NLO only (but α_s running at NNLO) because these data almost only affect the strange distribution, and the uncertainty on strangeness is rather larger than the size of the NNLO corrections, so we would rather keep the corresponding experimental information.

For inclusive jet production, we include all available data in the NLO fit, while in the NNLO fit we include only the data-points such that the relative difference between the exact and the

Experiment	N_{dat}	Exclusion regions in the (y, p_T) plane	
CDF Run-II k_t jets [83]	52	$1.1 < y < 1.6$ $ y > 1.6$	$p_T < 224 \text{ GeV}, p_T > 298 \text{ GeV}$ all p_T bins
ATLAS 2.76 TeV jets [63]	3	$0.0 < y < 0.3$ $ y > 0.3$	$p_T < 260 \text{ GeV}$ all p_T bins
ATLAS 7 TeV jets 2010 [50]	9	$0.0 < y < 0.3$ $0.3 < y < 0.8$ $ y > 0.8$	$p_T < 400 \text{ GeV}$ $p_T < 800 \text{ GeV}$ all p_T bins
CMS jets 2011 [62]	83	$1.0 < y < 1.5$ $ y > 1.5$	$p_T < 272 \text{ GeV}$ all p_T bins

Table 4: Summary of the exclusion regions in the jet transverse momentum p_T and rapidity $|y|$ used in the NNPDF3.0 NNLO fits for the inclusive jet production measurements. As explained in the text, these exclusion regions are determined from a cut-off in the relative difference between the exact and approximate threshold C -factors in the gluon-gluon channel [122]. N_{dat} in the second column is the number of experimental data points for these jet datasets that pass the selection cuts in the NNLO fits.

approximate gg -channel NNLO C -factors differ by less than 10%. As already mentioned in Sect. 2.3.2 above, a cut on the size of the C -factors, only accepting jet data for which the NNLO C -factor does not exceed 15%, results in approximately the same dataset. In Table 4 we summarize the ensuing exclusion regions in the (p_T, y) plane and the number of data-points N_{dat} which survive the cut for each experiment. As discussed in Ref. [122] and Sect. 2.3.2 above, this choice of cut ensures that NNLO corrections are included with an expected accuracy of a few percent for all the bins in (p_T, y) of the inclusive jet cross-section data.

For the ATLAS measurement of the W transverse momentum distribution, we include only those data points with $p_T^W > 25 \text{ GeV}$. This cut excludes the first two bins in p_T , and is motivated by the observation that at small p_T the perturbative series is not well-behaved and all-order resummation is needed (either analytically or by matching the fixed order calculation to a parton shower).

For the neutral-current Drell-Yan measurements from ATLAS and CMS, we include only data for which the dilepton invariant mass satisfies $M_{ll} < 200 \text{ GeV}$. This excludes the last six bins of the ATLAS DY invariant mass distribution and the 12 points in the rapidity distribution corresponding to the last bin of invariant mass, $M_{ll} \in [200, 1500] \text{ GeV}$, for the CMS measurement. The reason behind this cut is that the pure weak corrections that we include in our calculation, using FEWZ, do not include the QED subset of the total electroweak correction. In particular, including the photon-initiated contributions to the dilepton cross-section would require an initial photon PDF $\gamma(x, Q^2)$, which is not fitted in this analysis. Using the NNPDF2.3QED set, we have found that, while for most data-points this contribution is compatible with zero within the photon PDF uncertainty, for values of M_{ll} above M_Z such correction can become sizable, more than a few percent, and up to 20% for the highest invariant mass values. On this ground, we exclude all neutral-current Drell-Yan data with $M_{ll} > 200 \text{ GeV}$ from the fit, both at NLO and at NNLO.

The final cut is imposed, for the NLO fit only, to the lowest invariant-mass bin of the CMS Drell-Yan double differential distributions. As can be seen from Fig. 3 (right plot), for the bin with invariant mass $20 \leq M_{ll} \leq 30 \text{ GeV}$, the NNLO C -factors are large, around 10%, while experimental uncertainties are a few percent. Therefore, it is not possible to obtain a reasonable fit to these data-points in a NLO fit, and therefore the 24 points of the $20 \leq M_{ll} \leq 30 \text{ GeV}$ bin of the CMS data are excluded from the NLO analysis.

The number of data-points before and after cuts, both in the NLO and NNLO fit, are

summarized in Table 3. At LO we use the same cuts as in the NLO fit: this is a sensible choice, given that theory uncertainties at LO are larger than experimental uncertainties, it does not make much sense to attempt to devise a set of optimized kinematical cuts specifically for the LO fits.

2.4.2 Treatment of correlated systematic uncertainties

In NNPDF3.0 we include all available information on correlated systematic uncertainties. This information is summarized in Tables 1 and 2. All experiments provide uncorrelated statistical uncertainties; and most of them also a correlated normalization uncertainty. The exception of the latter case are ratio experiments such as NMC d/p , E866 d/p and the CMS W electron and muon asymmetries, and measurements which are normalized to the fiducial cross-section like the D0 Z rapidity distribution, the ATLAS W p_T distribution and the CMS 2D Drell-Yan data. On top of these two types of uncertainties, most experiments also provide information on correlated systematics, either as an overall covariance matrix, or as a set of nuisance parameters with the corresponding correlation. The last case is the most complete in that each individual source of systematic uncertainty can be in principle treated independently.

In Tables 1 and 2 we indicate in the 5th column whether the information on correlated systematics, on top of the normalization, is either not provided, provided as a covariance matrix, or fully provided with a complete breakdown of all the sources of systematic uncertainties. Furthermore, in the 6th column, datasets which are affected by common correlated systematics are flagged with the same letter. In Table 1, all the DIS datasets that are labelled with the same letter in this column share a common normalization uncertainty. In addition, the BCDMS, CHORUS, HERA-I, ZEUS HERA-II and H1 HERA-II datasets share other correlated systematic uncertainties. In Table 2, the CDF datasets share a normalization uncertainty, the ATLAS 7 TeV jets and W, Z rapidity distributions also have a correlated normalization uncertainty, and the ATLAS inclusive jet measurements at 2.76 TeV and 7 TeV share a number of systematic uncertainties, most importantly the jet energy scale. On the other hand, none of the CMS datasets has common correlated uncertainties among them.

Normalization uncertainties are special because the uncertainty is proportional to the measured value, i.e. they are “multiplicative”. This poses peculiar problems in including them in a fit, because their naive inclusion in the covariance matrix would lead to a systematically biased result [135]. A full solution to this problem based on an iterative procedure (the t_0 method) has been presented in Ref. [9] and adopted by NNPDF in all PDF determination from NNPDF2.0 onwards (previous approximate solutions such as adopted in the MSTW08 [136] PDF determination lead in practice to very similar results in most realistic situations, see Refs. [137, 138]).

In hadron collider experiments, not only normalization uncertainties but most, or perhaps all other correlated sources of uncertainty are also multiplicative. After checking with the respective experimental collaborations, we have thus concluded that the most accurate treatment of correlated systematics is obtained by treating all systematics as multiplicative for hadronic collisions. For deep-inelastic experiments, we treat all the correlated systematics of the HERA data as multiplicative, while for fixed-target experiments the systematics are treated as additive and only the normalization uncertainties as multiplicative. This information is also summarized in Tables 1 and 2, where in the 4th column experiments for which correlated systematics uncertainties are treated multiplicatively are denoted by M and others are denoted by A (additive uncertainties); normalization uncertainties are treated multiplicatively for all experiments. In Sect. 5 we will study the effect of a multiplicative vs. additive treatment of systematics, and find that it is small though perhaps not completely negligible. This will be specially true for

the gluon PDF at large- x , which is determined from collider jet data, for which an accurate treatment of experimental systematics is crucial.

As usual in the NNPDF methodology, we use the information on correlated systematics both in order to calculate the χ^2 used in fitting PDF to the Monte Carlo replicas, and also for generating the Monte Carlo replicas themselves. As explained in [6] (see also Eq. (21) in Sect. 4), the generation of the Monte Carlo replicas requires the knowledge of the individual sources of independent correlated systematic uncertainties. As can be seen from Table 2, for some LHC experiments this breakup is not provided and only the experimental covariance matrix is available. In those cases, we first create a set of artificial systematics consistent with the covariance matrix. These are obtained by diagonalizing the covariance matrix and then rescaling each eigenvector by the square root of its eigenvalue. This provides a set of N_{dat} artificial systematics which can be used to generate replica datasets in the same way as if we had the full breakup of systematics uncertainties, and which recombine correctly to give the original covariance matrix. This procedure effectively allows us to generate replica data according to the multivariate Gaussian distribution implied by the experimental covariance matrix.

3 Improved methodology

The inclusion of a substantial number of LHC datasets in a global PDF determination presents a variety of technical challenges, both in terms of the complexity of adding new hadronic observables into the fitting machinery and of the corresponding increase in the computational cost of running fits. The implementation of the enlarged LHC dataset is particularly demanding for NNPDF, which, being based upon a genetic algorithm minimization in a large parameter space, is already more computationally intensive than competing approaches. To provide an optimized solution to these challenges, for NNPDF3.0 we have developed a brand new fitting code based on C++ and Python, object-oriented languages that allow us to streamline the inclusion of new datasets and to achieve a highly efficient implementation of the minimization algorithms.

In what follows we first describe the modifications of the fitting code which make it more robust and efficient. Then we discuss the new PDF parametrization and the enforced positivity constraints. Finally, we discuss the improvements in the minimization strategy.

3.1 Code

The majority of the computer time spent in performing a PDF fit arises from the calculation of the physical observables for the comparison with experimental data. Indeed, any PDF determination involves an iterative procedure where all the data points included in the fit need to be recomputed a very large number of times for different functional forms of the input PDFs. The computation of physical observables in the NNPDF framework is based upon the FASTKERNEL method introduced in Refs. [10,13]. We refer the reader to the original publications for a detailed explanation of the method; here we simply recall a few basic facts that are necessary to explain the structure of the new code.

3.1.1 FastKernel methodology

For a suitable choice of interpolating functions, the parton distributions at a given scale Q^2 are represented on a grid of points in x :

$$\{f_i(x_\alpha, Q^2); \alpha = 1, \dots, N_x\}, \quad (4)$$

where the index i identifies the parton flavor, and the index α enumerates the points on the grid.

Deep-inelastic observables, which are linear in the PDFs, can be computed using a precomputed kernel $\hat{\sigma}_{\alpha j}^{I,J}$:

$$F_I(x_J, Q_J^2) = \sum_{j=1}^{N_{\text{pdf}}} \sum_{\alpha=1}^{N_x} \hat{\sigma}_{\alpha j}^{I,J} f_j(x_\alpha, Q_0^2). \quad (5)$$

The index I here labels the physical observable, x_J and Q_J are the corresponding kinematical variables for each particular experimental data point, α runs over the points in the interpolating grid and j is a parton flavor index. The kernel $\hat{\sigma}$ is referred to as an FK table. Similarly, hadronic observables, which are written as a convolution of two PDFs, can be computed in terms of an (hadronic) FK table $\hat{W}_{kl\gamma\delta}^{I,J}$; the final expression is:

$$F_I(x_J, Q_J^2) = \sum_{k,l=1}^{N_{\text{pdf}}} \sum_{\gamma,\delta=1}^{N_x} \hat{W}_{kl\gamma\delta}^{I,J} f_k(x_\gamma, Q_0^2) f_l(x_\delta, Q_0^2). \quad (6)$$

Once again we used the index I to identify the physical observable and J to label the specific experimental data point for this observable; the indices k, l run over the parton flavors, and the indices γ, δ count the points on the interpolating grids. In the fitting code, for each experimental dataset I we will have a separate FK table that encodes all the theory information.

The quantities $\hat{\sigma}$ and \hat{W} in Eqs. (5, 6) encode all the information about the theoretical description of the observables *e.g.* the perturbative order, the value of the strong coupling, the choice of scales, the QCD and electroweak perturbative corrections, or the prescription for the evolution. A modification in the theoretical formulation is reflected in a new FK table. The convolutions of the FK tables with the PDFs at the initial scale are generic, and do not require any knowledge about the theoretical framework. It is thus clear from their structure that the FK tables now implement a clean separation of the theory assumptions from the actual fitting procedure. During a fit, the PDFs at the initial scale are varied in order to minimize the χ^2 while the FK tables are always kept fixed and treated as an external input.

It should be emphasized that the main difference between fast NLO calculators such as FASTNLO, APPLGRID and AMCFast, and our FASTKERNEL method is that in the latter also the PDF evolution is included into the precomputed tables, while the other approaches require as input the PDFs evolved at the scales where experimental data is provided. This optimization step is essential to reduce the computational cost of running NNPDF fits. Note also that the generic structure Eqs. (5, 6) holds for any fast NLO calculator as well as for any PDF evolution code. In NNPDF3.0 we use our own internal Mellin-space FKGENERATOR code for PDF evolution and DIS observables, but using x -space codes such as HOPPET [139] or APFEL [102] should also be possible.

The fact that the expressions Eqs. (5, 6) of observables in terms of FK tables include the appropriate PDF evolution pre-computed in the FK tables $\hat{\sigma}$ (for DIS data) and \hat{W} (for hadronic observables) is particularly important for processes for which the precomputed fast NLO matrix elements, obtained for example from APPLGRID or FASTNLO, require the PDFs at a large number of Q^2 values. This is specially true for inclusive jet production, where for each experimental bin the range of Q^2 values for which PDFs need to be provided is different (since this range depends on the jet p_T and rapidity). In these cases, the associated acceleration due to the pre-computation of the PDF evolution makes all the difference between a process being practical or impossible to include in an NNPDF fit. The improvement due to this acceleration is quantified in Table 5 below.

3.1.2 NNPDF++

For the NNPDF3.0 determination, a new fitting toolchain has been developed from scratch with the FK procedure at its core in order to ensure that the most expensive part of any fit, the calculation of the physical observables, is performed in the most efficient manner possible. As discussed above the new fitting code has been designed with an explicit separation between experiment and theory, with all theory information contained within the FK tables, and all experimental information held within a new data file format, internally dubbed **CommonData**. These two data formats have an associated class structure built in modular C++ (which replaces the previous FORTRAN implementation), each retaining the following areas of jurisdiction:

- Theoretical treatment: FK tables
 - PDF evolution.
 - NLO QCD hard-scattering matrix elements.
 - Perturbative order, value of α_S/α_{EW} and values of the heavy quark masses.

- Heavy flavour renormalization scheme.
- NNLO QCD C-factors.
- NLO electroweak C-factors.
- Experimental information: `CommonData`
 - Process type information.
 - Experimental kinematics for each data-point.
 - Experimental central values.
 - Full breakdown of experimental systematic uncertainties.
 - Flexible choice of additive/multiplicative for the treatment systematic uncertainties.

The inclusion of a new experimental dataset into an NNPDF fit is therefore now a matter of converting the experimental results into the `CommonData` file format, and generating a set of FK tables each with the desired choices of theory parameters. Once generated, the `CommonData` and FK files are stored, and accessed thereon only via a set of classes used universally throughout the NNPDF++ fitting code.

These new code structures have enabled a number of methodological improvements by providing a more efficient and flexible treatment of experimental data. For example, performing kinematic cuts upon an experimental dataset can now be performed simply and algorithmically by selecting the points in the `CommonData` format which pass the required cuts according to their bundled kinematic information, and matching with the equivalent points in the FK `TABLE`. This is a considerable improvement over the earlier structure in `FORTRAN` whereby a variation of kinematical cuts necessitated a complete regeneration of the pre-computed theory tables due to the monolithic treatment of experimental data.

While the construction of this new fitting toolchain in C++ has made it possible to streamline any variation in the theory settings and experimental cuts, the new fitting code itself also achieves substantial improvements over the previous `FORTRAN` implementation. A modular treatment of most of the fitting procedure has been implemented, including PDF parametrization (both in terms of the flavor basis and of the specific functional form), minimization methods and stopping criteria. This allows for the rapid and safe replacement or improvement of separate modules of the NNPDF fitting framework without the need for a large amount of programming work, enabling methodological studies of a greater depth than those that we have been able to perform previously, as we shall report in later sections.

To ensure a fast and efficient minimization procedure, the calculation of observables in the FK convolution (the main bottleneck here) has been carefully optimized. In particular, the FK class that holds the hadronic table \hat{W} from Eq. (6) has been designed such that the \hat{W} table is stored with the optimal alignment in machine memory for use with `SIMD` (Single Instruction Multiple Data) instructions, which allow for an acceleration of the observable calculation by performing multiple numerical operations simultaneously. The large size of a typical FK product makes the careful memory alignment of the FK table and PDFs extremely beneficial. A number of `SIMD` instruction sets are available depending on the individual processor. By default we use a 16-byte memory alignment for suitability with Streaming `SIMD` Extensions (SSE) instructions, although this can be modified by a parameter to 32-bytes for use with processors enabled with Advanced Vector Extensions (AVX). The product itself is performed both with `SIMD` instructions and, where available, `OpenMP` is used to provide acceleration using multiple CPU cores. An implementation of the FK product for GPUs, while presenting no technical objections, has so far not been developed due to scalability concerns on available computing clusters.

Observable	APPLGRID	FK	optimized FK
W^+ production	1.03 ms	0.41 ms (2.5x)	0.32 ms (3.2x)
Inclusive jet production	2.45 ms	20.1 μ s (120x)	6.57 μ s (370x)

Table 5: Illustrative timings for the calculation of LHC observables using different methods to perform the convolution between the matrix elements and the PDFs. We provide results for two different observables: the total cross-section for W^+ production and for inclusive jet production for typical cuts of p_T and rapidity. In parenthesis we show the relative speed-up compared to the reference convolution based on APPLGRID. The main difference between the second and third columns is that in the latter case the PDF evolution is included in the precomputed FK table, while the differences between the third and fourth column arise from the use of explicit SSE acceleration in the product in the optimized FK convolution.

Thanks to the implementation of these various conceptual and technical optimizations in the calculation of physical observables, the NNPDF3.0 fitting code benefits from a substantial speed-up with respect to alternative calculations, using for example APPLGRID, where the PDF evolution needs to be performed separately from the calculation of hard-scattering matrix elements. This performance improvement is also clearly visible when comparing with the calculation of the hadronic convolution Eq. (6) using the optimized settings with that using non-optimized settings. To illustrate this point, in Table 5 we compare the timings, for a couple of representative LHC observables, for the convolution performed using APPLGRID, a naive FK implementation, and the optimized FK implementation using the SSE-accelerated calculation, for two representative observables. From Table 5 it is clear that a massive improvement in speed is achieved by precomputing the PDF evolution in the FK table, with further improvements obtained by the careful optimization of the FK product, and even further gains possible when combined with OpenMP on a multiprocessor platform.

3.2 PDF parametrization

As compared to the NNPDF2.3 analysis, there have been a number of substantial modifications in the way PDFs are parametrized and constrained. These include the choice of input scale, the parametrization basis, preprocessing, and the implementation of PDF positivity, which we now discuss in turn.

3.2.1 Parametrization basis

In previous NNPDF fits, PDFs were parametrized at a reference scale $Q_0^2 = 2 \text{ GeV}^2$. In NNPDF3.0 instead we choose $Q_0^2 = 1 \text{ GeV}^2$. Of course, this choice has no effect whatsoever on the results of the fit: indeed, changing the input scale amounts to a change in the input PDF parametrization, of which our fits are fully independent, as we have checked at length previously, and will verify again here (see in particular Sects. 4.5.3, 4.5.4, 5.3.5). The motivation for lowering Q_0^2 is to be able to span a wider range of possible values of the charm quark mass, without having to cross the initial evolution scale (*i.e.* while always having $Q_0 \leq m_c$), which will simplify future studies of heavy quark mass dependence.

At this reference scale, in previous NNPDF fits, including NNPDF2.3, PDFs were expressed

in terms of the following set of basis functions for quark and antiquark PDFs:

$$\begin{aligned}
\Sigma(x, Q_0^2) &= (u + \bar{u} + d + \bar{d} + s + \bar{s})(x, Q_0^2) \\
T_3(x, Q_0^2) &= (u + \bar{u} - d - \bar{d})(x, Q_0^2) \\
V(x, Q_0^2) &= (u - \bar{u} + d - \bar{d} + s - \bar{s})(x, Q_0^2) \\
\Delta_S(x, Q_0^2) &= (\bar{d} - \bar{u})(x, Q_0^2) \\
s^+(x, Q_0^2) &= (s + \bar{s})(x, Q_0^2) \\
s^-(x, Q_0^2) &= (s - \bar{s})(x, Q_0^2),
\end{aligned} \tag{7}$$

and then the gluon PDF $g(x, Q_0^2)$. This basis was chosen because it directly relates physical observables to PDFs, by making the leading order expression of some physical observables in terms of the basis functions particularly simple: for example, T_3 is directly related to the difference in proton and deuteron deep-inelastic structure functions $F_2^p - F_2^d$, and Δ_S is simply expressed in terms of Drell-Yan production in proton-proton and proton-deuteron collisions, for which there is data for example from the E866 experiment.

With the widening of the experimental dataset in NNPDF3.0, there is little reason to favor any particular PDF combination based on data, and thus we prefer to choose the basis that diagonalizes the DGLAP evolution equations. We emphasize that the only purpose of such choices is to speed up the minimization while leaving results unaffected: independence of our results of this basis change will be checked explicitly in Sects. 4.5.3 and 5 below. The default basis in the NNPDF3.0 fits is thus

$$\begin{aligned}
\Sigma(x, Q_0^2) &= (u + \bar{u} + d + \bar{d} + s + \bar{s})(x, Q_0^2) \\
T_3(x, Q_0^2) &= (u + \bar{u} - d - \bar{d})(x, Q_0^2) \\
T_8(x, Q_0^2) &= (u + \bar{u} + d + \bar{d} - 2s - 2\bar{s})(x, Q_0^2) \\
V(x, Q_0^2) &= (u - \bar{u} + d - \bar{d} + s - \bar{s})(x, Q_0^2) \\
V_3(x, Q_0^2) &= (u - \bar{u} - d + \bar{d})(x, Q_0^2) \\
V_8(x, Q_0^2) &= (u - \bar{u} + d - \bar{d} - 2s + 2\bar{s})(x, Q_0^2),
\end{aligned} \tag{8}$$

and of course the gluon. Here, as in previous NNPDF fits, we do not introduce an independent parametrization for the charm and anticharm PDFs (intrinsic charm). However we do plan to do in the near future.

As in all previous NNPDF fits, each basis PDF at the reference scale is parametrized in terms of a neural network (specifically a multi-layer feed-forward perceptron) times a preprocessing factor:

$$f_i(x, Q_0) = A_i \hat{f}_i(x, Q_0); \quad \hat{f}_i(x, Q_0) = x^{-\alpha_i} (1-x)^{\beta_i} \text{NN}_i(x) \tag{9}$$

where A_i is an overall normalization constant, and f_i and \hat{f}_i denote the normalized and un-normalized PDF respectively. The preprocessing term $x^{-\alpha_i} (1-x)^{\beta_i}$ is simply there to speed up the minimization, without biasing the fit. We now discuss the overall normalizations A_i , while the preprocessing will be addressed in Sect. 3.2.2 below.

Out of the seven normalization constants, A_i in Eq. (9), three can be constrained by the valence sum rules (for up, down and strange quarks) and another by the momentum sum rule. Which particular combinations depends of course of the choice of basis. With the default

NNPDF3.0 basis, Eq. (8), these constraints lead to

$$\begin{aligned} A_g &= \frac{1 - \int_0^1 dx x \Sigma(x, Q_0)}{\int_0^1 dx x \hat{g}(x, Q_0)}; & A_V &= \frac{3}{\int_0^1 dx \hat{V}(x, Q_0)}; \\ A_{V_3} &= \frac{1}{\int_0^1 dx \hat{V}_3(x, Q_0)}; & A_{V_8} &= \frac{3}{\int_0^1 dx \hat{V}_8(x, Q_0)}. \end{aligned} \quad (10)$$

The other normalization constants can be set arbitrarily to unity, that is $A_\Sigma = A_{T_3} = A_{T_8} = 1$: the overall size of these PDFs is then determined by the size of the fitted network. The finiteness of sum rule integrals Eq. (10) is enforced by discarding during the genetic algorithm minimization (see Sect. 3.3.1 below) any mutation for which the integrals would diverge. This condition, in particular, takes care of those NN configurations that lead to a too singular behavior at small- x .

Thanks to the flexibility of the new C++ fitting code, in NNPDF3.0 we support the option of using any arbitrary basis for the neural network parametrization of PDFs. This will allow us in particular to check explicitly basis independence. However, whenever we use a basis which differs from the evolution basis Eq. (8), the PDFs will be transformed back to the evolution basis before preprocessing and normalization are applied. This has the advantage of ensuring that the finite integrals Eq. (10) do not have to be constructed as the difference of divergent integrals, which would require large numerical cancellations, thereby potentially leading to numerical instabilities.

3.2.2 Effective preprocessing exponents

As mentioned, when parametrizing PDFs in terms of neural networks, Eq. (9), preprocessing is introduced as a means to speed up the minimization, by absorbing in a prefactor the bulk of the fitted behaviour so that the neural net only has to fit deviations from it. However, we must make sure that the choice of preprocessing function does not bias the result. As in previous NNPDF fits, starting with NNPDF1.2 [8] onwards, this is done by randomizing the preprocessing exponents, i.e. by choosing a different value for each replica within a suitable range. Unlike in previous NNPDF fits, where this range was determined based on a stability analysis of the results, we now determine the range self-consistently in a completely automatic way (the same methodology was already used in the NNPDFpol family of polarized PDF determinations [140, 141]).

This is done in the following way. First of all, we define effective asymptotic exponents as follows:

$$\alpha_{\text{eff},i}(x) = \frac{\ln f_i(x)}{\ln 1/x} \quad (11)$$

$$\beta_{\text{eff},i}(x) = \frac{\ln f_i(x)}{\ln(1-x)}. \quad (12)$$

Other definitions would be possible, such as

$$\alpha_{\text{eff},i}(x) = \frac{d \ln f_i(x)}{d \ln 1/x} \quad \beta_{\text{eff},i}(x) = \frac{d \ln f_i(x)}{d \ln(1-x)}. \quad (13)$$

We have checked that their use would not modify qualitatively our results. We choose a wide starting range for the preprocessing exponents for each PDF, and perform a fit. The effective exponents Eq. (11-12) are then computed for all replicas at $x = 10^{-6}$ and 10^{-3} for the low- x exponent α_i and at $x = 0.95$ and 0.65 for the large- x exponent β_i , for all PDFs (except for

PDF	NLO		NNLO	
	$[\alpha_{\min}, \alpha_{\max}]$	$[\beta_{\min}, \beta_{\max}]$	$[\alpha_{\min}, \alpha_{\max}]$	$[\beta_{\min}, \beta_{\max}]$
Σ	[1.06, 1.22]	[1.31, 2.68]	[1.02, 1.33]	[1.31, 2.74]
g	[0.96, 1.37]	[0.28, 5.45]	[1.05, 1.53]	[0.85, 5.20]
V	[0.54, 0.70]	[1.20, 2.91]	[0.54, 0.70]	[1.18, 2.80]
V_3	[0.29, 0.58]	[1.31, 3.42]	[0.29, 0.61]	[1.36, 3.73]
V_8	[0.54, 0.73]	[0.80, 3.09]	[0.55, 0.72]	[1.06, 3.07]
T_3	[-0.17, 1.36]	[1.58, 3.14]	[-0.25, 1.41]	[1.64, 3.20]
T_8	[0.54, 1.25]	[1.30, 3.42]	[0.54, 1.27]	[1.33, 3.23]

Table 6: Ranges from which the small- and large- x preprocessing exponents in Eq. 9 are randomly chosen for each PDF. For each replica, a value is chosen from these ranges assuming a flat probability distribution. The provide the results for the global NLO and NNLO NNPDF3.0 fits. The two sets of ranges, obtained at each perturbative order, are determined independently using an iterative procedure, as explained in the text.

the gluon and singlet small- x exponent, α_i , which is computed at $x = 10^{-6}$). The fit is then repeated by taking as new range for each exponent the envelope of twice the 68% confidence interval for each x value. The process is then iterated until convergence, i.e., until the output preprocessing exponents stop changing. Reassuringly, convergence is typically very fast: even in the cases where the fitted dataset is varied significantly, only one iteration is needed to achieve stability.

This procedure ensures that the final effective exponents are well within the range of variation both in the region of the smallest and largest x data points, and in the asymptotic region (these two regions coincide for the gluon and singlet at small x), thereby ensuring that the allowed range of effective exponents is not artificially reduced by the preprocessing, either asymptotically or sub-asymptotically.

The final output values of the preprocessing exponents for the central NLO and NNLO NNPDF3.0 fits are listed in Table 6. These ranges have been redetermined self-consistently for different fits: for example, for fits to reduced datasets, wider ranges are obtained due to the experimental information being less constraining.

3.2.3 Positivity constraints

As is well known [142], beyond leading-order parton distributions do not need to be positive definite. However, the requirement that measurable physical observables be positive still imposes a generalized positivity constraint on the PDFs. In previous NNPDF fits these constraints were enforced by imposing positivity of the deep-inelastic structure functions F_L , F_2^c and of the neutrino charm production (“dimuon”) cross-section at a scale of $Q_{\text{pos}}^2 = 5 \text{ GeV}^2$ in the range $x \in [10^{-5}, 1]$ (see in particular Sect. 4.5 and 5.5 of Ref. [10]). However, while these conditions were sufficient to guarantee the positivity of most physical observables in previous NNPDF fits, in order to ensure positivity of all observables, the number of independent positivity constraints must be at least equal to the number of independently parametrized PDFs.

In order to guarantee full positivity of physical observables we have thus enlarged the set of positivity constraints. In particular, we have chosen to impose positivity of some pseudo-observables which must respect positivity for reasons of principle, but which are not measurable in practice. We choose the three tagged deep-inelastic structure functions F_2^u , F_2^d and F_2^s , and the three flavor Drell-Yan rapidity distributions, $d\sigma_{u\bar{u}}^{\text{DY}}/dy$, $d\sigma_{d\bar{d}}^{\text{DY}}/dy$ and $d\sigma_{s\bar{s}}^{\text{DY}}/dy$, defined

respectively as the contribution to the structure function or the Drell-Yan rapidity distribution which is obtained when all quark electric charges are set to zero except that of the up, down or strange quark flavor, respectively. These six constraints enforce generalized positivity of the quark and antiquark PDFs. As in previous fits, our positivity conditions are imposed at the low scale of $Q_{\text{pos}}^2 = 5 \text{ GeV}^2$, and will then be also satisfied once the PDFs are evolved upwards in Q^2 .

Generalized positivity of the gluon is enforced by requiring positivity of the light component of the longitudinal structure function F_L^l , defined as the contribution to the structure function F_L when all quark electric charges are set to zero but those of the three lightest flavors. The use of F_L^l allows us to impose gluon positivity without having to make specific choices for the treatment of heavy quarks at the low scale where these conditions are imposed. This constraint is supplemented with that from the rapidity distribution $d\sigma_{gg}^H/dy$ for the production in gluon-gluon fusion of a Higgs-like scalar with mass $m_H^2 = 5 \text{ GeV}^2$ (such a constraint being much more stringent than that from production of a heavier Higgs). The use of these two observables ensures positivity of the gluon PDF both at small and at large x values.

In summary, the pseudo-observables used to enforce generalized PDF positivity, expressed schematically in terms of their underlying parton content, the following:

$$\begin{aligned}
F_2^u(x, Q^2) &\propto (u(x, Q^2) + \bar{u}(x, Q^2)) + \mathcal{O}(\alpha_s) \\
F_2^d(x, Q^2) &\propto (d(x, Q^2) + \bar{d}(x, Q^2)) + \mathcal{O}(\alpha_s) \\
F_2^s(x, Q^2) &\propto (s(x, Q^2) + \bar{s}(x, Q^2)) + \mathcal{O}(\alpha_s) \\
\frac{d^2\sigma_{u\bar{u}}^{\text{DY}}}{dM^2 dy} &\propto u(x_1, Q^2)\bar{u}(x_2, Q^2) + \mathcal{O}(\alpha_s) \\
\frac{d^2\sigma_{d\bar{d}}^{\text{DY}}}{dM^2 dy} &\propto d(x_1, Q^2)\bar{d}(x_2, Q^2) + \mathcal{O}(\alpha_s) \\
\frac{d^2\sigma_{s\bar{s}}^{\text{DY}}}{dM^2 dy} &\propto s(x_1, Q^2)\bar{s}(x_2, Q^2) + \mathcal{O}(\alpha_s) \\
F_L^l(x, Q^2) &\propto C_g \otimes g(x, Q^2) + C_q \otimes q(x, Q^2) + \mathcal{O}(\alpha_s^2) \\
\frac{d\sigma_{gg}^H}{dy} &\propto g(x_1, M_H^2)g(x_2, M_H^2) + \mathcal{O}(\alpha_s^3) \quad M_H \equiv Q_{\text{pos}}
\end{aligned} \tag{14}$$

All these positivity constraints are imposed at $Q_{\text{pos}}^2 = 5 \text{ GeV}^2$, and for $x \in [10^{-7}, 1]$, which, because of the structure of QCD evolution, ensures positivity at all higher scales (and explains our choice for the mass of the Higgs-like scalar). In practice we compute the observables at 20 points in the given x range, equally spaced on a log scale for $x < 0.1$ (ten points) and on a linear scale for $x \geq 0.1$.

During the minimization, the positivity constraints are imposed by adding a Lagrange multiplier, and then further discarding replicas for which any of the pseudo-observables is negative by more than 25% of its absolute value computed with a fixed reference PDF set (typically, the outcome of a previous fit). The latter condition is necessary for cross sections which are very close to zero (e.g. close to kinematic boundaries, like the rapidity tails of Drell-Yan distributions) where the Lagrange multiplier strategy is not effective. For DIS structure functions and Drell-Yan distributions, the pseudo-observables are computed using the internal NNPDF FASTKERNEL implementation, while the Higgs-like scalar production cross section is computed using MADGRAPH5_AMC@NLO [112] interfaced to AMCFast [111], using a tailored model which goes beyond the effective theory approximation for top quark mass effects [121, 143].

The strategy outlined above has been used to enforce generalized positivity both in the NLO and NNLO fits. The only difference is that in the NNLO fit we compute the pseudo-observables at NLO (with the PDFs evolved also at NLO for consistency) because at the low Q^2 values at which the positivity pseudo-observables are computed, large unresummed NNLO corrections lead to perturbatively unstable predictions at large and small x . In addition, there is some evidence that the resummed result is closer to NLO than to NNLO, see for example Ref. [144] for the case of deep-inelastic structure functions. The impact of the positivity constraints on the final PDFs is quite substantial, especially in the extrapolation regions, as one could expect. This statement will be quantified and discussed in Sect. 5.3, where we will compare two NNPDF3.0 NLO fits with and without the positivity constraints, and we will discuss further a posteriori checks of the implementation of the positivity conditions. In the results section, we also explore the impact of the improved positivity constraints on searches for high-mass new physics.

In the leading order fits, where PDFs should be strictly positive-definite, exactly the same strategy as in the NLO and NNLO fits is used (with pseudo-observables now computed at LO): we have verified that, as one could expect, using the conditions Eq. (14) expressed at LO leads to effectively positive-definite PDFs. To avoid the problem of the occasional replica of the LO fit that might become slightly negative, the LHAPDF6 NNPDF3.0LO grids provide by construction a positive-definite output. We have verified the robustness of our LO positivity implementation by comparing it to alternative strategies, such as using a neural network where the output of the last layer is squared.

3.3 Minimization algorithm

As in previous NNPDF fits, minimization is performed using genetic algorithms, which are especially suitable for dealing with very large parameter spaces. Because of the extreme flexibility of the fitting functions and the large number of parameters, the optimal fit is not necessarily the absolute minimum of the χ^2 (see in particular the discussion in Sect. 4 of Ref. [6]), which might correspond to an ‘overfit’ in which not only the desired best fit is reproduced, but also statistical fluctuation about it. As a consequence, a stopping criterion is needed on top of the minimization method. In NNPDF3.0 we have improved both the minimization strategy and the stopping criterion.

3.3.1 Genetic Algorithms

We have performed a comprehensive re-examination of the genetic algorithm minimization procedure utilized in previous NNPDF determinations. Our approach has been to take a minimal starting methodology with only a few basic features. New features were added in turn and only retained if they resulted in faster fitting.

The algorithm consists of three main steps: mutation, evaluation and selection. Firstly, a large number of mutant PDF sets are generated based on a parent set from the previous generation. The goodness of fit to the data for each mutant is then calculated. The best fit mutant is identified and passed on to the next generation, while the rest are discarded. The algorithm is then iterated until a set of stopping criteria are satisfied.

The number of mutants tested each generation is now set to 80 for all generations, removing the two GA ‘epochs’ used in previous determinations. However, this number is somewhat arbitrary, as the more significant quantity in terms of fit quality is the total number of mutants produced during the fit, i.e. the number of mutants multiplied by the number of generations. If one is increased while the other is decreased by the same factor the fit results are largely unchanged. All mutants are generated from the single best mutant from the previous generation.

To generate each mutant, the weights of the neural networks from the parent PDF set are altered by mutations. In previous NNPDF fits the mutations have consisted of point changes, where individual weights or thresholds in the networks were mutated at random. However, investigations of strategies for training neural networks [145] have found that employing coherent mutations across the whole network architecture instead leads to improved fitting performance. The general principle that explains this is that of changing multiple weights which are related by the structure of the network, leading to improvements in both the speed and quality of the training.

The neural networks used in NNPDF fits consist of connected nodes organized in layers. To get a value from the network, the nodes in the input layer are set with the required x and $\log x$ value and then the activations of nodes in successive layers are calculated according to

$$\xi_i^{(l)} = g \left(\sum_j w_{ij}^{(l)} \xi_j^{(l-1)} + \theta_i^l \right) \quad (15)$$

$$g(a) = \frac{1}{1 + e^{-a}} \quad (16)$$

where $\xi_i^{(l)}$ is the activation of the i th node in the l -th layer of the network, $w_{ij}^{(l)}$ are the weights from that node to the nodes in the previous layer and θ_i^l is the threshold for that node. The weights and the thresholds are the parameters in the fit, and so are the objects which are changed in the mutation. The exception to Eq. (16) is the last layer, where in order to allow for an unbounded output a linear activation function $g(a) = a$ is used instead. The flexibility of the fitting code allows us to easily explore other choices, for instance a quadratic output of the last layer, $g(a) = a^2$, has been used in studies of the PDF positivity in leading order fits.

In the NNPDF3.0 fits we use a nodal mutation algorithm, which gives for each node in each network an independent probability of being mutated. If a node is selected, its threshold and all of the weights are mutated according to

$$w \rightarrow w + \frac{\eta r_\delta}{N_{\text{ite}}^{r_{\text{ite}}}}, \quad (17)$$

where η is the baseline mutation size, r_δ is a uniform random number between -1 and 1 , different for each weight, N_{ite} is the number of generations elapsed and r_{ite} is a second uniform random number between 0 and 1 shared by all of the weights. An investigation performed on closure test fits found that the best value for η is 15 , while for the mutation probability the optimal value turns out to be around 5% , which corresponds to an average of 3.15 nodal mutations per mutant PDF set.

As with the removal of the fast- and slow-epochs and their replacement with a single set of GA parameters, the Targeted Weighted Training (TWT) procedure adopted in previous fits has also been dropped. This was originally introduced in order to avoid imbalanced training between datasets. With the considerably larger dataset of NNPDF3.0 along with numerous methodological improvements, such an imbalance is no longer observed even in fits without weighted training. Whereas previously the minimization was initiated with a TWT epoch in which the fit quality to individual datasets was minimized neglecting their cross-correlations, in NNPDF3.0 the minimization always includes all available cross-correlations between experimental datasets.

With the removal of the TWT mechanism, along with the consolidation of GA training into a single epoch with a unified set of mutation probabilities and sizes, the number of free parameters in the NNPDF minimization has been considerably reduced. In Table 7 we provide a comparison summarizing the relevant parameters in the NNPDF2.3 and NNPDF3.0 determinations.

	$N_{\text{gen}}^{\text{wt}}$	$N_{\text{gen}}^{\text{mut}}$	$N_{\text{gen}}^{\text{max}}$	E^{sw}	N_{mut}^a	N_{mut}^b
NNPDF 2.3	10000	2500	50000	2.3	80	30
NNPDF 3.0	—	-	30000	-	80	-

NNPDF2.3			NNPDF3.0		
Single Parameter Mutation			Nodal Mutation		
PDF	N_{mut}	η	PDF	P_{mut}	η
$\Sigma(x)$	2	10, 1	$\Sigma(x)$	5% per node	15
$g(x)$	3	10, 3, 0.4	$g(x)$	5% per node	15
$T_3(x)$	2	1, 0.1	$V(x)$	5% per node	15
$V(x)$	3	8, 1, 0.1	$V_3(x)$	5% per node	15
$\Delta_S(x)$	3	5, 1, 0.1	$V_8(x)$	5% per node	15
$s^+(x)$	2	5, 0.5	$T_3(x)$	5% per node	15
$s^-(x)$	2	1, 0.1	$T_8(x)$	5% per node	15

Table 7: Comparison of genetic algorithm parameter between the NNPDF2.3 and NNPDF3.0 fits. In the top table, parameters controlling the maximum fit length, number of mutants, and (for NNPDF2.3) target weighted training settings are shown. In the tables below, the mutation parameters are shown for the two determinations in terms of their respective fitting bases. For the NNPDF3.0 fit the mutation probability is now set at 5% per network node, and the mutation size is set to a consistent $\eta = 15$.

3.3.2 Stopping criterion

As in previous NNPDF fits, the optimal fit is determined using a cross-validation method. This is based on the idea of separating the data in two sets, a training set, which is fitted, and a validation set, which is not fitted. The genetic algorithm is used in order to minimize the χ^2 (or other figure of merit) of the training set, while the χ^2 of the validation set is monitored along the minimization, and the optimal fit is achieved when the validation χ^2 stops improving (see Ref. [6] for a more detailed discussion).

In previous NNPDF fits this stopping criterion was implemented by monitoring a moving average of the training and validation χ^2 , and stopping when the validation moving average increased while the training moving average decreased by an amount which exceeded suitably chosen threshold values. The use of a moving average and of threshold values was necessary in order to prevent the fit from stopping due to a statistical fluctuation, but introduced a certain arbitrariness since the value of these three parameters (the length of the moving average and the two thresholds) had to be tuned.

In NNPDF3.0 we have improved on this: we now simply stop all fits at the point in which the fit reaches the absolute minimum of the validation χ^2 within the maximum number of generations $N_{\text{gen}}^{\text{max}}$. In practice this is done in the following way. All replicas are minimized for $N_{\text{gen}}^{\text{max}}$ generations. In the beginning of the fit, both the validation χ^2 and the PDF configurations are stored. At the end of the each generation, the validation χ^2 is computed again, and if it is lower than the previous stored values, the PDFs are replaced with the current ones, if not the fit proceeds to the next iteration. At the end, the stored PDFs have the the lowest validation χ^2 seen during the fit, and they are taken as the global best fit. We call this the 'look-back' method, and it is clearly completely objective. The price to pay for this is that now all replicas have to be run up to $N_{\text{gen}}^{\text{max}}$.

In order for the look-back method to be effective, the value of $N_{\text{gen}}^{\text{max}}$ must be large enough that the PDFs at the minimum do not depend on it, i.e., such that the choice of a larger value

does not lead to a different minimum. We have verified this explicitly by checking that results are unchanged if the value of $N_{\text{gen}}^{\text{max}}$ is increased, see Sect. 4.5.1 below. We have furthermore checked that results are stable upon small variations of the position of the minimum: this guarantees that the choice of the absolute minimum (which could correspond to a local fluctuation) does not bias the result in any way. For this purpose, we have changed the look-back algorithm, by only updating the PDFs when the χ^2 decreases by more than some threshold: we tried increasingly large threshold values (0.1, 1, and 10 on the total χ^2 not divided by the number of data points) and verified that even though, of course, the stopping point changes, and happens at an earlier stage for large values of the threshold, the resulting PDFs are indistinguishable. This explicitly verifies that local minima whose validation χ^2 values differ from the absolute minimum by a small amount correspond to essentially the same PDFs, as one might expect.

Even so, there is still the possibility that occasionally for a particular replica the training length required to reach the global minimum would be exceptionally long, either because of an unusual fluctuation of the pseudodata, or because of fluctuations in the minimization. In that case, the look-back method would not be effective, because the best χ^2 within the maximal training length would still be far from the absolute minimum. In order to safeguard against this possibility, we consider the PDF arc-length defined as

$$L = \int \sqrt{1 + \left(\frac{df}{dx}\right)^2} dx, \quad (18)$$

suitably discretised. Use of an the arc-length penalty has been previously suggested as a way of penalizing functional forms that are too wiggly, see e.g. example the studies in Ref. [146].

We then take the value of L as an indicator of convergence: a PDF replica with a very unlikely value of L is assumed not to have converged, and thus discarded. This is implemented through an arc-length veto. The arc-length is computed for each PDF replica, and the average length and its standard deviation over the replica sample are determined for each PDF. We then discard PDF replicas for which at least one PDF the arc-length exceeds by more than four sigma the mean arc-length for that PDF.

Finally, at the end of the minimization, an a posteriori quality check on the resulting sample of Monte Carlo replicas is performed for each fit. These quality tests verify that the PDF generalized positivity has been successfully implemented (see discussion in Sect. 3.2.3), that no replica has a too unlikely arc-length, Eq. (18), and that likewise no replica has a too unlikely value of the χ^2 . As in the case of arc-length, if any given replica has a χ^2 whose value is more than four-sigma away the mean χ^2 , it is automatically replaced by another replica that instead satisfies this condition.

4 Closure testing

In this section we describe in detail the strategy that has been used to validate the new fitting methodology described in the previous section: closure testing. The benchmarking of fitting methodology is especially important due to the substantial increase in experimental data included in NNPDF3.0, and the increased precision of the resulting PDFs: as data become more precise and their kinematic coverage increases, it becomes more and more important to eliminate methodological uncertainties, so that the only uncertainties in our PDFs are experimental and theoretical. For example a basic requirement for a successful methodology is to be able to fit widely different datasets with the same methodology, and without having to tune it or modify it according to the dataset, or to the theory which is used to describe it (for example at different orders in perturbative QCD). As we will show below, these requirements can indeed be achieved in the NNPDF3.0 framework. NNPDF3.0 is in fact the first PDF determination for which the complete fitting methodology has been thoroughly tested and tuned in closure tests based on pseudo-data that have the same kinematical coverage and statistical properties as the experimental data included in the fit. The idea of using perfect pseudo-data to validate some specific aspects of a PDF fitting methodology has been previously explored in Ref. [14].

The basic idea of the closure test is simple [2]: we take a given assumed form for the PDFs (for example MSTW08), a given theoretical model (for example NLO pQCD), and with them generate a set of global pseudodata with known but realistic statistical properties (by using the covariance matrices of the real datasets that together make up, for example, the NNPDF3.0 dataset). These pseudodata are then ‘perfect’, in the sense that they have known statistical properties, no internal inconsistencies, and are also entirely consistent with the theoretical model used to produce them. Thus if we then use our fitting methodology to perform a fit to these pseudodata, we should reproduce the assumed underlying PDF, within the correct uncertainties. This latter point can be explored in some depth by changing by hand the level of uncertainties incorporated within the pseudodata.

We will first introduce the idea of closure testing in the context of the NNPDF fits, and discuss its practical implementation. Then we will quantify the efficiency of our neural network training methodology by performing closure tests to perfect pseudo-data, without any fluctuations, so that the fit quality can be arbitrarily good. We then perform fits where the pseudo-data is supplemented with different levels of statistical and systematic fluctuations. Finally, in the last part of this section we use closure tests to assess the robustness of our methodology against variations of some of its ingredients, like the choice of PDF fitting basis or the dependence on the size of the neural network. Some of these last checks have also been performed on real data, as will be reported in Sect. 5.

4.1 NNPDF closure testing

The framework used in NNPDF3.0 for the computation of observables, as presented in Sect. 3.1.1, provides us with the ideal tool to successfully implement closure tests. In particular, the clean separation between theoretical assumptions and input PDFs allows us to generate pseudo-data using a given set of PDFs and the experimental covariance matrix as an input, and to perform a fit to this pseudo-data using exactly the same theoretical settings (encoded in the FK tables) that were used for generating them.

Throughout this section we shall refer to the parton distributions used to generate the pseudo-data as the *input* PDFs, and denote them by f_{in} . Any PDF set available through the LHAPDF interface can be used as an input set to generate the pseudo-data: most of the closure

tests described here will be performed using MSTW08, though we will also later describe results using other input PDFs. We denote the set of pseudodata by $\mathcal{D} = \{D_I\}$: the dependence of the pseudo-data on the input PDF f_{in} and experimental covariance matrix will be left implicit.

The outcome of the closure test fits is then a set of fitted PDFs, f_{fit} , which we will compare to the input PDFs in order to study the statistical precision and possible systematic biases in the fitting methodology. For any PDF set f , whether input or fitted, the FASTKERNEL framework delivers a set of theoretical predictions, $\mathcal{T}[f] = \{T_I[f]\}$, based on a particular theoretical model, which for present purposes we take to be NLO perturbative QCD, precisely as implemented in the NNPDF3.0 fits to real data, with the same parameter choices and so on. The one exception to this statement is that for closure tests to MSTW08 we drop the generalized positivity constraints introduced in Sect. 3.2.3, since positivity was imposed differently by MSTW. We have checked however, by performing closure tests based on NNPDF3.0 as input PDFs, that none of the conclusions drawn in this section is affected once the generalized positivity is included in the fit.

Fitting is performed by minimizing a χ^2 function (denoted during the fitting as the ‘error function’): this is also used to assess the quality of the resulting fit. Such a χ^2 depends on the dataset, \mathcal{D} , and on the theoretical predictions of the PDFs f being fitted, $\mathcal{T}[f]$:

$$\chi^2[\mathcal{T}[f], \mathcal{D}] = \frac{1}{N_{\mathcal{D}}} \sum_{I,J} (T_I[f] - D_I) C_{IJ}^{-1} (T_J[f] - D_J). \quad (19)$$

In this expression, C_{IJ} is the covariance matrix of the data (here we always use the t_0 -covariance matrix in order to avoid bias in the inclusion of multiplicative uncertainties, see Sect. 2.4.2), and $N_{\mathcal{D}}$ is the total number of data points of the dataset. Note that when fitting the pseudo-data we thus use exactly the same procedure (in fact the same code) as we use in a fit to real data: we use the same error function, evaluated with the real data replaced by pseudo-data, the t_0 covariance matrix, and then the same fitting methodology (genetic algorithm, stopping criterion, etc). Since in the closure test fits the “correct” solution is known – it is given by the PDFs f_{in} used as an input – the result f_{fit} of the fit should then reproduce the input PDFs within the statistical uncertainties of f_{fit} as determined by the fit.

For the purposes of the studies that will be performed in this paper, we will introduce three distinct categories (levels) of closure tests depending on the amount of stochastic noise added to the pseudo-data points generated from the initial PDFs. In order to make these tests as realistic as possible, this stochastic noise is generated using the complete information in the experimental covariance matrix, so that the fluctuations and correlations of the pseudo-data reproduce precisely those of the real experimental data. For the baseline closure tests presented in this Section, the pseudo-data is in one-to-one correspondence with the experimental data used in the global fit, that is we have generated pseudo-data for every point in the NNPDF3.0 global dataset described in Tabs. 1 and 2.

The three levels of closure test that will we will study, which we call **Level 0**, **Level 1** and **Level 2** for reasons that will soon become apparent, are set up as follows:

- **Level 0.**

Using a given set of input PDFs, pseudo-data $\mathcal{D}_0 = \{D_I^0\}$ are generated using the FASTKERNEL convolution, Eqs. (5,6). In these Level 0 fits, no stochastic noise is added to the pseudo-data. Then we perform N_{rep} fits, each to exactly the same set of pseudo-data, minimizing the error function (which here is the same as the χ^2 per degree of freedom, i.e. $\chi^2[\mathcal{T}[f], \mathcal{D}_0]$, but using different seeds for the initialization of the random numbers used in the minimization. This yields an ensemble of PDF replicas $\{f_{\text{fit}}^k\}$, where $k = 1, \dots, N_{\text{rep}}$.

Note that the error function which we minimize is still computed using the covariance matrix of the data, even though now the pseudodata have zero uncertainty. While of course the overall normalization of the error function is immaterial (as the minimum is at zero), this has the advantage of reproducing the correlations in the underlying dataset, which means that the total amount of independent experimental information (the number of independent datapoints) is the same as in the original dataset.

It should be clear from its definition that in Level 0 closure tests, the fit quality can be arbitrarily good, provided we use a sufficiently flexible PDF parametrization and a sufficiently efficient minimization algorithm. Indeed, since, by construction, the pseudo-data does not have any stochastic noise, and there are no inconsistencies, there exist *perfect* fits to the Level 0 pseudo-data that have a vanishing χ^2 . We use the plural here because there clearly exists an infinity of fits which lead to vanishing χ^2 by going through all datapoints, but differ in the way they interpolate in between datapoints/ These optimal solutions to the minimization problem reproduce precisely the predictions of the set of PDFs used as input in the generation of the pseudo-data at each of the experimental data points. The Level 0 closure test is thus a highly non-trivial test of the efficiency of the minimization procedure: at Level 0 the value of χ^2 should decrease monotonically towards zero as the fit proceeds, if the functional form used for parametrizing the fitted PDFs is flexible enough.

- **Level 1.**

Now we add stochastic fluctuations on top of the pseudo-data generated for the Level 0 closure tests, D_I^0 , as follows:

$$D_I^1 = (1 + r_I^{\text{nor}} \sigma_I^{\text{nor}}) \left(D_I^0 + \sum_{p=1}^{N_{\text{sys}}} r_{I,p}^{\text{sys}} \sigma_{I,p}^{\text{sys}} + r_I^{\text{stat}} \sigma_I^{\text{stat}} \right), \quad (20)$$

where, as explained in [6], σ_I^{stat} , $\sigma_{I,p}^{\text{sys}}$ and σ_I^{nor} are the statistic, systematic and normalization uncertainties for each dataset, and the random numbers r_I^{nor} , $r_{I,p}^{\text{sys}}$ and r_I^{stat} are generated with the appropriate distribution to reproduce the experimental covariance matrix. These shifted datapoints represent the central values of a hypothetical experiment, for which the size of the statistical, systematic and normalization uncertainties are given in Eq. (20).

In Level 1 fits, the same underlying pseudo-data, generated by the random numbers in Eq. (20), is used for the fit of all the N_{rep} replicas, but as at Level 0 a different random seed is used to initialize the minimization of each replica. Therefore no additional stochastic fluctuations (“Monte Carlo replicas”) are added to Eq. (20), and thus, as we will confirm later, in a Level 1 fit the experimental uncertainties are not propagated into the uncertainties of the fitted PDFs. The ensemble of PDF replicas $\{f_{\text{fit}}^k\}$ resulting from the Level 1 fits is thus expected to underestimate the PDF uncertainties.

From its definition, given that the pseudo-data have fluctuated on average by one standard deviation away from the Level 0 value, we expect that in Level 1 closure tests the error function (which as at Level 0 coincides with the t_0 χ^2 per degree of freedom, i.e. $\chi^2[\mathcal{T}[f], \mathcal{D}_1]$) of the best fit will be around one. Moreover, also here there exist *perfect* solutions for the minimization, for which the χ^2 of the fitted PDFs coincides exactly with the χ^2 computed with the input PDFs used for the generation of the pseudo-data, i.e. such that $\chi^2[\mathcal{T}[f_{\text{fit}}], \mathcal{D}_1] = \chi^2[\mathcal{T}[f_{\text{in}}], \mathcal{D}_1]$. Nevertheless, not all PDFs produced by the fitting will be perfect, and the distribution of the fitted PDFs will be too narrow.

tag	f_{in}	level	stopping	training fraction	details
C1	MSTW08NLO	0	fixed length 1k	100%	NNPDF2.3 basis 2-20-15-1 NN arch with positivity
C2	MSTW08NLO	0	fixed length 3k	100%	
C3	MSTW08NLO	0	fixed length 5k	100%	
C4	MSTW08NLO	0	fixed length 10k	100%	
C5	MSTW08NLO	0	fixed length 30k	100%	
C6	MSTW08NLO	0	fixed length 50k	100%	
C7	MSTW08NLO	0	fixed length 100k	100%	
C8	MSTW08NLO	1	lookback 30k	50%	
C9	MSTW08NLO	2	lookback 30k	50%	
C10	CT10	2	lookback 30k	50%	
C11	MSTW08NLO	2	lookback 80k	50%	
C12	MSTW08NLO	2	lookback 30k	50%	
C13	MSTW08NLO	2	lookback 30k	25%	
C14	MSTW08NLO	2	lookback 30k	75%	
C15	MSTW08NLO	2	lookback 30k	50%	
C16	NNPDF3.0	2	lookback 30k	50%	
C17	NNPDF3.0	2	lookback 30k	50%	

Table 8: List of closure fits used in this Section. The tag in the first column is used to identify the fits in the text. The second column reports the choice of f_{in} for each fit. The level of the closure fit can be found in the third column, while the fourth column specifies the stopping method and the length of the fit, and the last columns provides additional relevant information.

- **Level 2.**

Now, starting from the shifted pseudo-data in Eq. (20), we generate N_{rep} Monte Carlo replicas as in the standard NNPDF procedure. Schematically, this means that we will have, for each replica k , a set of pseudo-data $\mathcal{D}_2^k = \{D_I^{2,k}\}$ with

$$D_I^{2,k} = \left(1 + r_I^{\text{nor},k} \sigma_I^{\text{nor}}\right) \left(D_I^1 + \sum_{p=1}^{N_{\text{sys}}} r_{I,p}^{\text{sys},k} \sigma_{I,p}^{\text{sys}} + r_I^{\text{stat},k} \sigma_I^{\text{stat}}\right), \quad (21)$$

for $k = 1, \dots, N_{\text{rep}}$, and the set of random numbers is different replica by replica. From the practical point of view, once we have a set of Level 1 pseudo-data Eq. (20), the Level 2 N_{rep} Monte Carlo pseudo-data replicas Eq. (21) are obtained using exactly the same code as is used for the fits to real data. Fits to the Level 2 pseudo-data replicas are also performed in exactly the same way as to real data replicas, using the error function, and with different random seeds to initialize the minimization fitting algorithm for each replica.

In Level 2 fits, each Monte Carlo replica represents a fluctuation around the Level 1 pseudo-data, and the procedure should correctly propagate the fluctuations in the pseudo-data, due to experimental statistical, systematic and normalization uncertainties, into the fitted PDFs. The fit to each data replica yields a PDF replica f_{fit}^k , and the ensemble of PDF replicas then contains all the information on PDF uncertainties and correlations. Hence we expect the final error function of a Level 2 fit (taken as $\chi^2[\mathcal{T}[f_{\text{fit}}^k], \mathcal{D}_2^k]$ just like in a fit to real data) to be close to two (since, in a sense, each replica contains two fluctuations), while the χ^2 per degree of freedom of the replica PDFs to the pseudo-data (i.e. $\chi^2[\mathcal{T}[f_{\text{fit}}^k], \mathcal{D}_1]$) will be close to one. Moreover we expect the input PDFs f_{in} to lie within the 1σ band of the fitted PDFs with a probability of around 68%.

In the following sections, we discuss in detail the results of these closure tests, starting from Level 0 and then moving to Level 1 and Level 2, including variations of some of the fitting

methodology settings such as the stopping or the fraction of trained data. For ease of reference, the closure fits discussed in this section have been collected in Table 8, where we collect the tag of the fit, the set of input PDFs used for the pseudo-data generation, the level of closure fit, the settings for the stopping and data training fraction, and other relevant information. All these fits are done using NLO theory, though of course given the use of pseudo-data this choice is immaterial in the interpretation of the closure tests.

4.2 Validation of the training efficiency: Level 0 closure tests

We present the results of a number of Level 0 closure tests, using them to assess the training efficiency of the NNPDF3.0 minimization. Indeed, in Level 0 fits there exists in principle an optimal solution for the minimization: that which reproduces exactly the input PDFs used for the pseudo-data generation. Therefore, we can study which minimization strategy gets closer to this optimal solution, with the smallest computational effort.

4.2.1 Training methodology efficiency

The two main ingredients of our fitting methodology that can be tested in Level 0 closure tests are the adequacy of the neural network architecture, ensuring that it provides a flexible enough parametrization to reproduce any kind of input PDFs, and the efficiency of the genetic algorithm minimization, by comparing different options and verifying which one gets to the optimal solution quickest. Since in this case no stochastic fluctuations are added in the generation of the pseudo-data, when the fitted PDFs are equal to the input ones, the χ^2 of the fit, as already mentioned, will exactly vanish: $\chi^2[\mathcal{T}[f_{\text{fit}}], \mathcal{D}_0] = 0$. As the length of the training is increased, the fitted PDFs are expected to get closer and closer to the input ones.

In order to verify this feature, we have performed a number of fixed length fits to the full data set, and studied the dependence of the χ^2 of the fitted PDFs, $\chi^2[\mathcal{T}[f_{\text{fit}}], \mathcal{D}_0]$, on the training length. The fits used for this study are the fits C1 to C7 in Table 8. All these fits use identical settings with the only difference being the training length, that ranges from short fits of 1k GA generations to very long fits of 100k GA generations. In all these Level 0 closure tests, the input PDF set is taken to be MSTW08 NLO.

The dependence of the χ^2 on the training length for these Level 0 closure tests is plotted in Fig. 4. We compare the results for the genetic algorithm used in NNPDF3.0, based in particular on the nodal mutations strategy discussed in Sect. 3.3.1, with the corresponding genetic algorithm used in the NNPDF2.3 fit. The results for each closure test are marked with crosses, with auxiliary lines joining these points to guide the eye. It is clear from the figure that as expected the χ^2 of the fit keeps decreasing as the fit length is increased, with a behaviour that is approximately described by a power law. Therefore, it is possible to achieve an arbitrarily small value of χ^2 by increasing the training length, though as we approach the minimum the number of generations needed will grow exponentially.¹

It is also clear from Fig. 4 that the NNPDF3.0 minimization, based on the nodal genetic algorithm, is substantially more efficient than the NNPDF2.3 one, reaching values of χ^2 which are up to one order of magnitude smaller for a given fixed training length. This of course will be of great importance when we come to the fits to experimental data, allowing a more efficient exploration of the minima of the parameter space.

¹ This is a well known behavior of genetic algorithms. In this particular case, once we are close enough to the absolute minimum, it might be more useful to switch to other strategies like steepest descent, though of course this is purely an academic issue.




Figure 4: The normalized χ^2 of Level 0 closure tests, Eq. (19), as a function of the length of the genetic algorithms minimization. We compare the results for the GA settings used in NNPf3.0, based in particular on the nodal mutations strategy, with the corresponding GA settings used in the NNPf2.3 fit. The results for each closure test are marked with crosses; lines joining these points are added to guide the eye.

Given that the χ^2 is essentially vanishing, we expect almost perfect agreement between the fitted and input PDFs. Indeed such an agreement can be found in the corresponding comparison in Figs. 5 and 6. In these plots we show, for the Level 0 fit with a training length of 100K (fit C7), the results of the fit and the input PDFs, in this case MSTW08. The central value of our fitted PDFs is computed as the average over replicas in the usual way:

$$\langle f_{\text{fit}} \rangle = \frac{1}{N_{\text{rep}}} \sum_{k=1}^{N_{\text{rep}}} f_{\text{fit}}^k, \quad (22)$$

the angled brackets denoting the average over replicas, and the variance as,

$$\sigma_{\text{fit}}^2 = \langle (f_{\text{fit}}^2 - \langle f_{\text{fit}} \rangle)^2 \rangle = \frac{1}{N_{\text{rep}}} \sum_{k=1}^{N_{\text{rep}}} \left(f_{\text{fit}}^k - \langle f_{\text{fit}} \rangle \right)^2. \quad (23)$$

It is clear from these plots that the neural network architecture is flexible enough to reproduce the input PDFs, and that the GA is working as expected. It is interesting to observe that PDFs for which experimental information is very dense, such as for example the up quark in the valence region, are perfectly reproduced with zero uncertainty. PDFs for which information is more sparse or indirect, such as for example the gluon, have an uncertainty even when the χ^2 at the datapoints is essentially zero, because there is still a certain freedom in interpolating between datapoints.

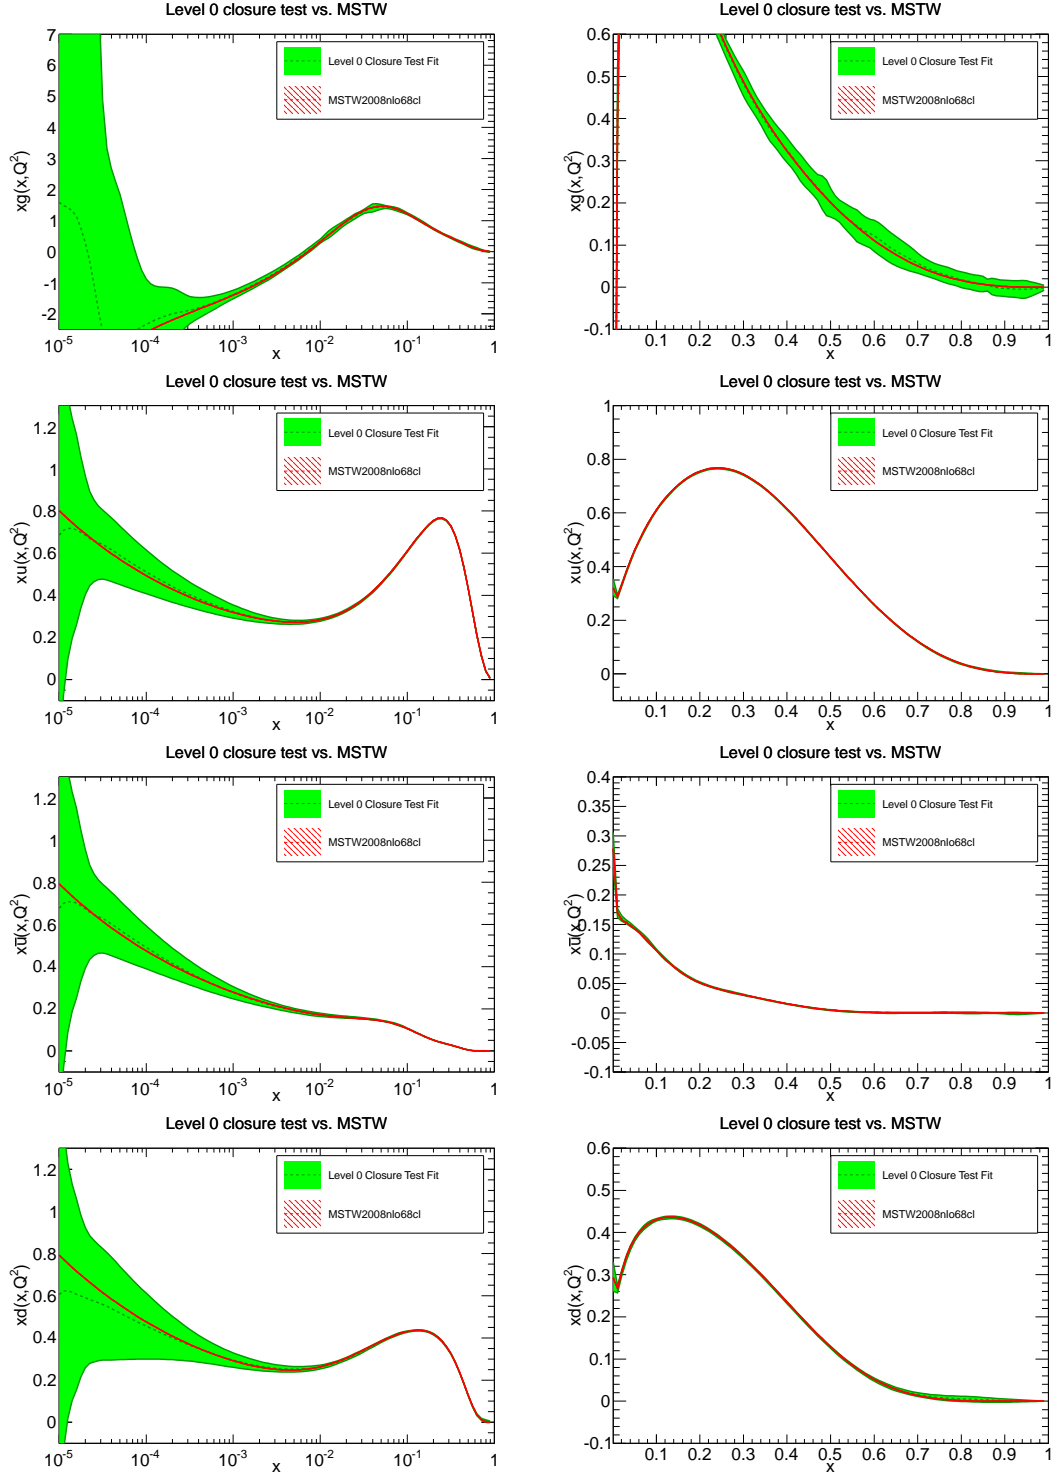


Figure 5: Comparison between the results of the Level 0 closure fit with 100K GA generations (fit C7) and the corresponding input PDF set, the central value of MSTW08 NLO PDF set. The green band shows the one-sigma interval computed over the sample of $N_{\text{rep}} = 100$ replicas, with the green dotted line showing the mean value. The plots show the gluon, u , \bar{u} and d PDFs on both linear (right hand side) and logarithmic (left) scales in x , at the scale $Q^2 = 1 \text{ GeV}^2$ where the PDFs are parametrized.

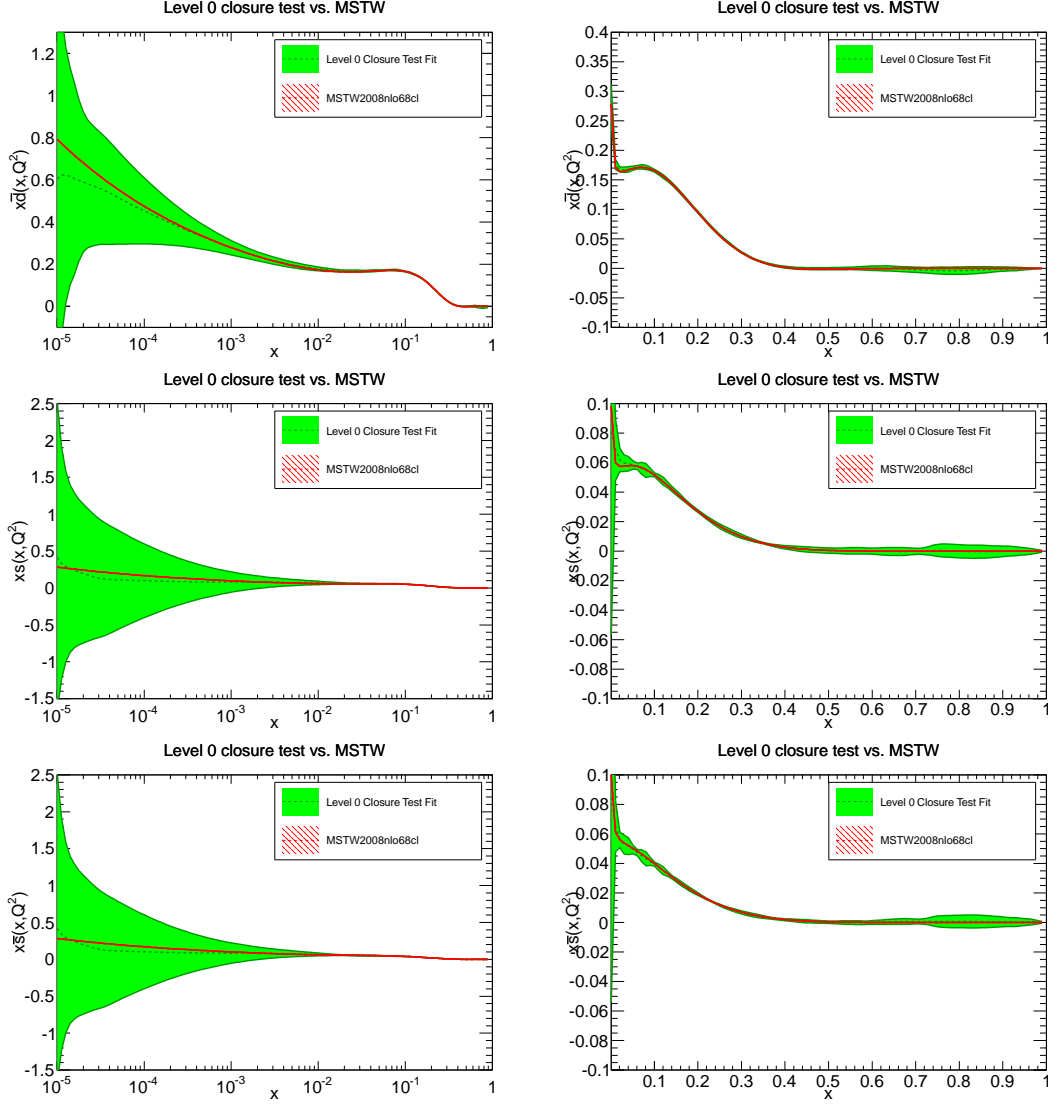


Figure 6: Same as Fig. 5 for the \bar{d} , s and \bar{s} PDFs.

Additional interesting information can be extracted from Level 0 fits by looking at the PDF uncertainties of the resulting fit, computed as usual as the standard deviation over the sample of $N_{\text{rep}} = 100$ fitted replicas, either at the level of parton distributions or at the level of physical observables. Recall that at Level 0 the only difference between each of the replicas is the random seed used to initialize the minimization, and therefore, their spread quantifies different possibilities of approaching the underlying input PDFs. Given that the Level 0 input pseudo-data do not fluctuate, the cross-sections computed from the fitted PDFs must converge to these input values for each replica as the training length is increased; i.e. the uncertainty on the predicted value for all the observables included in the fit must go to zero.

To verify this expectation, it is convenient to define an indicator which measures the standard deviation over the replica sample in units of the data uncertainty. This can be defined as follows:

$$\varphi_{\chi^2} \equiv \sqrt{\langle \chi^2[\mathcal{T}[f_{\text{fit}}], \mathcal{D}_0] \rangle - \chi^2[\langle \mathcal{T}[f_{\text{fit}}] \rangle, \mathcal{D}_0]}. \quad (24)$$

To see that this does the job, consider the first term in Eq. (24) in more detail: using the




Figure 7: The estimator φ_{χ^2} Eq. (24), as a function of the length of the genetic algorithms minimization. We compare the results for the GA settings used in NNPDF3.0, with the corresponding GA settings used in the NNPDF2.3 fit. The results for each closure test are marked with crosses, with auxiliary lines joining these points.

definition of the χ^2 in Eq. (19), we find that

$$\begin{aligned}
N_{\mathcal{D}} \langle \chi^2[\mathcal{T}[f], \mathcal{D}] \rangle &= \left\langle \sum_{I,J} (T_I[f] - D_I) C_{IJ}^{-1} (T_J[f] - D_J) \right\rangle \\
&= \sum_{I,J} \langle T_I[f] C_{IJ}^{-1} T_J[f] \rangle - \sum_{I,J} \langle T_I[f] \rangle C_{IJ}^{-1} D_J - \sum_{I,J} D_I C_{IJ}^{-1} \langle T_J[f] \rangle + \sum_{I,J} D_I C_{IJ}^{-1} D_J,
\end{aligned}$$

so that

$$\langle \chi^2[\mathcal{T}[f], \mathcal{D}] \rangle - \chi^2[\langle \mathcal{T}[f] \rangle, \mathcal{D}] = \frac{1}{N_{\mathcal{D}}} \sum_{I,J} (\langle T_I[f] C_{IJ}^{-1} T_J[f] \rangle - \langle T_I[f] \rangle C_{IJ}^{-1} \langle T_J[f] \rangle). \quad (25)$$

Thus in terms of the covariance matrix of the theoretical predictions

$$T_{IJ} \equiv (\langle T_I[f] T_J[f] \rangle - \langle T_I[f] \rangle \langle T_J[f] \rangle), \quad (26)$$

we have

$$\varphi_{\chi^2}^2 \equiv \frac{1}{N_{\mathcal{D}}} \sum_{I,J} C_{IJ}^{-1} T_{JI}, \quad (27)$$

i.e. the average over all the data points of the uncertainties and correlations of the theoretical predictions, T_{IJ} , normalized according to the corresponding uncertainties and correlations of the data as expressed through the covariance matrix C_{IJ} . If the covariance matrix was diagonal, i.e. in the absence of correlations, this would just be the variance of the predictions, divided by the experimental variance, averaged over datapoints; $\varphi_{\chi^2}^2$ is thus recognized to provide the generalization of this in the presence of correlations.

In Fig. 7 we show φ_{χ^2} , Eq. (24), for the Level 0 fits as a function of the length of the genetic algorithms minimization. As before, we compare the results for the GA settings used in NNP3.0, with the corresponding GA settings used in the NNP2.3 fit. We can see that indeed, as we increase the training length, the spread of the theoretical predictions at the data points for different replicas at the level of fitted cross-sections decreases monotonically, and again here we observe the improvement from the more efficient minimization strategy in NNP3.0.

On the other hand, the fitted PDFs themselves need not become identical replica by replica, even at Level 0, since different functional forms for the PDFs can yield the same predictions for the observables. Of course, in the kinematic regions where data are available, the fluctuations of the PDFs are limited by the fact that the neural networks provide a smooth interpolation between points that are constrained by data. In this respect, in regions where the PDFs are well constrained by experimental data, PDF uncertainties should be very small, and indeed they are.

This is however not necessarily the case in the extrapolation region, since the functional forms there can vary substantially without affecting the fitted figure of merit, the $\chi^2[\mathcal{T}[f_{\text{fit}}], \mathcal{D}_0]$, we expect large PDF uncertainties, which moreover are essentially independent of the training length. These two effects, very small PDF uncertainties in the data region, and large PDF uncertainties in the extrapolation regions, in particular at small and large x are still clearly visible in the plots in Fig. 5, even at the end of a 100k-generation training. These results provide a way of quantifying the *extrapolation* uncertainty on the PDFs due to the lack of direct constraints in these region: this is an irreducible source of PDF uncertainty that can only be reduced if new data is provided, and that accounts for most of the PDF uncertainties in the fits to real data in the extrapolation regions.

4.2.2 Effective preprocessing exponents

When using as an input to the closure test a PDF set based on a relatively simple functional form, such as for example MSTW08 or CT10, the exponents that characterize the asymptotic small- x and large- x behavior of the individual PDFs are known, since they are part of the set of parameters defining the PDF parametrization at the initial scale. For instance, in the MSTW08 fit, the PDFs are parametrized at the initial scale of $Q_0^2 = 1 \text{ GeV}^2$ using a generic functional form $Ax^{-\alpha}(1-x)^{\beta}(1+\epsilon x^{0.5}+\gamma x)$ and therefore, the asymptotic small- x and large- x behaviors will both be powerlike, parametrized by the exponents α and β . Note that in general one has to be careful with this naive interpretation since sub-asymptotic corrections can be numerically large.

In the context of NNP fits, no explicit assumption is made on the functional form of PDFs, which are instead parametrized using neural networks. On the other hand, preprocessing is introduced in order to absorb the dominant behavior of the fitted PDFs at small- and large- x and speed up the fitting. As explained in Sect 3.2.2, when fitting real data, the range in which preprocessing is varied for each PDF is determined dynamically through an iterative procedure.

In the context of a Level 0 closure test, when a PDF set based on a known functional form is used as input, it is interesting to verify that the small- x and large- x behavior of the fitted PDFs reproduce the ones of the input PDFs. This can be achieved by comparing the effective preprocessing exponents between the fitted and input PDFs. These effective exponents can be computed as in Eq. (11-12). The results of this comparison are shown in Fig. 8, for the Level 0 closure test with 100K generations (fit C7 in Table 8). A beautiful agreement is found, with the MSTW08 value always within in the one-sigma band of the fitted PDFs, thus showing that our methodology is capable of precise quantitative predictions on the behavior of the functions that we are trying to fit. Similar results are obtained for all the other PDF combinations.




Figure 8: Small- x (left) and large- x (right) behaviour of the fitted PDF for u_v compared to the behaviour of the MSTW08 input functions. In both cases the values are obtained using Eq. (11-12).

4.3 PDF uncertainties in closure tests

In the previous section we studied Level 0 closure tests, in which the fit quality can become arbitrarily good, and the PDF uncertainties arise purely due to the fact that the experimental data used in the fit has finite kinematical coverage. Now we turn to Level 1 and Level 2 closure tests, and in doing so we shed some light, in the cleanly controlled environment of closure testing, on the various different origins of PDF uncertainties: specifically those due to the fluctuations of the experimental data, to the choice of functional form, and to the extrapolation uncertainties due to the finite coverage of the data. Following this, we go on to quantify how well the closure test is passed, both in terms of central values and uncertainties, using various statistical estimators.

4.3.1 PDF uncertainties: data, functional and extrapolation components

A more sophisticated understanding of the various sources that form the total PDF uncertainties can be obtained in the context of closure tests by comparing Level 0, Level 1, and Level 2 fits. Indeed, in each of these fits the PDF uncertainty band has different components. In Level 0 fits, the only significant component is the interpolation and extrapolation uncertainty (which we will refer to as extrapolation uncertainty for short); in Level 1, one also has the uncertainty due to the choice of functional form; and in Level 2 finally one also adds the uncertainties due to the fluctuations of the experimental data. Therefore, by comparing Level 0, Level 1 and Level 2 closure fits we can analyze how the total PDF uncertainties are decomposed into data, functional and extrapolation uncertainties.

Let us begin with the extrapolation uncertainty. As discussed in the previous section, in a Level 0 closure test, the PDF uncertainty in the fitted datapoints should go to zero as the training length is increased. This implies that, in the experimental data region, PDF uncertainties should also decrease monotonically as a function of the training length wherever data are available. However, in between data (interpolation) and outside the data region (extrapolation) PDFs can fluctuate. This residual uncertainty, which remains even at infinite Level 0 training length, we refer to as the *extrapolation* uncertainty.

To illustrate this point, in Fig. 9 we show the results of a Level 0 closure tests for two fits with identical settings but different training length, 5k and 100k GA generations respectively. We can see that while the uncertainties at 100k generations are smaller than at 5k, they do not go to zero, and in fact for small values of x , where there is no experimental information,

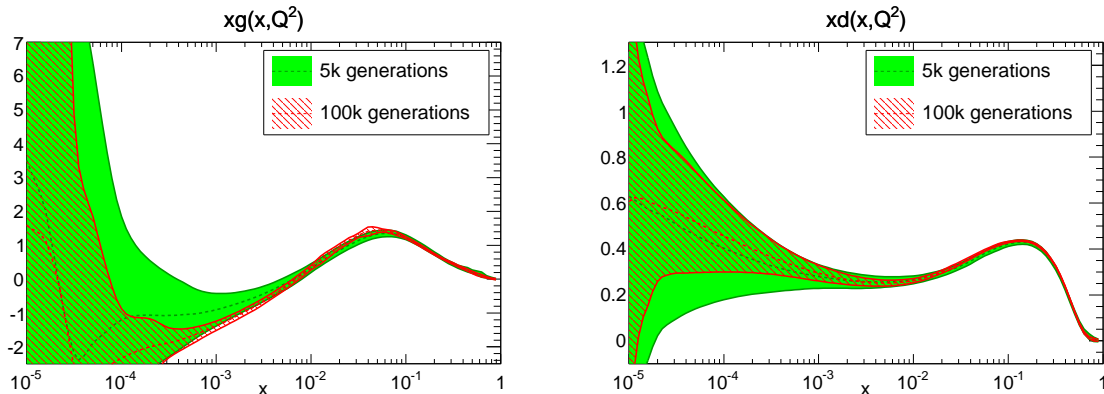


Figure 9: Comparison of the gluon (left plot) and down (right plot) PDFs for a Level 0 closure test with MSTW2008 NLO as input set, for two different training lengths, and all other fit settings identical. The green band shows the results with 5k GA generations and the red band those with 100k generations. PDFs are shown at the input parametrization scale of $Q^2 = 1 \text{ GeV}^2$.

they remain very large. This is a direct consequence of the lack of experimental data in this region, and illustrates this irreducible component of the PDF uncertainties that we denote as extrapolation uncertainty.

Note that, given the highly non-trivial dependence of PDFs on the measured cross-sections, and the wide range of observables included in the fit, it is very difficult to determine precisely how this extrapolation region is defined: while a non-negligible extrapolation component is expected for all PDFs at small enough and large enough values of x , significant uncertainties due to interpolation could also be present at intermediate x . Note that this also accounts for possible degeneracies between PDFs: for example, even in Level 0 closure fits some PDFs can compensate each other to produce exactly the same cross-section, and this effect will also be included in the extrapolation uncertainty.

Now we can see how the total PDF uncertainty also contains functional and data uncertainties by comparing the results of Level 0, Level 1 and Level 2 closure tests with otherwise all other fit settings identical. Fig. 10 shows the ratios of the uncertainty of the fitted PDFs to the respective central values in each case, for the C5, C8 and C9 closure test fits, see Table 8, corresponding to Level 0, Level 1 and Level 2 closure tests respectively. The MSTW08 NLO set is used as input PDF f_{in} , and a total training length of 30k generations is used in all cases. For the three fits, the PDF uncertainty bands are defined as the 68% confidence interval from the sample of $N_{\text{rep}} = 100$ fitted replicas. Results are provided for the PDFs in the flavor basis at the input parametrization scale of $Q^2 = 1 \text{ GeV}^2$, though the qualitative interpretation is the same in any other basis.

In order to understand these results, it is useful to review our expectations for the results of the Level 1 and Level 2 closure fits, as we have done above for Level 0. In a Level 1 fit, the central values of the data have been fluctuated around the theoretical prediction, and therefore $f_{\text{fit}} = f_{\text{in}}$ no longer provides an absolute minimum for the χ^2 . Indeed, provided the PDF parametrization is flexible enough, the minimization algorithm should find a large number of different functional forms that yield an equally good $\chi^2[\mathcal{T}[f_{\text{fit}}], \mathcal{D}_1] \approx 1$. Therefore, in Level 1 closure fits, on top of the extrapolation component, the total PDF uncertainty will include a new component, which we refer to as *functional* uncertainty.

This functional uncertainty is a consequence of the fact that, as discussed in Sect. 3.3.2, the optimal χ^2 in the presence of data fluctuations is not the absolute minimum of the χ^2 . In a

closure test, this is obvious: the optimal result corresponds to the true underlying functional form, and thus the optimal χ^2 is the χ^2 of the Level 1 pseudodata, whose value is of order one (and tends to one in the limit of infinite size of the data sample). For an infinite-dimensional space of functions, this χ^2 value can be obtained in an infinity of different ways, whose spread provides the functional uncertainty. The look-back method discussed in Sect. 3.3.2 ensures that regardless of the length of the fit, the final χ^2 does not decrease beyond the overlearning point, and thus again this functional uncertainty will survive even for infinite training length.

In a Level 2 fit, the starting point is again the Level 1 pseudo-data generated by adding a Gaussian fluctuation over the predictions obtained from the input PDFs. Now however there is a second step, which reproduces the actual fitting strategy used in the NNPDF framework: starting from these pseudo-data, a set of N_{rep} Monte Carlo replicas is generated that reflects the statistical and systematic errors given by the experimental measurements. Because each replica fluctuates around Level 1 data, the expectation for the minimized figure of merit in Level 2 fits is $\chi^2[\mathcal{T}[f_{\text{fit}}], \mathcal{D}_2] \approx 2$, although we still expect $\chi^2[\mathcal{T}[f_{\text{fit}}], \mathcal{D}_1] \approx 1$. In Level 2 closure tests, each replica is fitted with exactly the same algorithm, yielding an ensemble of fitted PDFs $\{f_{\text{fit}}^k\}$ whose statistical properties are a faithful propagation of the fluctuations in the underlying dataset. The increase in the uncertainty from Level 1 to Level 2 fits is the genuine *data* uncertainty.

Following these considerations, it is possible to understand quantitatively the features that are observed in Fig. 10. Firstly, we see that Level 0 uncertainties are smaller than Level 1 and in turn these are smaller than those at Level 2: this confirms the expectation that at each Level we are adding a new component of the total PDF uncertainty, extrapolation, functional and data components, respectively. We also observe that in the small- x and large- x regions it is the extrapolation uncertainty that dominates, that is, the Level 2 PDF uncertainties are already reasonably reproduced by those of Level 0 closure fits. However, as we explained, Level 0 extrapolation component can be not negligible also in some regions with experimental data, like medium x , due to interpolation and degeneracies. By comparing to the Level 1 results, we see that the functional uncertainty is generally sizable, especially in regions at the boundary between data and extrapolation, in particular at large x .

Interestingly, in regions where we have a rather reasonable coverage from available data, the three components are roughly of similar size. Take for example the gluon around $x \sim 10^{-3}$, which is well constrained by the high-precision HERA measurements. We see that the extrapolation, functional and data uncertainties are all of similar size, and thus a proper estimate of the total PDF uncertainty must include all three components. The same applies for other PDF flavors, such as for example strangeness for $x \gtrsim 0.01$ (with abundant constraints from neutrino DIS and LHC data) or the up and down quarks at medium and large- x (with many DIS and LHC datasets providing information on them).

An important general conclusion is thus that data uncertainties are not dominant anywhere, and thus a PDF determination that does not include the extrapolation and functional components will underestimate the overall PDF uncertainty. This conclusion is consistent with that of previous rather less sophisticated, NNPDF studies such as those of Ref. [147]. It is natural to conjecture that the tolerance method [148] which is used in Hessian fits, provides an effective way of supplementing the data uncertainty which the Hessian method gives with these extra necessary components.

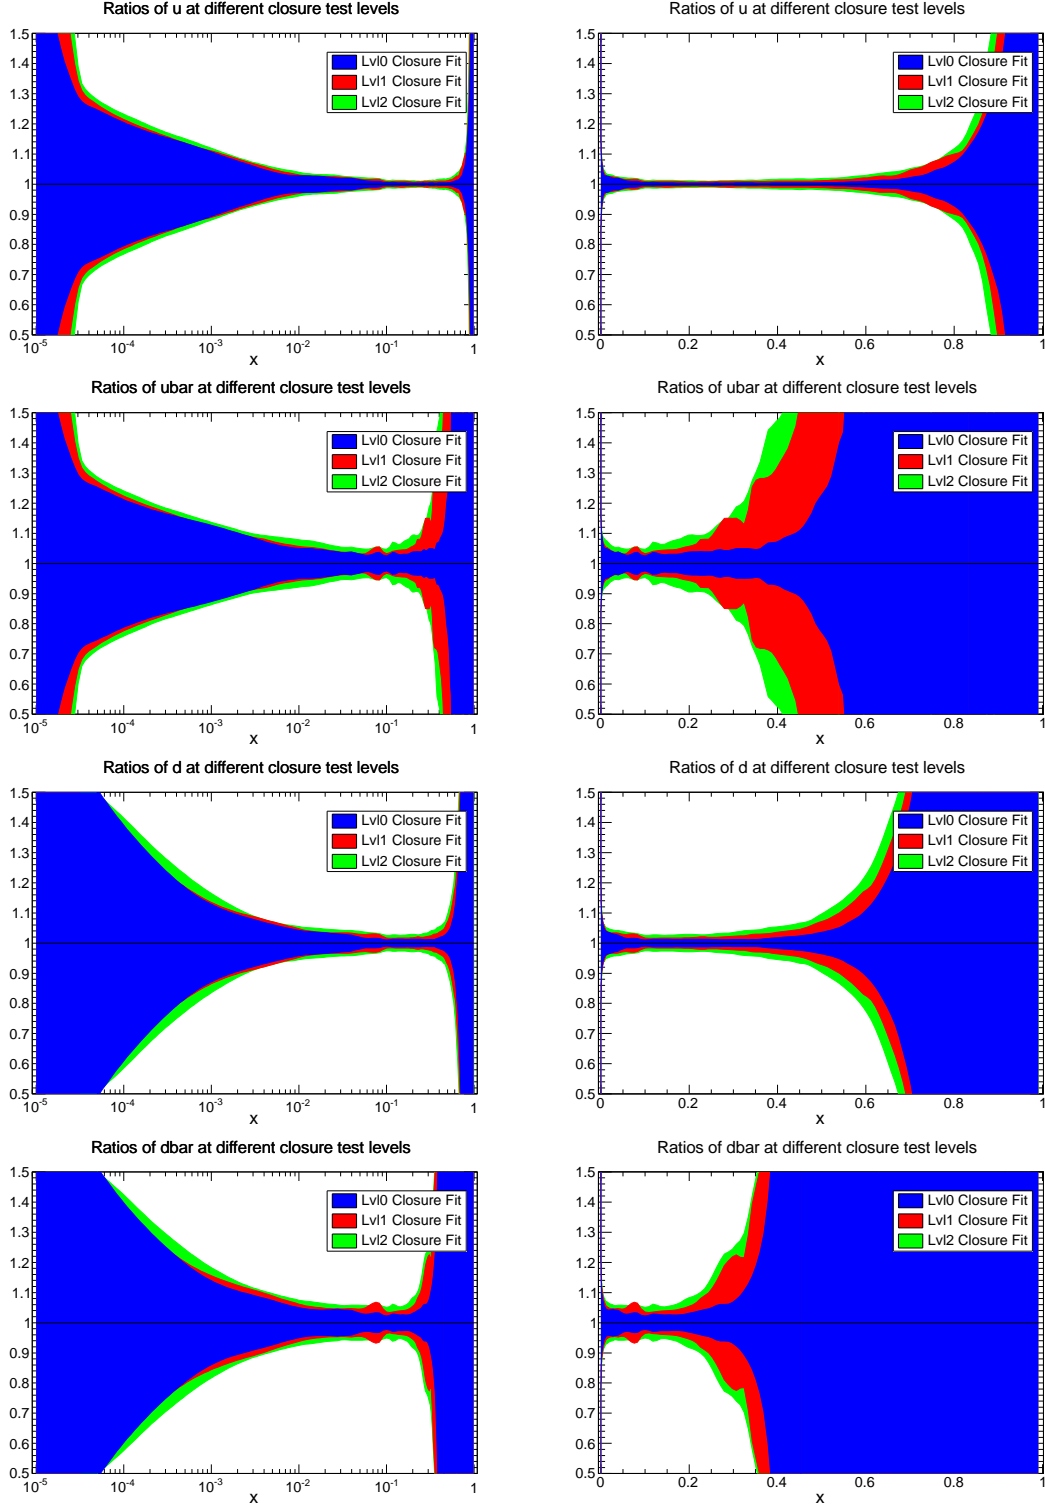


Figure 10: Comparison of the relative PDF uncertainties obtained from Level 0, Level 1 and Level 2 closure test fits. The plots show the 68% confidence level PDF uncertainty band for each of the fits, normalized to the corresponding central value of each fit. The results for the u, \bar{u}, d and \bar{d} PDFs are shown at the initial parametrization scale $Q^2 = 1 \text{ GeV}^2$, both in a logarithmic (left) and in a linear (right) scales.

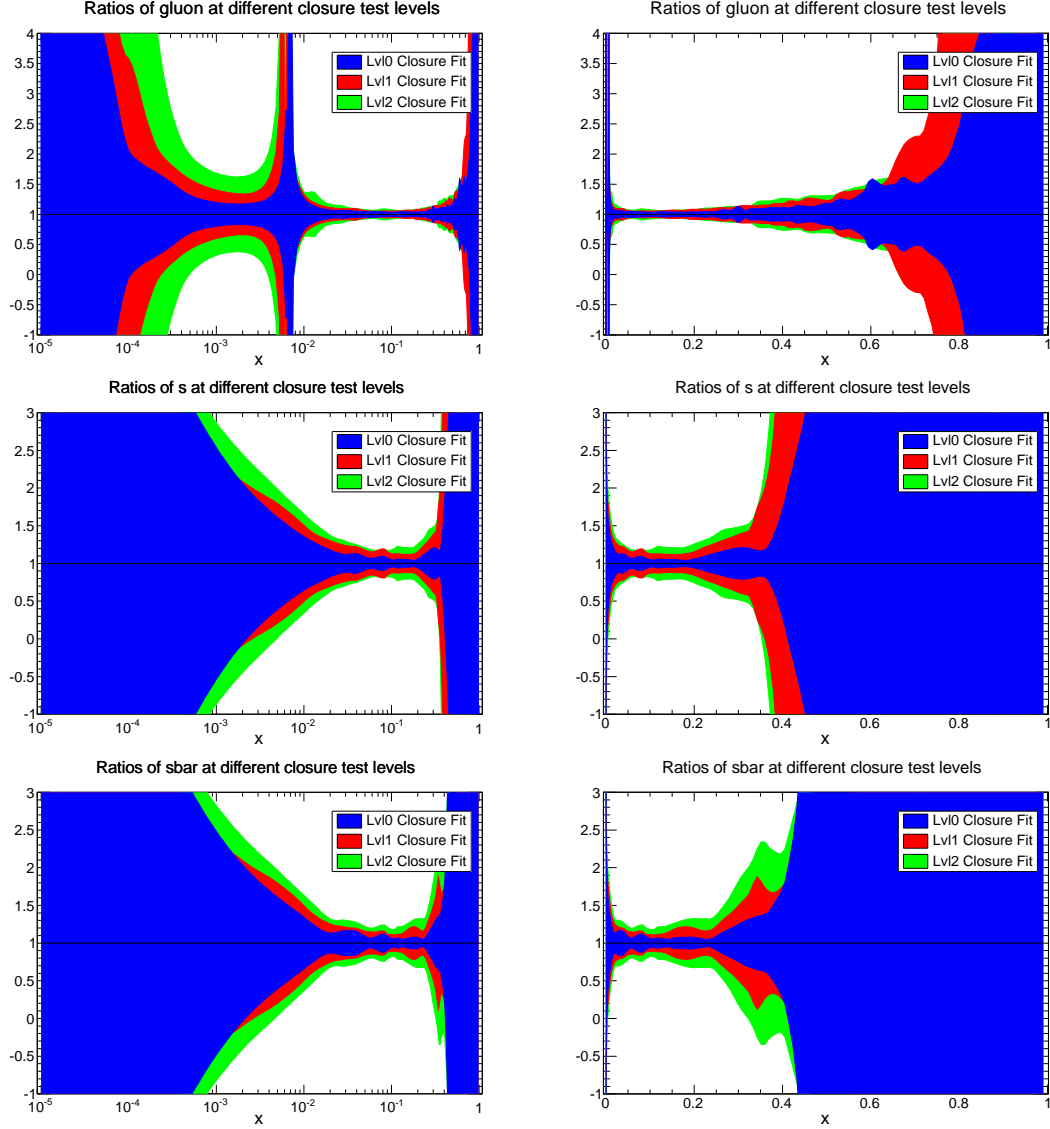


Figure 10 (Cont.): Same as above for the g , s and \bar{s} PDFs. The spike on the gluon PDF at $x \sim 0.005$ is caused by the fact that the gluon from the input PDF set, MSTW08, has a node in this region, and thus the relative PDF uncertainty blows up.

4.3.2 Quantitative validation of PDF uncertainties in closure tests

In the previous section we discussed the various components of the total PDF uncertainty in a purely qualitative way, by comparing the results of Level 0, Level 1 and Level 2 closure tests. In this section we provide some quantitative estimators of the successfulness and effectiveness of the closure tests, which will be used in Sect. 4.4 to validate the closure test fits. which we will use in the remainder of this section.

In Level 1 and Level 2 closure fits, we expect that the central χ^2 obtained from the average of the fitted PDFs should reproduce the one computed using the input PDFs, that is $\chi^2[\langle \mathcal{T}[f] \rangle, \mathcal{D}_1] \approx \chi^2[\mathcal{T}[f_{\text{in}}], \mathcal{D}_1]$ where by \mathcal{D}_1 we indicate that that we use the Level 1 pseudo-data, which include the fluctuations.

We can test for a particular fit whether this is the case by defining the following statistical

estimator

$$\Delta_{\chi^2} = \frac{\chi^2[\langle \mathcal{T}[f] \rangle, \mathcal{D}_1] - \chi^2[\mathcal{T}[f_{\text{in}}], \mathcal{D}_1]}{\chi^2[\mathcal{T}[f_{\text{in}}], \mathcal{D}_1]}, \quad (28)$$

that is, the difference between the central χ^2 of the closure test fit, computed with an average over replica PDFs, and the χ^2 of the input PDF set, both computed with respect to the same closure test dataset. This estimator is therefore a measure of how close the closure test fit reproduces the theoretical predictions of the input PDF, and therefore it provides a quantitative measure of the success of the test.

In particular, noting that in genetic algorithm minimization χ^2 is a decreasing function along the training, $\Delta_{\chi^2} > 0$ corresponds to underlearning (the optimal χ^2 has not been reached yet), $\Delta_{\chi^2} = 0$ corresponds to perfect learning of the underlying law, while in fixed-length genetic minimization $\Delta_{\chi^2} < 0$ would correspond to overlearning, i.e., the fit is learning the noise in the data. Of course in a closure test the true underlying χ^2 is known and thus one could stop at the optimal point by comparing to it, which if course not possible in a realistic situation. The effectiveness of the method that is actually used — in our case, the look-back method described in Sect. 3.3.1 — can then be gauged by simply evaluating how close Δ_{χ^2} Eq. 28 is at the actual stopping point.

The estimator Eq. (28) thus quantifies how successful the closure test fit is in terms of reproducing the central values of the input PDFs. We now introduce an estimator which allows for an assessment of the accuracy with which PDF uncertainties are reproduced. To this purpose, we first recall that for an unbiased estimator, assuming gaussianity, the n -sigma intervals about the prediction can be interpreted as confidence levels for the true value (see e.g. Ref. [15]): this means that the true value must fall on average within the one sigma uncertainty band in 68.3% of cases, within the two sigma band in 95.5% of cases, and so on.

We thus define the estimator

$$\xi_{n\sigma} = \frac{1}{N_{\text{PDF}}} \frac{1}{N_x} \frac{1}{N_{\text{fits}}} \sum_{i=1}^{N_{\text{PDF}}} \sum_{j=1}^{N_x} \sum_{l=1}^{N_{\text{fits}}} I_{[-n\sigma_{\text{fit}}^{i(l)}(x_j), n\sigma_{\text{fit}}^{i(l)}(x_j)]} \left(\langle f_{\text{fit}}^{i(l)}(x_j) \rangle - f_{\text{in}}^i(x_j) \right), \quad (29)$$

where n is a positive integer, and N_{PDF} , N_x and N_{fits} are the number of PDF flavors, x values and fits respectively, over which averages are performed. For the sampling of the PDFs in x , we use 20 points between 10^{-5} and 1, half of them log spaced below 0.1 and the rest linearly spaced. In Eq. (29), $I_A(x)$ denotes the indicator function of the interval A , that is, it is only non-zero if its argument lies in the interval A , and one otherwise. Finally, $\langle f_{\text{fit}}^{i(l)} \rangle$ and $\sigma_{\text{fit}}^{i(l)}$ are the average PDFs and the corresponding standard deviation of the i PDF flavor for fit l , computed over the sample of $N_{\text{rep}} = 100$ replicas of fit l .

The estimators $\xi_{1\sigma}$, $\xi_{2\sigma}$, ... provide the fraction of those fits for which the input PDF falls within one sigma, two sigma, etc of the central PDF $\bar{f}_{\text{fit}}^{i(l)}$, averaged over PDF flavors and values of x . In a successful closure test we must thus have that $\xi_{1\sigma} \approx 0.68$, $\xi_{2\sigma} \approx 0.95$, etc.

In principle, in order calculate $\xi_{1\sigma}$, Eq. (29), we would need to, for instance, generate 100 closure test fits each one with $N_{\text{rep}} = 100$ replicas, following the procedure explained above. However performing this very large number of fits is very computationally expensive, and we would instead like to obtain an estimate of $\xi_{1\sigma}$ which involves fewer fits. To achieve this, we can approximate the mean of each fit, $\langle f_{\text{fit}}^{i(l)} \rangle$, by fitting a single replica to each set of closure test data at Level 1, i.e. without additional replica fluctuations. We can also replace the individual values of $\sigma^{i(l)}$ in Eq. (29) with the corresponding values taken from a single 100 replica fit, making use of the fact that the variation in the PDF uncertainties between different closure test fits is small.

4.4 Validation of the closure test fits

Using the statistical estimators introduced in the previous section, we now move to validate quantitatively the results of the closure fits. First, we show how close the central values and χ^2 values of the input and fitted PDFs are to each other, both for the total dataset and for individual experiments. Then we discuss a quantitative validation of the PDF uncertainties obtained in the closure tests, using the estimators defined in Sec. 4.3.2 for this purpose. Finally in this section we will show how one can also use the Bayesian reweighting procedure [18,149] to provide further evidence that the closure tests are working as expected. Indeed, it turns out that reweighting provides the most stringent validation test, since to be successful it is necessary to reproduce not only central values and uncertainties, but also higher moments and correlations.

4.4.1 Central values

A first indicator of the quality of the closure test is provided by the values of the central χ^2 obtained using the fit, $\chi^2[\langle\mathcal{T}[f_{\text{fit}}]\rangle, \mathcal{D}_1]$. This should reproduce the values of the χ^2 obtained using the generating PDFs, that is $\chi^2[\mathcal{T}[f_{\text{in}}], \mathcal{D}_1]$. For our baseline Level 2 closure test, using pseudo-data based on MSTW08, and performed with the look-back stopping criterion rather than at fixed length (fit C9 in Table 8), we obtain the following result for Eq. (28),

$$\Delta_{\chi^2} = -0.011, \quad (30)$$

which shows that the fitted PDFs reproduce the χ^2 of the generating PDFs at the 1% level, with a small amount of overlearning, which must be viewed as an inefficiency of the look-back method. The result for the corresponding Level 1 fit (C8) is very similar: $\Delta_{\chi^2} = -0.015$.

This level of agreement is achieved not only for the total χ^2 , but also for the individual experiments included in the global fit. This is important to test, since each experiment fluctuates by a different amount, and we want the closure test to reproduce these fluctuations. Fig. 11 shows the contributions in the Level 2 fit C9 to $\chi^2[\langle\mathcal{T}[f_{\text{fit}}]\rangle, \mathcal{D}_1]$ from the pseudo-data generated for each individual experiment, compared to the corresponding contributions to $\chi^2[\mathcal{T}[f_{\text{in}}], \mathcal{D}_1]$. The horizontal bars show the total χ^2 for the two PDF sets, and are computed as the weighted average of the χ^2 of each individual dataset. The datasets shown are the same used in the baseline NNPDF3.0 global fit, see Tables 1 and 2 in Sect. 2.

We see from Fig. 11 that the closure test successfully reproduces the χ^2 of the input PDFs, MSTW08, not only for the total dataset but also experiment by experiment. Note that specially for experiments with only a small number of points the fluctuations added to the pseudo-data can lead to a χ^2 quite different from one, and this is exactly reproduced by the closure test results. Fig. 11 is thus a strong check that, at least at the level of central values, the Level 2 closure fits are successful.

4.4.2 PDF uncertainties: qualitative validation

We can get a first assessment of the uncertainty of the Level 2 closure fit by means of the estimator φ_{χ^2} Eq. (24). For the Level 2 closure fit C9 we get

$$\varphi_{\chi^2} = 0.254, \quad (31)$$

a result that indicates that the PDF uncertainty on the datapoints from the Level 2 closure fits is roughly a factor of four smaller than that the original experimental uncertainty, due to the combination of all the data into the PDFs. For the Level 1 fit C8 this number is rather less:

Distribution of χ^2 for experiments




Figure 11: Comparison of the χ^2 to the closure test data, $\chi^2[\mathcal{D}_1]$, obtained with the input (red) and with the fitted (green) PDFs, for the Level 2 closure test fit based on MSTW08 pseudo-data, fit C8 in Table 8. The horizontal bars show the total χ^2 for the two PDF sets, and are computed as the weighted average of the χ^2 of each individual dataset. The datasets shown are the same used in the baseline NNPDF3.0 global fit, see Tables 1 and 2 in Sect. 2.

$\varphi_{\chi^2} = 0.173$, confirming that at Level 1 the data uncertainty is still missing. Since contributions to φ_{χ^2} are entirely from the data points, it is reasonable to assume that the contribution to it from extrapolation uncertainties is relatively small, so we can use it to compare the functional and data uncertainties. Assuming further that these are combining in quadrature, this gives a figure of 0.186 for the data uncertainty, indicating that, at least on average, the data and functional uncertainties are roughly equal.

However, we must now verify that this error reduction is real, namely that uncertainties are correctly estimated by the closure test. We do this first in a more qualitative way, and then more quantitatively in Sects. 4.4.3-4.4.4 below.

To this purpose, we first look at the distance between the fitted PDF and the underlying “truth” in units of the standard deviation. This is most efficiently done using the distance estimators, as defined in Appendix B, between the closure test fit PDFs and the input MSTW08 PDFs. The distances between fitted and input PDFs in closure tests are normalized such that $d(x, Q) \sim 1$ correspond to agreement at the one-sigma level (in units of the uncertainties of the closure test fitted PDF), $d(x, Q) \sim 2$ correspond to agreement at the two-sigma level and so on, that is,

$$d_{\sigma} [f_{i,\text{fit}}, f_{i,\text{in}}] (x, Q) \equiv \sqrt{\frac{(\bar{f}_{i,\text{fit}}(x, Q) - f_{i,\text{in}}(x, Q))^2}{\sigma^2 [f_{i,\text{fit}}] (x, Q)}}, \quad (32)$$

where i stands for the PDF flavor. In the following, distances are computed at the initial parametrization scale of $Q^2 = 1\text{GeV}^2$.

Using the definition Eq. (32), the distances between the central values of the fitted and the input PDFs (MSTW08) in the Level 2 closure test C9 are shown in Fig. 12. The distances in Fig. 12 show that the fitted and input PDFs are in good agreement, at the level of one sigma or better (in units of the uncertainty of the fitted PDF), with some PDFs for some points in x differing by an amount between one and two sigma, as one would expect if the underlying distribution was roughly Gaussian. In the extrapolation regions, small- x and large- x , the distances between input and fitted PDFs become smaller because of the large PDF uncertainties in these regions.

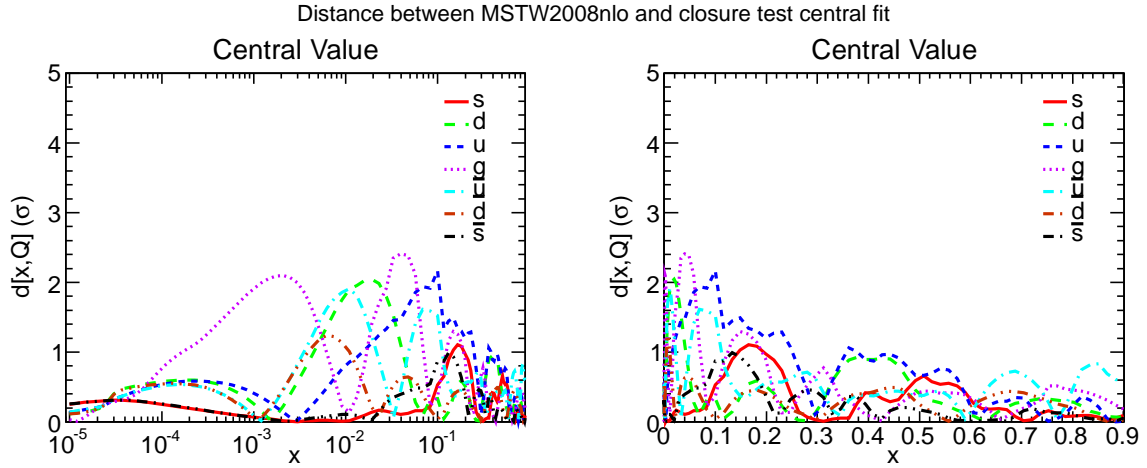


Figure 12: Distances, Eq. (32), between the central values of the fitted PDFs from the Level 2 closure test C8 and the MSTW08 PDFs, which are used as input to generate the theory predictions of the closure test. These distances are computed using Eq. (32), that is, normalized to the standard deviation of the fitted PDFs. Distances are computed at the input parametrization scale of $Q^2 = 1 \text{ GeV}^2$, see text for more details.

From the distance comparisons of Fig. 12 we can see that, at the qualitative level, the closure test is successful since the fitted PDFs fluctuate around the truth by an amount which is roughly compatible with statistical expectations. More insight on this comparison is provided by plotting, for all PDF flavors, the ratio $f_{\text{fit}}/f_{\text{in}}$ between the fitted and input PDFs, including both the fitted PDF central values and uncertainties. This comparison is shown in Fig. 13 for each parton flavor. It is clear from these plots that the NNPDF methodology reproduces successfully the input PDFs, with deviations from the input functions usually by two standard deviations at most. This comparison provides further evidence that PDF uncertainties are properly estimated in Level 2 closure tests, in that the central value of the fitted PDFs fluctuates around the truth by an amount which is consistent with the size of the PDF errors.




Figure 13: Ratio of the PDFs obtained from the Level 2 closure test C8, which uses MSTW08 as input, with respect to the input MSTW2008 PDFs themselves. The green band shows the one sigma interval of the fitted PDFs, while the green dotted line is the corresponding mean. The plots show the up and down PDFs on both linear (right hand side) and logarithmic (left) scales in x . The corresponding comparison in terms of distances, Eq. (32), is shown in Fig. 12. The comparison is performed at the fitting scale of $Q^2 = 1 \text{ GeV}^2$.




Figure 13 (Cont.): Same as above for gluon and strange PDFs.

A final qualitative comparison can be performed by using the arc-length, which, as discussed in Sect. 3.3.2, provides a measure of the complexity of a function defined over a finite interval. A function that has a more complicated structure is expected to have a larger arc-length, and vice-versa. In the context of our fits, PDFs with too large arc-lengths might be an indication of over-fitting, with a PDF trying to use a contrived shape in order to fit statistical fluctuations. For this reason, as explained in Sect. 3.3.2, replicas that at the end of the fit have some PDFs with unnaturally large arc-lengths are discarded from the final sample.

In the context of closure tests, comparing the arc-lengths of the fitting PDFs with those of the input PDFs provide an integrated comparison, rather than point by point, that the fitted PDFs reproduce the input ones. The arc-lengths of the Level 2 fit C9, using MSTW08 as input, are shown in Fig. 14. Arc-lengths are computed at the input parametrization scale of $Q^2 = 1 \text{ GeV}^2$. As expected from the comparison at the level of PDFs, there is also good agreement between the input and fitted PDFs at the level of arc-lengths, the arc-length of the input PDF typically lying within the 68% confidence level of the fitted PDFs.




Figure 14: Comparison of the arc-lengths of the fitted and input PDFs, calculated using Eq. (18) for the C8 Level 2 closure test, using MSTW08 NLO as input. Arc-lengths are computed at the input parametrization scale of $Q^2 = 1 \text{ GeV}^2$. For the fitted PDFs, we show the central value and the 68% confidence interval.

4.4.3 PDF uncertainties: quantitative validation

The results of Sect. 4.4.2 provide evidence the Level 2 closure tests are successful, both in terms of central values and of PDF uncertainties. We now provide a quantitative assessment using the indicators of Sect. 4.3.2. more quantitative confirmation of this success.

In order to validate quantitatively central values we compute the $\xi_{n\sigma}$ estimators Eq. (29) for Level 2 fit C9. We obtain the following results:

$$\xi_{1\sigma}^{(12)} = 0.699, \quad \xi_{2\sigma}^{(12)} = 0.948, \quad (33)$$

to be compared with the theoretical expectations of 0.68 and 0.95. This excellent agreement confirms that the PDF replicas obtained by our fitting methodology can indeed be interpreted as a representation of the probability distribution for the PDFs given the data used in the fit.

To verify that this agreement is not accidental, or a fluke of the definition of the estimator $\xi_{n\sigma}$, but rather a robust feature of our analysis, we have redone the whole procedure but this time based on Level 1 closure fits, set up in the same way as C8, and computed again the $\xi_{n\sigma}$ estimators. We already know from the qualitative comparisons between Level 1 and Level 2 uncertainties from Fig. 10 that in Level 1 closure tests PDF uncertainties are underestimated, and therefore in this case the central value of the fitted PDF will fluctuate more than its estimated uncertainty would suggest.

This implies that in Level 1 closure tests we expect the $\xi_{n\sigma}$ estimators to be smaller than the theoretical expectations above. Indeed, computing $\xi_{1\sigma}$ and $\xi_{2\sigma}$ at Level 1 this is precisely what we find:

$$\xi_{1\sigma}^{(11)} = 0.512, \quad \xi_{2\sigma}^{(11)} = 0.836, \quad (34)$$

which show that indeed the Level 1 closure tests fail, in the sense that Level 1 fits underestimate PDF uncertainties. This confirms that the Level 2 step, that is, the generation of N_{rep} Monte Carlo replicas on top of the pseudo-data with the fluctuations, is essential to obtain the correct PDF uncertainties.

The fact that the Level 2 closure test lead to a correct estimate of the PDF uncertainties, while the Level 1 fails, can be tested in more detailed by looking at the distribution of the mean of our fit for different closure test datasets: this tests not only the one and two sigma confidence intervals, but the shape of the whole distribution of deviations between the prediction and the truth. Fig. 15 shows the histograms of the differences between $\langle f_{\text{fit}} \rangle$ obtained using different closure test datasets (that is, pseudo-data generated with different random seeds) and the central value f_{in} of the MSTW input PDFs, in units of the standard deviation of a standard Level 1 or Level 2 closure test. The histogram is generated using the values at $x = 0.05, 0.1$ and 0.2 for each PDF, as a representative sampling. The resulting distribution is very close to a Gaussian distribution with standard deviation 1 when scaled with the Level 2 uncertainties (also shown in Fig. 15), and is considerably wider using the Level 1 uncertainties. This is consistent with the results found above that Level 2 uncertainties reproduce the correct fluctuations of the central values while Level 1 overestimates them.

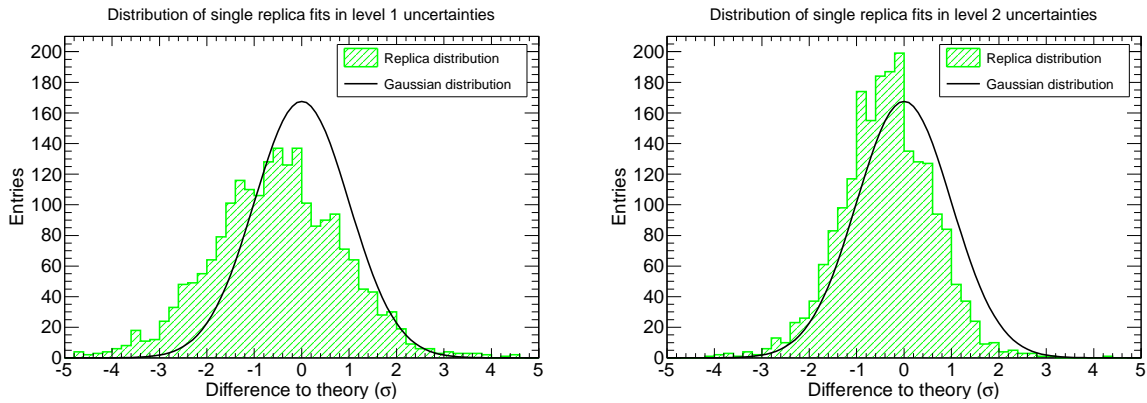


Figure 15: Histograms for the difference between the input PDF and the singlet replica fits obtained from different sets of closure test data, in units of the standard deviation of standard Level 1 (left) and Level 2 (right) fits. An appropriately scaled Gaussian distribution is shown for comparison.

4.4.4 Closure test validation using Bayesian reweighting

The results above show that our Level 2 closure tests are successful in terms of both central values and uncertainties, both at the qualitative and quantitative levels. However, it is still conceivable that they could fail for higher moments of the fitted distribution, including the correlations. In this respect, the most stringent validation of our fitting methodology, also in the context of closure tests, is provided by exploiting the Bayesian reweighting method [149]. Reweighting allows one to determine a new set of PDFs from a prior set when new data are added. Since reweighting is analytic, in the sense that the weights are computed using the χ^2 to the new data without the need for refitting, by comparing the reweighted distribution to a refitted distribution in a closure test, one can test rather precisely the fitting methodology.

With this motivation, we have performed, based on closure tests, two reweighting analyses, one at Level 1 and the other at Level 2. The priors consist in each case of sets of 1000 PDF replicas, produced using the same dataset as was used in the NNPDF2.3-like fits described in Sect. 5, in particular without the CDF and ATLAS inclusive jet production measurements. The

pseudo-data were generated using MSTW08 as the input PDF, NLO QCD, and so on, just as for the C8 and C9 closure tests. They were then fitted in the usual way. The CDF and ATLAS jet data were then included by reweighting the prior in the usual way [149], and the results compared with the corresponding closure test PDFs in which instead the CDF and ATLAS jet data were included by generating pseudo-data, adding them to the prior pseudo-data and then refitting.

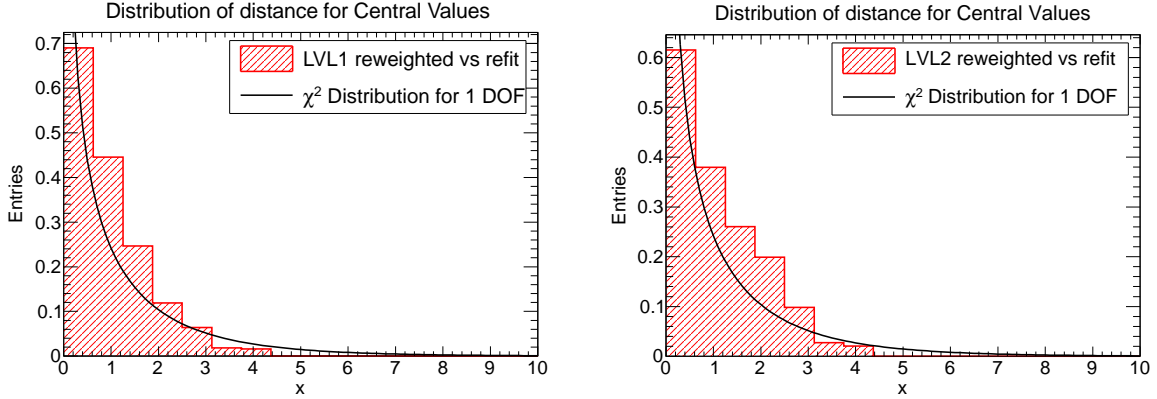


Figure 16: Distribution of the distances, for the central values, between the reweighted and the refitted PDFs for the Level 1 (left) and the Level 2 (right) closure test fits. See Appendix B for the definition of the distances. Also shown is a χ^2 distribution with one degree of freedom.

First, we compute the distance between the central values of the reweighted and refitted PDFs. The distribution of these distances, in both the Level 1 and Level 2 cases, is shown in Fig. 16. These distributions are obtained from a sampling of produced using 100 points in x for each PDF in the flavor basis, in the range between 10^{-5} and 1. We expect this distribution of distances, for a successful reweighting test, to follow a χ^2 distribution with one degree of freedom.

From Fig. 16 we see that the distribution of distances follows the statistical expectation of a χ^2 distribution with one degree of freedom, as seen by comparing with the superimposed curve representing the latter, both in Level 1 and Level 2. This means that both Level 1 and Level 2 closure tests faithfully reproduce central values.

A more stringent test is obtained when computing uncertainties. The comparison between the reweighted and refitted results including uncertainties for Level 1 and Level 2 closure tests, are shown in Fig. 17. Several interesting conclusions can be drawn from this comparison. First, it is clear that at Level 2 the reweighted and refitted results are in perfect agreement, thus validating the procedure. Second, at Level 1 the refitted uncertainties are rather smaller than the reweighted ones, thus signaling a breakdown of the procedure.

The more interesting observation however is that in fact the Level 1 reweighted and Level 2 reweighted results essentially coincide. This, upon reflection, is to be expected: indeed, if the new data bring in sufficient new information, Bayesian reweighting produces results which are largely independent of the prior. Because the jet data are controlling the large x gluon uncertainty, and these are introduced by reweighting in both cases, this uncertainties comes out to be the same even when the uncertainties in the prior are not accurately estimated, as it is the case for the Level 1 fit. If instead the jet data are introduced by Level 1 refitting, the ensuing uncertainty comes out to be too small, because Level 1 does not properly account for the data uncertainty as discussed in Sect. 4.3.1.




Figure 17: The large- x gluon PDF at the initial parametrization scale of $Q^2 = 1 \text{ GeV}^2$, comparing the results of adding the CDF and ATLAS inclusive jet data by reweighting (green solid band) and by refitting (red hatched band). Results are shown for a Level 1 closure test fit (left) and for a Level 2 closure test fit (right). See text for more details.

4.5 Robustness of the fitting methodology

We now turn to study the robustness of the fitting methodology against changes in some of its parameters, in the context of the closure test fits. The main advantage of closure tests, as compared to attempting to address these issues in fits to real experimental data, is that any deviation from expectations must be attributed uniquely to deficiencies in the methodology, rather than blamed on spurious reasons like inconsistencies in the data, or failings of theoretical description: in a closure test data and theory are by construction ‘perfect’.

First, we study how the fit results change when the training length is modified, then when the fraction of data included in the training set is varied, then we examine the effects of modifying the basis in which the PDFs are parametrized, and finally we verify that our baseline neural network architecture is redundant enough by comparing with the results of a fit with a huge neural network. Most of the closure tests in this section are performed using MSTW08 as input PDF, but we finally also explore the robustness of the closure testing against variations of this choice, in particular using CT10 and NNPDF3.0 itself as input PDFs.

4.5.1 Independence of the maximum number of GA generations

In NNPDF3.0, the cross-validation strategy that determines the optimal stopping point for the neural network training is based on the look-back method, described in Sect. 3.3.2. Provided that the maximum number of genetic algorithm generations N_{gen} is large enough, we expect results based on the look-back method to be independent of the value of N_{gen} . The condition for this to take place is that N_{gen} must be sufficiently large so that the look-back algorithm is able to identify an optimal stopping point by finding somewhere close to the absolute minimum of the validation χ^2 .

To verify this expectation, in Fig. 18 we show the distances for two Level 2 closure test fits with the maximum numbers of generations set to 30k and 80k (fits C9 and C11 in Table 8), all other settings being identical. Both fits use the look-back procedure to determine the ideal stopping point, and only differ in the maximum number of generations that are analyzed by the look-back algorithm. Since the random seed used for the fit is the same in the two cases, this is completely equivalent to looking at the same fit had it been left to run for a larger number of generations.

Both the distances between central values, and the distances between uncertainties, are less than one for all the PDFs in the flavor basis and for the whole range of x values, showing that the two fits are statistically indistinguishable. We can therefore rule out any sizable dependence on the total training length in our current results, and we can stick to a baseline maximum number of generations of 30k in our fits to real data.

4.5.2 Dependence on the training fraction

As discussed in Sect. 3.3.2, the look-back cross-validation algorithm requires us to separate the fitted dataset into two disjoint subsets, the training set and the validation set, and in the standard fits each contains 50% of the total dataset. During the minimization, only the training set is seen by the neural network, with the validation being used only to determine the optimal stopping point. While the exact value of the training fraction would be irrelevant in the limit of infinite statistics, one could argue that for our large, but finite, dataset, results might change substantially if this 50% is varied to some other value. In particular, we must be sure that a training fraction of 50% is enough to retain all the relevant information contained in the original dataset.

Distances between 30k and 80k generation fits

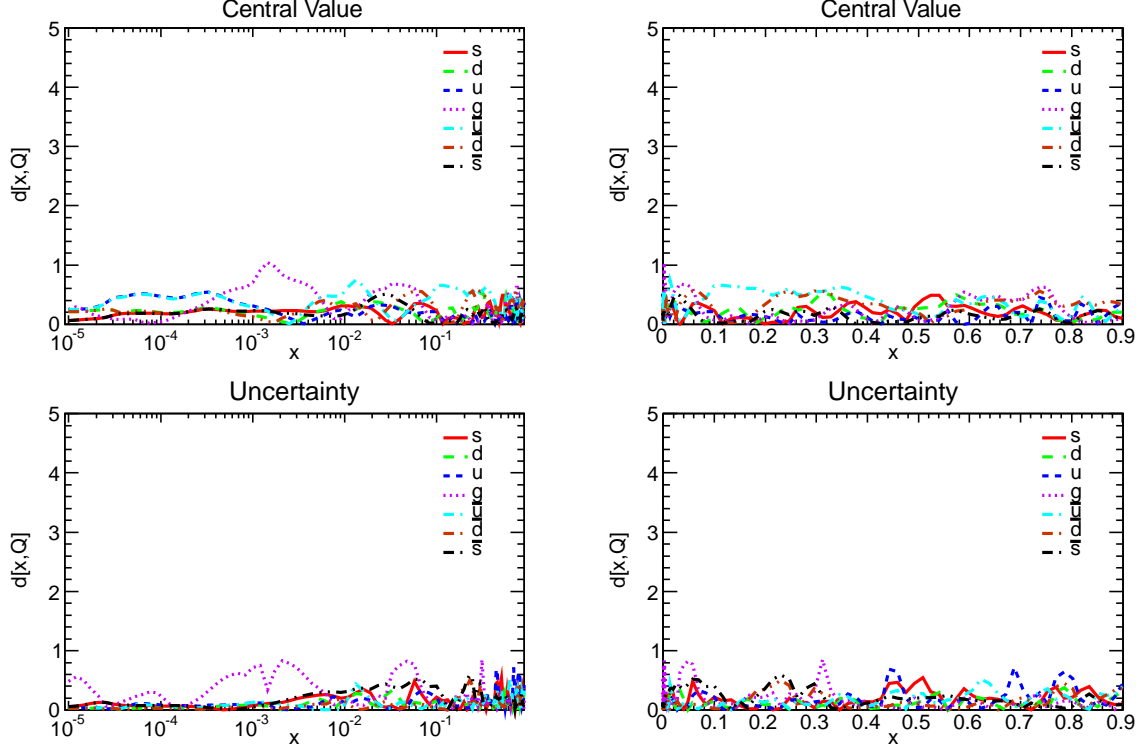


Figure 18: Distances between central values and uncertainties of two Level 2 closure test fits with maximum numbers of generations set to 30k and 80k, fits C9 and C11 in Table 8, with all other fit setting identical. Distances are calculated at the input parametrization scale of $Q^2 = 1 \text{ GeV}^2$. Left plots are in a log scale while right plots are in a linear scale.

In order to study the impact on the fitted PDFs of the use of a different training fraction, we have produced Level 2 closure fits with identical settings, the only difference being the value of this training fraction. Compared to the baseline fit with a training fraction of 50% (fit C9 in Table 8)), we have produced a fit with a smaller training fraction, 25% (fit C13), and another one with a larger training fraction, 75% (fit C14). The comparison between these three values of the training fraction, quantified by the distance plots of Fig. 19 and Fig. 20, indicate that the resulting uncertainties on the fitted PDFs increase when the training fraction is reduced to 25%, showing that, for the NNPDF3.0 dataset, some of the information which is lost when using the smaller training fraction is not redundant. This effect can be quantified more directly by looking at the PDFs in the two fits, shown in Fig. 21. Is clear that the reduction of the training fractions entails a considerable increase of the PDF uncertainties.

On the other hand, as can be seen from the distances in Fig. 20, the fits with training fractions of 50% and 75% are statistically indistinguishable. We thus conclude that the loss of information due to the training-validation splitting of the dataset required by the cross-validation procedure is small provided the training fraction is above 50%, but would be problematic for smaller training fractions.

Distances between fits with training fractions of 0.5 and 0.25

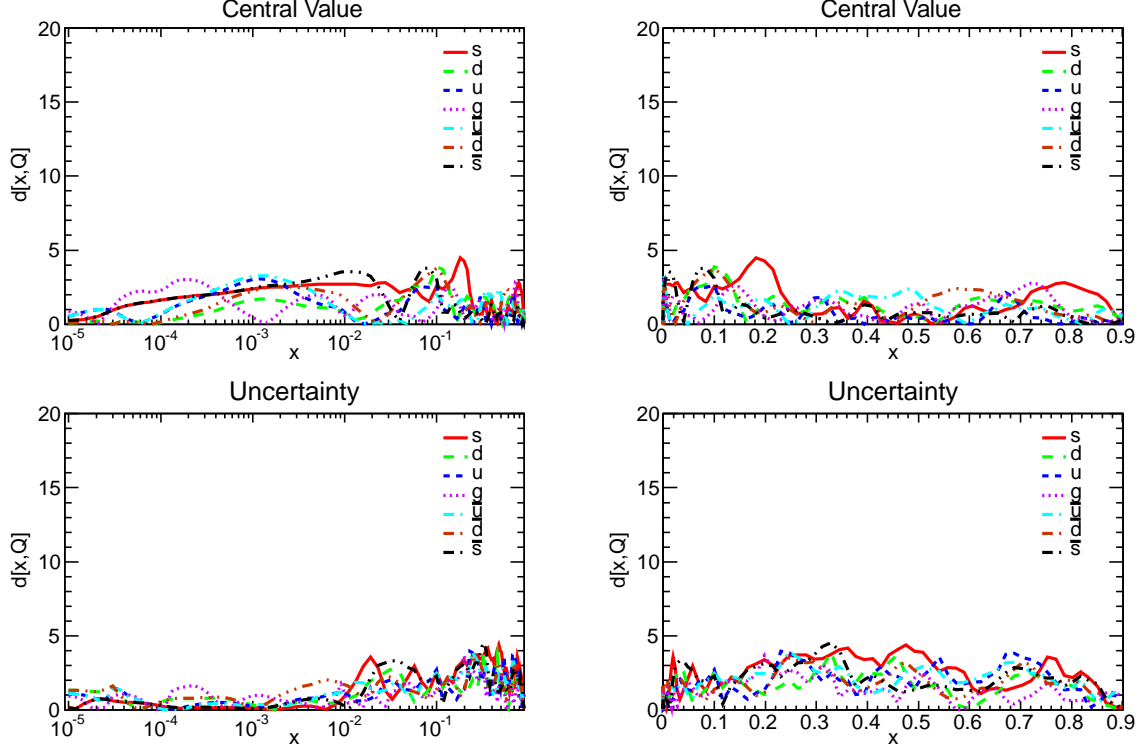


Figure 19: Same as Fig. 18 for the Level 2 closure fits based on training fractions of 25% and 50%, fits C9 and C13 in Table 8.

4.5.3 Parametrization basis independence

One of the most attractive features of the new NNPDF3.0 methodology is that thanks to the new flexible C++ fitting code, it is possible to change the basis for the PDF parametrization rather easily, as discussed in Sect. 3.2. In turn, this allows to perform detailed tests of the robustness of the whole fitting procedure: indeed, for a truly unbiased fit, results should be independent of the parametrization basis, since they are all related to each other by linear transformation. Note that in the standard Hessian approach to PDF fit even adding a few extra parameters to the input parametrization can be quite complicated, let alone changing the parametrization basis altogether.

In order to test for this, we have produced a closure test Level 2 fit, C12 in Table 8, which is otherwise identical to fit C9 but uses the NNPDF2.3 parametrization basis, Eq. (7), rather than the new NNPDF3.0 parametrization basis, Eq. (8). The distances between fits C9 and C12, that is, between fits in the NNPDF3.0 and NNPDF2.3 parametrization basis, are shown in Fig. 22. We see that distances are quite small, though slightly larger than what is found for statistically indistinguishable fits. Distances are typically between $d \sim 1$ and $d \sim 3$, so some slight differences due to the choice of basis are found, which are however much smaller than the corresponding PDF uncertainties: this means that basis independence is satisfied up to inefficiencies of the algorithm at the quarter-sigma level. We can conclude that our results are independent for all practical purposes of the choice of fitting basis. As will be shown in Sect. 5.3.5, similar conclusions will be derived of the corresponding tests in the fits to real experimental data.

Distances between fits with training fractions of 0.5 and 0.75

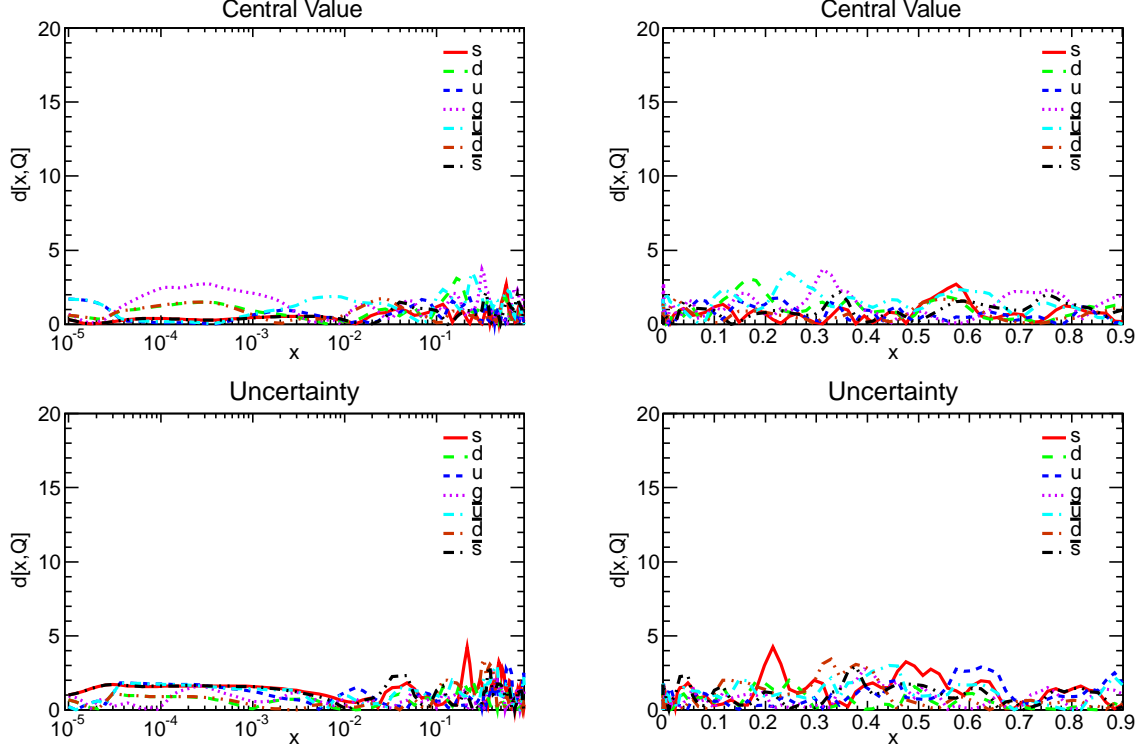


Figure 20: Same as Fig. 18 for the Level 2 closure fits based on training fractions of 50% and 25%, fits C9 and C14 in Table 8.

4.5.4 Redundancy of the neural network architecture

To be used as reliable unbiased interpolants, neural networks have to be characterized by an architecture redundant enough for the problem at hand: adding or removing nodes, even a substantial number of them, should leave the final results unchanged. The choice of a 2-5-3-1 architecture as baseline for the NNPfD fits can be justified precisely by showing its redundancy, as we did in previous work [6]. Using the closure testing technology, we have revisited for the current fits the stability of the fit results with respect to this choice of the architecture of the neural networks.

Starting from our baseline architecture, 2-5-3-1, which has around 37 free parameters for each PDF, or 259 in total, we have performed a closure fit with instead a huge neural network, with architecture 2-20-15-1, which therefore increases by more than a factor 10 the number of free parameters in the fit. Other than this modification, exactly the same fit settings were used in the two cases.

The distances between the two fits can be seen in Fig. 23. Remarkably, even increasing the architecture to have a factor 10 more free parameters, the fit using the huge networks is still quite close to the fit with standard sized networks. Again, we do not have perfect statistical equivalence, but inefficiencies are well below the half-sigma level.

This test also provides another demonstration of the resilience to over-fitting of the fitting methodology. Indeed, one might fear that such a large increase in architecture would lead to massive overlearning, while in actual fact results are virtually unchanged. This is thus a significant test of the effectiveness of the look-back method.

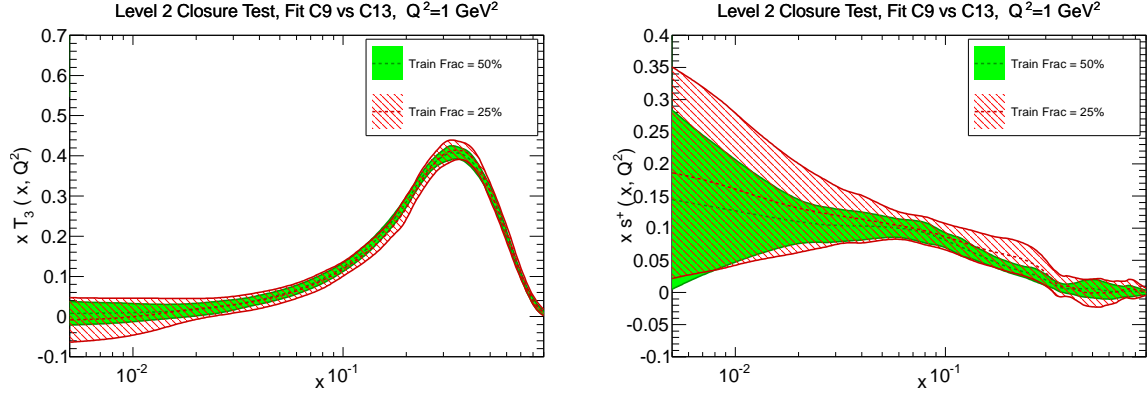


Figure 21: Comparison between PDFs in the Level 2 closure fits based on two different training fractions, 50% and 75%, corresponding to fits C9 and C14 in Table 8. The show the quark triplet $xT_3(x, Q_0^2)$ (left plot) and the total strangeness $xs^+(x, Q_0^2)$ (right plot) at the initial parametrization scale of $Q_0^2 = 1 \text{ GeV}^2$. The increase in PDF uncertainties when the training fraction is decreased to 25% is a result of the loss of non-redundant information as compared to the baseline fit with 50% training fraction

Interestingly, we also observe a very moderate increase in CPU time when using the huge network as compared to our baseline architecture. This is because when a network is very redundant, it often makes little difference which parameters the algorithm changes.

Finally, let us mention that several other tests of neural network architecture were also performed, including changing the number of layers in the networks and the number of inputs. In most cases the results were either very similar to those from the standard structure or (for some smaller networks) noticeably worse.

4.5.5 Robustness with respect to the choice of input PDF set

So far in this section we have shown results fitted to the pseudo-data generated using MSTW08 as input PDF, and studied the effects of changing some of the settings of the closure test fits. However, it is important to verify that there is nothing special in using pseudo-data generated with MSTW08, and that our methodology is flexible enough so that similar successful closure fits are achieved if other PDFs are used as input. In particular, we want to explicitly verify that self-closure is successful: using NNPDF3.0 as input PDF, and checking that it is correctly reproduced by the closure fit. Note that the MSTW08 and CT10 parametrizations are relatively simple, so we need to verify that everything works fine even when a very flexible PDF parametrization is used as input in the closure test.

As a preliminary test, we have verified that Level 0 tests works regardless of the input PDFs, by producing a number of fixed length Level 0 fits for different input PDFs, and checking found that all the conclusions of Sect. 4.2 are unchanged if input PDFs other than MSTW08 are used to generate the pseudo-data. We have then moved to Level 2 tests: we have performed closure test fits using exactly the same methodology as those based on MSTW08, but now using pseudo-data generated using different input PDF sets, in particular with the CT10 and NNPDF3.0 NLO sets. In the latter case, we have performed closure tests with and without the inclusion of the generalized positivity constraints described in Sect. 3.2.3. Given that NNPDF3.0 satisfies these constraints, this is a useful cross-check. These closure tests are labeled in Table 8 as C10 (for CT10) and C16 and C17 (for NNPDF3.0 without and with positivity constraints).

First, Fig. 24 shows the distances between the fitted and the input PDFs, Eq. (32), for

Distances between fits in evolution and NNPDF2.3 bases

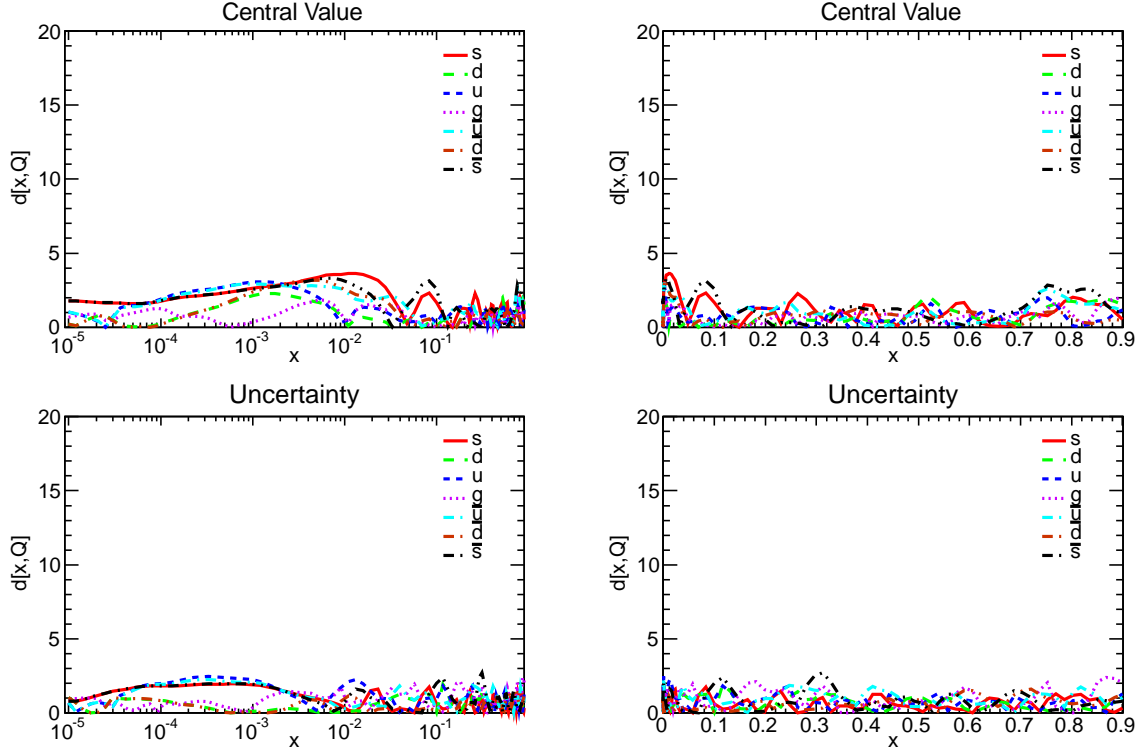


Figure 22: Same as Fig. 18 for the Level 2 closure fits based on the NNPDF3.0 and NNPDF2.3 input parametrization basis, fits C9 and C112 in Table 8.

the closure test fit which uses the CT10 NLO PDF set to generate the data. Let us recall, as discussed above, that these distances are normalized to the standard deviation of the fitted PDFs, and thus a value $d_\sigma \sim 1$ indicates that input and fitted PDFs agree at the one sigma level, and so on. These distances should be compared with the corresponding results obtained using MSTW08 as input, Fig. 12. We observe that, just as with MSTW08, the fitted PDFs are mostly within one sigma of the input PDFs, and never more than around two sigma away. In this respect, the closure test based on CT10 is as successful as that based on MSTW08. We have verified that this is also the case using some of the other estimators considered in this section.

The distances Eq. (32) are shown again in Fig. 25 now for a closure test fit using as input the NNPDF3.0 NLO PDFs. As in other closure tests, no generalized positivity is required here. This fit is therefore the ultimate closure test, providing evidence of self-closure, where the PDF set used to generate the pseudo-data has been determined using the same methodology as the closure test fit. Again, the agreement using NNPDF3.0 as input is as good as that obtained using other PDFs with a less flexible parametrization, showing that the closure tests works also in the case of a initial condition with substantial structure.

All of the closure test fits shown so far have been performed without the positivity constraints used in the fits to real data, described in Sect. 3.2.3. The motivation for this is that for some of the input PDFs used in the closure tests, in particular MSTW08, generalized positivity is not satisfied and therefore including such constraints would result in inconsistencies with the generated pseudo-data in the closure test, potentially biasing the results. On the other hand, as shown in Sect. 5.3.2, the NNPDF3.0 fits satisfy the generalized positivity constraints by

Distances between fits with standard and huge neural networks

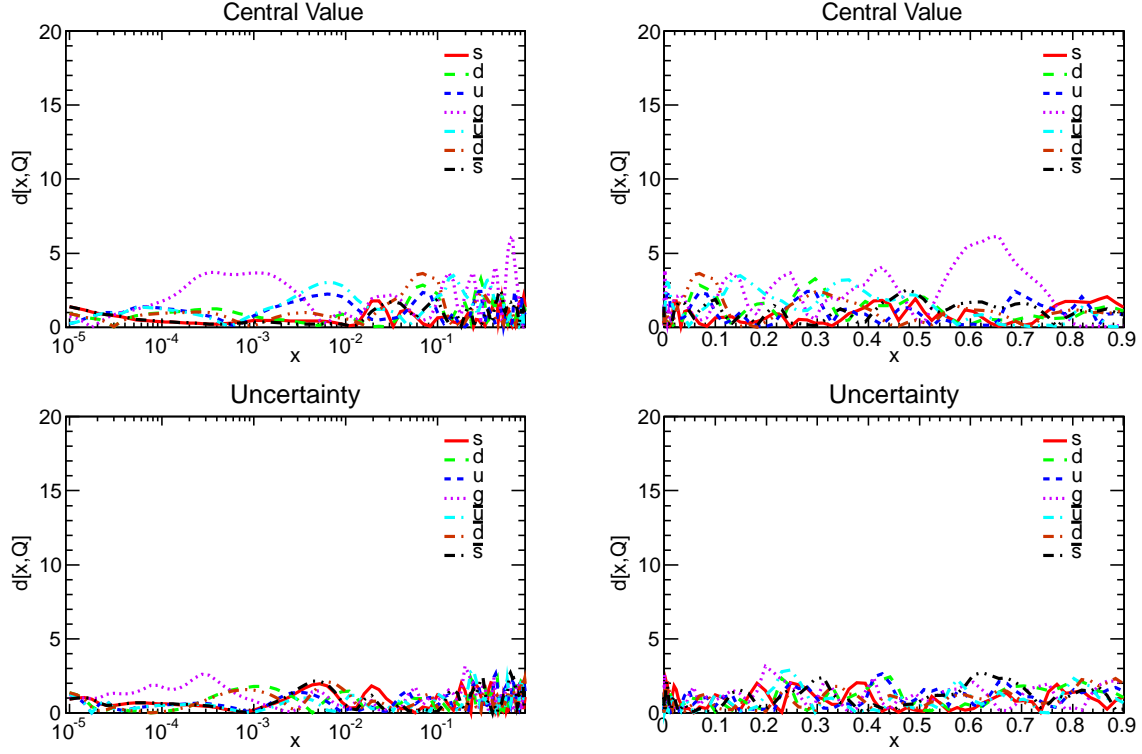


Figure 23: Same as Fig. 18 for the Level 2 closure fits based on 2-5-3-1 and 2-20-15-1 architectures, fits C9 and C15 in Table 8.

construction, and therefore if we use NNPDF3.0 as input PDF we can include generalized positivity in the closure test, expecting to find no differences with respect to the previous case. Fig. 26 shows the distances, Eq. (32), for the closure test using as input the NNPDF3.0 NLO PDFs and now also with the generalized positivity constraints imposed during the closure test fit. We indeed find that the level of agreement is similar as that of the closure test when no positivity constraints were imposed during the minimization.




Figure 24: Distances, Eq. (32), between the central values of the PDFs from the closure test fit and the CT10 PDFs, in units of the standard deviation of the fit PDFs.

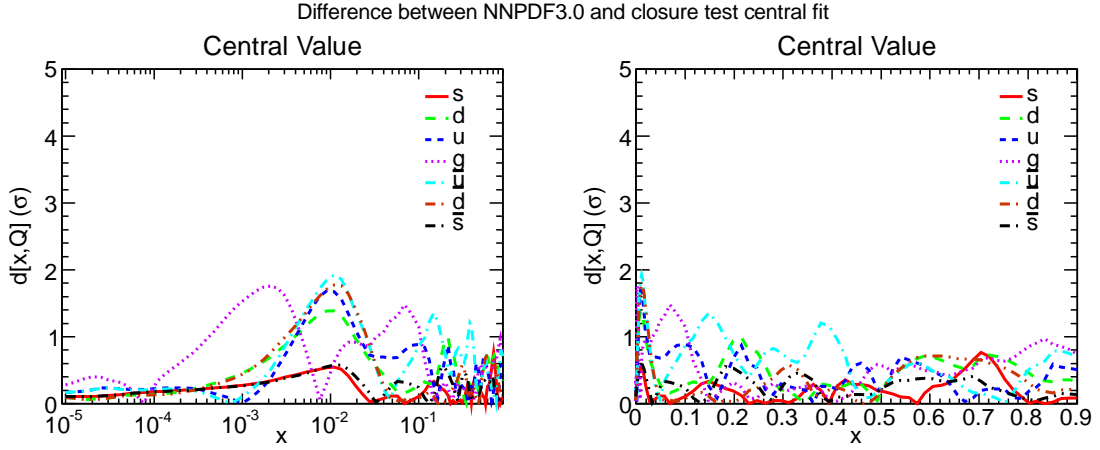


Figure 25: Same as Fig. 24 for the closure test based on NNPDF3.0 as input PDFs, without positivity constraints.




Figure 26: Same as Fig. 25, this time including the generalized positivity constraints in the closure test fit. See text for more details.

5 The NNPDF3.0 PDF set

In this Section we present the NNPDF3.0 LO, NLO and NNLO global fits. First, we discuss fit quality and the dependence of the χ^2 on its exact definition and on the treatment of systematic and normalization uncertainties. Then we show results for PDFs, and compare the new sets with NNPDF2.3 and with other existing PDF sets.

Next, we explore the dependence of the NNPDF3.0 partons on the choice of dataset. We propose a new definition of conservative PDF set, based on an estimator which allows for an assessment of the consistency of an individual dataset with the global fit, and compare these conservative partons with the global fit results. We then study a wide range of variations of the fitted dataset, including fits without LHC data, fits with only HERA and LHC data, and fits based on the NNPDF2.3 dataset. Fits to reduced datasets are also used to study the impact of jet data on the global and their stability, and the size of the strange PDF, which has been the object of various recent studies.

We then turn to an assessment of the stability of the NNPDF3.0 results upon variations in the fitting methodology. These include the impact of positivity constraints, the stability upon change of fitting basis, and the dependence on the additive or multiplicative (recall Sect. 2.4.2) treatment of systematic errors.

Finally, we study the implications of NNPDF3.0 for LHC phenomenology. We compare NLO and NNLO PDF luminosities at $\sqrt{s}=13$ TeV with NNPDF2.3 and with CT10 and MSTW08. We then provide predictions for the LHC at 13 TeV, using the MADGRAPH5_AMC@NLO program [112], for a number of representative processes like vector boson production, top quark production and Higgs production. In addition, we also discuss, using the iHixs [150] code, the implications for the dominant Higgs production channel at the LHC, gluon-fusion, and compute NNLO cross-sections. In this case we also study the dependence of $\sigma(gg \rightarrow h)$ on the choice of fitted dataset. To complete our brief exploration of the implications of NNPDF3.0 at the LHC, we study the production of systems of very high invariant masses, close to the kinematic threshold, which are relevant for searches of massive New Physics at the energy frontier.

5.1 The NNPDF3.0 set of parton distributions

In this Section first of all we discuss the quality of the LO, NLO and NNLO fits and the dependence of the χ^2 on its definition, and then we compare the NNPDF3.0 PDFs to the previous NNPDF2.3 set. A comparison of PDFs at different perturbative orders follows, and we discuss how this information can be used to estimate theoretical uncertainties on PDFs.

5.1.1 Fit quality

In Tab. 9 we present the results for the fit quality of the global LO, NLO and NNLO sets. The comparison is performed for a common value of $\alpha_s(M_Z) = 0.118$ (but PDFs in a wide range of values of $\alpha_s(M_Z)$ are also available, see Sect. 6). We show the results obtained both using the experimental and the t_0 χ^2 definition (see Refs. [9, 137]). Note that the t_0 definition varies with the perturbative order, as it depends on the theoretical values of the cross-sections included in the fit. The t_0 χ^2 is used for minimization as it corresponds to an unbiased maximum-likelihood estimator even in the presence of multiplicative uncertainties [9] (see Sect. 2.4.2). The experimental definition, which is based on the experimental covariance matrix, cannot be used for minimization as it would lead to biased results, but it is best suited for benchmarking as it only depends on publicly available results (the final PDFs and the experimental covariance matrix).

	LO			NLO			NNLO		
	N_{dat}	χ^2_{exp}	$\chi^2_{t_0}$	N_{dat}	χ^2_{exp}	$\chi^2_{t_0}$	N_{dat}	χ^2_{exp}	$\chi^2_{t_0}$
Total	4258	2.42	2.17	4276	1.23	1.25	4078	1.29	1.27
NMC d/p	132	1.41	1.09	132	0.92	0.92	132	0.93	0.93
NMC	224	2.83	3.3	224	1.63	1.66	224	1.52	1.55
SLAC	74	3.29	2.96	74	1.59	1.62	74	1.13	1.17
BCDMS	581	1.78	1.78	581	1.22	1.27	581	1.29	1.35
CHORUS	862	1.55	1.16	862	1.11	1.15	862	1.09	1.13
NuTeV	79	0.97	1.03	79	0.70	0.66	79	0.86	0.81
HERA-I	592	1.75	1.51	592	1.05	1.16	592	1.04	1.12
ZEUS HERA-II	252	1.94	1.44	252	1.40	1.49	252	1.48	1.52
H1 HERA-II	511	3.28	2.09	511	1.65	1.65	511	1.79	1.76
HERA σ_{NC}^e	38	1.80	2.69	47	1.27	1.12	47	1.28	1.20
E886 d/p	15	2.04	1.10	15	0.53	0.54	15	0.48	0.48
E886 p	184	0.98	1.64	184	1.19	1.11	184	1.55	1.17
E605	119	0.67	1.07	119	0.78	0.79	119	0.90	0.72
CDF Z rapidity	29	2.02	3.88	29	1.33	1.55	29	1.53	1.62
CDF Run-II k_t jets	76	1.51	2.12	76	0.96	1.05	52	1.80	1.20
D0 Z rapidity	28	1.35	2.48	28	0.57	0.68	28	0.61	0.65
ATLAS W, Z 2010	30	5.94	3.20	30	1.19	1.25	30	1.23	1.18
ATLAS 7 TeV jets 2010	90	2.31	0.62	90	1.07	0.52	9	1.36	0.85
ATLAS 2.76 TeV jets	59	3.88	0.61	59	1.29	0.65	3	0.33	0.33
ATLAS high-mass DY	5	13.0	15.6	5	2.06	2.84	5	1.45	1.81
ATLAS W p_T	-	-	-	9	1.13	1.28	-	-	-
CMS W electron asy	11	10.9	0.95	11	0.87	0.79	11	0.73	0.70
CMS W muon asy	11	76.8	2.25	11	1.81	1.80	11	1.72	1.72
CMS jets 2011	133	1.83	1.74	133	0.96	0.91	83	1.9	1.07
CMS $W + c$ total	5	11.2	25.8	5	0.96	1.30	5	0.84	1.11
CMS $W + c$ ratio	5	2.04	2.17	5	2.02	2.02	5	1.77	1.77
CMS 2D DY 2011	88	4.11	12.8	88	1.23	1.56	110	1.36	1.59
LHCb W rapidity	10	3.17	4.01	10	0.71	0.69	10	0.72	0.63
LHCb Z rapidity	9	5.14	6.17	9	1.10	1.34	9	1.59	1.80
$\sigma(t\bar{t})$	6	42.1	115	6	1.43	1.68	6	0.66	0.61

Table 9: The values of the χ^2 per data point for the LO, NLO and NNLO central fits of the NNPDF3.0 family with $\alpha_s(M_Z) = 0.118$, obtained using both the experimental and the t_0 definitions.

The overall fit quality is good, with an experimental χ^2 value of 1.23 at NLO and 1.29 at NNLO. The LO fit is characterized of course by a much poorer fit quality, due to the missing NLO corrections. While for some experiments like CHORUS, SLAC, ATLAS high-mass Drell-Yan, the W lepton asymmetry or top quark pair production, the χ^2 improves when going from NLO to NNLO, for most of the experiments it remains either very similar or gets slightly worse. This is also the case for the new HERA-II datasets. For the jet data the fit quality is quite similar at NLO and NNLO using the t_0 definition, but note that the kinematical cuts in the two cases are different (see Sect. 2.4.1). This is also the case for the CMS Drell-Yan data: the χ^2 is slightly worse at NNLO but only because at NLO we impose kinematical cuts that remove the region with large NNLO corrections: without such cuts, the χ^2 is much poorer at NLO.

Another interesting feature that one can observe from Tab. 9 is that, even for the same underlying PDFs, the numerical differences between the two definitions of the χ^2 can be substantial. This effect is particularly acute for experiments where systematic uncertainties dominate over statistical ones, and emphasizes the crucial role of a careful estimation of systematic errors for PDF fitting. One such example is provided by the NNLO fit, where for the CMS inclusive jet data the best-fit χ^2 changes from 1.90 (experimental definition) to 1.07 (t_0 definition). Reassur-

ingly, as we will show in Sect. 5.3.3 below, these differences in the value of the χ^2 do not have a large impact on the PDFs, which are rather stable upon changes of the χ^2 definition. The dependence of the χ^2 on its definition is weaker for fixed target experiments and DIS data, for which statistical uncertainties are dominant.

5.1.2 Parton distributions

We now compare the NNPDF3.0 LO, NLO and NNLO partons, with $\alpha_s(M_Z) = 0.118$, with the corresponding NNPDF2.3 sets and with each other. In Fig. 27 we show the distances between the parton distributions in the NNPDF3.0 and NNPDF2.3 sets for the three perturbative orders, LO, NLO and NNLO. We recall that when comparing two sets of $N_{\text{rep}} = 100$ replicas, $d \sim 1$ means that the two sets are statistically equivalent (they cannot be distinguished from replica sets extracted from the same underlying probability distribution), while $d \sim 10$ means that the sets correspond to PDFs that agree at the one-sigma level. A full discussion of distances is given in Appendix B; note that in comparison to previous NNPDF papers we have slightly changed the algorithm used in computing distances (in particular by removing an averaging and smoothing procedure), without changing their statistical interpretation. Distances are computed at a scale of $Q^2 = 2 \text{ GeV}^2$. At LO, when comparing to NNPDF2.3 we use the set with $\alpha_s(M_Z) = 0.119$, since $\alpha_s(M_Z) = 0.118$ is not available for NNPDF2.3 LO. This has a minor effect on the comparison.

The sizes of the distances shown in Fig. 27, vary significantly with the perturbative order. At LO, the gluon in the two sets is in very good agreement for $x \lesssim 0.01$. This suggests that Monte Carlo tunes (which strongly depend on the small- x gluon) based on NNPDF2.3LO, such as the Monash 2013 tune of PYTHIA8 [151] should also work reasonably well with NNPDF3.0LO. On the other hand, larger differences, between one and two-sigma, are found at medium and large x , both for the quarks and the gluon. Note that, however, at LO theory uncertainties dominate over PDF uncertainties.

At NLO and NNLO, NNPDF2.3 and NNPDF3.0 are typically in agreement at the one-sigma level, with occasionally somewhat larger distances, of order 1.5-sigma. In particular, while the total quark singlet PDF is relatively stable, there are larger differences for individual quark flavors, especially at medium and large- x . Significant differences are also found for the gluon PDF, especially at NLO, where however it should be kept in mind that NNPDF2.3 used the FONLL-A treatment of heavy quarks, while NNPDF3.0 uses FONLL-B (see Sect. 2.3.4). This comparison also shows that PDF uncertainties change at the level of one-sigma: this is to be expected, as a consequence of the constraints coming from new data, and the improved fitting methodology.

Now we turn to the direct comparison of NNPDF2.3 and NNPDF3.0 NLO PDFs: in Fig. 28 the gluon, singlet PDF, isospin triplet and total valence PDFs are shown, with $\alpha_s(M_Z) = 0.118$ at $Q^2 = 2 \text{ GeV}^2$. We can see that in the NNPDF3.0 NLO set, the central value of the gluon never turns negative, even at small- x : it is flat down to $x \sim 10^{-4}$ and then it begins to grow, within its large uncertainty, always remaining above its NNPDF2.3 counterpart. The difference can be understood as a consequence of moving to the FONLL-B heavy quark scheme, and due to the more stringent positivity constraints that are imposed now (see Sect. 3.2.3). For the total quark singlet there is good agreement between 2.3 and 3.0. For the quark triplet we see two interesting features: first at large x the result in 3.0 is larger than in 2.3, especially in the region where the PDF peaks, and also the small- x uncertainties are substantially larger, suggesting that they were somewhat underestimated in NNPDF2.3, presumably due to the less flexible methodology and less efficient treatment of preprocessing (see Sect. 3.2.2). As will be shown

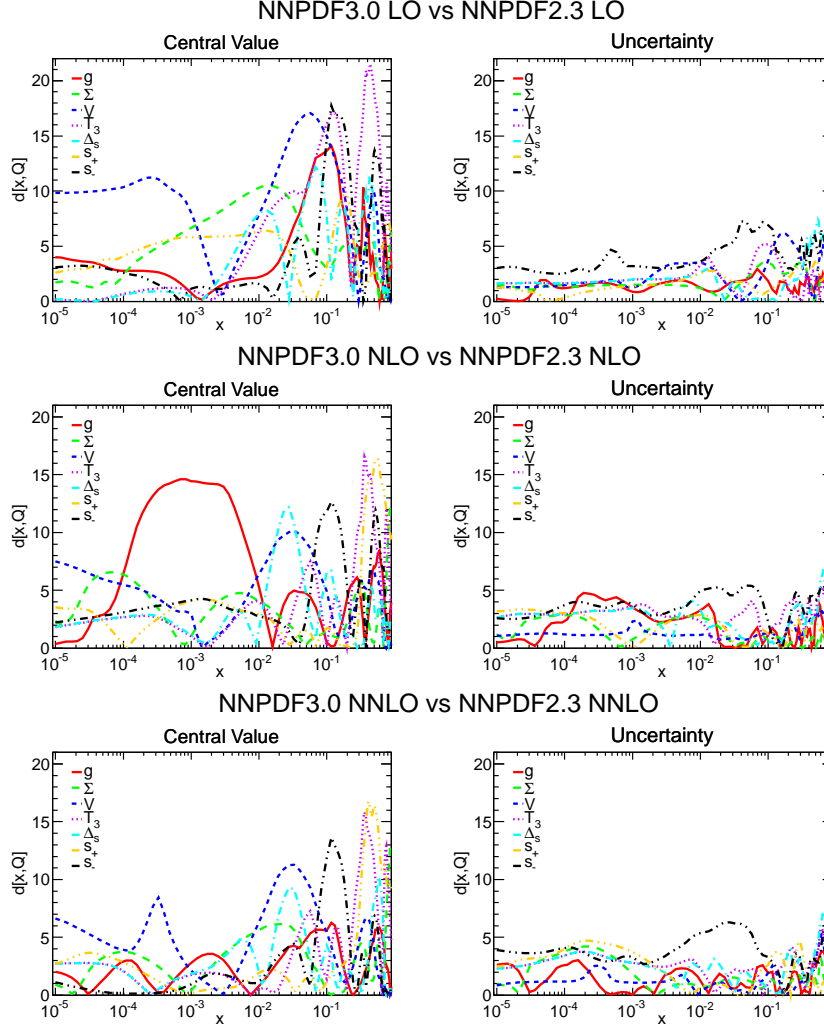


Figure 27: Distances between NNPDF2.3 and NNPDF3.0 at LO (top), NLO (center) and NNLO (bottom) PDFs, with $\alpha_s(M_Z) = 0.118$. computed between sets of $N_{\text{rep}} = 100$ replicas at $Q^2 = 2 \text{ GeV}^2$. Note that at LO, the NNPDF2.3 set has $\alpha_s(M_Z) = 0.119$.

below, the small- x uncertainties in the triplet between NNPDF2.3 and NNPDF3.0 become more similar when the comparison is performed in terms of 68% confidence level intervals, see Fig. 31.

The same comparison is performed at NNLO in Fig. 29. In this case, we observe good consistency for the gluon PDF, with a reduction in the PDF uncertainties at small- x . Note that in both NNLO fits the same FONLL-C GM-VFN scheme is being used. The agreement is also good for the quark singlet, as it was at NLO. For the triplet PDF T_3 , the small- x uncertainties, as at NLO, are larger in 3.0 than in 2.3, though as in the case of NLO the two PDF uncertainty bands are in rather better agreement if 68% CL contours instead of one-sigma uncertainties are compared.

It is also interesting to compare the NNPDF2.3 and 3.0 sets at the higher scale $Q^2 = 10^4 \text{ GeV}^2$, typical of LHC processes. Results for this comparison, at NNLO, are shown Fig. 30, as ratios to the NNPDF2.3 central value. We see that the two PDF sets agree typically at the one-sigma level or better, with some exceptions. The NNPDF3.0 gluon is somewhat softer than




Figure 28: Comparison of NNPDF2.3 and NNPDF3.0 NLO PDFs at $Q^2 = 2 \text{ GeV}^2$ with $\alpha_s(M_Z) = 0.118$. From top to bottom and from left to right the gluon, singlet, isospin triplet and total valence are shown.

in NNPDF2.3, in particular in the region around $x \sim 0.01$ which is important for the Higgs cross-section in gluon fusion. The total quark singlet is also harder for $x \leq 10^{-3}$, where the two error bands do not overlap, except for $x \lesssim 10^{-4}$ where the two sets agree again at one-sigma. For the triplet, there is good agreement, except near $x \sim 0.3$ where the NNPDF2.3 and 3.0 fits disagree at the two-sigma level; for the total valence PDF there is a reasonable agreement for all values of x .

Finally, we compare one-sigma uncertainty bands to 68% confidence level intervals. For this comparison we use the NNPDF3.0 NLO fit with $N_{\text{rep}} = 1000$ replicas; the conclusions would be qualitatively the same for the NNLO fit. Of course if the PDF probability distribution is Gaussian the one-sigma and 68% intervals coincide. While this is usually the case, for some PDFs in specific x regions there are significant deviations from gaussianity that can be quantified by this type of comparison: typically, this happens in extrapolation regions where there are no direct experimental constraints, especially if positivity constraints, which are asymmetric, play a significant role.

This comparison is shown in Fig. 31: for most PDFs there is a good agreement, the only exception being the small- x extrapolation regions, $x \lesssim 10^{-4}$, where 68% CL intervals are typically smaller than the one-sigma bands, indicating the presence of non-gaussian outliers. For the triplet, differences are already significant for $x \lesssim 10^{-3}$. Note that because of the generalized positivity constraints, the lower limit of 68% CL interval is typically just above zero.




Figure 29: Same as Fig. 28 but at NNLO.

5.1.3 Perturbative stability and theoretical uncertainties

We now compare the NNPDF3.0 LO, NLO and NNLO sets: this is a meaningful comparison because the same methodology is used at all orders, with only the underlying QCD theory (and to a small extent the dataset) changing from one order to the next. In Fig. 32 we show the distances between the NNPDF3.0 pairs of fits at two consecutive orders: LO vs. NLO, and NLO vs. NNLO. In the former case, the main variation is of course seen in the gluon PDF, which as well known is very different at LO; there are also significant differences in the large- x quarks. Of course, at LO theory uncertainties completely dominate over the PDF uncertainty, which depends on the data and is roughly the same at all orders, as this comparison clearly shows.

Distances become much smaller when comparing NLO to NNLO: for central values, the main differences are in the gluon PDF, both at small x and at large x , and in the medium- and large- x quarks, in particular the total quark singlet. Uncertainties are again quite stable, with the exception on the large- x gluon, where the PDF uncertainties are larger at NNLO because of the additional cuts applied to the jet data (recall Sect. 2.3.2). These differences in the fitted jet dataset also impact the central values of the two fits.

Next, in Fig. 33 we compare the LO, NLO and NNLO NNPDF3.0 parton distributions at $Q^2 = 2 \text{ GeV}^2$. The large shift in the gluon between LO and NLO and its subsequent stability at NNLO is clearly seen. Specifically, the LO gluon is very large, compensating for missing NLO terms in the DIS splitting functions and anomalous dimensions. This is a crucial ingredient for the tunes to semi-hard and soft data in Monte Carlo parton shower programs (see e.g. Ref. [151]). At NLO, the small- x gluon is rather flatter than the NNLO one, which tends to go negative at small- x , being prevented to do so by positivity bounds. This relatively unstable perturbative behaviour of the small- x gluon might be related to unresummed small- x




Figure 30: Same as Fig. 29, but at $Q^2 = 10^4 \text{ GeV}^2$, and with results shown as ratios to the NNPDF2.3 central value.

perturbative corrections [152]. Quark PDFs are generally quite stable, with NNLO and NLO always in agreement at the one-sigma level, and sizable shifts only seen when going from LO to NLO, especially in the region around $x \sim 0.1$.

The comparison of parton distributions at different perturbative orders provides a way of estimating the uncertainty related to missing higher-order corrections in the computations used for PDF determination, using a method suggested in Refs. [153, 154] and applied to PDFs in Ref. [155]. This source of theoretical uncertainty is currently not part of the total PDF errors, which only includes uncertainties due to the data and the methodology (recall Sect. 4.3, especially Fig. 10). As these become increasingly small, however, an estimate of the theoretical uncertainty related to missing higher perturbative orders becomes more and more important. While a full study of these uncertainties is beyond the scope of this work, we may provide a first estimate by simply studying the variation of each individual PDF when going up one order. In the remainder of this Section, we will refer to the uncertainty related to missing higher orders as “theory uncertainty”, while we will call PDF uncertainty the standard uncertainty as discussed on Sect. 4.3, which only includes the uncertainty from the data and the methodology.

In Fig. 34 the relative shift in the central values of the NNPDF3.0 PDFs at $Q^2 = 10^4 \text{ GeV}^2$ when going from LO to NLO (normalized to the NLO) is compared to the PDF uncertainty at the two order. This shift may be viewed as the LO theory uncertainty, and thus likely an upper bound to the NLO theory uncertainty. It is clear that at LO theory uncertainties are dominant essentially everywhere, and especially for the small- x gluon and the medium and small- x quarks, as also apparent from Fig. 33.

The shifts when going from NLO to NNLO (normalized to the NNLO) are shown in Fig. 35. At this order, the theory uncertainty becomes smaller than the PDF uncertainty, thereby suggesting that their current neglect is mostly justified. However, is also apparent that theory and




Figure 31: Comparison of one-sigma uncertainty bands and 68% confidence level intervals for the NNPDF3.0 NLO set with $\alpha_s(M_z) = 0.119$ at $Q^2 = 2 \text{ GeV}^2$. The set with $N_{\text{rep}} = 1000$ replicas has been used. From top to bottom and from left to right the gluon, singlet, isospin triplet and total valence are shown.

PDF uncertainties may become comparable in some important cases, like the large- x gluon, relevant for example for top quark pair production, or the medium- x quark singlet PDF, which is relevant for LHC electroweak boson production. Note that for the gluon at $x \sim 10^{-2}$, relevant for Higgs production in gluon-fusion, the perturbative convergence is very good, as already highlighted in Ref. [155].

The results of Figs. 34 and 35 suggest that as PDF uncertainties decrease, either by the addition of new experimental constraints or by refinements in the fitting methodology, a careful estimate of theory uncertainties will become mandatory. In the case of NNPDF3.0 this is especially true now given that, thanks to the closure test validation, methodological uncertainties are under full control.

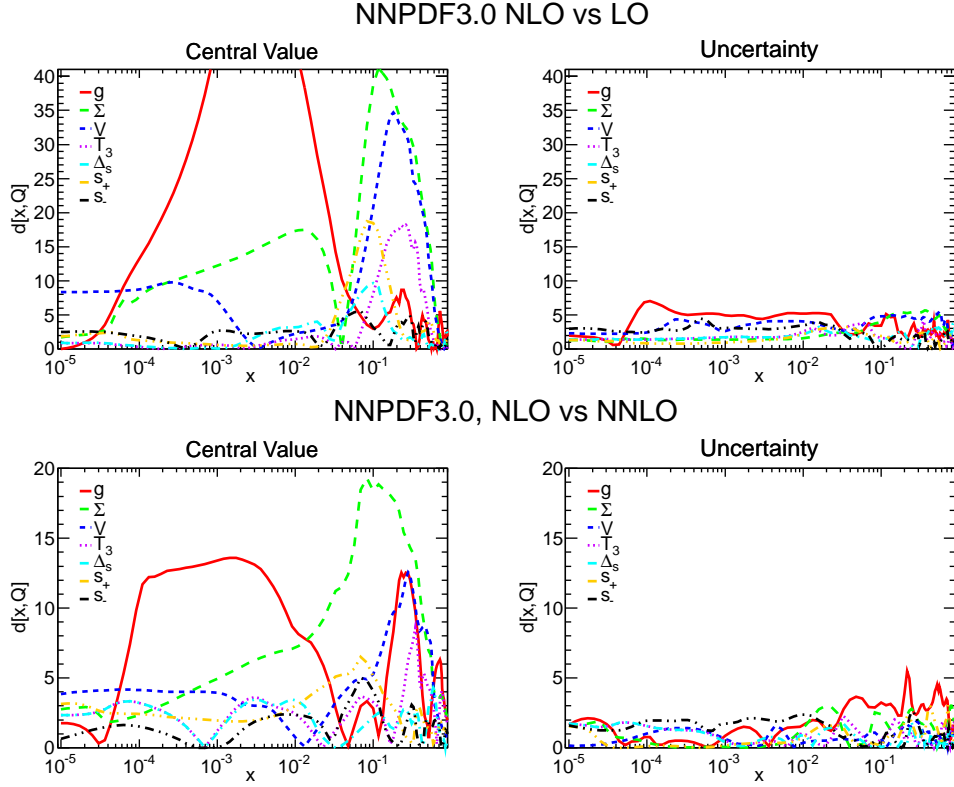


Figure 32: Same as Fig. 27, but now comparing NNPDF3.0 LO vs NLO (top) and NLO vs. NNLO (bottom).




Figure 33: Same as Fig. 28, but now comparing NNPDF3.0 LO, NLO and NNLO PDFs.

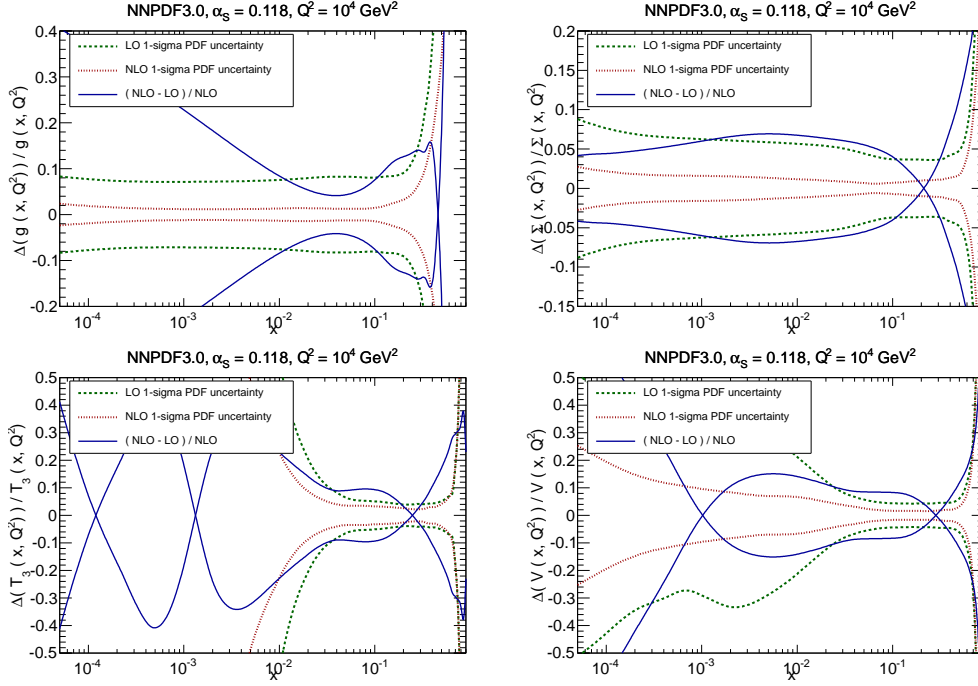


Figure 34: The shift in NNPDF3.0 PDFs with $\alpha_s(M_Z) = 0.118$ at $Q^2 = 10^4 \text{ GeV}^2$ when going from LO to NLO, normalized to the NLO central value, compared to the LO and NLO PDF uncertainties. From top to bottom and from left to right the gluon, singlet, isospin triplet and total valence are shown.

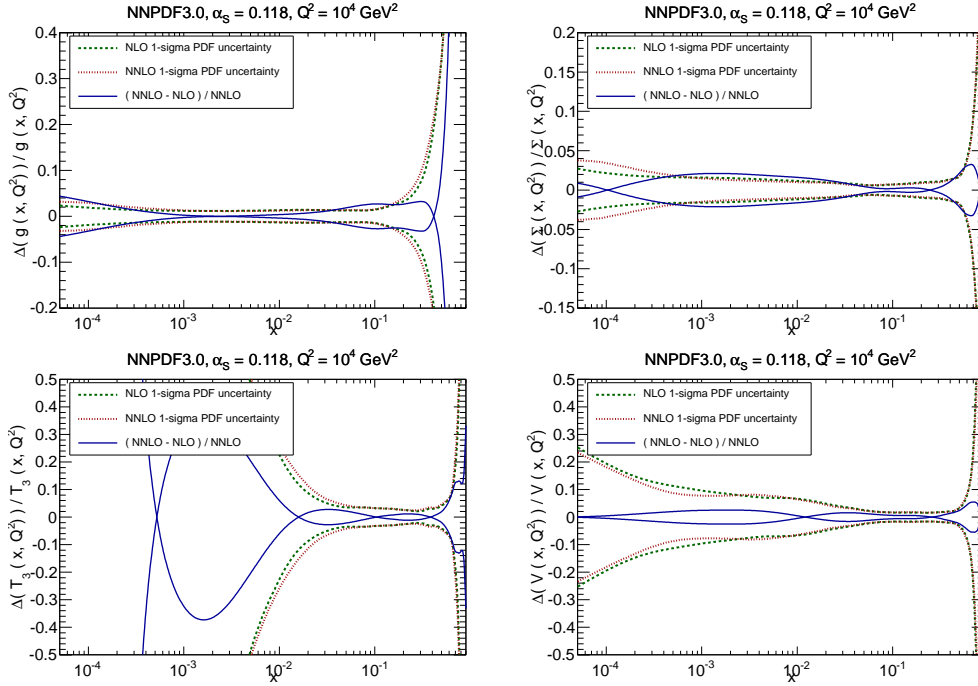


Figure 35: Same as Fig. 34, but now comparing NLO and NNLO.

5.2 Dependence on the dataset

We now study the dependence of the NNPDF3.0 PDFs on the choice of dataset, exploring a wide range of variations as compared to the dataset used in the reference global fit. First, we propose a new definition of conservative PDFs based on a maximally consistent dataset, along the lines of previous proposals [156], but now using an objective criterion, rather than theoretical expectations. Then we explore systematically the impact of the new data in the NNPDF3.0 framework, by comparing fits to various subsets of data both to the global fit, and to a minimal fit based on HERA data only; in this context we also study the impact of the new HERA-II data. We also construct sets where the HERA data is supplemented by all available data by either ATLAS or CMS; these are specially useful for comparisons with related studies by the LHC collaborations. We specifically study in detail the impact of jet data, the associated theoretical uncertainties, and their impact on the uncertainty on the large- x gluon, and the impact of $W+c$ data on strangeness and their relative weight in comparison to neutrino data.

5.2.1 Conservative PDFs from a consistent dataset

Inconsistencies between data which enter a global PDF determination can distort the statistical interpretation of PDF uncertainties. Inconsistency of any individual dataset with the bulk of the global fit may suggest that its understanding, either from the theoretical or experimental point of view, is not complete, and that its exclusion from the fit might be advantageous. In order to minimize such inconsistencies, “conservative” partons have been suggested, for example by introducing restrictive kinematic cuts which remove potentially dangerous regions [156], or by picking data which one might expect to be more reliable: for example, the NNPDF2.3 collider-only fit [13], based on the expectation that collider data, because of their higher energy, should be more reliable than fixed-target data.

We propose a new objective definition of a conservative set of PDFs based on a measure of consistency between datasets introduced in Ref. [18, 149]. This is based on observing that lack of compatibility can always be viewed as an underestimate of the covariance matrix: if the covariance matrix is inflated by a factor α^2 , then compatibility can always be attained if α^2 is large enough (crudely speaking, if uncertainties are all multiplied by a factor α). It is then possible to measure compatibility by assuming that the prior knowledge is given by all experiments in the global dataset but the given one, and using Bayes’ theorem to study how this prior is modified when adding the experiment whose consistency is under investigation. One may then compute the a posteriori probability $P(\alpha)$ that the covariance matrix of the given experiment should be rescaled by a factor α . Compatibility corresponds to the case in which $P(\alpha)$ peaks around $\alpha \sim 1$, while if the most likely value is at $\alpha_0 > 1$, this means that compatibility is only achieved when uncertainties are inflated by α_0 (see Ref. [18, 149] for a more detailed discussion and definition). The t_0 definition of the χ^2 , which is used for minimization (see Sect. 5.1.1), is also used in the determination of $P(\alpha)$.

We then proceed as follows. We compute the probability distribution of the rescaling variable α , $P(\alpha)$ for each dataset included in the global fit, and we determine the mean, the median and the mode of the corresponding $P(\alpha)$ distribution. We then exclude from the conservative fit all experiments for which at least two of these three quantities are above some threshold value, denoted by α_{max} . We discard all datasets for which the criterion fails either at NLO or at NNLO (or both), which corresponds to the most conservative choice of only retaining experiments which are well described at all perturbative orders, and has the obvious advantage that the dataset does not depend on the perturbative order, thereby keeping PDF uncertainties separate from the theory uncertainties discussed in Sect. 5.1.3. In practice, for simplicity we compute the

Experiment	NLO global fit			NNLO global fit		
	mean	mode	median	mean	mode	median
NMC d/p	1.04	1.01	1.03	1.04	1.01	1.03
NMC $\sigma_{\text{NC},p}^c$	1.32	1.31	1.27	1.27	1.26	1.27
SLAC	1.31	1.27	1.30	1.13	1.09	1.12
BCDMS	1.17	1.16	1.17	1.20	1.19	1.20
CHORUS	1.11	1.10	1.11	1.10	1.09	1.09
NuTeV	1.04	0.90	0.98	1.06	0.92	1.00
HERA-I	1.09	1.09	1.10	1.10	1.09	1.09
ZEUS HERA-II	1.23	1.22	1.23	1.25	1.24	1.25
H1 HERA-II	1.30	1.3	1.31	1.35	1.34	1.34
HERA σ_{NC}^c	1.10	1.06	1.09	1.14	1.11	1.13
E886 d/p	1.00	0.88	0.96	1.01	0.88	0.96
E886 p	1.13	1.11	1.12	1.15	1.14	1.15
E605	0.97	0.94	0.96	0.94	0.91	0.93
CDF Z rapidity	1.34	1.28	1.32	1.39	1.32	1.36
CDF Run-II k_t jets	1.09	1.06	1.08	1.15	1.12	1.14
D0 Z rapidity	1.34	1.28	1.32	0.86	0.82	0.85
ATLAS W, Z 2010	1.20	1.15	1.18	1.17	1.12	1.15
ATLAS 7 TeV jets 2010	0.76	0.74	0.75	1.09	0.92	1.02
ATLAS 2.76 TeV jets	0.86	0.83	0.85	1.07	0.57	0.83
ATLAS high-mass DY	2.22	1.68	2.03	1.82	1.34	1.63
CMS W electron asy	1.05	0.91	0.99	1.00	0.87	0.95
CMS W muon asy	1.62	1.42	1.54	1.60	1.40	1.53
CMS jets 2011	1.01	0.97	0.99	1.09	1.07	1.08
CMS $W+c$ total	1.60	1.17	1.42	1.50	1.09	1.33
CMS $W+c$ ratio	1.93	1.43	1.74	1.88	1.39	1.69
CMS 2D DY 2011	1.27	1.25	1.27	1.28	1.27	1.28
LHCb W, Z rapidity	1.10	1.02	1.07	1.20	1.12	1.17
$\sigma(tt)$	1.65	1.24	1.49	1.09	0.75	0.95

Table 10: The mean, mode and median of the $P(\alpha)$ distributions [18,149] (see text) for all the experiments in the NNPDF3.0 global fits, both at NLO (left) and at NNLO (right).

probability $P(\alpha)$ without excluding the given experiment from the global fit: this provides a conservative estimate of the compatibility (which is clearly increased by including the experiment under investigation in the prior) without requiring us to construct a new set of 1000 replicas when each of the experiments is excluded in turn. The values of the mean, median and mode thus computed for all the experiments in the NNPDF3.0 global fits at NLO and NNLO are collected in Tab. 10.²

In the following we present results for partons obtained by fitting to datasets constructed by only including datasets from Tab. 10. which pass the conservative corresponding to three different values of α_{max} , namely 1.1, 1.2 and 1.3. In Tab. 11 we provide the χ^2 (for ease of comparison we show results obtained using the experimental definition, see Sect. 5.1.1) for the PDF fits to these datasets. To facilitate the comparison with the global fit, we also provide its χ^2 values in the same table, taken from Tab. 9.

The improvement in global fit quality as α_{max} is lowered is apparent, with the most conservative option leading to an essentially perfect χ^2 of order one. It is interesting to observe that NMC proton data, which are known to have internal inconsistencies [3], as well as other

²We notice that the ATLAS W p_T data, which are included in the NNPDF3.0 global fit, are not included in Tab. 12. This is due to the fact that we construct NLO and NNLO conservative PDF sets based on the same datasets. Since the ATLAS W p_T data are excluded from the global NNLO fit because of the lack of availability of a NNLO prediction, they will not be included in our conservative set.

	$\alpha_{\max} = 1.1$		$\alpha_{\max} = 1.2$		$\alpha_{\max} = 1.3$		Global fit	
	χ^2_{nlo}	χ^2_{nnlo}	χ^2_{nlo}	χ^2_{nnlo}	χ^2_{nlo}	χ^2_{nnlo}	χ^2_{nlo}	χ^2_{nnlo}
Total	0.96	1.01	1.06	1.10	1.12	1.16	1.23	1.29
NMC d/p	0.91	0.91	0.89	0.89	0.88	0.89	0.92	0.93
NMC $\sigma^{\text{NC},p}$	-	-	-	-	-	-	1.63	1.52
SLAC	-	-	-	-	1.77	1.19	1.59	1.13
BCDMS	-	-	1.11	1.15	1.12	1.16	1.22	1.29
CHORUS	-	-	1.06	1.02	1.09	1.07	1.11	1.09
NuTeV	0.35	0.34	0.62	0.64	0.70	0.70	0.70	0.86
HERA-I	0.97	0.98	1.02	1.00	1.02	0.99	1.05	1.04
ZEUS HERA-II	-	-	-	-	1.41	1.48	1.40	1.48
H1 HERA-II	-	-	-	-	-	-	1.65	1.79
HERA $\sigma^{\text{c}}_{\text{NC}}$	-	-	1.21	1.32	1.20	1.31	1.27	1.28
E886 d/p	0.30	0.30	0.43	0.40	0.44	0.46	0.53	0.48
E886 p	-	-	1.18	1.40	1.27	1.53	1.19	1.55
E605	1.04	1.10	0.74	0.83	0.75	0.88	0.78	0.90
CDF Z rapidity	-	-	-	-	-	-	1.33	1.53
CDF Run-II k_t jets	-	-	1.01	2.01	1.04	1.84	0.96	1.80
D0 Z rapidity	0.56	0.61	0.62	0.71	0.60	0.69	0.57	0.61
ATLAS W, Z 2010	-	-	1.19	1.13	1.19	1.17	1.19	1.23
ATLAS 7 TeV jets 2010	0.96	1.65	1.08	1.58	1.10	1.54	1.07	1.36
ATLAS 2.76 TeV jets	1.03	0.38	1.38	0.36	1.35	0.35	1.29	0.33
ATLAS high-mass DY	-	-	-	-	-	-	2.06	1.45
ATLAS W p_T	-	-	-	-	-	-	1.13	-
CMS W electron asy	0.98	0.84	0.82	0.72	0.85	0.73	0.87	0.73
CMS W muon asy	-	-	-	-	-	-	1.81	1.72
CMS jets 2011	0.90	2.09	0.96	2.09	0.99	2.10	0.96	1.90
CMS $W + c$ total	-	-	-	-	-	-	0.96	0.84
CMS $W + c$ ratio	-	-	-	-	-	-	2.02	1.77
CMS 2D DY 2011	-	-	-	-	1.20	1.30	1.23	1.36
LHCb W rapidity	-	-	0.69	0.65	0.74	0.69	0.71	0.72
LHCb Z rapidity	-	-	1.23	1.78	1.11	1.58	1.10	1.59
$\sigma(t\bar{t})$	-	-	-	-	-	-	1.43	0.66

Table 11: The experimental χ^2 values at NLO and NNLO for NNPDF3.0 fits to conservative datasets corresponding to three different values of the threshold α_{\max} (see text). In each case, the χ^2 is shown for the datasets which pass the conservative cut. The values for the global fit (same as in Tab. 9) are also shown for ease of comparison.

datasets such as the H1 HERA-II data, the ATLAS high-mass Drell-Yan data, and the CMS $W+c$ data are excluded even from the least conservative set, the one with $\alpha_{\max} = 1.3$. On the other hand, the CMS inclusive jet data is included for all values of α_{\max} ; note that for this dataset the experimental χ^2 shown in Tab. 11 is significantly worse than the t_0 value used for minimization and the determination of $P(\alpha)$.

The maximally consistent dataset, found with $\alpha_{\max} = 1.1$, includes the NMC d/p data, the NuTeV and HERA-I DIS data, the Drell-Yan data from E866 and E605, the D0 Z rapidity, the ATLAS and CMS inclusive jets and the CMS W electron asymmetry.

In Tab. 12, we furthermore compare the mean, mode and median of the $P(\alpha)$ distributions for the experiments excluded from the NNLO conservative fit with $\alpha_{\max} = 1.1$ when the global fit is used as prior (i.e. the same numbers for the corresponding entries in Tab. 10), to the same quantities computed using as a prior the conservative fit itself. All the peak values of $P(\alpha)$ deteriorate when using the conservative set as a prior, as they ought to. Clearly, this deterioration will be maximal for datasets which are internally consistent, but inconsistent with

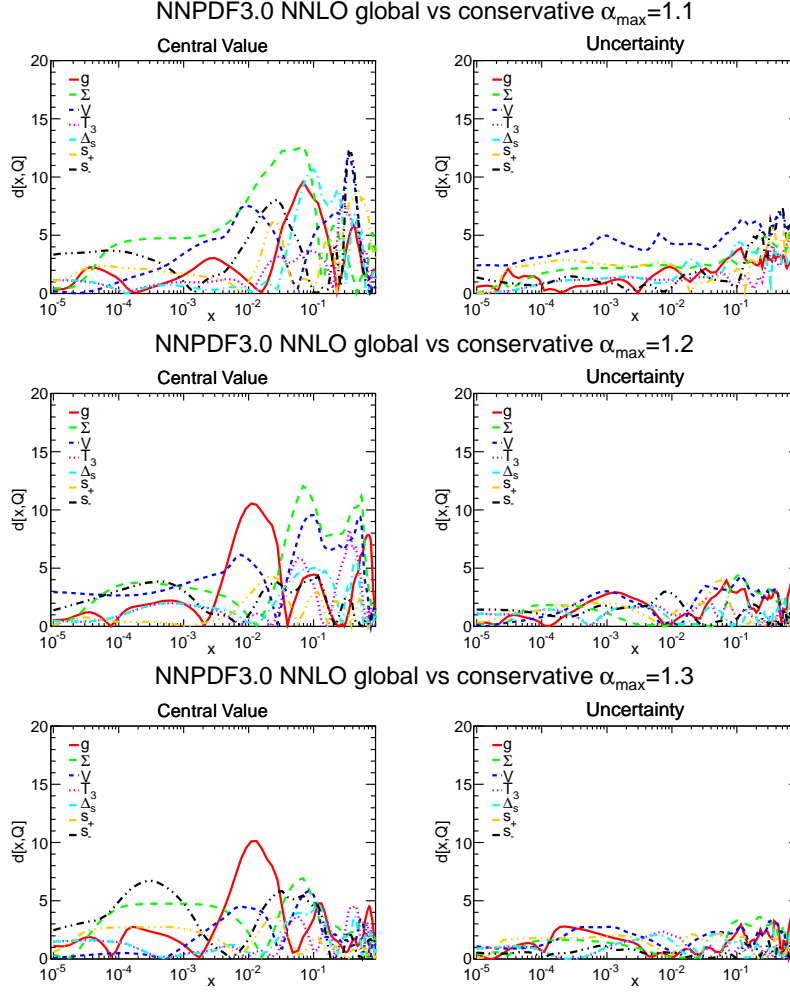


Figure 36: Same as Fig. 27, but now comparing the baseline NNPDF3.0 global fit to the conservative fits obtained using the three values of α_{\max} of Tab. 11: $\alpha_{\max} = 1.1$ (top), $\alpha_{\max} = 1.2$ (center), $\alpha_{\max} = 1.3$ (bottom).

the rest, and more moderate for experiments which are affected by internal inconsistencies, so that a rescaling of uncertainties is needed in order to describe them, regardless of what one takes as a prior. This is the case for instance for the NMC $\sigma^{\text{NC,p}}$ which are affected by internal inconsistencies as already mentioned.

The distance between the conservative sets and the baseline NNPDF3.0 NNLO global fit are shown in Fig. 36, while PDFs are compared directly in Fig. 37, where the NNLO conservative fits with $\alpha_{\max} = 1.1$ and 1.2 and the reference fit are shown. All sets are consistent with the global fit, with PDFs that differ at most at the one-sigma level, thereby confirming the consistency of the procedure, though of course PDF uncertainties are larger in the fits to reduced datasets. The small- x gluon is similar in all cases because it is driven by the HERA-I data, while there is more dependence on the choice of α_{\max} at medium and large x : interestingly, in the region relevant for Higgs production in gluon fusion the gluon is significantly affected by the choice of α_{\max} , though not beyond the one-sigma level. The quarks are also in good agreement, with the main differences seen at medium x . The set with $\alpha_{\max} = 1.1$ has of course the largest PDF

Experiment	NNLO global fit			NNLO cons. fit $\alpha_{\max} = 1.1$		
	mean	mode	median	mean	mode	median
NMC $\sigma^{\text{NC,p}}$	1.27	1.26	1.27	1.50	1.45	1.48
SLAC	1.13	1.09	1.12	1.61	1.37	1.48
BCDMS	1.20	1.19	1.20	2.02	1.86	1.92
CHORUS	1.10	1.09	1.09	2.55	1.69	2.32
ZEUS HERA-II	1.25	1.24	1.25	1.38	1.33	1.36
H1 HERA-II	1.35	1.34	1.34	1.51	1.47	1.49
HERA σ_{NC}^e	1.14	1.11	1.13	1.13	1.09	1.12
E886 p	1.15	1.14	1.15	2.18	1.62	2.03
CDF Z rapidity	1.39	1.32	1.36	1.56	1.40	1.50
CDF Run-II k_t jets	1.15	1.12	1.14	1.25	1.18	1.22
ATLAS W, Z 2010	1.17	1.12	1.15	1.38	1.25	1.32
ATLAS high-mass DY	1.00	1.34	1.63	1.63	1.19	1.45
CMS W muon asy	1.60	1.40	1.53	2.90	2.48	2.81
CMS $W+c$ total	1.50	1.09	1.33	1.85	1.37	1.67
CMS $W+c$ ratio	2.00	1.39	1.69	2.12	1.58	1.94
CMS 2D DY 2011	1.28	1.27	1.28	1.29	1.28	1.29
LHCb	1.20	1.12	1.17	1.58	1.22	1.48

Table 12: The mean, mode and median of the $P(\alpha)$ distributions at NNLO for the experiments excluded from the conservative fit with $\alpha_{\max} = 1.1$, either when the prior is the global fit (same as Tab. 10) or when using as prior the conservative set itself.

uncertainties, though even with the correspondingly restricted dataset of Tab. 11 they are not much worse than those of the global fit.

These conservative parton sets may be used for studies aimed at assessing how individual datasets affect LHC observables, by studying their effect on a maximally self-consistent dataset, such as performed in Ref. [157]. In the future, as more and more data will become available, this approach might also be used in deciding for an optimal dataset on which a global fit should be based.

5.2.2 Impact of the new HERA and LHC data

We now examine in detail the impact of the new HERA and LHC data in NNPDF3.0. Results will be shown for NNLO fits, by investigating the impact of the new data both on the global fit, and also on a HERA-only fit: while the former is more realistic, the latter allows for an assessment of the specific impact of each individual piece of data (though of course it overestimates their impact in a realistic setting). Jet data will be specifically discussed in Sect. 5.2.3 below. In all these fits, exactly the same theory and the same methodology of the default set will be used, with only the dataset changing, so that the impact of the dataset is specifically assessed. This will eventually allow us to provide a quantitative assessment of the dependence on the dataset of the uncertainty on our prediction.

In order to have a first overall assessment, we have produced a variant of the NNPDF3.0 fit using the same methodology, but using an NNPDF2.3-like dataset. We include in the dataset for this fit all, and only, the data from Tables 1-2 which were included in the NNPDF2.3 dataset. This is not quite identical to the NNPDF2.3 dataset, because we include these data with the same cuts as in NNPDF3.0 (which are sometimes slightly different than those of NNPDF2.3, as discussed in Sect. 2) and not the data of Tables 3, despite the fact that these were included in NNPDF2.3. It is however adequate to assess the (moderate) impact of the new data as we now see.

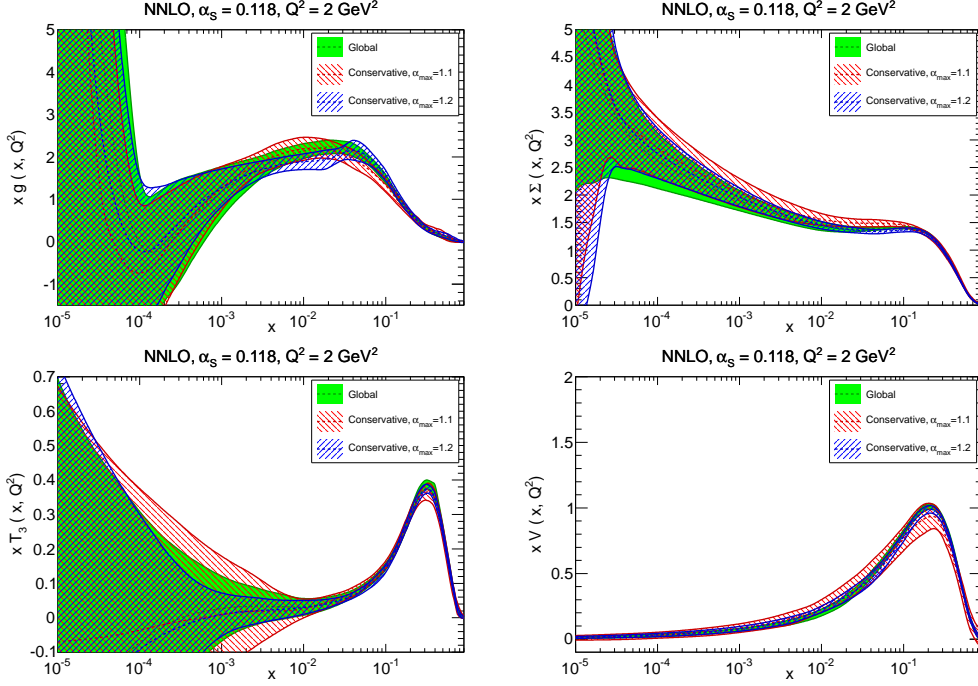


Figure 37: Same as Fig. 34, but now comparing the default global NNLO fit to the two conservative fits with $\alpha_{\max} = 1.1$ and $\alpha_{\max} = 1.2$.

The distances between PDFs from this fit and their NNPDF3.0 counterparts are shown in Fig. 38, while the PDFs at $Q = 100$ GeV are compared in Fig. 39. It is clear that the new data affects moderately all PDFs: central values vary within half a sigma of the PDF errors at most. This was to be expected, since the NNPDF2.3 PDFs already described rather well all the new experimental data that has been added in NNPDF3.0, so the main impact of the new data is to reduce uncertainties. Indeed, the PDF comparison shows that the change in uncertainties, seen in the distance plot again at a half-sigma level, always corresponds to a reduction in uncertainty.

The largest effect on central values is seen for the large- and medium- x quarks, followed by the gluon in the same region. The small- x gluons and quarks are quite stable since there is no new data that affects them in this region. Uncertainties mostly improve for the gluon PDF, both at large x thanks to the LHC jet and top quark data, and at medium and small x from the new HERA-II data. The new data favor a rather softer gluon at large x in comparison to the NNPDF2.3-like dataset, though differences are always within the PDF uncertainties. Also for the antiquark sea there is a visible improvement, especially at medium x , where the bulk of the LHC electroweak vector boson production data is. Finally, there are some improvements in strangeness; the role of the LHC data in pinning down $s(x, Q)$ is discussed in more detail in Sect. 5.2.4.

We now focus specifically on the impact of LHC data: to this purpose, we produce a fit excluding all the LHC data from the dataset, including those which were already in NNPDF2.3, and keeping all the other data in the NNPDF3.0 fit (including the HERA-II data). The corresponding distances are shown in Fig. 40, while PDFs are compared in Fig. 41.

The impact is seen to be moderate, at a half-sigma level, both for central values and for uncertainties, but it always leads to an improvement in uncertainties. Central values are mostly affected for quarks at medium and large x , at to a lesser extent for the gluon.




Figure 38: Same as Fig. 27, but now comparing the default NNLO set to a set obtained using the same methodology but an NNPDF2.3-like dataset.




Figure 39: Comparison of NNPDF3.0 NNLO PDFs at $Q^2 = 10^4 \text{ GeV}^2$ to PDFs obtained using an NNPDF2.3-like dataset. Results are shown as ratio to the default set. From top to bottom and from left to right the gluon, anti-up, anti-down quarks and total strangeness are shown.

NNPDF3.0 NNLO, Global vs noLHC

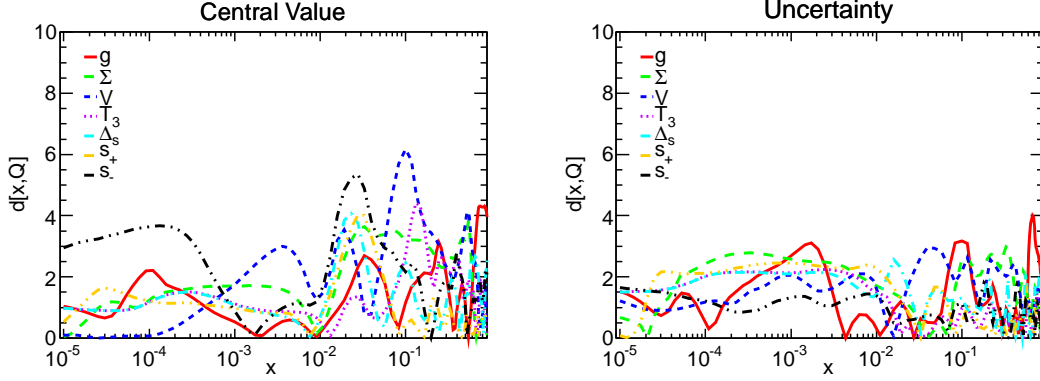


Figure 40: Same as Fig. 27, but now comparing the default NNLO set to PDFs obtained using the same dataset, but with all LHC data excluded.

Reassuringly, PDFs without LHC data are always within the one sigma uncertainty bands of the global fit PDFs, confirming the consistency of LHC data with the previous dataset. The gluon at medium and small x is already well constrained by HERA and Tevatron data, but the LHC improves uncertainties for $x \geq 0.02$, thanks to the ATLAS and CMS inclusive jet data and top quark production data. The down quark and strange PDFs are also affected, especially in the small- x region, but also at medium x .

Next, we construct a PDF set based on HERA data only. This is then used to further assess the impact of the ATLAS and CMS data. These HERA-only PDFs are compared to the global fit at $Q^2 = 2 \text{ GeV}^2$ in Fig. 42. Clearly, most PDFs, with the partial exception of the small- x gluon, PDFs have much larger uncertainties than the global fit: specifically, the quark flavor separation and the large- x gluon are very poorly constrained in a HERA-only fit, which is thus not competitive with a global fit for phenomenology applications.

The HERA dataset has widened considerably with the addition of the complete HERA-II inclusive data from H1 and ZEUS and combined HERA charm production data. In order to study the impact of this data, we have produced a version of the HERA-only fit in which we have kept only the combined HERA-I data, i.e. a HERA-I-only fit. The NNLO PDFs of the HERA-only and HERA-I-only fits are compared at $Q^2 = 10^4 \text{ GeV}^2$ in Fig. 43. The additional information provided by HERA-II apparently has a moderate impact: the gluon is mostly unchanged, while the PDF uncertainties on the medium- and large- x up antiquarks and (to a lesser extent) on the down antiquarks are moderately reduced. We conclude that, while certainly beneficial, the new HERA-II data does not change substantially the known fact that HERA-only fits are affected by large PDF uncertainties.

We now study the response of HERA-only fit of Fig. 37 upon addition of various pieces of data. In particular, we produce two fits, respectively from a dataset obtained by adding to the HERA data all the ATLAS data or all the CMS data. Specifically (see Tab. 1) in the HERA+CMS fit, the HERA data is supplemented with data on jet production, W asymmetries, Drell-Yan differential distributions, $W+c$ production and top quark total cross-sections, while in the HERA+ATLAS fit, the HERA measurements are supplemented with W, Z rapidity distributions from the 2010 dataset, inclusive jet data at 7 TeV and 2.76 TeV, and high-mass Drell-Yan production.

The distances between the HERA-only and HERA+ATLAS, or HERA-only vs. HERA+CMS




Figure 41: Same as Fig. 39 but now comparing the default NNLO set to PDFs obtained using the same dataset, but with all LHC data excluded.




Figure 42: Same as Fig. 37 but now comparing the default NNLO set to PDFs obtained using only HERA data.




Figure 43: Same as Fig. 39 but now comparing HERA-only and HERA-I-only PDFs (see text).

fits are shown in Fig. 44, while the gluon and \bar{d} PDFs are shown in Fig. 45, with the default global fit also shown for reference. The impact of the LHC data is now apparent, in particular for PDF combinations which are poorly constrained by a fit to HERA data only, specifically the large- x gluon and the quark flavor separation. Note that the CMS data provides more stringent constraints on the gluon at large x since it uses the 2011 inclusive jet data, which for ATLAS is still not available. ATLAS and CMS have a similar constraining power for the medium and large- x quarks, with CMS slightly superior for the strangeness PDFs thanks to the availability of the $W+c$ measurements, and also for flavor separation (and thus for \bar{d}) due to the fact that the CMS electroweak dataset is somewhat more extensive.

On the other hand, comparison to the global fit shows that neither of these HERA+ATLAS or HERA+CMS fits is yet competitive. Also, comparison of these result to our previous assessment of the overall impact of LHC data on the global fit (Figs. 40 and 41) shows that gauging the impact of LHC data on PDFs by their effect on a HERA-only fit might be somewhat misleading: because of the good consistency of LHC data with the pre-LHC global dataset, their impact in the global fit is rather less pronounced.

We finally come to a global assessment of the dependence of the PDF uncertainty on the dataset. This can be done by using the estimator φ_{χ^2} , Eq. (24), which, as discussed in Sect. 4, is essentially the ratio of the PDF uncertainty on the prediction to the original experimental uncertainties, averaged over all data, and keeping into account all correlations. Its value thus measures by how much the data uncertainty, propagated to the PDFs, and then propagated back to the data, is reduced due to the fact that the PDF fits combines many data points into a single set of PDFs which are constrained both by theory (such as perturbative evolution, which connects different values of Q^2 , and sum rules, which relates different values of x) and also by smoothness (the value of any PDF at neighboring points in x cannot differ by an arbitrarily

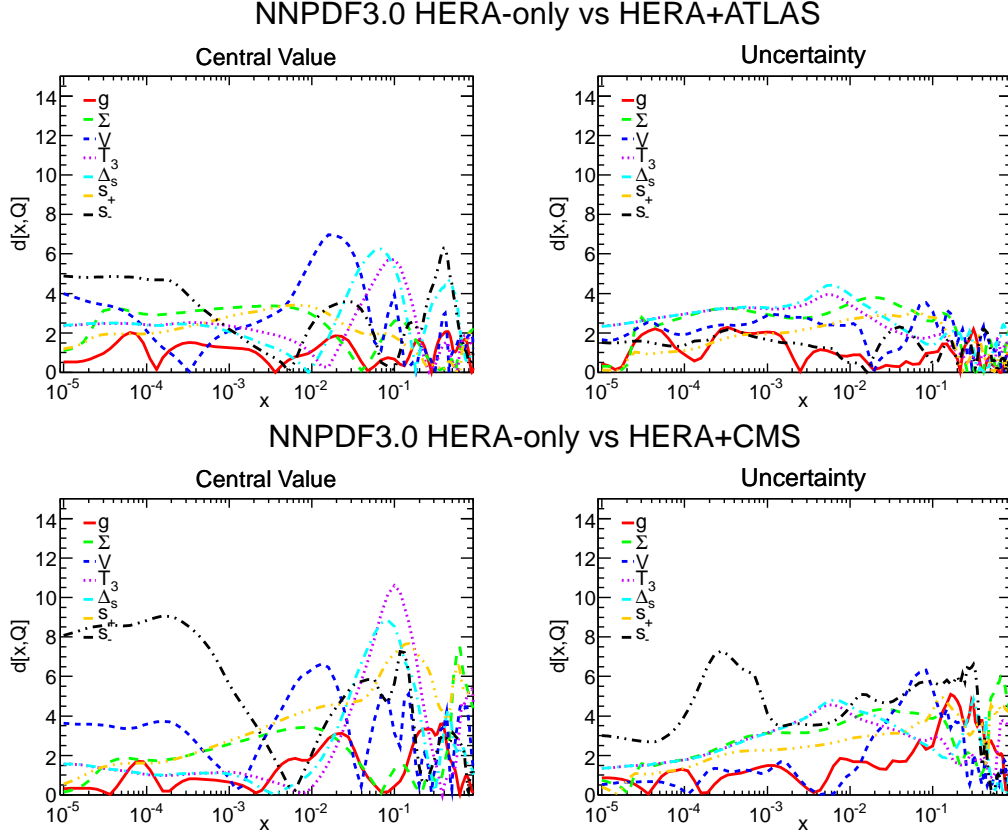


Figure 44: Same as Fig. 27, but now comparing the NNLO HERA-only and HERA+ATLAS PDFs (top) or the HERA-only and HERA+CMS PDFs (bottom) (see text).

large amount). For a single data point, or if datapoints were independent, its value would be one by definition, while its deviation from one estimates the uncertainty reduction (or increase, which could happen in principle when combining inconsistent data).

In Tab. 13 we show φ_{χ^2} for the global NNPDF3.0 NLO and NNLO fits, as well as for the fits based on reduced datasets, starting from the HERA-I only fit and including the fit without LHC data, in increasing order of the size of the fitted dataset. The result for the conservative set refers to the fit with threshold $\alpha_{\max} = 1.1$.

For the global fits, we find $\varphi_{\chi^2} = 0.291$ and 0.302 for the NLO and NNLO sets respectively, to be compared with the corresponding value at LO, $\varphi_{\chi^2} = 0.512$. The improvement between LO and NLO, almost by a factor of two in terms of the reduction of the PDF uncertainties to the fitted datapoints, is clear evidence of the better consistency of the NLO fit in comparison to the LO one. On the other hand, the NNLO fit is very similar to the NLO one in this respect (perhaps marginally worse), consistent with the observation that the quality of the NNLO fit is not better than that of the NLO fit, as also measured by the value of the χ^2 , see Tab. 9.

The monotonic decrease of the values of φ_{χ^2} for the fits to reduced datasets, from HERA-I to HERA-all, to HERA+ATLAS or HERA+CMS, to the global fit, shows the constant uncertainty reduction as more data are combined. The uncertainty on the conservative fit is larger than that of any fit except the HERA-I fit. This provides evidence in favor of using our current default as reference, rather than the conservative set, as the substantial decrease in φ_{χ^2} in the global




Figure 45: Comparison of the gluon and antidown NNLO PDFs at $Q^2 = 10^4 \text{ GeV}^2$ of the HERA-only and HERA+ATLAS sets (top) or the HERA-only and HERA+CMS sets (bottom), shown as ratios to the HERA-only PDFs. For reference, the PDFs from the default NNPDF3.0 global set are also shown.

fit in comparison to the conservative one suggests that the overall consistency of the global fit is still quite good. Finally, the uncertainty on the no-LHC fit is larger than that of a global fit, and in fact comparable to that of the HERA+CMS fit: this shows that even though the impact of the LHC data is moderate, it is visible, and it is comparable to the impact of all the other non-HERA data.

Dataset	φ_{χ^2} NLO	φ_{χ^2} NNLO
Global	0.291	0.302
HERA-I	0.453	0.439
HERA all	0.375	0.343
HERA+ATLAS	-	0.330
HERA+CMS	-	0.315
Conservative	0.422	0.478
no LHC	0.312	0.316

Table 13: The value of the fractional uncertainty φ_{χ^2} , Eq. (24), for the default NNPDF3.0 NLO and NNLO fits compared to that obtained in various fits to reduced datasets. For the LO global fit, we find $\varphi_{\chi^2} = 0.512$. The result for the conservative set refers to the fit with $\alpha_{\text{max}} = 1.1$.




Figure 46: Comparison of the gluon in a fit to a dataset without jet data and in the global fit at NLO (top) and NNLO (bottom), plotted at $Q^2 = 2 \text{ GeV}^2$ vs. x on a logarithmic (left) and linear (right) scale.

5.2.3 Impact of jet data on the global fit

We now explore the impact of jet data in the NLO and NNLO NNPDF3.0 fits, with the motivation of making sure that theoretical limitations in the description of jet data, and in particular the current lack of full knowledge of NNLO corrections, does not bias the fit results.

To this purpose, we have produced versions of the NNPDF3.0 PDF fit in which all jet data are removed from the global dataset: the gluon from these sets is compared to that from the default global fit at $Q^2 = 2 \text{ GeV}^2$ in Fig. 46: Other PDFs are essentially unchanged upon removing jet data. It is clear that removing jet data from the global fit leads to a substantial increase of the PDF uncertainties on the gluon at medium- and large- x . However, when jet data are included, the uncertainties are very similar at NLO and NNLO, despite the fact that at NNLO the jet dataset is significantly smaller due to the more restrictive cuts which we have introduced in order to account for the incomplete knowledge of NNLO corrections to jet production (see Sect. 2.3.2). This is reassuring in that it is consistent with the expectation that no instabilities are introduced by jet data in the NNLO fit despite potentially large perturbative corrections.

In Tab. 14 we compare at NLO and NNLO the χ^2 to the collider jet data, both in the reference NNPDF3.0 fit and in the fit without jet data. We provide the results using both the experimental and the t_0 χ^2 definitions, whose values can differ significantly, especially at NNLO. The description of jet data turns out to be reasonably good even when they are not included in the fit. This is evidence for consistency, and it explains why they help in reducing the gluon uncertainty. We also show the value of the χ^2 for top pair production, which is also sensitive to the gluon. The fact that this value changes very little upon inclusion of jet data is also evidence for general consistency.

Dataset	NLO			
	Exp χ^2		$t_0 \chi^2$	
	Global	No Jets	Global	No Jets
CDF Run II	0.95	<i>1.51</i>	1.05	<i>1.62</i>
ATLAS 7 TeV + 2.76 TeV	1.58	<i>1.88</i>	0.86	<i>0.96</i>
CMS 7 TeV 2011	0.96	<i>1.32</i>	0.90	<i>1.17</i>
Top quark pair-production	1.43	1.26	1.67	1.49

Dataset	NNLO			
	Exp χ^2		$t_0 \chi^2$	
	Global	No Jets	Global	No Jets
CDF Run II	1.84	<i>1.85</i>	1.20	<i>1.58</i>
ATLAS 7 TeV + 2.76 TeV	1.17	<i>1.00</i>	0.72	<i>0.65</i>
CMS 7 TeV 2011	1.91	<i>2.23</i>	1.07	<i>1.37</i>
Top quark pair-production	0.73	0.43	0.61	0.42

Table 14: The χ^2 to jet data, computed using the default NNPDF3.0 PDFs, or PDFs based on a dataset without jet data; values in italics correspond to data which have not been included in the fit. Values are provided using both the experimental and t_0 definition of the χ^2 (see Sect. 5.1.1). The value for top data (included in all fits) is also shown.

We conclude that based on all evidence not including jet data (or not including them at NNLO) would not lead to a significant change of the central value of the extracted gluon distribution but it would lead to a deterioration of its uncertainty. Given our conservative treatment of NNLO perturbative corrections, described in Sect. 2.3.2, and in the absence of indications of instability or inconsistency, we believe that the determination of the gluon is most reliable when jet data are kept in the data set, as we do for our default fit.

5.2.4 Nucleon strangeness

Recently the size of the strange distribution has been the object of experimental and phenomenological debate. In global fits, the strange PDF is mostly constrained by the neutrino-induced deep-inelastic scattering data, such as CHORUS, NuTeV and NOMAD [36, 37, 158, 159]. While also inclusive data is sensitive to strangeness, the strongest constraint come from the so-called dimuon process, charm production in charged-current DIS. However, the theoretical treatment of this data is affected by various sources of theoretical uncertainty, such as the need to model charm fragmentation, the treatment of charm quark mass effects at low scales, and effects related to the use of nuclear targets. Recently, LHC data which constrain the strange PDF have become available: on top of inclusive W and Z production on- and off-shell production, W production in association with charm quarks which directly probes strangeness at leading order.

In PDF global fits, with strangeness mostly based on neutrino data, the strange sea is typically smaller than the up and down quark sea by a factor of order $\sim \frac{1}{2}$. In 2012, a QCD analysis of the ATLAS data on W and Z rapidity distributions at 7 TeV [160] suggested that the size of the strange and up and down sea is comparable, at least for $x \sim 0.01$. This analysis was revisited in the NNPDF2.3 framework [13], with the conclusion that while the ATLAS data in isolation do favor a central value of $s(x, Q^2)$ similar in size to $\bar{u}(x, Q^2)$ and $\bar{d}(x, Q^2)$, the uncertainties involved are so large that it is difficult to make a clear-cut statement, and in particular the central value of the strangeness fraction in the global NNPDF2.3 fit is compatible

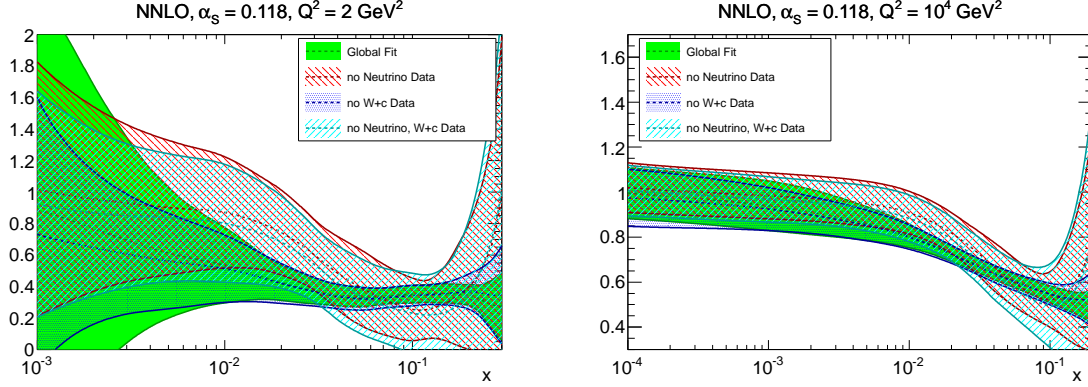


Figure 47: The strangeness ratio r_s Eq. (35), at NNLO sets with $\alpha_s(M_Z) = 0.118$ plotted vs. x at $Q^2 = 2 \text{ GeV}^2$ (left) and $Q^2 = 10^4 \text{ GeV}^2$ (right) for the default NNPDF3.0 PDF set compared to set obtained excluding from the fitted dataset either neutrino data, or $W+c$, or both neutrino and $W+c$ data. and a fit with neither of the two datasets.

with that of a HERA+ATLAS fit at the one-sigma level. Also, it was found that when including the ATLAS data in the global fit they would have little impact, and strangeness would still be suppressed.

As discussed in Sect. 2.1, in NNPDF3.0 we have also included (both at NLO and NNLO, see Sect. 2.4.1) the CMS data [60] for $W+c$ which have become available, and which directly constrain the strange distribution. The CMS $W+c$ data have been recently used in a QCD analysis [58], together with HERA data, to show that the strange PDF $s(x, Q^2)$ from collider-only data can be determined with a precision comparable to that of global fits which include neutrino data. The CMS data favors a suppressed strangeness, consistent with the indications from the neutrino data. ATLAS $W+c$ data (not included in NNPDF3.0 because they are only available at the hadron level) seem instead to favor a less suppressed strangeness [92]. Fits including LHC W, Z and $W+c$ data along with fixed target deep-inelastic scattering and Drell-Yan data have also been studied in Ref. [75], with the conclusion that a good fit to all these datasets can be obtained and finding again a suppressed strangeness.

We now study this issue in light of the NNPDF3.0 global PDF determination, by constructing PDF sets based on data which include or exclude in turn various pieces of experimental information which are sensitive to strangeness. Specifically we have produced PDF sets based on the following modifications of the NNPDF3.0 default dataset of Tables 1-2:

- all neutrino data (CHORUS and NuTeV) removed, CMS $W+c$ data included;
- all neutrino data (CHORUS and NuTeV) included, CMS $W+c$ data removed;
- both neutrino data (CHORUS and NuTeV) and CMS $W+c$ data removed.

We now compare results obtained in each case, specifically for the strangeness fraction r_s , defined as

$$r_s(x, Q^2) = \frac{s(x, Q^2) + \bar{s}(x, Q^2)}{d(x, Q^2) + \bar{u}(x, Q^2)}. \quad (35)$$

In Fig. 47 r_s is shown for the default NNPDF3.0 fit and the three above fits, plotted as a function of x , for $Q^2 = 2 \text{ GeV}^2$ and $Q^2 = 10^4 \text{ GeV}^2$.

	χ^2_{exp}			
	Global	No neutrino	No $W+c$	No neutrino/ $W+c$
CHORUS	1.13	<i>3.87</i>	1.09	<i>3.45</i>
NuTeV	0.62	<i>4.31</i>	0.66	<i>6.45</i>
ATLAS W, Z 2010	1.21	1.05	1.24	1.08
CMS $W+c$ 2011	0.86	0.50	<i>0.90</i>	<i>0.61</i>

Table 15: Values of the χ^2 (experimental definition, see Sect. 5.1.1) to different data with sensitivity to strangeness, using as input PDFs obtained from fits in which these data are included or excluded in turn; values in italics denotes cases in which the given data was not included in the fit. See text for more details.

First, we observe the remarkable compatibility of the various fits (with, as usual, smaller uncertainty at a higher scale due to asymptotic freedom): for all fits and all x the one-sigma PDF uncertainty bands overlap. The global fit is always the most accurate, though at very low $x \lesssim 10^{-3}$ and high scale the uncertainty is of similar size in all fits. While when removing neutrino data the uncertainty blows up, removing the $W+c$ has very little impact, leading to a moderate uncertainty reduction for $x \geq 0.05$ when added to the fit without neutrino data. The fits without neutrino data do lead to a slightly higher central value of r_s in the region of $x \sim 0.01$, but the effect is completely negligible on the scale of the uncertainty on these fits, though it would be significant on the scale of the uncertainty of the global fit. It must therefore be considered a statistical fluctuation due to the large uncertainty in the fit without neutrino data.

The χ^2 for the relevant experiments in these various fits are collected in Tab. 15, thereby allowing us to compare how well each experiment is described when included or excluded from the fit. We see that $W+c$ data are well described regardless of whether they are included in the fit or not, while the neutrino data are not well described unless they are included in the fit. This again shows that the impact of the $W+c$ can only be moderate.

We conclude that the $W+c$ data alone are not yet competitive with the neutrino data in determining strangeness, and that their inclusion does not modify significantly the assessment of the size of the nucleon strangeness in previous global fits, for which there is thus clear evidence for suppression in comparison to light quarks by a factor of between two and three at low scales.

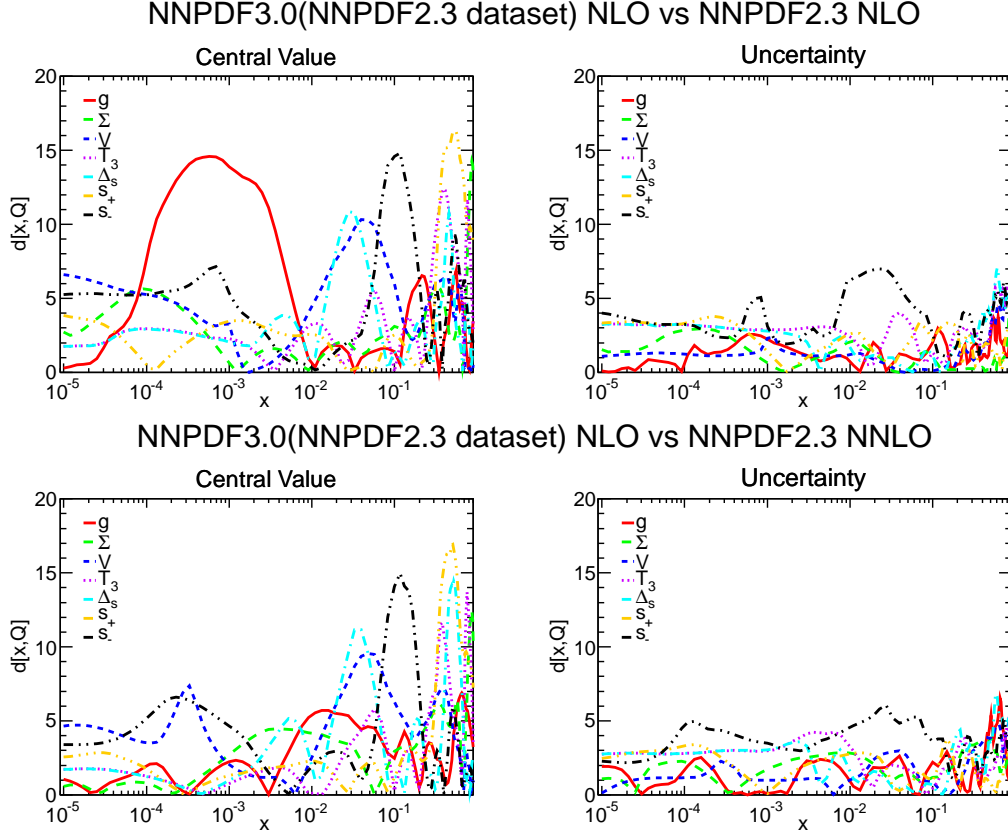


Figure 48: Same as Fig. 38, but now comparing the PDFs obtained from an NNPDF2.3-like dataset with NNPDF3.0 methodology and theory to the published [13] NNPDF2.3 sets at NLO(top) and NNLO (bottom).

5.3 Stability

We will now check the dependence and stability of our results upon our methodology and its variations. Some of the issues which we address here were already studied in Sect. 4.5 in the context of closure tests, but others are specific to the fit to actual data.

Specifically, we will study the impact on the NNPDF3.0 results of the new minimization and stopping methodology discussed in Sect. 3.3 in comparison to that previously used in NNPDF2.3, the impact of the improved treatment of positivity discussed in Sect. 3.2.3, the impact of a multiplicative vs. additive treatment of systematic uncertainties (see Sect. 2.4.2), our improved self-consistent determination of preprocessing ranges presented in Sect. 3.2.2, and finally the independence of fitting basis (already tested in closure tests in Sect. 4.5.3).

5.3.1 Impact of the NNPDF3.0 methodology

First, we discuss the impact of the NNPDF3.0 methodology, by comparing the NNPDF2.3 global fit with the fit discussed in Sect. 5.2.2 and shown in Fig. 39, based on an NNPDF2.3-like dataset, but with NNPDF3.0 methodology and theory settings (specifically the use of FONLL-B), i.e. by varying the methodology with fixed dataset. This is the complementary comparison to that which we performed in Sect. 5.2.2, where instead we varied the dataset with fixed methodology.




Figure 49: Same as Fig. 39, but now comparing the PDFs obtained from an NNPDF2.3-like dataset with NNPDF3.0 methodology and theory to the published [13] NNPDF2.3 sets at NLO(top) and NNLO (bottom).

With this comparison we can fully disentangle the effects of the new experimental data in NNPDF3.0 from that of the improved fitting methodology and the new theoretical settings.

The distances between the original NNPDF2.3 PDFs of Ref. [13] and the NNPDF3.0 fit with NNPDF2.3 data are shown in Fig. 48 both at NLO and NNLO, while the NNLO PDFs are compared in Fig. 49.

In the NLO fit the new methodology and theory settings have an impact on the small- x gluon and large- x quarks at the one and a half sigma level. The differences in the gluon can be understood as a consequence of having switched from the FONLL-A heavy quark scheme used in NNPDF2.3 to the more accurate FONLL-B adopted in NNPDF3.0, while the differences seen for quarks are necessarily a consequence of the more efficient methodology and extended positivity constraints (see Sect. 5.3.2 below). At NNLO the non-insignificant differences seen in all PDFs reflect the improved methodology and positivity, as the NNLO theory used in 2.3 and 3.0 is the same. At high scale the most noticeable difference is the softening of the small- x gluon seen in Fig. 49.

The main conclusion of this comparison is that a significant part of the differences between the final NNPDF2.3 and NNPDF3.0 sets, as seen specifically at high scale in Fig. 30 and at low scale in Fig. 29, are due to the improved methodology (minimization and generalized positivity). This is consistent with the conclusion of Sect. 5.2.2 (see in particular Figs. 38-39) that the new data added in NNPDF3.0 have generally a moderate impact.




Figure 50: Same as Fig. 38, but now comparing the default set to its counterpart obtained without imposing positivity.

5.3.2 Constraints from positivity

As explained in Sect. 3.2.3, in NNPDF3.0 we adopt a more extensive set of positivity constraints: specifically now the number of constraints is equal to the number of PDFs (in fact, it exceeds them because the gluon is constrained by two different pseudo-observables), thereby ensuring not only positivity of the observables used in PDF fitting, but also of potential new observables such as cross-sections for new physics processes used in searches (see Sect. 5.4.4 below).

In order to quantify the impact of these positivity constraint, we have produced an unphysical variant of the NNPDF3.0 NNLO in which positivity constraints are removed (so in particular physical cross-sections could become negative). The distances between the default fit and the fit without positivity are shown in Fig. 50, while some of the PDFs which change most are compared in Fig. 51.

The impact of positivity is relatively mild except for the small- x gluon and the large- x strangeness, for which there is little direct experimental information. For all other PDFs and x ranges the impact of positivity is below the one sigma level. Note that even so, strict positivity is necessary if one wishes to obtain meaningful predictions, e.g. in new physics searches. For the gluon, the effects of the positivity can be noticed already at $x \lesssim 10^{-3}$, while at even smaller x the gluon would become much more negative if positivity were not imposed: so the impact of positivity is to make the gluon less negative at small x . For the strangeness asymmetry, interestingly, the dip-bump structure seen in the global fit is seen to be a consequence of positivity.

As a test of the efficiency of the Lagrange multiplier method we use to impose positivity, we have explicitly checked a posteriori that physical cross sections at NLO and NNLO are indeed not negative (note that because positivity is always imposed at NLO, as discussed in Sect. 3.2.3, small violations of positivity at NNLO are in principle possible). This is illustrated in Fig. 52, where we plot two pseudo-observables which are used to impose positivity, namely the light component of F_L , and the $s\bar{s}$ Drell-Yan rapidity distribution: individual replicas are shown (green dashed curves) as well as the reference set in the positivity implementation (see Sect. 3.2.3), and the effectiveness of positivity (especially for the Drell-Yan distribution) is clearly seen.

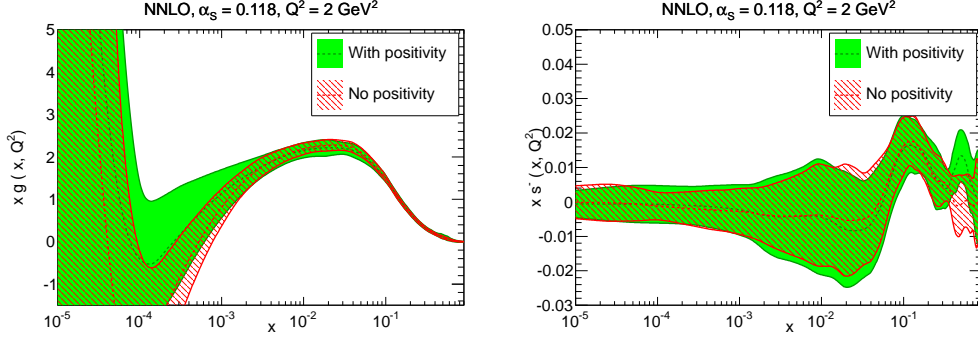


Figure 51: Comparison of the default NNPDF3.0 NNLO PDFs at $Q^2 = 2 \text{ GeV}^2$ with $\alpha_s(M_Z) = 0.118$ to their counterpart obtained without imposing positivity; the gluon (left) and strangeness asymmetry (right) are shown.

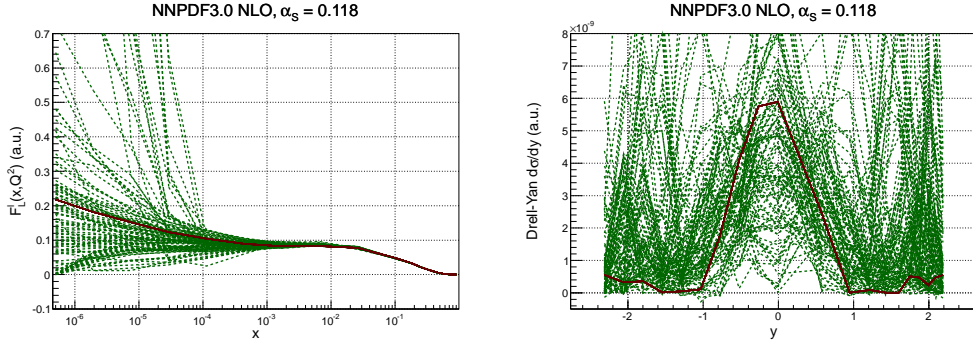


Figure 52: The light quark contribution to F_L (left), and the $s\bar{s}$ Drell-Yan rapidity distribution (right) plotted in arbitrary units at $Q^2 = 5 \text{ GeV}^2$ for individual replicas in the NNPDF3.0 NLO set (dashed green lines). The reference set used in the positivity implementation (see Sect. 3.2.3) is also shown (red line).

5.3.3 Additive versus multiplicative systematics

As discussed in Sect. 2.4.2, there is a certain latitude in the treatment of correlated systematics, in particular whether they are treated additively or multiplicatively (only normalization uncertainties being certainly multiplicative). In order to test the impact of the additive vs. multiplicative treatment of systematics, we have produced two modified version of the NNPDF3.0 fit, which only differ in the treatment of the systematics: in the first one we treat all systematics but the normalization as additive, and in the second, all systematics but the normalization is randomized, and treated as either additive or multiplicative at random for each replica. The default treatment (multiplicative or additive) of systematics is given in Tables 1 and 2 (fourth column).

The distances between these two fits and the default are shown in Fig. 53, while two PDFs for which the effect of the change is most noticeable are shown in Fig. 54 for the additive case, compared to the NNPDF3.0 default. The impact of the treatment of systematics turns out to be essentially indistinguishable from statistical fluctuations for all PDFs except the large- x gluon, for which it can have an effect at the half-sigma level if all systematics is treated as additive. This can be understood as a consequence of the fact that the gluon depends strongly on jet data,




Figure 53: Same as Fig. 38, but now comparing the default set to its counterpart in which systematics (except normalization) is treated as additive (top) or in which the treatment of systematics (except normalization) is randomized (bottom).

which are affected by large systematics. The impact on the gluon is explicitly shown in Fig. 54. The singlet is also shown: in this case, the change in uncertainty at small x seen in the plot is actually compatible with a statistical fluctuations, as the distance plot Fig. 53 demonstrates. When systematics are randomized the effect is diluted and the changes are always compatible with statistical fluctuations.

We conclude that the treatment of systematics, while an issue in principle, in practice affects results at a level which is compatible with fluctuations: in the absence of information on how to treat the systematics this would be randomized, but then results are compatible with the default. Even when all systematics are treated as additive, which is an extreme case, only the gluon changes significantly: but then the gluon is mostly affected by HERA and jet data, for both of which the preferred treatment of systematics is multiplicative [161, 162]. The default treatment of systematics in the NNPDF3.0 fit thus appears to be both reliable and robust.

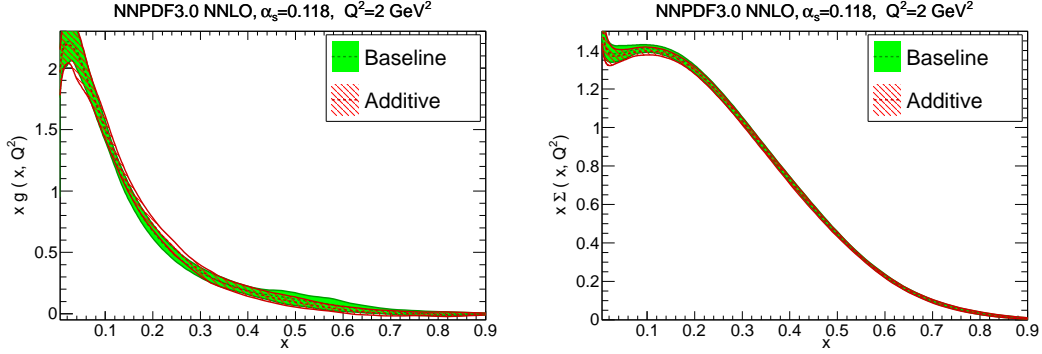


Figure 54: Comparison of the NNLO gluon (left) and singlet (right) in a fit in which all systematic uncertainties but normalization are treated as additive, to the baseline NNPDF3.0 fit, where systematic uncertainties are treated as specified in Tables 1–2. The PDFs are plotted at $Q^2 = 2 \text{ GeV}^2$ with a linear scale, since the difference between baseline and additive is most important at large- x .

5.3.4 Independence of preprocessing

As discussed in Sect. 3.2.2, in NNPDF3.0 the preprocessing exponents which characterize the function that relates PDFs to their neural network parametrization according to Eq. (9) are determined self-consistently. This method has been already used in NNPDF polarized PDF fits [140, 141]. We now check that indeed the method works, and specifically that the ranges of the effective exponents Eq. (11-12) are well within the preprocessing ranges, and thus not biased by it.

The effective exponents and preprocessing ranges for the gluon and singlet in the global NNPDF3.0 NLO fit are shown in Fig. 55: the green solid band is the 68% confidence level interval for each effective exponent as a function of x , for the $N_{\text{rep}} = 1000$ replica fit, and the green dashed line is twice this interval. The red hatched band (and the red dashed lines) are the corresponding results obtained from the $N_{\text{rep}} = 100$ replica fit. In each case, the black solid horizontal lines are the extremes of the preprocessing range, determined from the effective exponents of a previous fit according to the procedure discussed in Sect. 3.2.2. The preprocessing range is self-consistent in the sense that if the range is determined again from the values of the effective exponents shown in Fig. 55 it is essentially unchanged, as it can be seen from visual inspection of the plot.

It is furthermore clear that indeed the effective range is always well within the allowed preprocessing range, thereby ensuring that the former is not influenced or distorted by the latter. We have checked that this is the case for all PDFs and all PDF sets discussed in this paper. Of course, when changing substantially the dataset, such as, for example, when constructing the HERA-only PDFs of Sect. 5.2, the preprocessing ranges may change significantly (with less data the acceptable ranges are generally wider) and thus the procedure has to be iterated again for convergence.

5.3.5 Independence of the PDF fitting basis

We finally wish to test for independence of the fit results of the choice of PDF basis. This was already tested explicitly in Sect. 4 in the framework of a closure test, and it is now verified for the fit to actual data, which is possibly more complicated, both because the final NNPDF3.0 PDFs have sometimes more structure than the MSTW08 PDFs used in that closure test, and

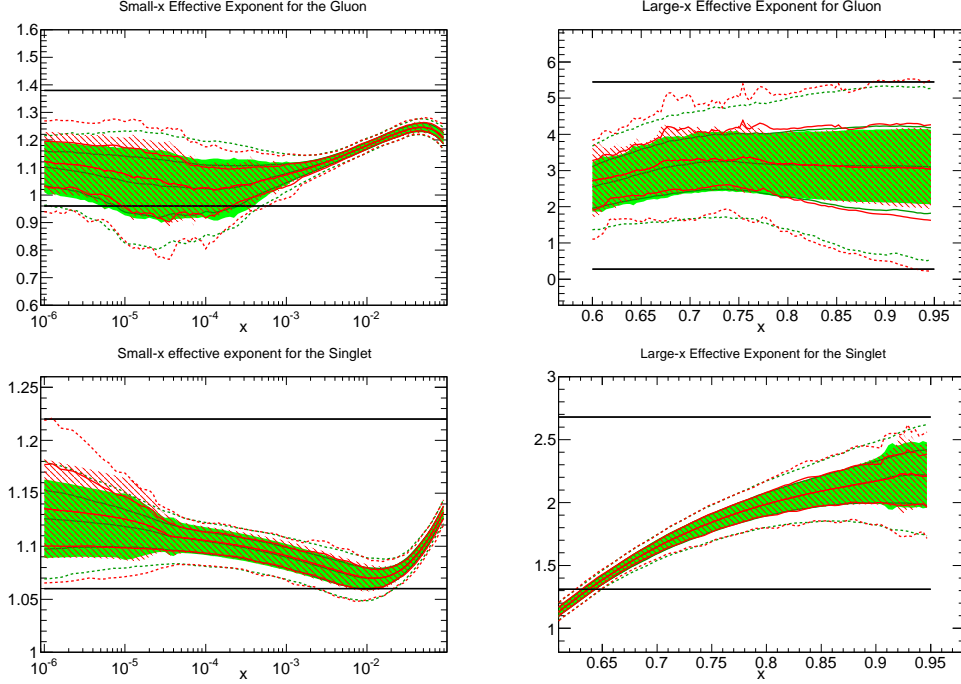


Figure 55: The small- x and large- x effective asymptotic exponents, α_{eff} , Eq. (11) (left) and β_{eff} , Eq. (12) (right), for the gluon (top) and singlet (bottom) in the NNPDF3.0 NLO set. The green solid band is the 68% confidence level interval for each effective exponent as a function of x , for the $N_{\text{rep}} = 1000$ replica fit, and the green dashed line is twice this interval. The red hatched band (and the red dashed lines) are the corresponding results obtained from the $N_{\text{rep}} = 100$ replica fit. In each case, the black solid horizontal lines provide the corresponding range for the preprocessing exponents in the input PDF parametrization Eq. (9).

also because of potential inconsistencies between data which are by construction absent in a closure test. It is worth pointing out that this independence is usually not achieved, or at least not completely, by PDF determinations based on a standard functional forms, as shown in Ref. [163].

In Fig. 56 we show the distances between the default NNPDF3.0 NNLO PDFs and the same fit but using the NNPDF2.3 fitting PDF basis Eq. (7) instead of the default NNPDF3.0 basis Eq. (8) for the parametrization of input PDFs. Results are consistent to what was found in the closure test: distances are mostly compatible with statistical equivalence, with only strangeness at the valence peak deviating at the half-sigma level (which is above the threshold of statistical indistinguishability). Note that the dip-bump structure in s^- seen in Fig. 51 (and related to positivity) arises due to interference of different PDFs in the NNPDF2.3 basis, yet it is perfectly reproduced, thereby confirming the remarkable stability of the NNPDF3.0 results.

5.4 Implications for LHC phenomenology

We now turn to a brief preliminary investigation of the implications of NNPDF3.0 PDFs for LHC phenomenology. We start by comparing the NNLO PDF luminosities at 13 TeV, both to NNPDF2.3, with CT10 and MSTW08. Then we move to predictions for a variety of LHC cross-sections at 13 TeV, which we compute at NLO using the automated MADGRAPH5_AMC@NLO program [112], and for which we compare results obtained using NNPDF2.3 and NNPDF3.0

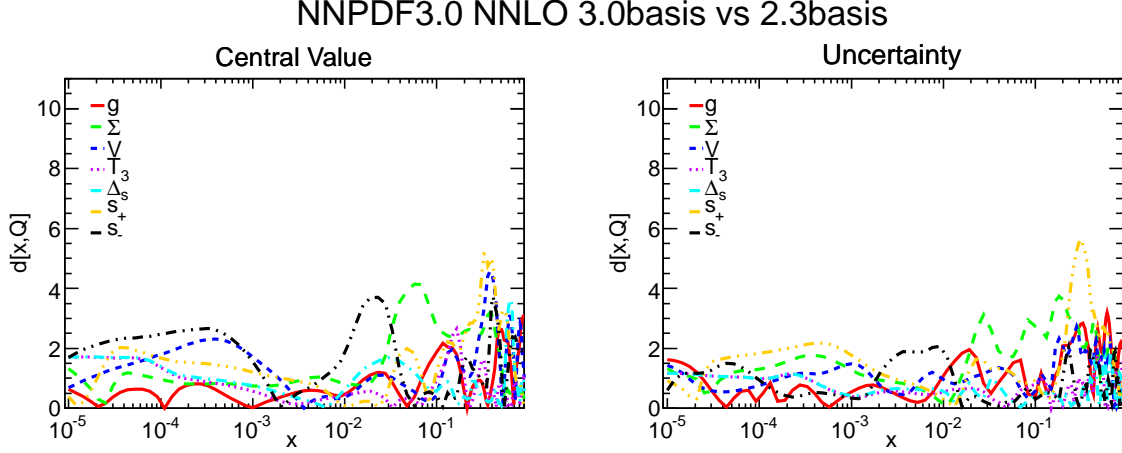


Figure 56: Same as Fig. 38, but now comparing the default NNPDF3.0 fit to one in which the NNPDF2.3 basis Eq. (7) instead of the basis Eq. (8) has been used for PDF parametrization.

PDFs: specifically, vector boson, top, and Higgs production. Then we turn to the implications of NNPDF3.0 PDFs for the dominant Higgs production channel at the LHC, gluon-fusion, and provide NNLO cross-sections computed with the `iHixs` code [150], including a study the dependence of results on the dataset used for PDF determination. Finally, we study the production of high-mass states, close to the LHC kinematic threshold, as relevant for searches of massive New Physics at the energy frontier.

5.4.1 PDF luminosities

Parton luminosities provide a simple way to compare PDFs and estimate the PDF dependence of collider observables. Following Ref. [164], we define the parton luminosity for the ij initial state as

$$\Phi_{ij}(M_X^2) = \frac{1}{s} \int_{\tau}^1 \frac{dx_1}{x_1} f_i(x_1, M_X^2) f_j(\tau/x_1, M_X^2) , \quad (36)$$

where $f_i(x, M_X^2)$ is the PDF for the i -th parton, $\tau \equiv M_X^2/s$ and M_X is the invariant mass of the final state.

We compare gg , qq , $q\bar{q}$ and qg luminosities obtained using NLO and NNLO PDF sets for $\sqrt{s}=13$ TeV and $\alpha_s(M_Z) = 0.118$ (where for quarks a sum over light flavors is understood). In Figs. 57-58 we compare NNPDF2.3 and NNPDF3.0. At NLO, we generally find agreement at the one-sigma level, with some differences in the qq channel for $M_X \sim 500$ GeV, where the error bands of the two PDF sets barely overlap, and in the qg channel above 1 TeV, where the luminosity is rather larger in NNPDF3.0 than in NNPDF2.3. Note that in the gg channel in the region around 100-200 GeV the NNPDF3.0 luminosity is somewhat softer than in NNPDF2.3, though always in agreement within PDF uncertainties.

At NNLO, in the qq and $q\bar{q}$ channels there is generally good agreement, with differences well within one sigma: for $q\bar{q}$, the NNPDF3.0 luminosity tends to be larger at high invariant masses, while for qq around 500 GeV NNPDF3.0 is somewhat lower, with barely overlapping error bands. More significant differences are found in the gg channel, where the luminosity at medium invariant masses is smaller by about one sigma in NNPDF3.0 than in NNPDF2.3: in particular, for $30 \text{ GeV} \leq M_X \leq 300 \text{ GeV}$, the gg one sigma bands barely overlap. This has




Figure 57: Parton luminosities, Eq. (36) computed using NNPDF2.3 and NNPDF3.0 NLO PDFs with $\alpha_s(M_Z) = 0.118$, as a function of the invariant mass of the final state M_X . Results are shown as ratios to NNPDF2.3. From top to bottom and from left to right the $q\bar{q}$, qq , qg and gg luminosities are shown.

important consequences for gluon-initiated processes such as inclusive Higgs production, see Sect. 5.4.3 below. As discussed in Sect. 5.1.2, these differences stem from a combination of the improved fitting methodology and the new constraints from HERA and LHC data.

We then compare NNPDF3.0 to CT10 and MSTW08: results are shown in Fig. 59. Note however that this comparison will become soon obsolete with the upcoming updates from the MSTW and CT groups. The three sets are consistent within their uncertainties. Quite in general NNPDF3.0 has smaller uncertainties in the data region, but larger uncertainties in the extrapolation regions. For the gg luminosity in the region relevant for Higgs production, the agreement between the three sets has improved in comparison to previous benchmarks using NNPDF2.3 [137]. Large differences are found for large values of M_X , relevant for the production of very massive New Physics particles.

5.4.2 Implications for $\sqrt{s}=13$ TeV LHC processes

An extensive phenomenological comparison of NNPDF3.0 for all processes and a wide range of PDF sets is beyond the scope of this work. We will thus only compare NNPDF2.3 and NNPDF3.0, for a limited set of LHC observables. Unless otherwise stated, we use the NLO sets with $\alpha_s = 0.118$ and work in the $N_f = 5$ variable-flavor-number scheme with massless bottom quarks. Theoretical predictions at NLO are computed using the MADGRAPH5_AMC@NLO program [112], version 2.1.2, interfaced to LHAPDF6. The NLO results are sufficient to assess the PDF dependence of these observables, as typically the NNLO/NLO K -factor have a weak PDF dependence. In addition to NNPDF3.0 global fit results, we also provide predictions using the conservative parton set with $\alpha_{\max} = 1.1$ as an illustration of results found using a maximally




Figure 58: Same as Fig. 57 at NNLO.

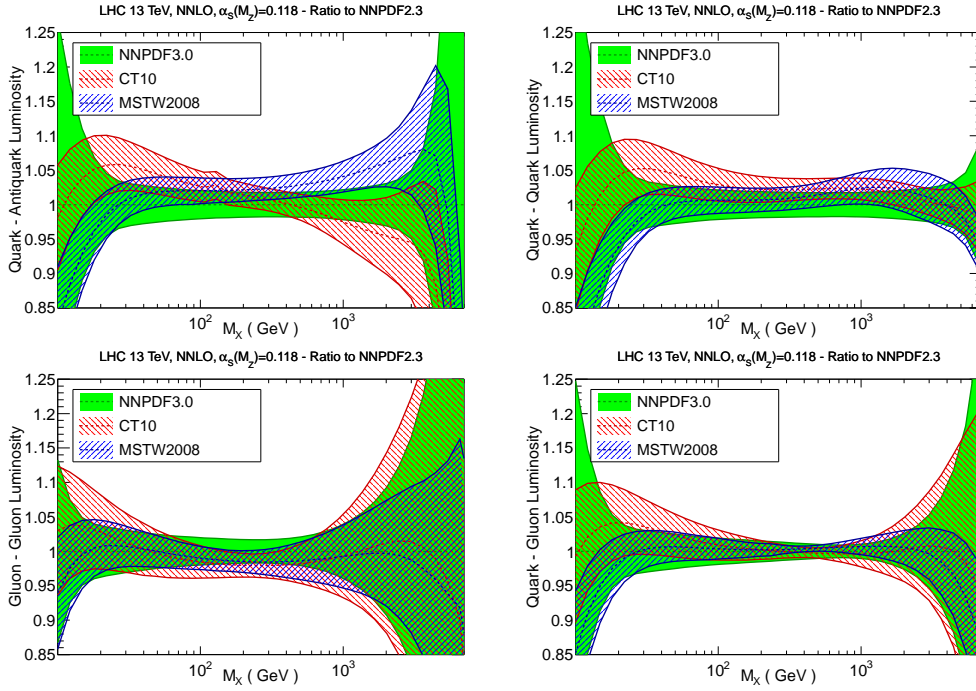


Figure 59: Same as Fig. 58 but now comparing NNPDF3.0, MSTW08 and CT10 NNLO (all with $\alpha_s(M_Z) = 0.118$). Results are shown as ratios to NNPDF3.0.

Process	NNPDF2.3	NNPDF3.0	RelDiff	NNPDF3.0 $\alpha_{\max} = 1.1$
p p > l+ l- [QCD]	1.403 nb ($\pm 1.5\%$)	1.404 nb ($\pm 2.0\%$)	+0.1%	1.45 nb ($\pm 2.0\%$)
p p > l+ vl [QCD]	10.30 nb ($\pm 1.3\%$)	10.21 nb ($\pm 1.9\%$)	-0.9%	10.29 nb ($\pm 2.3\%$)
p p > l- vl [QCD]	7.67 nb ($\pm 1.3\%$)	7.75 nb ($\pm 1.9\%$)	+1.1%	7.96 nb ($\pm 1.9\%$)
p p > w+ c [QCD]	2.665 nb ($\pm 3.5\%$)	2.680 nb ($\pm 4.2\%$)	+0.56%	2.807 nb ($\pm 8.8\%$)
p p > l+ vl j [QCD]	2.353 nb ($\pm 1.2\%$)	2.332 nb ($\pm 1.5\%$)	-0.9%	2.325 nb ($\pm 1.6\%$)
p p > a j [QCD]	62.24 nb ($\pm 1.2\%$)	63.85 ($\pm 1.8\%$)	+2.6%	61.51 ($\pm 1.9\%$)
p p > t t [QCD]	678 pb ($\pm 1.7\%$)	672 pb ($\pm 1.6\%$)	-0.9%	655 pb ($\pm 3.3\%$)
p p > h l+ l- [QCD]	26.48 fb ($\pm 1.4\%$)	26.58 fb ($\pm 1.5\%$)	+0.4%	27.07 fb ($\pm 2.3\%$)
p p > h l+ vl [QCD]	0.134 pb ($\pm 1.6\%$)	0.131 pb ($\pm 1.6\%$)	-2.2%	0.137 pb ($\pm 2.6\%$)
p p > h t t [QCD]	0.458 pb ($\pm 2.2\%$)	0.4595 pb ($\pm 1.7\%$)	+0.6%	0.459 pb ($\pm 4.0\%$)

Table 16: Cross-sections for LHC at 13 TeV, computed at NLO using e MADGRAPH5_AMC@NLO using NNPDF2.3 and NNPDF3.0 NLO PDFs, with $N_f = 5$ and $\alpha_s(M_Z) = 0.118$. In each case, central values and the one-sigma PDF uncertainty (in parenthesis) are given. We also show the percentage difference between the central values using the two PDF sets and the prediction using conservative partons with $\alpha_{\max} = 1.1$.

consistent dataset (see [165]).

Even though we do not aim to a direct comparison to the raw data, cross-sections have been computed at the fiducial level, including resonance decays for some processes, and using realistic generation cuts. Jets are reconstructed with the anti- k_T algorithm [166] with radius $R = 0.5$, and the following cuts are applied to all jets in the final state:

$$|\eta_{\text{jet}}| \leq 4.5, \quad p_{T,\text{jet}} \geq 25 \text{ GeV}. \quad (37)$$

For final-state leptons, the following cuts are applied:

$$|\eta_l| \leq 2.5, \quad p_{T,l} \geq 25 \text{ GeV}, \quad m_{l+l-} \geq 30 \text{ GeV}. \quad (38)$$

Finally, for photons we impose

$$|\eta_\gamma| \leq 2.5, \quad p_{T,\gamma} \geq 25 \text{ GeV}, \quad (39)$$

and use the Frixione isolation criterion [167], with $\epsilon_\gamma = 1.0$ and $n = 1$ and an isolation cone radius $R_0 = 0.4$. No further analysis cuts are applied. Renormalization and factorization scales are set dynamically event by event to $\mu_f = \mu_r = H_T/2$, with H_T the scalar sum of the transverse energies of all the final-state particles. Within each run, PDF and scale uncertainties in MADGRAPH5_AMC@NLO are obtained at no extra cost using the reweighting technique introduced in Ref. [168]: we thus provide in each case both the central value and the PDF uncertainty.

Results are collected in Tab. 16, where processes are grouped into three subsets: processes which are sensitive to quark and antiquarks, processes which are sensitive to the gluon PDF, and Higgs production processes, except gluon fusion which is discussed in the next Section. The results of Tab. 16 are also represented in Fig. 60, normalized to NNPDF2.3.

For all these cross-section a remarkable stability between NNPDF2.3 and NNPDF3.0 is observed, with all results varying by no more than the size of the corresponding PDF uncertainty. For top-quark pair production, going from NNPDF2.3 to NNPDF3.0 the cross-section increases by about 1%, about half the PDF uncertainty. This can be understood recalling that the NNPDF3.0 gg luminosity is harder than is NNPDF2.3 counterpart for $M_X \sim 400 \text{ GeV}$; this also explains the behaviors of the closely related ttH production process. Note that NNPDF2.3

LHC 13 TeV, $\alpha_s=0.118$, MadGraph5_aMC@NLO fNLO




Figure 60: Graphical comparison the results of Tab 16. Results are shown normalized to the NNPDF2.3 central value.

already gave a very good description of all available ATLAS and CMS 7 TeV and 8 TeV production data [76], even though they were not included in the fit. For W production in association with charm quarks, we use a $N_F = 3$ scheme in order to retain the full charm mass dependence. For this observable, results are very stable when moving from NNPDF2.3 to NNPDF3.0.

Coming now to Higgs production, for ttH the NNPDF3.0 result is about 2.3% larger than the NNPDF2.3 prediction, consistent with the expectation from the gg luminosity comparisons in Fig. 57 for $M_X \sim 500 = 700$ GeV (recall that the calculation uses a dynamical setting of the factorization scale). For the associate production channels, hW and hZ , driven by the $q\bar{q}$ luminosities, differences are well within one sigma, as expected from the luminosities of Fig. 57.

Coming finally to the comparison with the conservative PDFs, we note that prediction obtained using the latter are generally consistent at the one sigma level with the default ones, occasionally with differences at the two sigma level, such as for example for $hl^+\nu$.

Of course, predictions from conservative partons are generally affected by larger PDF uncertainties because of the smaller dataset, though in some cases they are only slightly less accurate than the global fit, such as for example for inclusive W and Z production. On the other hand, for processes that depend on strangeness (like $W+c$) or that are gluon-driven (like $t\bar{t}$ and $t\bar{t}h$) the PDF uncertainties are substantially larger in the conservative partons than in the global fit.

	σ_{ggh} (pb) NLO	Pull	σ_{ggh} (pb) NNLO	Pull
NNPDF2.3	34.72 ± 0.33	-	46.39 ± 0.46	-
NNPDF3.0 with 2.3 data	34.06 ± 0.57	1.0	45.14 ± 0.74	1.4
NNPDF3.0 global	33.96 ± 0.61	1.1	45.01 ± 0.72	1.6
NNPDF3.0 conservative $\alpha_{\max} = 1.1$	33.31 ± 0.54	2.2	43.70 ± 1.12	2.2
NNPDF3.0 no Jets	34.56 ± 1.04	0.2	45.32 ± 0.92	1.0
NNPDF3.0 noLHC	34.12 ± 0.80	0.7	45.10 ± 0.91	1.3
NNPDF3.0 HERA-only	31.96 ± 3.03	0.9	43.02 ± 2.21	1.5

Table 17: The total cross-section for Higgs production in gluon fusion at the LHC 13 TeV at NLO (left) and NNLO (right) for $\alpha_s(M_Z) = 0.118$. The pull P Eq. (40) is also given.

5.4.3 Higgs production in gluon fusion

We now focus specifically on Higgs production in gluon fusion, the dominant channel at the LHC, for which theoretical uncertainties are a limiting factor for the determination of Higgs properties. We provide predictions for the total cross-section at NLO and NNLO for the LHC 13 TeV, comparing the default NNPDF3.0 set to NNPDF2.3 and to the various sets based on alternative datasets discussed in Sect. 5.2, with the main goal of studying the dependence of this prediction on the underlying PDFs, along the lines of the study presented in Ref. [157]. The uncertainties shown are pure PDF uncertainties, with $\alpha_s(M_Z) = 0.118$, i.e. the α_s uncertainty is not included. The inclusive cross-section is computed using iHixs 1.3.3, with $m_h = 125$ GeV, renormalization and factorization scales set to $\mu_r = \mu_f = m_h$ and the infinite top mass (effective theory) approach. Clearly, our predictions are not meant to be realistic, however, the effects which we do not include (such as for instance electroweak corrections, or finite top, bottom and charm mass contributions) have a negligible PDF dependence, while α_s uncertainties are completely independent of the PDF uncertainty, given that the PDF and α_s uncertainties combine in quadrature in the gaussian case even when correlated [169]. Hence our results do provide an accurate assessment of the PDF dependence of the cross-section and its uncertainty.

Results are collected in Tab. 17, and summarized graphically in Figs. 61-62. We show results obtained using NNPDF2.3, and the following NNPDF3.0 sets: 2.3-like dataset, default, conservative with $\alpha_{\max} = 1.1$, no jet data, no LHC data and HERA-only, all of which have been discussed in detail in Sect. 5.2. In Tab. 17 we also show the pull of each prediction compared to the NNPDF2.3 result, defined as

$$P \equiv \frac{(\sigma_{ggh}(2.3) - \sigma_{ggh}(3.0))}{\sqrt{\Delta\sigma_{ggh}^2(3.0) + \Delta\sigma_{ggh}^2(2.3)}}, \quad (40)$$

where $\Delta\sigma_{ggh}$ is the one-sigma PDF uncertainty.

As expected from the comparison of the gluon-gluon luminosities in Fig. 58, the NNLO cross-section decreases by about 2-sigma when going from NNPDF2.3 to NNPDF3.0, while the PDF uncertainty increases substantially. At NLO the effect is less marked, with the NNPDF2.3 and NNPDF3.0 in agreement at the one-sigma level. Because a similar result is found using NNPDF3.0 PDFs based on a 2.3-like dataset we must conclude that the change is mostly due to methodological improvements, rather than the underlying data. Because of the validation from closure testing, the NNPDF3.0 results are more reliable, both in terms of central values and uncertainties.

Results obtained using NNPDF3.0 sets based on alternative datasets are always mutually consistent at the one sigma level, with the conservative partons leading to a lower result and




Figure 61: Graphical representation of the NLO results of Tab. 17.

the fit with no jets to a slightly higher one but with significantly larger uncertainty. The lowest central value is found using the HERA-only set, which however is affected by a PDF uncertainty which is by a factor three larger than the default.

Despite these variations in results, all pulls are of similar size, between 0.7 and 1.1 at NLO and between 1.3 and 1.6 at NNLO, with the only exception of conservative partons, which have a larger pull of 2.2 both at NLO and NNLO. This means that the dependence of the results on the dataset appears to be consistent with statistical fluctuations.

5.4.4 New Physics particle production at high masses

A wide variety of scenarios of physics beyond the standard model include heavy particles at the TeV scale that could be within the reach of the LHC. Production of very massive particles probes PDFs at large x , where they are poorly known due to the lack of direct experimental information, and thus the corresponding predictions are affected by substantial PDF uncertainties (see e.g. Refs. [170,171]). Consequently, PDFs can be a limiting factor in the determination of exclusion regions, and improving their knowledge can lead to an increase in the search potential. In this context an accurate assessment of their uncertainties is therefore crucial. The unbiased NNPDF approach is advantageous in this respect in that it leads to uncertainty estimates which are not biased by assumptions on the functional form of PDFs. The only significant constraint on PDFs close to threshold comes from positivity, which is now implemented in an optimal way as we discussed in Sects. 3.2.3-5.3.2 above.

As an example, we consider high-mass Drell-Yan production and the pair production of supersymmetric particles. High-mass dilepton production is frequently used to search for new physics that couples to the electroweak sector, and thus it is important to provide precise predictions for the SM production mechanisms.

We have computed the dilepton invariant mass distribution in $pp \rightarrow \gamma^*/Z \rightarrow l^+l^-$ events at the LHC 14 TeV with NNPDF3.0 at NLO and NNLO using FEWZ. Recall from Sect. 3.2.3 that positivity is always imposed at NLO, so an explicit check of positivity of the NNLO result




Figure 62: Graphical representation of the NNLO results of Tab. 17.

is nontrivial. Results are shown in Fig. 63, in different M_{ll} bins: each of the $N_{\text{rep}} = 100$ NLO (left) or NNLO (right) Monte Carlo replicas is shown as a green dashed line, together with the corresponding central values and one-sigma intervals. All MC replicas are positive up to the highest invariant mass bins.

As a second example, we provide predictions for the pair production of supersymmetric particles at the LHC 14 TeV. The computation has been performed using PROSPINO [172, 173] with NNPDF3.0 NLO, using settings as close as possible to those of Refs. [170, 171], though the only relevant physical input for our illustrative study are the sparticle masses. Note that for this processes NNLO calculations are not available. We have produced results for squark-squark, squark-antisquark and gluino-gluino production, for three different values of the sparticle masses, namely 1 TeV, 2 TeV and 3 TeV. We show the predictions for the $N_{\text{rep}} = 100$ Monte Carlo replicas of NNPDF3.0 for the squark-antisquark and gluino-gluino channels; the squark-squark cross-section (not shown) is always positive in the mass range that we are considering. In each case, we also provide the average result and the 68% confidence level interval. For comparison, predictions using the NNPDF2.3 NLO PDFs are also shown. Results are shown in Fig. 64. In the case of gluino-gluino production (right), all replicas are strictly positive up to $m_{\tilde{g}} = 2$ TeV. At 3 TeV, some replicas lead to slightly negative cross-sections: 15 in NNPDF2.3, and only 3 in NNPDF3.0. In both cases, the 68% confidence levels are always positive; we conclude that the occasional negative replica can be set to zero with no impact on the central value or uncertainty. For squark-antisquark production (left) all cross sections are positive up to $m_{\tilde{q}} = 2$ TeV, while for $m_{\tilde{q}} = 3$ TeV some cross-sections are slightly negative. For NNPDF2.3, the central value was negative, while now for NNPDF3.0 the central value is positive and only a small part of the 68% confidence level range is in the negative region. Setting these negative cross-sections to zero would of course not modify the upper uncertainty range, implying that for this particular subprocesses and sparticle masses arbitrarily small production cross-sections are allowed within the large PDF uncertainties. Note however the very large PDF uncertainties at the largest masses: for example, for squark-antisquark production, the 68% CL range for the cross-section for $m_{\tilde{q}} = 3$ TeV is $\sigma(\tilde{q}\tilde{q}) \in [0, 15 \cdot 10^{-5}]$ pb, with a central value of $\sim 5 \cdot 10^{-5}$ pb.




Figure 63: The dilepton invariant mass distribution in $pp \rightarrow \gamma^*/Z \rightarrow l^+l^-$ at the LHC 14 TeV with NNPDF3.0 at NLO (left) and NNLO (right) using FEWZ. Each (green) dashed line is one of the $N_{\text{rep}} = 100$ Monte Carlo replicas, while the solid (red) line is the average and the outer solid (blue) lines give the edges of the one-sigma interval. Both the absolute result (top) and the ratio to the central value (bottom) are shown.

We conclude that, thanks to the improved implementation of positivity in NNPDF3.0, the cross-section for high-mass particle production is positive, thereby improving over NNPDF2.3. Occasional replicas leading to negative cross-sections very close to threshold can be set to zero without significantly affecting central values and uncertainties.

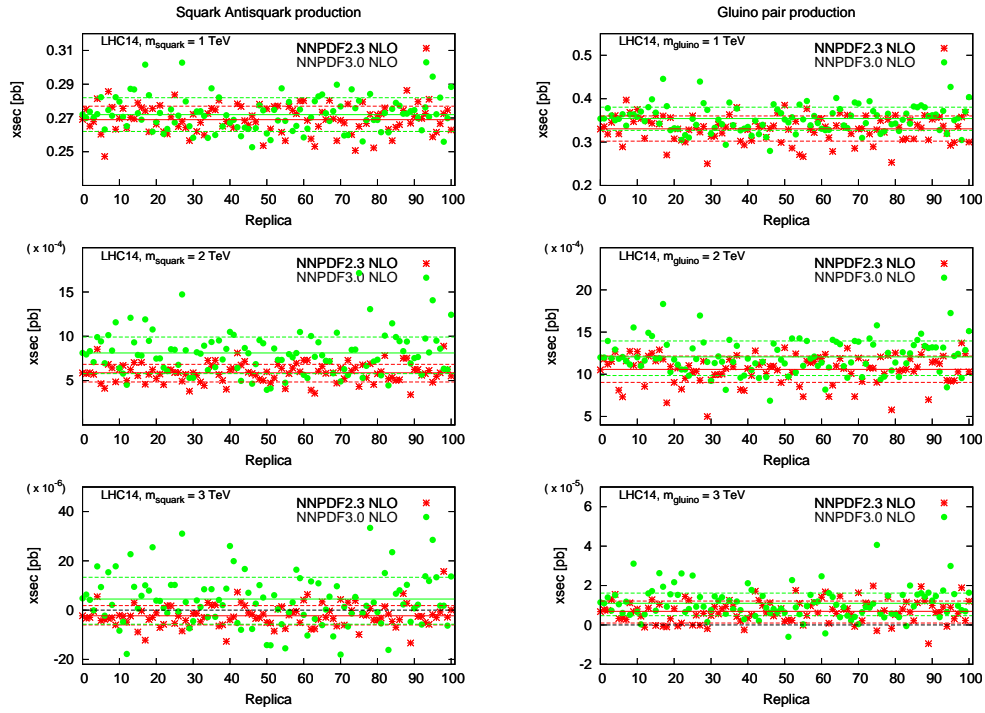


Figure 64: Cross-sections for NLO squark-antisquark (left) and gluino-gluino (right) pair production at the LHC 14 TeV with NNPDF3.0 and NNPDF2.3, for sparticle masses of 1 TeV, 2 TeV and 3 TeV. In each case, we show the predictions for the $N_{\text{rep}} = 100$ Monte Carlo replicas as well as by the average result and the 68% confidence level interval.

6 Summary and outlook

We conclude by summarizing the NNPDF3.0 fits that will be made public, their use in specific phenomenological application, and future developments.

All PDF sets are made available via the LHAPDF6 library [174], from version 6.1.4 onward,

<http://lhpdf.hepforge.org/>

For the global fits, we also provide the corresponding grids in LHAPDF5 format via the the NNPDF HEPFORGE webpage,

<http://nnpdf.hepforge.org/>

though whenever possible, the modern v6 interface should be used.

The interpolation used in the delivery of NNPDF sets has been substantially improved in the transition from LHAPDF5 to LHAPDF6.

- The default interpolation algorithms, of LHAPDF6 are now used, instead of the in-house interpolation used by NNPDF sets in LHAPDF5. These algorithms are faster and more accurate, but we have checked that in most of the relevant range of x and Q^2 the LHAPDF6 and LHAPDF5 algorithms agree at the permille level.

- A wider range for the x and Q grids is now adopted: the NNPDF3.0 sets can be used for

$$10^{-9} \leq x \leq 1, \quad 1 \text{ GeV} \leq Q \leq 100 \text{ TeV}, \quad (41)$$

which is adequate for future phenomenological applications, including calculations of cross-section of ultra-high energy neutrinos, which require PDFs down to $x \sim 10^{-8}$, and studies of a Future Circular Collider with a center-of-mass energy of 100 TeV. Outside the range in Eq. (41) the values of the PDFs are frozen.

- Thanks to the novel functionalities in LHAPDF6, for all the NNPDF3.0 sets the Q^2 interpolation grid has been divided into different subgrids separated by heavy quark thresholds. This ensures that the correct threshold behavior of the heavy quark PDFs is satisfied also by the interpolation, and in particular that the heavy quark PDFs are exactly zero below the threshold.

The NNPDF3.0 PDF sets to be made available via LHAPDF6 are the following:

- Global NNPDF3.0 sets - baseline fits

The baseline LO, NLO and NNLO NNPDF3.0 sets are based on the global dataset, with $\alpha_s(M_Z) = 0.118$ and a variable-flavor number with up to five active flavors, the convention for the corresponding LHAPDF6 grid files is

```
NNPDF30_lo_as_0118
NNPDF30_nlo_as_0118
NNPDF30_nnlo_as_0118
```

The leading order NNPDF3.0 allows, among other applications, to compare theoretical predictions for cross-sections computed at different perturbative orders. in LHAPDF6, the output of the NNPDF3.0 LO PDFs is forced to be positive using the suitable meta-data instructions. For uses in leading-order Monte Carlo event generators, we still recommend the use NNPDF2.3LO, since, as part of the updated Monash 2013 Tune [151] of PYTHIA8 [175], it has been shown to provide an excellent description of a variety of soft and semi-hard data.

- Global NNPDF3.0 sets - fits with α_s variations

Using the same global dataset as above, for the NLO and NNLO fits we have produced fits with different values of $\alpha_s(M_Z)$, that can be used to evaluate the combined PDF+ α_s in cross-sections:

NNPDF30_nlo_as_0115, NNPDF30_nlo_as_0117, NNPDF30_nlo_as_0119,
 NNPDF30_nlo_as_0121
 NNPDF30_nnlo_as_0115, NNPDF30_nnlo_as_0117, NNPDF30_nnlo_as_0119,
 NNPDF30_nnlo_as_0121

The procedure to combine PDF and α_s uncertainties in a Monte Carlo set like NNPDF3.0 is explained in [133] (see also Ref. [176]). The choice of the uncertainty $\delta\alpha_s$ to be assigned to the strong coupling is left to the PDF users. For LO, we also provide a fit with $\alpha_s(M_Z) = 0.130$, as this is the value typically required in LO calculations:

NNPDF30_lo_as_0130

Note that from the NLO and NNLO range of α_s values, the predictions for any other value of α_s not included in the list can be easily obtained from either interpolation or extrapolation.

- Global NNPDF3.0 - fits with different maximum number of active flavors

For various important phenomenological applications, PDF sets in which either the charm quark or both the charm and bottom quarks are treated as massive are required. In the following we denote these two schemes as the $N_f = 3$ and $N_f = 4$ schemes. In order to obtain $N_f = 3$ and $N_f = 4$ versions of the NNPDF3.0 sets, we follow the procedure outlined in [11]. We start from the baseline NNPDF3.0 fits, where the charm and bottom quarks are massless parton ($N_f = 5$ scheme), at the input scale, $Q_0^2 = 1 \text{ GeV}^2$. This boundary condition is then evolved upwards in Q^2 , but using the $N_f = 3$ and $N_f = 4$ schemes, to produce the desired sets with different maximum number of active partons. This has been done for a reduced range of $\alpha_s(M_Z)$ values, to be able to compute the combined PDF+ α_s uncertainties. These fits are:

NNPDF30_nlo_as_0117_nf3, NNPDF30_nlo_as_0118_nf3, NNPDF30_nlo_as_0119_nf3
 NNPDF30_nnlo_as_0117_nf3, NNPDF30_nnlo_as_0118_nf3, NNPDF30_nnlo_as_0119_nf3
 NNPDF30_nlo_as_0117_nf4, NNPDF30_nlo_as_0118_nf4, NNPDF30_nlo_as_0119_nf4
 NNPDF30_nnlo_as_0117_nf4, NNPDF30_nnlo_as_0118_nf4, NNPDF30_nnlo_as_0119_nf4

In addition, using a similar procedure we have provided sets in the $N_f = 6$ scheme, that is, treating the top quark as a massless parton. While treating the top quark as massless is not justified at the LHC, it might become appropriate at higher energy colliders, see for example Ref. [177]. These sets that include the top PDF are denoted as

NNPDF30_nlo_as_0117_nf6, NNPDF30_nlo_as_0118_nf6, NNPDF30_nlo_as_0119_nf6
 NNPDF30_nnlo_as_0117_nf6, NNPDF30_nnlo_as_0118_nf6, NNPDF30_nnlo_as_0119_nf6

Likewise, we also provide similar fits with different maximum number of active flavors for the LO fits, for the two values of α_s used:

NNPDF30_lo_as_0118_nf3, NNPDF30_lo_as_0118_nf4, NNPDF30_lo_as_0118_nf6,
 NNPDF30_lo_as_0130_nf3, NNPDF30_lo_as_0130_nf4, NNPDF30_lo_as_0130_nf6,

- Fits based on reduced datasets

We have also released fits based on reduced datasets, as discussed in detail in Sect. 5.2. First, conservative partons, for which the fit with the choice of threshold value of $\alpha_{\max} = 1.1$ is used. We provide results in the restricted range of α_s values:

NNPDF30_nlo_as_0117_cons, NNPDF30_nlo_as_0118_cons, NNPDF30_nlo_as_0119_cons
 NNPDF30_nnlo_as_0117_cons, NNPDF30_nnlo_as_0118_cons, NNPDF30_nnlo_as_0119_cons

Then, HERA-only PDFs, which may be useful to gauge the impact of new LHC measurements:

NNPDF30_nlo_as_0117_hera, NNPDF30_nlo_as_0118_hera, NNPDF30_nlo_as_0119_hera
 NNPDF30_nnlo_as_0117_hera, NNPDF30_nnlo_as_0118_hera, NNPDF30_nnlo_as_0119_hera

In view of possible reweighting applications, for the central value of $\alpha_s(M_Z) = 0.118$ also sets $N_{\text{rep}} = 1000$ are provided,

NNPDF30_nlo_as_0118_hera_1000
 NNPDF30_nnlo_as_0118_hera_1000

Next, HERA+ATLAS and HER+CMS fits. These can be useful to compare with analogous fits presented by the LHC collaborations, and are labeled as

NNPDF30_nlo_as_0117_atlas, NNPDF30_nlo_as_0118_atlas, NNPDF30_nlo_as_0119_atlas
 NNPDF30_nnlo_as_0117_atlas, NNPDF30_nnlo_as_0118_atlas, NNPDF30_nnlo_as_0119_atlas
 NNPDF30_nlo_as_0117_cms, NNPDF30_nlo_as_0118_cms, NNPDF30_nlo_as_0119_cms
 NNPDF30_nnlo_as_0117_cms, NNPDF30_nnlo_as_0118_cms, NNPDF30_nnlo_as_0119_cms

Then, sets without LHC data, sometimes used for theory comparisons with LHC measurements. Also here, in view of possible reweighting exercises with LHC data, we provide for the central value of $\alpha_s(M_Z) = 0.118$ a set of $N_{\text{rep}} = 1000$ replicas:

NNPDF30_nlo_as_0117_nolhc, NNPDF30_nlo_as_0118_nolhc, NNPDF30_nlo_as_0119_nolhc
 NNPDF30_nlo_as_0118_nolhc_1000
 NNPDF30_nnlo_as_0117_nolhc, NNPDF30_nnlo_as_0118_nolhc, NNPDF30_nnlo_as_0119_nolhc
 NNPDF30_nnlo_as_0118_nolhc_1000

Finally, a fit excluding all jet data, which avoids the use of approximate NNLO results. Again, these are provided in a range of α_s values, and denoted by

NNPDF30_nlo_as_0117_nojet, NNPDF30_nlo_as_0118_nojet, NNPDF30_nlo_as_0119_nojet
 NNPDF30_nnlo_as_0117_nojet, NNPDF30_nnlo_as_0118_nojet, NNPDF30_nnlo_as_0119_nojet

Note that for all fits for which $N_{\text{rep}} = 1000$ replicas are available, the corresponding $N_{\text{rep}} = 100$ replica sets have been obtained using independent random seeds, so that the $N_{\text{rep}} = 1000$ and $N_{\text{rep}} = 100$ replica sets are fully independent.

Now that PDFs validated by a closure test, and fitted to a wide variety of LHC data are available, three main directions of progress are foreseen. From a methodological point of view, some optimization is still possible: specifically by producing sets with a reduced number of PDF replicas which retain as much as possible of the information from the full dataset. From the phenomenological point of view, new data will be added to the dataset as soon as they become available. These include the final combined HERA dataset, jet data from ATLAS and CMS from the 2011 and 2012 runs, and electroweak gauge boson production at 8 TeV, and, in the longer term, LHC data from Run II will also be included. Finally, from the theory point of view we plan to improve the treatment of heavy quarks by introducing an intrinsic charm PDFs, to provide PDF sets based on resummed QCD theory, and to produce QCD+QED sets including a full fit of the photon PDF.

With data, methodology and theory under control the main outstanding issue is a full characterization of theoretical uncertainties on PDFs, which remains an open problem. It will have to be addressed in order for parton distributions to become a controlled tool for precision physics at present and future colliders.

Acknowledgments

We thank L. Lyons for stressing the importance of closure testing of parton distributions, and A. de Roeck, A. David, J. Huston, M. Mangano, A. Mitov, P. Nadolsky, G. Passarino, G. Salam, R. Tanaka, R. Thorne and G. Watt for many illuminating discussions. We thank F. Petriello for his help in running FEWZ and comparing the results and G. Ferrera for his help in setting up DYNNLO. We are grateful to S. Glazov for support with the HERA-II H1 data; A. Cooper-Sarkar and I. Abt for support with the ZEUS HERA-II data; S. Glazov, C. Gwenlan and A. Tricoli for information about the ATLAS measurements; J. Alcaraz, J. Berryhill, M. Cepeda, M. Gouzevitch, I. Josa, K. Lipka and K. Rabbertz for assistance with the CMS data; and D. Ward, J. Anderson, K. Mueller, R. McNulty and T. Shears for help with the LHCb electroweak measurements.

V. B. is supported by the ERC grant 291377, LHCtheory: Theoretical predictions and analyses of LHC physics: advancing the precision frontier. S. F. and S. C. are supported in part by an Italian PRIN2010 grant, by a European Investment Bank EIBURS grant, and by the European Commission through the HiggsTools Initial Training Network PITN-GA-2012-316704. J. R. is supported by an STFC Rutherford Fellowship ST/K005227/1. J. R. and N. H. are supported by an European Research Council Starting Grant "PDF4BSM". R. D. B. and L. D. D. are funded by an STFC Consolidated Grant ST/J000329/1.

A QCD and weak corrections to vector boson production data

In this appendix we provide a detailed overview of the theoretical calculations that we have used to include LHC vector boson production data in the NNPDF3.0 analysis. For each datapoint, we provide the size of the NNLO QCD and of the NLO weak corrections, comparing with the total experimental uncertainty. Results for the QCD and weak corrections are given for all the LHC neutral current Drell-Yan measurements included in NNPDF3.0, namely for the CMS double differential distributions from the 2011 run [59], the ATLAS high mass Drell-Yan distributions from the 2011 run [56], the ATLAS Z boson rapidity distribution from the 2010 run [47], and for LHCb Z rapidity distributions in the forward region from the 2011 run [61]. Using these calculations, here we also provide, for the neutral-current Drell-Yan data, the NNLO and weak C -factors, defined as in Sect. 2.3.

As we will show now, the effect of the pure weak corrections is negligible for neutral-current Drell-Yan data around Z peak region, and they become numerically important for small and large values of the dilepton mass M_{ll} . For charged-current Drell-Yan productions, weak corrections are small outside the W peak region (where all the available data lies) and thus can be safely neglected in the present analysis. For charged-current DY, NNLO QCD corrections are moderate, they have been included in the NNPDF3.0 analysis but are not discussed explicitly in this Appendix.

In this work, theoretical predictions are obtained with FEWZ3.1 [118] and have been cross-checked against DYNLO1.3 [113]. The NNPDF2.3 NLO set with $\alpha_s = 0.118$ is used as input for the NLO theoretical predictions, and NNPDF2.3 NNLO is used for the corresponding NNLO calculations, as well as to determine the C -factors. The G_μ scheme is adopted for all electroweak computations, where the input parameters are M_Z , M_W and G_F , while α_e and s_W may be obtained from these using:

$$s_W = 1 - \frac{M_W^2}{M_Z^2} \quad \alpha_e = R_t^2 G_F M_W^2 \frac{s_W}{\pi}, \quad (42)$$

where $R_t^2 = 1.4142135624$. In this analysis we use $G_F = 1.16637 \cdot 10^{-5} \text{ GeV}^{-2}$, $M_W = 80.398 \text{ GeV}$, $M_Z = 91.1876 \text{ GeV}$ and do not adopt the narrow-width approximation.

A.1 CMS double differential distributions

The CMS experiment has measured [59] the lepton pair rapidity distributions y_{ll} for the Drell-Yan process in six bins of the invariant mass of the final state lepton pair, M_{ll} . For each bin in M_{ll} , the y_{ll} distribution is divided in 24 equally spaced bins in rapidity, up to $y_{ll} < 2.4$, except for the last bin in M_{ll} that is divided in 12 bins. In the theoretical calculations, we set μ_F and μ_R to the average value $\langle M_{ll} \rangle$ of each dilepton invariant bin. The cuts on the final state leptons are the following:

$$\begin{aligned} p_T^{\min} &> 9 \text{ GeV} \quad \wedge \quad p_T^{\max} > 14 \text{ GeV}, \\ 20 \text{ GeV} &< M_{ll} < 1500 \text{ GeV}, \\ |\eta_l^{1,2}| &< 2.4, \end{aligned}$$

with p_T^{\min} (p_T^{\max}) being the transverse momentum of the softer (harder) lepton, and $|\eta_l|$ are the rapidities of the two leptons. In Tabs. 18-23 we provide the results of our calculations, and compare the total experimental uncertainty with the corresponding NNLO QCD and the NLO pure weak corrections, for the six bins in invariant mass of this measurement.

The NNLO corrections to this measurement are found to be substantial, specially in the bins with M_{ll} below the Z peak region. For example, the average NNLO correction is around 10% in the first bin, $\sim 6\%$ in the second bin, and $\sim 4\%$ in the Z peak region. These corrections are also found to be relatively constant as a function of the dilepton rapidity $|y_{ll}|$, except for the last bin of the distribution, where the correction is substantially larger. The comparison between data and theory predictions for the first bin of this measurement was shown in Fig 3. In the region with dilepton mass $M_{ll} \leq 45$ GeV, the NNLO QCD corrections are larger than the experimental uncertainties, and therefore their inclusion is essential to achieve a good fit quality.

Concerning the pure weak corrections, 6th column of Tabs. 18-23, we note that it is found to be at most at the few percent level, for example, the average over all rapidity bins is $\sim 3\%$ for the first invariant mass bin and $\sim 4\%$ for the second. Above the Z peak mass we find these weak corrections to be below the 1% range. Note that in the highest invariant mass bin, with $200 \leq M_{ll} \leq 1500$ GeV, the cross-sections are dominated by the region around $M_{ll} \sim 200$ GeV, while the contribution from higher masses, where weak effects are known to be more substantial, has less weight.

20 GeV < M_{ll} < 30 GeV					
$ y_{ll} $	$d\sigma^{\text{exp}}/dy_{ll}/dM_{ll}$ (pb)	$d\sigma^{\text{NLO}}/dy_{ll}/dM_{ll}$ (pb)	Δ_{exp} (%)	Δ_{NNLO} (%)	Δ_{pureEW} (%)
[0.0, 0.1]	17.84	14.83	6.5	8.1	0.4
[0.1, 0.2]	17.68	14.78	6.2	8.9	0.6
[0.2, 0.3]	17.21	14.80	6.4	9.0	0.3
[0.3, 0.4]	17.63	14.78	6.0	8.9	0.4
[0.4, 0.5]	17.84	14.70	5.6	8.8	0.4
[0.5, 0.6]	18.10	14.71	5.2	8.8	0.2
[0.6, 0.7]	18.41	14.59	4.8	8.9	0.4
[0.7, 0.8]	18.16	14.54	4.3	9.2	0.2
[0.8, 0.9]	18.05	14.46	3.8	9.4	0.6
[0.9, 1.0]	17.84	14.36	3.5	9.6	0.5
[1.0, 1.1]	17.52	14.31	3.3	9.7	0.3
[1.1, 1.2]	17.47	14.16	3.3	9.8	0.4
[1.2, 1.3]	16.90	14.07	3.4	9.9	0.4
[1.3, 1.4]	16.95	13.93	4.0	10.1	0.1
[1.4, 1.5]	16.32	13.75	4.5	10.4	0.5
[1.5, 1.6]	16.48	13.60	4.5	10.7	0.1
[1.6, 1.7]	15.48	13.25	4.7	10.7	0.2
[1.7, 1.8]	15.22	12.84	5.2	10.3	0.4
[1.8, 1.9]	14.06	12.25	5.2	9.2	0.1
[1.9, 2.0]	12.59	11.24	5.4	7.6	0.8
[2.0, 2.1]	11.02	9.58	5.2	6.1	0.8
[2.1, 2.2]	8.39	7.28	5.6	6.0	0.6
[2.2, 2.3]	5.46	4.58	7.7	9.5	0.3
[2.3, 2.4]	2.05	1.48	10.3	13.4	0.4
Average			5.2	9.3	0.6

Table 18: Experimental measurements and theoretical predictions for the CMS Drell-Yan double differential distributions in the dilepton invariant mass bin with 20 GeV < M_{ll} < 30 GeV, as a function of the dilepton rapidity $|y_{ll}|$. For each rapidity bin we provide the experimental central value (2nd column), the NLO theoretical predictions obtained with NNPDF2.3NLO with $\alpha_s = 0.118$ (3rd column), the total percentage experimental uncertainty (4th column), the NNLO QCD correction (5th column) and the NLO pure weak correction (6th column). In the last row the average over all the rapidity bins is provided. See text for more details.

30 GeV < M_{ll} < 45 GeV					
$ y_{ll} $	$d\sigma^{\text{exp}}/dy_{ll}/dM_{ll}$ (pb)	$d\sigma^{\text{NLO}}/dy_{ll}/dM_{ll}$ (pb)	Δ_{exp} (%)	Δ_{NNLO} (%)	Δ_{pureEW} (%)
[0.0, 0.1]	26.87	25.36	3.5	6.7	3.1
[0.1, 0.2]	25.66	25.42	3.3	6.4	3.0
[0.2, 0.3]	25.87	25.37	3.0	6.2	2.7
[0.3, 0.4]	25.66	25.45	3.1	6.1	2.6
[0.4, 0.5]	26.13	25.28	2.8	6.1	3.0
[0.5, 0.6]	27.08	25.22	2.7	6.0	2.9
[0.6, 0.7]	26.18	25.29	2.8	5.8	2.7
[0.7, 0.8]	25.66	25.19	3.1	5.6	2.8
[0.8, 0.9]	26.29	25.06	3.0	5.3	2.8
[0.9, 1.0]	25.61	24.95	2.9	5.2	3.0
[1.0, 1.1]	25.71	24.93	2.9	5.2	2.6
[1.1, 1.2]	25.55	24.72	2.7	5.3	2.7
[1.2, 1.3]	24.50	24.33	2.8	5.5	3.6
[1.3, 1.4]	25.76	24.21	2.6	5.8	2.4
[1.4, 1.5]	24.08	23.30	2.6	5.8	3.0
[1.5, 1.6]	23.35	22.19	2.7	5.6	2.9
[1.6, 1.7]	21.67	20.58	2.7	4.9	3.8
[1.7, 1.8]	19.26	18.63	2.7	3.6	2.8
[1.8, 1.9]	16.58	16.13	3.2	1.6	3.1
[1.9, 2.0]	13.80	13.21	3.4	0.9	3.7
[2.0, 2.1]	11.02	10.30	4.3	3.7	3.4
[2.1, 2.2]	7.87	7.41	5.3	5.4	3.7
[2.2, 2.3]	4.67	4.43	6.7	3.2	4.3
[2.3, 2.4]	1.63	1.49	9.6	17.5	6.2
Average			3.5	5.6	3.2

Table 19: Same as Table 18 for the dilepton invariant mass bin with 30 GeV < M_{ll} < 45 GeV.

45 GeV < M_{ll} < 60 GeV					
$ y_{ll} $	$d\sigma^{\text{exp}}/dy_{ll}/dM_{ll}$ (pb)	$d\sigma^{\text{NLO}}/dy_{ll}/dM_{ll}$ (pb)	Δ_{exp} (%)	Δ_{NNLO} (%)	Δ_{pureEW} (%)
[0.0, 0.1]	11.23	10.51	2.8	5.7	2.9
[0.1, 0.2]	11.18	10.51	2.8	4.6	3.3
[0.2, 0.3]	11.07	10.51	3.3	3.9	2.7
[0.3, 0.4]	11.33	10.50	3.2	3.4	3.3
[0.4, 0.5]	10.65	10.46	3.9	3.1	2.8
[0.5, 0.6]	11.70	10.50	2.7	3.0	2.4
[0.6, 0.7]	11.49	10.41	2.7	3.0	3.4
[0.7, 0.8]	11.28	10.39	3.3	3.0	3.0
[0.8, 0.9]	11.28	10.30	2.8	3.0	3.4
[0.9, 1.0]	11.18	10.21	2.8	2.9	2.9
[1.0, 1.1]	10.76	10.02	2.9	2.8	2.7
[1.1, 1.2]	10.28	9.73	3.1	2.6	3.0
[1.2, 1.3]	9.76	9.25	3.2	2.2	3.3
[1.3, 1.4]	9.81	8.69	2.7	1.8	3.5
[1.4, 1.5]	8.92	8.15	2.9	1.2	3.0
[1.5, 1.6]	8.40	7.48	3.8	0.5	3.3
[1.6, 1.7]	7.35	6.79	2.9	0.5	2.6
[1.7, 1.8]	6.40	5.99	3.3	1.6	3.0
[1.8, 1.9]	5.93	5.14	4.4	3.1	3.2
[1.9, 2.0]	4.93	4.28	5.3	4.8	3.0
[2.0, 2.1]	4.04	3.34	6.5	6.9	3.6
[2.1, 2.2]	2.57	2.42	6.1	9.4	2.5
[2.2, 2.3]	1.78	1.44	11.8	12.7	3.7
[2.3, 2.4]	0.63	0.50	16.7	18.1	0.4
Average			4.4	4.3	3.9

Table 20: Same as Table 18 for the dilepton invariant mass bin with 45 GeV < M_{ll} < 60 GeV.

60 GeV < M_{ll} < 120 GeV					
$ y_{ll} $	$d\sigma^{\text{exp}}/dy_{ll}/dM_{ll}$ (pb)	$d\sigma^{\text{NLO}}/dy_{ll}/dM_{ll}$ (pb)	Δ_{exp} (%)	Δ_{NNLO} (%)	Δ_{pureEW} (%)
[0.0, 0.1]	317.44	294.30	1.3	2.2	1.1
[0.1, 0.2]	315.35	291.35	1.3	2.1	2.1
[0.2, 0.3]	315.87	291.02	1.3	2.0	1.6
[0.3, 0.4]	312.72	289.59	1.3	1.9	1.3
[0.4, 0.5]	312.72	286.10	1.3	1.8	1.6
[0.5, 0.6]	309.57	282.63	1.4	1.7	1.8
[0.6, 0.7]	306.95	276.41	1.4	1.5	1.6
[0.7, 0.8]	299.60	270.50	1.4	1.4	2.1
[0.8, 0.9]	292.26	264.98	1.4	1.2	0.9
[0.9, 1.0]	285.96	255.89	1.5	1.0	1.8
[1.0, 1.1]	275.99	245.52	1.5	0.8	2.3
[1.1, 1.2]	264.45	237.54	1.6	0.6	0.9
[1.2, 1.3]	249.76	223.67	1.7	0.3	1.4
[1.3, 1.4]	234.02	208.77	1.6	0.1	2.2
[1.4, 1.5]	219.33	193.84	1.7	0.2	1.5
[1.5, 1.6]	196.76	176.08	1.6	0.5	2.1
[1.6, 1.7]	177.87	156.71	1.8	0.9	2.3
[1.7, 1.8]	153.74	139.25	2.0	1.3	1.2
[1.8, 1.9]	131.70	117.22	2.0	1.8	1.8
[1.9, 2.0]	106.51	96.52	2.5	2.4	1.4
[2.0, 2.1]	82.38	74.07	3.2	3.3	2.5
[2.1, 2.2]	57.19	52.84	3.7	4.8	0.3
[2.2, 2.3]	33.58	31.05	6.3	8.5	1.9
[2.3, 2.4]	10.49	10.40	5.0	44.5	0.1
Average			2.1	3.6	1.6

Table 21: Same as Table 18 for the dilepton invariant mass bin with 60 GeV < M_{ll} < 120 GeV.

120 GeV < M_{ll} < 200 GeV					
$ y_{ll} $	$d\sigma^{\text{exp}}/dy_{ll}/dM_{ll}$ (pb)	$d\sigma^{\text{NLO}}/dy_{ll}/dM_{ll}$ (pb)	Δ_{exp} (%)	Δ_{NNLO} (%)	Δ_{pureEW} (%)
[0.0, 0.1]	3.47	3.24	4.7	1.0	0.8
[0.1, 0.2]	3.61	3.22	4.5	0.5	0.3
[0.2, 0.3]	3.54	3.21	4.8	0.0	0.8
[0.3, 0.4]	3.52	3.17	4.5	0.3	0.0
[0.4, 0.5]	3.38	3.15	4.2	0.5	0.7
[0.5, 0.6]	3.52	3.10	4.2	0.5	0.2
[0.6, 0.7]	3.31	3.05	4.6	0.4	0.4
[0.7, 0.8]	3.32	2.99	4.1	0.2	0.4
[0.8, 0.9]	3.22	2.92	4.6	0.1	0.1
[0.9, 1.0]	3.05	2.83	4.1	0.0	0.4
[1.0, 1.1]	3.03	2.74	4.0	0.1	0.7
[1.1, 1.2]	2.79	2.61	4.1	0.2	0.2
[1.2, 1.3]	2.71	2.49	4.2	0.5	0.4
[1.3, 1.4]	2.52	2.34	4.4	0.8	0.4
[1.4, 1.5]	2.44	2.16	4.5	1.3	0.0
[1.5, 1.6]	2.19	1.98	5.0	1.7	0.5
[1.6, 1.7]	1.95	1.76	5.4	2.2	0.5
[1.7, 1.8]	1.58	1.54	6.3	2.4	0.4
[1.8, 1.9]	1.25	1.29	7.5	2.4	0.6
[1.9, 2.0]	1.16	1.05	8.6	1.8	0.4
[2.0, 2.1]	0.73	0.80	10.8	0.4	0.2
[2.1, 2.2]	0.53	0.56	13.9	2.3	0.5
[2.2, 2.3]	0.28	0.32	18.7	6.8	2.1
[2.3, 2.4]	0.11	0.10	24.8	13.4	0.6
Average			6.9	1.7	0.5

Table 22: Same as Table 18 for the dilepton invariant mass bin with 120 GeV < M_{ll} < 200 GeV.

200 GeV < M_{ll} < 1500 GeV					
$ y_{ll} $	$d\sigma^{\text{exp}}/dy_{ll}/dM_{ll}$ (pb)	$d\sigma^{\text{NLO}}/dy_{ll}/dM_{ll}$ (pb)	Δ_{exp} (%)	Δ_{NNLO} (%)	Δ_{pureEW} (%)
[0.0, 0.2]	0.530	0.533	13.8	1.3	0.0
[0.2, 0.4]	0.609	0.527	10.3	1.6	0.0
[0.4, 0.6]	0.651	0.518	8.0	1.5	0.1
[0.6, 0.8]	0.522	0.501	9.2	2.0	0.0
[0.8, 1.0]	0.534	0.476	7.7	2.5	0.1
[1.0, 1.2]	0.527	0.439	7.4	0.8	0.2
[1.2, 1.4]	0.476	0.388	6.9	3.4	0.1
[1.4, 1.6]	0.323	0.319	9.0	2.6	0.1
[1.6, 1.8]	0.247	0.238	10.5	0.5	0.3
[1.8, 2.0]	0.152	0.154	14.5	2.2	0.0
[2.0, 2.2]	0.097	0.078	20.6	4.4	0.5
[2.2, 2.4]	0.022	0.021	36.4	11.4	0.6
Average			12.9	2.9	0.2

Table 23: Same as Table 18 for the dilepton invariant mass bin with 200 GeV < M_{ll} < 1200 GeV.

A.2 ATLAS high-mass Drell-Yan differential cross section

Now we turn to present the results for the ATLAS high-mass Drell-Yan production data from the 2011 run [56], based on an integrated luminosity of 4.9 fb^{-1} . This measurement is presented in terms of the invariant mass of the electron pairs produced in the range $116 \text{ GeV} < M_{ll} < 1.5 \text{ TeV}$. Theoretical predictions are computed using the same codes and settings described above, we have used as renormalization and factorization scales $\mu_F = \mu_R = M_Z$ and the following cuts to the final states leptons have been applied:

$$p_T^l > 25 \text{ GeV} , \\ |\eta_l^{1,2}| < 2.5 .$$

In Tab. 24, with the same structure as Table 18, we provide the experimental central value, the NLO theoretical predictions obtained with NNPDF2.3NLO with $\alpha_s = 0.118$, the total percentage experimental uncertainty, the NNLO QCD correction, and the NLO pure weak correction. In the last row, the average over all the data points in the measurement is given.

As expected, the size of the pure weak corrections increases monotonically with the value of M_{ll} , reaching up to $\sim 7\%$ in the highest mass bin. While the uncertainties in experimental data are still rather higher due to the limited statistics, is clear that for 8 TeV data and even more for Run II measurements, including electroweak corrections in Pdf fits will be mandatory. On the other hand, QCD NNLO corrections are found to be at the few percent level, typically smaller than the total experimental uncertainty.

ATLAS 2011 DY invariant mass distribution					
M_{ll}	$d\sigma^{\text{exp}}/dM_{ll} \text{ (pb)}$	$d\sigma^{\text{NLO}}/dM_{ll} \text{ (pb)}$	$\Delta_{\text{exp}} \text{ (%)}$	$\Delta_{\text{NNLO}} \text{ (%)}$	$\Delta_{\text{pureEW}} \text{ (%)}$
[116, 130]	224.00	212.78	4.3	0.9	0.7
[130, 150]	102.00	91.60	4.5	1.3	0.4
[150, 170]	51.20	45.15	5.0	1.5	0.1
[170, 190]	28.40	25.55	5.4	2.8	0.1
[190, 210]	18.70	15.76	6.1	3.4	0.2
[210, 230]	10.70	10.28	7.5	1.6	0.1
[230, 250]	8.23	7.22	7.9	0.4	0.2
[250, 300]	4.66	3.93	7.2	0.0	0.0
[300, 400]	1.70	1.40	7.8	3.4	0.3
[400, 500]	0.47	0.45	11.3	4.7	1.0
[500, 700]	0.15	0.12	12.4	5.0	1.5
[700, 1000]	0.0221	0.0196	25.1	9.0	3.4
[1000, 1500]	0.0029	0.0022	51.0	1.0	6.8
Average			12.0	2.7	1.2

Table 24: Same as Table 18 for the ATLAS measurement of the high-mass Drell-Yan invariant dilepton mass distribution from the 2011 dataset.

A.3 ATLAS Z rapidity distribution

Now we consider the ATLAS measurements of the Z rapidity distribution from an integrated luminosity of 36 pb^{-1} [47]. This dataset was already included in our previous NNPDF2.3 analysis [13]. In the theoretical calculations, we have used the following cuts for the final-state

lepton kinematics:

$$\begin{aligned}
p_T^l &\geq 20 \text{ GeV}, \\
66 \text{ GeV} &\leq M_{ll} \leq 116 \text{ GeV}, \\
|\eta_l^{1,2}| &\leq 4.9.
\end{aligned}$$

In Tab. 25, with again the same structure as Table 18, we provide for this measurement the experimental central value, the NLO theoretical predictions obtained with NNPDF2.3NLO with $\alpha_s = 0.118$, the total percentage experimental uncertainty, the NNLO QCD correction, and the NLO pure weak correction.

From Tab. 25 we see that in the Z peak region, both the NNLO QCD and NLO weak corrections are small, of the order 1% at most, and slowly varying with the lepton rapidity. On the other hand, the precision of the data is high and therefore it is necessary to include these corrections in the theory calculation, as has been done in this work. This is specially true in the central region, with $|y_{ll}| \lesssim 1.6$.

ATLAS Z rapidity distribution					
$ y_{ll} $	$d\sigma^{\text{exp}}/dy_{ll}$ (pb)	$d\sigma^{\text{NLO}}/dy_{ll}$ (pb)	Δ_{exp} (%)	Δ_{NNLO} (%)	Δ_{pureEW} (%)
[0.0, 0.4]	129.27	123.44	1.9	0.7	0.9
[0.4, 0.8]	129.44	122.22	1.9	0.9	0.7
[0.8, 1.2]	125.81	119.98	1.8	0.8	0.7
[1.2, 1.6]	118.23	116.58	1.9	0.7	0.8
[1.6, 2.0]	113.37	112.07	2.3	0.6	0.6
[2.0, 2.4]	105.26	105.32	3.7	0.6	0.9
[2.4, 2.8]	92.18	92.63	6.3	0.9	0.7
[2.8, 3.6]	53.38	54.93	10.1	2.2	0.7
Average			3.7	0.9	0.8

Table 25: Same as Table 18 for the ATLAS measurement of the rapidity distributions of Z bosons from the 2010 dataset.

A.4 LHCb Z forward rapidity distribution

Finally, we present here the theoretical predictions for the LHCb $Z \rightarrow ee$ rapidity distributions in the forward region from the 2011 dataset [61]. In the calculation, the following cuts on the lepton kinematics have been used,

$$\begin{aligned}
p_T^l &\geq 20 \text{ GeV}, \\
60 \text{ GeV} &\leq M_{ll} \leq 120 \text{ GeV}, \\
2 &\leq \eta_l^{1,2} \leq 4.5.
\end{aligned}$$

Results are shown in Tab. 26, with the usual structure. As in the previous case of the ATLAS measurement of the Z rapidity distribution, and as we expect for calculations in the Z peak region, the NLO pure weak corrections are found to be quite small, of order 1% at most, and slowly varying with the lepton rapidity. On the the hand the NNLO QCD corrections are more important, and interestingly they increase monotonically with the rapidity of the dilepton system, reaching up to $\sim 4\%$ in the most forward bin. The experimental uncertainties for this measurement in any case are larger than the NNLO QCD corrections.

LHCb $Z \rightarrow e^+e^-$ forward rapidity distribution					
y_{ll}	$d\sigma^{\text{exp}}/dy_{ll}$ (pb)	$d\sigma^{\text{NLO}}/dy_{ll}$ (pb)	Δ_{exp} (%)	Δ_{NNLO} (%)	Δ_{pureEW} (%)
[2.00, 2.25]	13.6	13.2	6.6	1.2	1.2
[2.25, 2.50]	39.4	36.7	3.6	0.8	0.4
[2.50, 2.75]	56.7	51.8	3.4	0.2	0.6
[2.75, 3.00]	63.2	61.0	3.3	0.5	0.7
[3.00, 3.25]	59.9	56.1	3.8	1.2	0.6
[3.25, 3.50]	43.8	38.0	4.3	1.9	0.3
[3.50, 3.75]	20.5	17.3	6.8	2.6	0.6
[3.75, 4.00]	5.9	4.6	16.9	3.0	0.1
[4.00, 4.25]	0.66	0.5	80.3	4.0	0.1
Average			14.3	1.7	0.5

Table 26: Same as Table 18 for the LHCb measurement of the rapidity distributions of Z bosons in the forward region from the 2011 dataset.

B Distance estimators

In previous publications we have extensively used, for comparisons between PDFs, the distance between two NNPDF fits, each represented by a sample of Monte Carlo replicas. This distance estimator was first introduced in Ref. [10], and can be used both to assess the compatibility between two PDFs sets, and to test whether two PDF sets are statistically equivalent. In the course of the present work we have updated the definition of this distance estimator, as we discuss the details of this new definition below. We also explain how the distance estimator can be suitably modified to be used in the validation of the closure test fits.

Given a Monte Carlo sample of N_{rep} replicas representing the probability distribution of a given PDF set, $\{f^{(k)}\}$, the expectation value of the distribution as a function of x and Q^2 is given by

$$\bar{f}(x, Q^2) \equiv \langle f(x, Q^2) \rangle_{\text{rep}} = \frac{1}{N_{\text{rep}}} \sum_i^{N_{\text{rep}}} f^{(k)}(x, Q^2), \quad (43)$$

where the index (k) runs over all the replicas in the sample. The variance of the sample is estimated as

$$\sigma^2[f(x, Q^2)] = \frac{1}{N_{\text{rep}} - 1} \sum_i^{N_{\text{rep}}} \left(f^{(k)}(x, Q^2) - \langle f(x, Q^2) \rangle_{\text{rep}} \right)^2. \quad (44)$$

The variance of the mean is, in turn, defined in terms of the variance of the sample by

$$\sigma^2[\langle f(x, Q^2) \rangle_{\text{rep}}] = \frac{1}{N_{\text{rep}}} \sigma^2[f(x, Q^2)]. \quad (45)$$

The variance of the variance itself can be computed using

$$\sigma^2[\sigma^2[f(x, Q^2)]] = \frac{1}{N_{\text{rep}}} \left[m_4[f(x, Q^2)] - \frac{N_{\text{rep}} - 3}{N_{\text{rep}} - 1} (\sigma^2[f(x, Q^2)])^2 \right], \quad (46)$$

where $m_4[f(x, Q^2)]$ denotes the fourth moment of the probability distribution for $f(x, Q^2)$, namely

$$m_4[f(x, Q^2)] = \frac{1}{N_{\text{rep}}} \left[\sum_{k=1}^{N_{\text{rep}}} \left(f^{(k)}(x, Q^2) - \langle f(x, Q^2) \rangle_{\text{rep}} \right)^4 \right]. \quad (47)$$

Given the definitions above, the distance between two sets of PDFs, each characterized by a given distribution of the Monte Carlo replicas, denoted by $\{f^{(k)}\}$ and $\{g^{(k)}\}$, can be defined as the square root of the square difference of the PDF central values in units of the uncertainty of the mean, that is

$$d_{\bar{f}, \bar{g}}(x, Q^2) = \sqrt{\frac{(\bar{f} - \bar{g})^2}{\sigma^2[\bar{f}] + \sigma^2[\bar{g}]}}. \quad (48)$$

In Eq. (48), the denominator uses the variance of the mean of the distribution, defined as in Eq. (45). An analogous distance can be defined for the variances of the two samples:

$$d_{\sigma[f], \sigma[g]}(x, Q^2) = \sqrt{\frac{(\sigma^2[f] - \sigma^2[g])^2}{\sigma^2[\sigma^2[f]] + \sigma^2[\sigma^2[g]]}}. \quad (49)$$

where now in the denominator we have the variance of the variance, Eq. (46).

The distances for the central values and for the variances defined in Eqs. (48) and (49) test whether the underlying distributions from which the two Monte Carlo samples $\{f^{(k)}\}$ and $\{g^{(k)}\}$ are drawn have respectively the same mean and the same standard deviation. In particular, it is possible to show that one expects these distances to fluctuate around $d \sim 1$ if the two samples do indeed come from the same distribution. On the other hand, values of the distances around $d \sim \sqrt{N_{\text{rep}}}$ indicates that the central values (the variances) of the two PDF sets differ by one standard deviation in units of the variance of the distribution Eq. (44) (in units of the variance of the variance Eq. (46)).

In this paper, when producing the distances between two NNPDF sets, Eqs. (48) and (49), we always compare sets of $N_{\text{rep}} = 100$ replicas, and therefore a value of the distance around $d \sim 10$ indicates that two sets differ by one standard deviation (this applies both to central values and to variances). The PDFs are sampled at the scale $Q_0^2 = 2 \text{ GeV}^2$ for 50 points in x equally spaced in logarithmic scale in the interval $[10^{-5}, 0.1]$ and then 50 more points linearly spaced over the interval $[0.1, 0.95]$.

The main difference between the new definition of the distance estimators and the previous one introduced in Ref. [10], and used in previous NNPDF publications, is that we have now removed the additional bootstrap sampling of the distances distributions. For this reason, with the new definition, distances can become arbitrarily small, for instance when the central values of the two distributions coincide. Other than this, the interpretation of the distance plots is very similar both with the old and new definitions.

The new definition for the distance introduced above is used in this paper whenever we are comparing two NNPDF fits. In the context of closure tests described in Sect. 4, it is useful to use a slightly different definition. The motivation is that we want to compare the closure test fitted PDFs with the input PDF, and therefore we need to remove the scaling factor $1/\sqrt{N_{\text{rep}}}$, which is only required when comparing two Monte Carlo sets. In addition, in the denominator we should include only the variance of the fitted PDFs, since for the input PDF only the central value is used in the definition of the pseudo data. Therefore, when comparing closure test fitted PDFs with the corresponding input PDF in Sect. 4, we define the distance of the fitted PDF set, f_{fit} , with respect to the initial PDF set used for generating the pseudodata, f_{in} , as follows

$$d_{\text{ct}} [f_{i,\text{fit}}, f_{i,\text{in}}] (x, Q) \equiv \sqrt{\frac{(\bar{f}_{i,\text{fit}}(x, Q) - f_{i,\text{in}}(x, Q))^2}{\sigma^2 [f_{i,\text{fit}}] (x, Q)}}. \quad (50)$$

With this definition, the distance Eq. (50) between the closure test PDF and the underlying theory is measured in units of the standard deviation of the fitted PDFs. In this work, the distances between closure test fits will be always calculated at the initial parametrisation scale of $Q^2 = 1 \text{ GeV}^2$.

References

- [1] M. Diemoz, F. Ferroni, E. Longo, and G. Martinelli, *Parton Densities from Deep Inelastic Scattering to Hadronic Processes at Super Collider Energies*, *Z. Phys.* **C39** (1988) 21.
- [2] L. Demortier, *Proceedings, PHYSTAT 2011 Workshop on Statistical Issues Related to Discovery Claims in Search Experiments and Unfolding, CERN, Geneva, Switzerland 17-20 January 2011*, ch. Open Issues in the Wake of Banff 2011. 2011.
- [3] S. Forte, L. Garrido, J. I. Latorre, and A. Piccione, *Neural network parametrization of deep-inelastic structure functions*, *JHEP* **05** (2002) 062, [[hep-ph/0204232](#)].
- [4] **The NNPDF** Collaboration, L. Del Debbio, S. Forte, J. I. Latorre, A. Piccione, and J. Rojo, *Unbiased determination of the proton structure function $f_2(p)$ with estimation*, *JHEP* **03** (2005) 080, [[hep-ph/0501067](#)].
- [5] **The NNPDF** Collaboration, L. Del Debbio, S. Forte, J. I. Latorre, A. Piccione, and J. Rojo, *Neural network determination of parton distributions: The nonsinglet case*, *JHEP* **03** (2007) 039, [[hep-ph/0701127](#)].
- [6] **The NNPDF** Collaboration, R. D. Ball et al., *A determination of parton distributions with faithful uncertainty estimation*, *Nucl. Phys.* **B809** (2009) 1–63, [[arXiv:0808.1231](#)].
- [7] **The NNPDF** Collaboration, J. Rojo et al., *Update on Neural Network Parton Distributions: NNPDF1.1*, [arXiv:0811.2288](#).
- [8] **The NNPDF** Collaboration, R. D. Ball et al., *Precision determination of electroweak parameters and the strange content of the proton from neutrino deep-inelastic scattering*, *Nucl. Phys.* **B823** (2009) 195–233, [[arXiv:0906.1958](#)].
- [9] **The NNPDF** Collaboration, R. D. Ball et al., *Fitting Parton Distribution Data with Multiplicative Normalization Uncertainties*, *JHEP* **05** (2010) 075, [[arXiv:0912.2276](#)].
- [10] **The NNPDF** Collaboration, R. D. Ball et al., *A first unbiased global NLO determination of parton distributions and their uncertainties*, *Nucl. Phys.* **B838** (2010) 136–206, [[arXiv:1002.4407](#)].
- [11] **The NNPDF** Collaboration, R. D. Ball et al., *Impact of Heavy Quark Masses on Parton Distributions and LHC Phenomenology*, *Nucl. Phys.* **B849** (2011) 296–363, [[arXiv:1101.1300](#)].
- [12] **The NNPDF** Collaboration, R. D. Ball et al., *Unbiased global determination of parton distributions and their uncertainties at NNLO and at LO*, *Nucl. Phys.* **B855** (2012) 153–221, [[arXiv:1107.2652](#)].
- [13] **NNPDF** Collaboration, R. D. Ball, V. Bertone, S. Carrazza, C. S. Deans, L. Del Debbio, et al., *Parton distributions with LHC data*, *Nucl. Phys.* **B867** (2013) 244–289, [[arXiv:1207.1303](#)].
- [14] G. Watt and R. Thorne, *Study of Monte Carlo approach to experimental uncertainty propagation with MSTW 2008 PDFs*, *JHEP* **1208** (2012) 052, [[arXiv:1205.4024](#)].

- [15] G. Cowan, *Statistical data analysis*. Oxford University Press, 2002.
- [16] **H1 and ZEUS** Collaboration, V. Radescu, *Combination and QCD analysis of the HERA inclusive cross sections*, *PoS ICHEP2010* (2010) 168.
- [17] **H1 and ZEUS** Collaboration, A. Cooper-Sarkar, *PDF Fits at HERA*, *PoS EPS-HEP2011* (2011) 320, [[arXiv:1112.2107](#)].
- [18] **The NNPDF** Collaboration, R. D. Ball et al., *Reweighting NNPDFs: the W lepton asymmetry*, *Nucl. Phys.* **B849** (2011) 112–143, [[arXiv:1012.0836](#)].
- [19] **New Muon** Collaboration, M. Arneodo et al., *Accurate measurement of F_2^d/F_2^p and $R_d - R_p$* , *Nucl. Phys.* **B487** (1997) 3–26, [[hep-ex/9611022](#)].
- [20] **New Muon** Collaboration, M. Arneodo et al., *Measurement of the proton and deuteron structure functions, F_2^p and F_2^d , and of the ratio σ_L/σ_T* , *Nucl. Phys.* **B483** (1997) 3–43, [[hep-ph/9610231](#)].
- [21] **BCDMS** Collaboration, A. C. Benvenuti et al., *A high statistics measurement of the proton structure functions $f_2(x, q^2)$ and r from deep inelastic muon scattering at high q^2* , *Phys. Lett.* **B223** (1989) 485.
- [22] **BCDMS** Collaboration, A. C. Benvenuti et al., *A high statistics measurement of the deuteron structure functions $f_2(x, q^2)$ and r from deep inelastic muon scattering at high q^2* , *Phys. Lett.* **B237** (1990) 592.
- [23] L. W. Whitlow, E. M. Riordan, S. Dasu, S. Rock, and A. Bodek, *Precise measurements of the proton and deuteron structure functions from a global analysis of the SLAC deep inelastic electron scattering cross-sections*, *Phys. Lett.* **B282** (1992) 475–482.
- [24] **H1 and ZEUS** Collaboration, F. Aaron et al., *Combined Measurement and QCD Analysis of the Inclusive $e^\pm p$ Scattering Cross Sections at HERA*, *JHEP* **1001** (2010) 109, [[arXiv:0911.0884](#)].
- [25] **H1** Collaboration, F. D. Aaron et al., *Measurement of the Proton Structure Function F_L at Low x* , *Phys. Lett.* **B665** (2008) 139–146, [[arXiv:0805.2809](#)].
- [26] **ZEUS** Collaboration, J. Breitweg et al., *Measurement of $D^{*\pm}$ production and the charm contribution to F_2 in deep inelastic scattering at HERA*, *Eur. Phys. J.* **C12** (2000) 35–52, [[hep-ex/9908012](#)].
- [27] **ZEUS** Collaboration, S. Chekanov et al., *Measurement of $D^{*\pm}$ production in deep inelastic $e^\pm p$ scattering at HERA*, *Phys. Rev.* **D69** (2004) 012004, [[hep-ex/0308068](#)].
- [28] **ZEUS** Collaboration, S. Chekanov et al., *Measurement of D^\pm and D^0 production in deep inelastic scattering using a lifetime tag at HERA*, *Eur. Phys. J.* **C63** (2009) 171–188, [[arXiv:0812.3775](#)].
- [29] **ZEUS** Collaboration, S. Chekanov et al., *Measurement of charm and beauty production in deep inelastic ep scattering from decays into muons at HERA*, *Eur. Phys. J.* **C65** (2010) 65–79, [[arXiv:0904.3487](#)].
- [30] **H1** Collaboration, C. Adloff et al., *Measurement of $D^{*\pm}$ meson production and F_2^c in deep inelastic scattering at HERA*, *Phys. Lett.* **B528** (2002) 199–214, [[hep-ex/0108039](#)].

- [31] **H1** Collaboration, F. D. Aaron et al., *Measurement of the D^* Meson Production Cross Section and F_2^c , at High Q^2 , in ep Scattering at HERA*, *Phys. Lett.* **B686** (2010) 91–100, [[arXiv:0911.3989](#)].
- [32] **H1** Collaboration, F. D. Aaron et al., *Measurement of the Charm and Beauty Structure Functions using the H1 Vertex Detector at HERA*, *Eur. Phys. J.* **C65** (2010) 89–109, [[arXiv:0907.2643](#)].
- [33] **ZEUS** Collaboration, S. Chekanov et al., *Measurement of high- Q^2 neutral current deep inelastic e^-p scattering cross sections with a longitudinally polarised electron beam at HERA*, *Eur. Phys. J.* **C62** (2009) 625–658, [[arXiv:0901.2385](#)].
- [34] **ZEUS** Collaboration, S. Chekanov et al., *Measurement of charged current deep inelastic scattering cross sections with a longitudinally polarised electron beam at HERA*, *Eur. Phys. J.* **C61** (2009) 223–235, [[arXiv:0812.4620](#)].
- [35] **CHORUS** Collaboration, G. Onengut et al., *Measurement of nucleon structure functions in neutrino scattering*, *Phys. Lett.* **B632** (2006) 65–75.
- [36] **NuTeV** Collaboration, M. Goncharov et al., *Precise measurement of dimuon production cross-sections in $\nu_\mu Fe$ and $\bar{\nu}_\mu Fe$ deep inelastic scattering at the Tevatron*, *Phys. Rev.* **D64** (2001) 112006, [[hep-ex/0102049](#)].
- [37] D. A. Mason, *Measurement of the strange - antistrange asymmetry at NLO in QCD from NuTeV dimuon data*, . FERMILAB-THESIS-2006-01.
- [38] G. Moreno et al., *Dimuon production in proton - copper collisions at $\sqrt{s} = 38.8\text{-GeV}$* , *Phys. Rev.* **D43** (1991) 2815–2836.
- [39] **NuSea** Collaboration, J. C. Webb et al., *Absolute Drell-Yan dimuon cross sections in 800-GeV/c $p p$ and $p d$ collisions*, [hep-ex/0302019](#).
- [40] J. C. Webb, *Measurement of continuum dimuon production in 800-GeV/c proton nucleon collisions*, [hep-ex/0301031](#).
- [41] **FNAL E866/NuSea** Collaboration, R. S. Towell et al., *Improved measurement of the anti-d/anti-u asymmetry in the nucleon sea*, *Phys. Rev.* **D64** (2001) 052002, [[hep-ex/0103030](#)].
- [42] **CDF** Collaboration, T. Aaltonen et al., *Direct Measurement of the W Production Charge Asymmetry in $p\bar{p}$ Collisions at $\sqrt{s} = 1.96\text{ TeV}$* , *Phys. Rev. Lett.* **102** (2009) 181801, [[arXiv:0901.2169](#)].
- [43] **CDF** Collaboration, T. A. Aaltonen et al., *Measurement of $d\sigma/dy$ of Drell-Yan e^+e^- pairs in the Z Mass Region from $p\bar{p}$ Collisions at $\sqrt{s} = 1.96\text{ TeV}$* , *Phys. Lett.* **B692** (2010) 232–239, [[arXiv:0908.3914](#)].
- [44] **D0** Collaboration, V. M. Abazov et al., *Measurement of the shape of the boson rapidity distribution for $p\bar{p} \rightarrow Z/\gamma^* \rightarrow e^+e^- + X$ events produced at $\sqrt{s}=1.96\text{-TeV}$* , *Phys. Rev.* **D76** (2007) 012003, [[hep-ex/0702025](#)].
- [45] **CDF** Collaboration, T. Aaltonen et al., *Measurement of the Inclusive Jet Cross Section at the Fermilab Tevatron p - p bar Collider Using a Cone-Based Jet Algorithm*, *Phys. Rev.* **D78** (2008) 052006, [[arXiv:0807.2204](#)].

- [46] **D0** Collaboration, V. M. Abazov et al., *Measurement of the inclusive jet cross-section in $p\bar{p}$ collisions at $\sqrt{s}=1.96\text{-TeV}$* , *Phys. Rev. Lett.* **101** (2008) 062001, [[arXiv:0802.2400](#)].
- [47] **ATLAS** Collaboration, G. Aad et al., *Measurement of the inclusive W^\pm and Z/γ^* cross sections in the electron and muon decay channels in pp collisions at $\sqrt{s}=7\text{ TeV}$ with the ATLAS detector*, *Phys.Rev.* **D85** (2012) 072004, [[arXiv:1109.5141](#)].
- [48] **CMS** Collaboration, S. Chatrchyan et al., *Measurement of the electron charge asymmetry in inclusive W production in pp collisions at $\sqrt{s}=7\text{ TeV}$* , *Phys.Rev.Lett.* **109** (2012) 111806, [[arXiv:1206.2598](#)].
- [49] **LHCb Collaboration** Collaboration, R. Aaij et al., *Inclusive W and Z production in the forward region at $\sqrt{s}=7\text{ TeV}$* , *JHEP* **1206** (2012) 058, [[arXiv:1204.1620](#)].
- [50] **ATLAS** Collaboration, G. Aad et al., *Measurement of inclusive jet and dijet production in pp collisions at $\sqrt{s}=7\text{ TeV}$ using the ATLAS detector*, *Phys. Rev.* **D86** (2012) 014022, [[arXiv:1112.6297](#)].
- [51] **H1** Collaboration, F. Aaron et al., *Inclusive Deep Inelastic Scattering at High Q^2 with Longitudinally Polarised Lepton Beams at HERA*, *JHEP* **1209** (2012) 061, [[arXiv:1206.7007](#)].
- [52] **H1** Collaboration, F. Aaron et al., *Measurement of the Inclusive $e^\pm p$ Scattering Cross Section at High Inelasticity y and of the Structure Function F_L* , *Eur.Phys.J.* **C71** (2011) 1579, [[arXiv:1012.4355](#)].
- [53] **ZEUS** Collaboration, A. Cooper Sarkar, *Measurement of high- Q^2 neutral current deep inelastic $e+p$ scattering cross sections with a longitudinally polarised positron beam at HERA*, [arXiv:1208.6138](#).
- [54] **ZEUS** Collaboration, H. Abramowicz et al., *Measurement of high- Q^2 charged current deep inelastic scattering cross sections with a longitudinally polarised positron beam at HERA*, *Eur.Phys.J.* **C70** (2010) 945–963, [[arXiv:1008.3493](#)].
- [55] **H1 , ZEUS** Collaboration, H. Abramowicz et al., *Combination and QCD Analysis of Charm Production Cross Section Measurements in Deep-Inelastic ep Scattering at HERA*, *Eur.Phys.J.* **C73** (2013) 2311, [[arXiv:1211.1182](#)].
- [56] **ATLAS** Collaboration, G. Aad et al., *Measurement of the high-mass Drell–Yan differential cross-section in pp collisions at $\sqrt{s}=7\text{ TeV}$ with the ATLAS detector*, *Phys.Lett.* **B725** (2013) 223–242, [[arXiv:1305.4192](#)].
- [57] **ATLAS** Collaboration, G. Aad et al., *Measurement of the Transverse Momentum Distribution of W Bosons in pp Collisions at $\sqrt{s}=7\text{ TeV}$ with the ATLAS Detector*, *Phys.Rev.* **D85** (2012) 012005, [[arXiv:1108.6308](#)].
- [58] **CMS** Collaboration, S. Chatrchyan et al., *Measurement of the muon charge asymmetry in inclusive pp to WX production at $\sqrt{s}=7\text{ TeV}$ and an improved determination of light parton distribution functions*, *Phys.Rev.* **D90** (2014) 032004, [[arXiv:1312.6283](#)].
- [59] **CMS Collaboration** Collaboration, S. Chatrchyan et al., *Measurement of the differential and double-differential Drell–Yan cross sections in proton-proton collisions at $\sqrt{s}=7\text{ TeV}$* , *JHEP* **1312** (2013) 030, [[arXiv:1310.7291](#)].

- [60] **CMS** Collaboration, S. Chatrchyan et al., *Measurement of associated $W + \text{charm}$ production in pp collisions at $\sqrt{s} = 7 \text{ TeV}$* , [arXiv:1310.1138](#).
- [61] **LHCb** Collaboration, R. Aaij et al., *Measurement of the cross-section for $Z \rightarrow e^+e^-$ production in pp collisions at $\sqrt{s} = 7 \text{ TeV}$* , *JHEP* **1302** (2013) 106, [[arXiv:1212.4620](#)].
- [62] **CMS** Collaboration, S. Chatrchyan et al., *Measurements of differential jet cross sections in proton-proton collisions at $\sqrt{s} = 7 \text{ TeV}$ with the CMS detector*, *Phys.Rev.* **D87** (2013) 112002, [[arXiv:1212.6660](#)].
- [63] **ATLAS** Collaboration, G. Aad et al., *Measurement of the inclusive jet cross section in pp collisions at $\sqrt{s}=2.76 \text{ TeV}$ and comparison to the inclusive jet cross section at $\sqrt{s}=7 \text{ TeV}$ using the ATLAS detector*, *Eur.Phys.J.* **C73** (2013) 2509, [[arXiv:1304.4739](#)].
- [64] G. P. Salam and G. Soyez, *A practical Seedless Infrared-Safe Cone jet algorithm*, *JHEP* **05** (2007) 086, [[arXiv:0704.0292](#)].
- [65] **ATLAS** Collaboration, G. Aad et al., *Measurement of the cross section for top-quark pair production in pp collisions at $\sqrt{s} = 7 \text{ TeV}$ with the ATLAS detector using final states with two high-pt leptons*, *JHEP* **1205** (2012) 059, [[arXiv:1202.4892](#)].
- [66] **ATLAS** Collaboration, G. Aad et al., *Measurement of the $t\bar{t}$ production cross-section in pp collisions at $\sqrt{s} = 7 \text{ TeV}$ using kinematic information of lepton+jets events*, *ATLAS-CONF-2011-121*, *ATLAS-COM-CONF-2011-132*.
- [67] **ATLAS** Collaboration, G. Aad et al., *Measurement of the $t\bar{t}$ production cross-section in pp collisions at $\sqrt{s} = 8 \text{ TeV}$ using $e\mu$ events with b -tagged jets*, *ATLAS-CONF-2013-097*, *ATLAS-COM-CONF-2013-112*.
- [68] **CMS** Collaboration, S. Chatrchyan et al., *Measurement of the $t\bar{t}$ production cross section in the dilepton channel in pp collisions at $\sqrt{s} = 8 \text{ TeV}$* , *JHEP* **1402** (2014) 024, [[arXiv:1312.7582](#)].
- [69] **CMS** Collaboration, S. Chatrchyan et al., *Measurement of the $t\bar{t}$ production cross section in the dilepton channel in pp collisions at $\sqrt{s} = 7 \text{ TeV}$* , *JHEP* **1211** (2012) 067, [[arXiv:1208.2671](#)].
- [70] **CMS** Collaboration, S. Chatrchyan et al., *Measurement of the $t\bar{t}$ production cross section in pp collisions at $\sqrt{s} = 7 \text{ TeV}$ with lepton + jets final states*, *Phys.Lett.* **B720** (2013) 83–104, [[arXiv:1212.6682](#)].
- [71] S. Forte, *Parton distributions at the dawn of the LHC*, *Acta Phys.Polon.* **B41** (2010) 2859–2920, [[arXiv:1011.5247](#)].
- [72] S. Forte and G. Watt, *Progress in the Determination of the Partonic Structure of the Proton*, *Ann.Rev.Nucl.Part.Sci.* **63** (2013) 291–328, [[arXiv:1301.6754](#)].
- [73] M. L. Mangano and J. Rojo, *Cross Section Ratios between different CM energies at the LHC: opportunities for precision measurements and BSM sensitivity*, *JHEP* **1208** (2012) 010, [[arXiv:1206.3557](#)].
- [74] W. Stirling and E. Vryonidou, *Charm production in association with an electroweak gauge boson at the LHC*, *Phys.Rev.Lett.* **109** (2012) 082002, [[arXiv:1203.6781](#)].

- [75] S. Alekhin, J. Bluemlein, L. Caminadac, K. Lipka, K. Lohwasser, et al., *Determination of Strange Sea Quark Distributions from Fixed-target and Collider Data*, [arXiv:1404.6469](#).
- [76] M. Czakon, M. L. Mangano, A. Mitov, and J. Rojo, *Constraints on the gluon PDF from top quark pair production at hadron colliders*, *JHEP* **1307** (2013) 167, [[arXiv:1303.7215](#)].
- [77] M. Beneke, P. Falgari, S. Klein, J. Piclum, C. Schwinn, et al., *Inclusive Top-Pair Production Phenomenology with TOPIXS*, *JHEP* **1207** (2012) 194, [[arXiv:1206.2454](#)].
- [78] R. Gauld, *Feasibility of top quark measurements at LHCb and constraints on the large- x gluon PDF*, *JHEP* **1402** (2014) 126, [[arXiv:1311.1810](#)].
- [79] S. A. Malik and G. Watt, *Ratios of W and Z cross sections at large boson p_T as a constraint on PDFs and background to new physics*, *JHEP* **1402** (2014) 025, [[arXiv:1304.2424](#)].
- [80] M. Czakon and A. Mitov, *NNLO corrections to top pair production at hadron colliders: the quark-gluon reaction*, *JHEP* **1301** (2013) 080, [[arXiv:1210.6832](#)].
- [81] M. Czakon, P. Fiedler, and A. Mitov, *The total top quark pair production cross-section at hadron colliders through $O(\alpha_S^4)$* , *Phys.Rev.Lett.* **110** (2013) 252004, [[arXiv:1303.6254](#)].
- [82] A. Gehrmann-De Ridder, T. Gehrmann, E. Glover, and J. Pires, *Second order QCD corrections to jet production at hadron colliders: the all-gluon contribution*, *Phys.Rev.Lett.* **110** (2013) 162003, [[arXiv:1301.7310](#)].
- [83] **CDF - Run II** Collaboration, A. Abulencia et al., *Measurement of the Inclusive Jet Cross Section using the k_T algorithm in $p\bar{p}$ Collisions at $\sqrt{s}=1.96$ TeV with the CDF II Detector*, *Phys. Rev.* **D75** (2007) 092006, [[hep-ex/0701051](#)].
- [84] **H1 Collaboration** Collaboration, V. Andreev et al., *Measurement of inclusive ep cross sections at high Q^2 at $\sqrt{s} = 225$ and 252 GeV and of the longitudinal proton structure function F_L at HERA*, *Eur.Phys.J.* **C74** (2014) 2814, [[arXiv:1312.4821](#)].
- [85] **ZEUS Collaboration** Collaboration, S. Chekanov et al., *Measurement of the Longitudinal Proton Structure Function at HERA*, *Phys.Lett.* **B682** (2009) 8–22, [[arXiv:0904.1092](#)].
- [86] **ZEUS** Collaboration, H. Abramowicz, *Measurement of high- Q^2 neutral current deep inelastic e^+p scattering cross sections with a longitudinally polarised positron beam at HERA*, [arXiv:1208.6138](#).
- [87] V. Bertone and J. Rojo, *Parton Distributions with the Combined HERA Charm Production Cross Sections*, *AIP Conf.Proc.* **1523** (2012) 51–54, [[arXiv:1212.0741](#)].
- [88] S. Alekhin, J. Blmlein, K. Daum, K. Lipka, and S. Moch, *Precise charm-quark mass from deep-inelastic scattering*, *Phys.Lett.* **B720** (2013) 172–176, [[arXiv:1212.2355](#)].
- [89] J. Gao, M. Guzzi, and P. M. Nadolsky, *Charm quark mass dependence in a global QCD analysis*, *Eur.Phys.J.* **C73** (2013) 2541, [[arXiv:1304.3494](#)].
- [90] **NNPDF** Collaboration, R. D. Ball et al., *Parton distributions with QED corrections*, *Nucl.Phys.* **B877** (2013) 290–320, [[arXiv:1308.0598](#)].

- [91] **ATLAS** Collaboration, G. Aad et al., *Measurement of the transverse momentum distribution of Z/γ^* bosons in proton-proton collisions at $\sqrt{s} = 7$ TeV with the ATLAS detector*, *Phys.Lett.* **B705** (2011) 415–434, [[arXiv:1107.2381](#)].
- [92] **ATLAS** Collaboration, G. Aad et al., *Measurement of the production of a W boson in association with a charm quark in pp collisions at $\sqrt{s} = 7$ TeV with the ATLAS detector*, *JHEP* **1405** (2014) 068, [[arXiv:1402.6263](#)].
- [93] **LHCb** Collaboration, *Measurement of the cross-section for $Z \rightarrow \mu\mu$ production with 1 fb^{-1} of pp collisions at $\sqrt{s}=7$ TeV*, LHCb-CONF-2013-007.
- [94] **LHCb** Collaboration, *Inclusive low mass Drell-Yan production in the forward region at $\sqrt{s} = 7$ TeV*, LHCb-ANA-2012-029.
- [95] **CMS** Collaboration, S. Chatrchyan et al., *Measurement of the Inclusive Jet Cross Section in pp Collisions at $\sqrt{s} = 7$ TeV*, *Phys.Rev.Lett.* **107** (2011) 132001, [[arXiv:1106.0208](#)].
- [96] B. Watt, P. Motylinski, and R. Thorne, *The Effect of LHC Jet Data on MSTW PDFs*, *Eur.Phys.J.* **C74** (2014) 2934, [[arXiv:1311.5703](#)].
- [97] **CMS Collaboration** Collaboration, V. Khachatryan et al., *Constraints on parton distribution functions and extraction of the strong coupling constant from the inclusive jet cross section in pp collisions at $\sqrt{s} = 7$ TeV*, [arXiv:1410.6765](#).
- [98] **ATLAS Collaboration** Collaboration, G. Aad et al., *Measurement of dijet cross sections in pp collisions at 7 TeV centre-of-mass energy using the ATLAS detector*, *JHEP* **1405** (2014) 059, [[arXiv:1312.3524](#)].
- [99] **CMS** Collaboration, S. Chatrchyan et al., *Top pair cross section in dileptons*, CMS-PAS-TOP-12-007.
- [100] J. Currie, A. Gehrmann-De Ridder, E. Glover, and J. Pires, *NNLO QCD corrections to jet production at hadron colliders from gluon scattering*, *JHEP* **1401** (2014) 110, [[arXiv:1310.3993](#)].
- [101] D. de Florian, P. Hinderer, A. Mukherjee, F. Ringer, and W. Vogelsang, *Approximate next-to-next-to-leading order corrections to hadronic jet production*, *Phys.Rev.Lett.* **112** (2014) 082001, [[arXiv:1310.7192](#)].
- [102] V. Bertone, S. Carrazza, and J. Rojo, *APFEL: A PDF Evolution Library with QED corrections*, *Comput.Phys.Commun.* **185** (2014) 1647–1668, [[arXiv:1310.1394](#)].
- [103] J. Campbell, K. Hatakeyama, J. Huston, F. Petriello, J. R. Andersen, et al., *Working Group Report: Quantum Chromodynamics*, [arXiv:1310.5189](#).
- [104] **The NNPDF** Collaboration, R. D. Ball et al., *Theoretical issues in PDF determination and associated uncertainties*, *Phys.Lett.* **B723** (2013) 330–339, [[arXiv:1303.1189](#)].
- [105] T. Carli, D. Clements, A. Cooper-Sarkar, C. Gwenlan, G. P. Salam, et al., *A posteriori inclusion of parton density functions in NLO QCD final-state calculations at hadron colliders: The APPLGRID Project*, *Eur.Phys.J.* **C66** (2010) 503–524, [[arXiv:0911.2985](#)].

- [106] J. Campbell and R. K. Ellis, *Next-to-leading order corrections to $W + 2\text{jet}$ and $Z + 2\text{jet}$ production at hadron colliders*, *Phys. Rev.* **D65** (2002) 113007, [[hep-ph/0202176](#)].
- [107] *Mcfm*, <http://mcfm.fnal.gov>.
- [108] Z. Nagy, *Next-to-leading order calculation of three-jet observables in hadron hadron collision*, *Phys. Rev.* **D68** (2003) 094002, [[hep-ph/0307268](#)].
- [109] T. Kluge, K. Rabbertz, and M. Wobisch, *Fast pQCD calculations for PDF fits*, [hep-ph/0609285](#).
- [110] **fastNLO** Collaboration, M. Wobisch, D. Britzger, T. Kluge, K. Rabbertz, and F. Stober, *Theory-Data Comparisons for Jet Measurements in Hadron-Induced Processes*, [arXiv:1109.1310](#).
- [111] V. Bertone, R. Frederix, S. Frixione, J. Rojo, and M. Sutton, *aMCfast: automation of fast NLO computations for PDF fits*, *JHEP* **1408** (2014) 166, [[arXiv:1406.7693](#)].
- [112] J. Alwall, R. Frederix, S. Frixione, V. Hirschi, F. Maltoni, et al., *The automated computation of tree-level and next-to-leading order differential cross sections, and their matching to parton shower simulations*, *JHEP* **1407** (2014) 079, [[arXiv:1405.0301](#)].
- [113] *Dynnlo*, <http://theory.fi.infn.it/grazzini/dy.html>.
- [114] S. Catani, G. Ferrera, and M. Grazzini, *W Boson Production at Hadron Colliders: The Lepton Charge Asymmetry in NNLO QCD*, *JHEP* **1005** (2010) 006, [[arXiv:1002.3115](#)].
- [115] S. Catani and M. Grazzini, *An NNLO subtraction formalism in hadron collisions and its application to Higgs boson production at the LHC*, *Phys.Rev.Lett.* **98** (2007) 222002, [[hep-ph/0703012](#)].
- [116] S. Catani, L. Cieri, G. Ferrera, D. de Florian, and M. Grazzini, *Vector boson production at hadron colliders: a fully exclusive QCD calculation at NNLO*, *Phys.Rev.Lett.* **103** (2009) 082001, [[arXiv:0903.2120](#)].
- [117] R. Gavin, Y. Li, F. Petriello, and S. Quackenbush, *W Physics at the LHC with FEWZ 2.1*, *Comput.Phys.Commun.* **184** (2013) 208–214, [[arXiv:1201.5896](#)].
- [118] Y. Li and F. Petriello, *Combining QCD and electroweak corrections to dilepton production in FEWZ*, *Phys.Rev.* **D86** (2012) 094034, [[arXiv:1208.5967](#)].
- [119] M. Czakon and A. Mitov, *Top++: a program for the calculation of the top-pair cross-section at hadron colliders*, [arXiv:1112.5675](#).
- [120] R. Frederix, S. Frixione, V. Hirschi, F. Maltoni, O. Mattelaer, et al., *Higgs pair production at the LHC with NLO and parton-shower effects*, *Phys.Lett.* **B732** (2014) 142–149, [[arXiv:1401.7340](#)].
- [121] P. Artoisenet, P. de Aquino, F. Demartin, R. Frederix, S. Frixione, et al., *A framework for Higgs characterisation*, *JHEP* **1311** (2013) 043, [[arXiv:1306.6464](#)].
- [122] S. Carrazza and J. Pires, *Perturbative QCD description of jet data from LHC Run-I and Tevatron Run-II*, [arXiv:1407.7031](#).

- [123] C. Carloni Calame, G. Montagna, O. Nicrosini, and A. Vicini, *Precision electroweak calculation of the production of a high transverse-momentum lepton pair at hadron colliders*, *JHEP* **0710** (2007) 109, [[arXiv:0710.1722](#)].
- [124] F. A. Berends and R. Kleiss, *Hard Photon Effects in W^\pm and Z^0 Decay*, *Z.Phys.* **C27** (1985) 365.
- [125] F. A. Berends, R. Kleiss, J. Revol, and J. Vialle, *QED Radiative Corrections and Radiative Decays of the Intermediate Weak Bosons Produced in Proton - Anti-proton Collisions*, *Z.Phys.* **C27** (1985) 155.
- [126] S. Dittmaier and M. Huber, *Radiative corrections to the neutral-current Drell-Yan process in the Standard Model and its minimal supersymmetric extension*, *JHEP* **1001** (2010) 060, [[arXiv:0911.2329](#)].
- [127] U. Baur, S. Keller, and W. Sakumoto, *QED radiative corrections to Z boson production and the forward backward asymmetry at hadron colliders*, *Phys.Rev.* **D57** (1998) 199–215, [[hep-ph/9707301](#)].
- [128] U. Baur, O. Brein, W. Hollik, C. Schappacher, and D. Wackeroth, *Electroweak radiative corrections to neutral current Drell-Yan processes at hadron colliders*, *Phys.Rev.* **D65** (2002) 033007, [[hep-ph/0108274](#)].
- [129] R. Boughezal, Y. Li, and F. Petriello, *Disentangling radiative corrections using high-mass Drell-Yan at the LHC*, *Phys.Rev.* **D89** (2014) 034030, [[arXiv:1312.3972](#)].
- [130] S. Forte, E. Laenen, P. Nason, and J. Rojo, *Heavy quarks in deep-inelastic scattering*, *Nucl. Phys.* **B834** (2010) 116–162, [[arXiv:1001.2312](#)].
- [131] J. Ablinger, A. Behring, J. Blumlein, A. De Freitas, A. von Manteuffel, et al., *The 3-Loop Pure Singlet Heavy Flavor Contributions to the Structure Function $F_2(x, Q^2)$ and the Anomalous Dimension*, [arXiv:1409.1135](#).
- [132] S. Alekhin and S. Moch, *Heavy-quark deep-inelastic scattering with a running mass*, *Phys. Lett.* **B699** (2011) 345–353, [[arXiv:1011.5790](#)].
- [133] F. Demartin, S. Forte, E. Mariani, J. Rojo, and A. Vicini, *The impact of PDF and α_s uncertainties on Higgs Production in gluon fusion at hadron colliders*, *Phys. Rev.* **D82** (2010) 014002, [[arXiv:1004.0962](#)].
- [134] F. Cascioli, P. Maierhoefer, N. Moretti, S. Pozzorini, and F. Siegert, *NLO matching for $ttbb$ production with massive b -quarks*, [arXiv:1309.5912](#).
- [135] G. D’Agostini, *On the use of the covariance matrix to fit correlated data*, *Nucl.Instrum.Meth.* **A346** (1994) 306–311.
- [136] A. D. Martin, W. J. Stirling, R. S. Thorne, and G. Watt, *Parton distributions for the LHC*, *Eur. Phys. J.* **C63** (2009) 189–285, [[arXiv:0901.0002](#)].
- [137] R. D. Ball, S. Carrazza, L. Del Debbio, S. Forte, J. Gao, et al., *Parton Distribution Benchmarking with LHC Data*, *JHEP* **1304** (2013) 125, [[arXiv:1211.5142](#)].

- [138] J. Butterworth, G. Dissertori, S. Dittmaier, D. de Florian, N. Glover, et al., *Les Houches 2013: Physics at TeV Colliders: Standard Model Working Group Report*, [arXiv:1405.1067](#).
- [139] G. P. Salam and J. Rojo, *A Higher Order Perturbative Parton Evolution Toolkit (HOPPET)*, *Comput. Phys. Commun.* **180** (2009) 120–156, [[arXiv:0804.3755](#)].
- [140] **The NNPDF Collaboration**, R. D. Ball et al., *Unbiased determination of polarized parton distributions and their uncertainties*, *Nucl.Phys.* **B874** (2013) 36–84, [[arXiv:1303.7236](#)].
- [141] **NNPDF Collaboration**, E. R. Nocera, R. D. Ball, S. Forte, G. Ridolfi, and J. Rojo, *A first unbiased global determination of polarized PDFs and their uncertainties*, *Nucl.Phys.* **B887** (2014) 276–308, [[arXiv:1406.5539](#)].
- [142] G. Altarelli, S. Forte, and G. Ridolfi, *On positivity of parton distributions*, *Nucl. Phys.* **B534** (1998) 277–296, [[hep-ph/9806345](#)].
- [143] F. Demartin, F. Maltoni, K. Mawatari, B. Page, and M. Zaro, *Higgs characterisation at NLO in QCD: CP properties of the top-quark Yukawa interaction*, *Eur.Phys.J.* **C74** (2014), no. 9 3065, [[arXiv:1407.5089](#)].
- [144] G. Altarelli, R. D. Ball, and S. Forte, *Small x Resummation with Quarks: Deep-Inelastic Scattering*, *Nucl. Phys.* **B799** (2008) 199–240, [[arXiv:0802.0032](#)].
- [145] D. J. Montana and L. Davis, *Training Feedforward Neural Networks Using Genetic Algorithms*, in *Proceedings of the 11th International Joint Conference on Artificial Intelligence - Volume 1, IJCAI’89*, (San Francisco, CA, USA), pp. 762–767, Morgan Kaufmann Publishers Inc., 1989.
- [146] A. Glazov, S. Moch, and V. Radescu, *Parton Distribution Uncertainties using Smoothness Prior*, *Phys.Lett.* **B695** (2011) 238–241, [[arXiv:1009.6170](#)].
- [147] **NNPDF Collaboration**, R. D. Ball et al., *Parton Distributions: Determining Probabilities in a Space of Functions*, [arXiv:1110.1863](#).
- [148] J. Pumplin, *Parametrization dependence and $\Delta\chi^2$ in parton distribution fitting*, *Phys.Rev.* **D82** (2010) 114020, [[arXiv:0909.5176](#)].
- [149] R. D. Ball, V. Bertone, F. Cerutti, L. Del Debbio, S. Forte, et al., *Reweighting and Unweighting of Parton Distributions and the LHC W lepton asymmetry data*, *Nucl.Phys.* **B855** (2012) 608–638, [[arXiv:1108.1758](#)].
- [150] C. Anastasiou, S. Buehler, F. Herzog, and A. Lazopoulos, *Total cross-section for Higgs boson hadroproduction with anomalous Standard Model interactions*, *JHEP* **1112** (2011) 058, [[arXiv:1107.0683](#)].
- [151] P. Skands, S. Carrazza, and J. Rojo, *Tuning PYTHIA 8.1: the Monash 2013 Tune*, *European Physical Journal* **74** (2014) 3024, [[arXiv:1404.5630](#)].
- [152] F. Caola, S. Forte, and J. Rojo, *HERA data and DGLAP evolution: Theory and phenomenology*, *Nucl.Phys.* **A854** (2011) 32–44, [[arXiv:1007.5405](#)].

- [153] M. Cacciari and N. Houdeau, *Meaningful characterisation of perturbative theoretical uncertainties*, *JHEP* **1109** (2011) 039, [[arXiv:1105.5152](#)].
- [154] E. Bagnaschi, M. Cacciari, A. Guffanti, and L. Jenniches, *An extensive survey of the estimation of uncertainties from missing higher orders in perturbative calculations*, [arXiv:1409.5036](#).
- [155] S. Forte, A. Isgur, and G. Vita, *Do we need N^3 LO Parton Distributions?*, *Phys.Lett.* **B731** (2014) 136–140, [[arXiv:1312.6688](#)].
- [156] A. Martin, R. Roberts, W. Stirling, and R. Thorne, *MRST partons and uncertainties*, [hep-ph/0307262](#).
- [157] S. Forte and J. Rojo, “Section II.2 in: J. Butterworth et al., ”Les Houches 2013: Physics at TeV Colliders: Standard Model Working Group Report”.” [arXiv:1405.1067](#), 2014.
- [158] D. Mason et al., *Measurement of the Nucleon Strange-Antistrange Asymmetry at Next-to-Leading Order in QCD from NuTeV Dimuon Data*, *Phys. Rev. Lett.* **99** (2007) 192001.
- [159] **NOMAD** Collaboration, O. Samoylov et al., *A Precision Measurement of Charm Dimuon Production in Neutrino Interactions from the NOMAD Experiment*, *Nucl.Phys.* **B876** (2013) 339–375, [[arXiv:1308.4750](#)].
- [160] **ATLAS** Collaboration, G. Aad et al., *Determination of the strange quark density of the proton from ATLAS measurements of the W, Z cross sections*, *Phys.Rev.Lett.* (2012) [[arXiv:1203.4051](#)].
- [161] A. Glazov, “Private communication, on behalf of the H1-ZEUS combination and ATLAS.”.
- [162] M. Gouzevitch, “Private communication, on behalf of CMS.”.
- [163] P. Jimenez-Delgado, *The role of the input scale in parton distribution analyses*, *Phys.Lett.* **B714** (2012) 301–305, [[arXiv:1206.4262](#)].
- [164] J. M. Campbell, J. W. Huston, and W. J. Stirling, *Hard interactions of quarks and gluons: A primer for LHC physics*, *Rept. Prog. Phys.* **70** (2007) 89, [[hep-ph/0611148](#)].
- [165] J. Rojo, *Parton distributions based on a maximally consistent dataset*, [arXiv:1409.3029](#).
- [166] M. Cacciari, G. P. Salam, and G. Soyez, *The Anti- $k(t)$ jet clustering algorithm*, *JHEP* **0804** (2008) 063, [[arXiv:0802.1189](#)].
- [167] S. Frixione, *Isolated photons in perturbative QCD*, *Phys.Lett.* **B429** (1998) 369–374, [[hep-ph/9801442](#)].
- [168] R. Frederix, S. Frixione, V. Hirschi, F. Maltoni, R. Pittau, et al., *Four-lepton production at hadron colliders: aMC@NLO predictions with theoretical uncertainties*, *JHEP* **1202** (2012) 099, [[arXiv:1110.4738](#)].
- [169] H.-L. Lai et al., *Uncertainty induced by QCD coupling in the CTEQ global analysis of parton distributions*, *Phys. Rev.* **D82** (2010) 054021, [[arXiv:1004.4624](#)].

- [170] M. Kramer, A. Kulesza, R. van der Leeuw, M. Mangano, S. Padhi, et al., *Supersymmetry production cross sections in pp collisions at $\sqrt{s} = 7$ TeV*, [arXiv:1206.2892](#).
- [171] C. Borschensky, M. Krmer, A. Kulesza, M. Mangano, S. Padhi, et al., *Squark and gluino production cross sections in pp collisions at $\sqrt{s} = 13, 14, 33$ and 100 TeV*, [arXiv:1407.5066](#).
- [172] W. Beenakker, R. Hopker, M. Spira, and P. Zerwas, *Squark and gluino production at hadron colliders*, *Nucl.Phys.* **B492** (1997) 51–103, [[hep-ph/9610490](#)].
- [173] T. Plehn. Publicly available from <http://www.thphys.uni-heidelberg.de/~plehn>.
- [174] D. Bourilkov, R. C. Group, and M. R. Whalley, *LHAPDF: PDF use from the Tevatron to the LHC*, [hep-ph/0605240](#).
- [175] T. Sjostrand, S. Mrenna, and P. Z. Skands, *A Brief Introduction to PYTHIA 8.1*, *Comput. Phys. Commun.* **178** (2008) 852–867, [[arXiv:0710.3820](#)].
- [176] S. Alekhin, S. Alioli, R. D. Ball, V. Bertone, J. Blumlein, et al., *The PDF4LHC Working Group Interim Report*, [arXiv:1101.0536](#).
- [177] S. Dawson, A. Ismail, and I. Low, *A Redux on "When is the Top Quark a Parton?"*, *Phys.Rev.* **D90** (2014) 014005, [[arXiv:1405.6211](#)].